Natural Language Processing And Understanding

110 notes across 14 week(s)

Contents

Week 0

Natural Language Processing and Understanding: Course Overview

Natural Language Processing and Understanding: Course Overview

Why This Course Exists

Machines that read, interpret, and generate human language are no longer research curiosities — they are production infrastructure. Search engines rank billions of documents, translation systems bridge languages in milliseconds, keyboards predict the next word, and large language models draft code, summarise reports, and power conversational agents.

Natural Language Processing (NLP) sits at the centre of this shift. It is the discipline that turns unstructured text into actionable signals and fluent output. For postgraduate ML practitioners, NLP literacy is among the highest-leverage skills in modern AI engineering.

---

1. NLP as the Interface Layer of Modern AI

NLP has evolved from a narrow AI subfield into the **primary interface** between humans and intelligent systems.

Application Area	NLP Role
Markets and analytics	Sentiment analysis on news and social data drives trading and brand strategy
Software development	LLMs co-author code, documentation, and test cases
Customer experience	Chatbots, ticket routing, and personalised responses at scale
Knowledge work	Summarisation, Q&A, and semantic search over enterprise documents

The ability to **process text computationally** is now a core engineering competency — not an optional specialisation.

---

2. Course Learning Arc

The curriculum progresses from linguistic foundations to state-of-the-art generative systems:

Phase 1 — Foundations

- What NLP, NLU, and NLG are and how they relate
- Linguistic concepts: morphology, ambiguity, polysemy, synonymy
- Text preprocessing: tokenisation, noise removal, stemming, lemmatisation

Phase 2 — Representing Words as Numbers

- **Sparse methods:** one-hot encoding, Bag of Words, TF-IDF

- **Dense embeddings:** Word2Vec, GloVe, static vs contextual representations
- Visualising and interpreting embedding spaces

Phase 3 — Sequential Modelling

- Sequence-to-sequence framing for language tasks
- Recurrent Neural Networks (RNNs) and their limitations
- **LSTM** and **GRU** — gated architectures that preserve long-range context
- Vanishing gradient problem and why gates matter

Phase 4 — The Transformer Revolution

- Self-attention and the *Attention Is All You Need* architecture
- Encoder-decoder and decoder-only model families (GPT, BERT)
- Hugging Face ecosystem for pretrained model access
- BERT: pre-training objectives, fine-tuning, and variants (RoBERTa, ALBERT, DistilBERT)

Phase 5 — Applied NLP

- Text classification and sentiment analysis (NLTK, BERT, Flair, spaCy)
- Topic modelling: LDA, BERTopic, GSDMM
- Large Language Models: how they generate text, output control, prompting strategies

Phase 6 — Capstone Project

- End-to-end LLM-powered application (quiz generation)
- Prompt templates, API integration, backend/frontend architecture

---

3. Hands-On Projects

Theory is paired with implementation throughout:

Project	Skills Practised
Semantic search engine	Embeddings, similarity, retrieval
Sentiment classifier	Feature engineering, transformer fine-tuning, tool comparison
Quiz app powered by LLMs	Prompt engineering, API integration, full-stack deployment

Hands-on work solidifies the **why** (linguistic and architectural reasoning) alongside the **how** (library usage and pipeline construction).

---

4. Learning Outcomes

By course completion, a student should be able to:

- Explain why human language is hard for machines and how preprocessing addresses it
- Choose between sparse and dense representations for a given task
- Describe RNN → LSTM → Transformer evolution and when each paradigm applies
- Fine-tune or prompt pretrained models for classification and generation tasks
- Design and deploy a practical NLP application using modern LLM APIs

The trajectory moves from **AI consumer** to **AI engineer** — understanding mechanisms, not just invoking APIs.

---

5. Technology Stack Preview

Layer	Tools / Concepts
Preprocessing	NLTK, spaCy, Flair, regex
Classical ML features	scikit-learn, TF-IDF, BoW
Embeddings	Gensim (Word2Vec), GloVe
Deep learning	TensorFlow/Keras (RNN, LSTM, GRU)
Transformers	Hugging Face Transformers, BERT
Topic modelling	LDA (Gensim), BERTopic
LLMs	Gemini, OpenAI APIs, Google AI Studio
Capstone	Python backend + frontend, prompt templates

Common Pitfalls / Exam Traps

- Treating NLP as only "ChatGPT usage" — classical preprocessing and representations remain foundational and examinable
- Skipping early modules to jump to transformers — attention mechanisms assume understanding of sequences, embeddings, and task framing
- Confusing **tool familiarity** with **conceptual mastery** — knowing Hugging Face API syntax ≠ explaining why BERT uses masked language modelling
- Underestimating **linguistic ambiguity** — it motivates the entire arc from rule-based systems to contextual embeddings

---

Quick Revision Summary

- NLP is the interface layer of modern AI — search, translation, sentiment, LLMs
- Course arc: foundations → sparse/dense representations → RNNs/LSTMs → transformers/BERT → applications → LLM capstone
- Sparse: one-hot, BoW, TF-IDF; Dense: Word2Vec, GloVe; Contextual: BERT, LLMs
- Sequential models (RNN, LSTM, GRU) precede transformers; vanishing gradients motivate gated architectures
- Hands-on: semantic search, sentiment classifier, LLM quiz application
- Goal: understand both the *why* behind language AI and the *how* of implementation

---

← **Previous** · **Index** · **Next** →
Practitioner-Led NLP Education: What to Expect

Practitioner-Led NLP Education: What to Expect

Intuition First

Academic NLP curricula often emphasise theory and benchmark datasets. Industry NLP work demands something additional: shipping pipelines that handle messy production text, integrating pretrained models with business APIs, and iterating under latency and cost constraints. A course led by an experienced **AI engineer** bridges that gap — grounding concepts in patterns seen across real product deployments.

This note frames what practitioner-led instruction typically emphasises and how students should engage to maximise transfer to professional work.

---

1. Industry Experience vs Pure Academia

Dimension	Academic Focus	Practitioner Focus
Data	Clean benchmark corpora (IMDB, CoNLL)	Noisy user-generated text, domain jargon, PII
Success metric	Accuracy, F1, BLEU on test sets	Latency, cost per request, maintainability, user satisfaction
Model choice	State-of-the-art novelty	Fit-for-purpose: VADER vs BERT vs LLM by constraint
Delivery	Notebook experiments	APIs, monitoring, versioning, failure handling

A practitioner-led NLP course does not replace theoretical rigour — it **contextualises** it. Every architecture choice (TF-IDF vs embedding vs fine-tuned BERT vs prompted LLM) maps to a deployment trade-off students will face in industry.

---

2. Cross-Industry Exposure

Experienced NLP practitioners typically work across multiple domains:

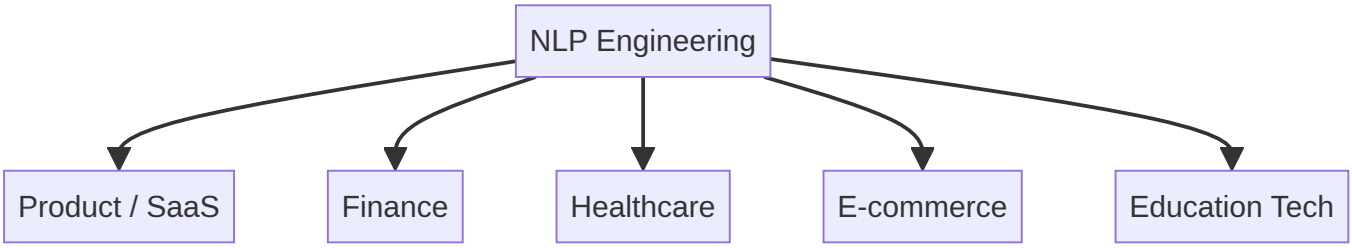

Each domain imposes different constraints:

- **Finance** — compliance, audit trails, low tolerance for hallucination
- **Healthcare** — specialised vocabulary, HIPAA-style privacy, high precision for NER
- **Product/SaaS** — scale, multilingual support, A/B testing of model versions
- **E-commerce** — real-time sentiment, search relevance, personalisation

Exposure to varied industries teaches **transferable patterns** (preprocessing pipelines, evaluation discipline) while highlighting where domain adaptation is non-negotiable.

---

3. Mentorship-Oriented Learning

Effective practitioner instruction often includes:

- **Worked examples** from production-adjacent scenarios, not only textbook datasets
- **Tool literacy** — NLTK, spaCy, Flair, Hugging Face, LLM APIs — with guidance on when each fits
- **Debugging mindset** — why a sentiment model fails on sarcasm, why LDA topics are incoherent, why prompts drift
- **Career-relevant framing** — how NLP skills compose with MLOps, data engineering, and full-stack delivery

Students benefit by treating labs as **portfolio artefacts**, not checkbox exercises.

---

4. How to Engage With This Course

Practice	Rationale
Implement, don't only read	NLP intuition comes from seeing tokenisation edge cases and embedding failures firsthand
Compare approaches	Week 10's NLTK vs BERT vs Flair vs spaCy comparison mirrors real tool-selection decisions
Build the capstone end-to-end	QuizGenius exercises prompt design, API integration, and architecture — skills directly transferable to LLM product work
Document trade-offs	For each technique, note accuracy, speed, cost, and interpretability — the dimensions practitioners optimise

---

5. From Student to Practitioner

The course trajectory aligns with a common professional path:

Practitioner-led delivery accelerates the final steps — students see how pretrained models and APIs compress time-to-value while still requiring solid fundamentals for debugging and customisation.

Common Pitfalls / Exam Traps

- Assuming industry experience means skipping theory — exams test **conceptual understanding** (LSTM gates, attention, LDA assumptions), not just API usage
- Believing one tool wins everywhere — practitioners constantly trade off VADER speed vs BERT accuracy vs LLM flexibility
- Ignoring **mentorship value** — asking why a design choice was made beats memorising code snippets
- Conflating years of experience with infallibility — validate recommendations against documentation and reproducible experiments

---

Quick Revision Summary

- Practitioner-led NLP education emphasises production trade-offs alongside theory
- Industry work spans domains with different compliance, vocabulary, and latency needs
- Mentorship adds worked examples, tool selection guidance, and debugging mindset
- Engage by implementing, comparing approaches, and building the capstone as a portfolio piece
- Course path mirrors professional growth: fundamentals → pipelines → models → transformers → LLM deployment
- Balance practitioner insights with exam-focused conceptual mastery

---

← **Previous** · **Index** · **Next** →

Week 1

Week 1

Introduction to Natural Language Processing

Introduction to Natural Language Processing

The Core Question

How do machines work with human language? Text and speech are everywhere — search queries, support tickets, clinical notes, social posts — yet computers natively operate on numbers, not words. **Natural Language Processing (NLP)** is the discipline that bridges that gap: it gives machines the ability to read, interpret, and respond in human language.

This opening module establishes the vocabulary, scope, and linguistic foundations needed for every technique that follows — from tokenisation and tagging to transformers and large language models.

---

1. What This Module Covers

The module is organised around three pillars:

Pillar	Focus
Core NLP concepts	What NLP is, what problems it solves, and where it appears in production systems
NLP, NLU, and NLG	How understanding and generation relate under the NLP umbrella
Linguistic foundations	Morphology, ambiguity, polysemy, and synonymy — why language is hard for machines

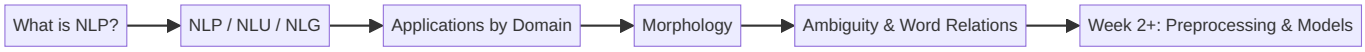

---

2. Why Language Is a Distinct Engineering Problem

Unlike structured database fields, natural language is:

- **Unstructured** — no fixed schema; meaning emerges from word order, punctuation, and context
- **Ambiguous** — the same surface form can carry multiple interpretations
- **Context-dependent** — word meaning shifts with domain, culture, and surrounding text

NLP systems must therefore handle noise, variation, and implicit meaning — not just pattern-match strings.

---

3. Module Learning Path

1. Define NLP and its high-level goals (read, extract, generate)
2. Distinguish **NLU** (understanding) from **NLG** (generation)
3. Survey real-world applications across search, healthcare, finance, and education
4. Introduce **morphology** — how words are built from roots and affixes
5. Examine **ambiguity**, **polysemy**, and **synonymy** — core reasons embedding-based methods eventually replace naive dictionary lookups

---

Common Pitfalls / Exam Traps

- Treating NLP as only "chatbots" — it spans search ranking, fraud monitoring, clinical summarisation, and more
- Assuming NLP = NLG — **NLU** (classification, NER, sentiment) is equally central and often the upstream step
- Believing rule-based dictionary lookup suffices — **context** is required to resolve polysemy and domain-specific senses
- Ignoring linguistic foundations — morphology directly motivates stemming, lemmatisation, and vocabulary reduction in preprocessing pipelines

Quick Revision Summary

- NLP enables machines to understand, process, and interact using human language
- Language is unstructured, ambiguous, and context-dependent — unlike typical tabular data
- This module covers NLP scope, NLU/NLG relationship, applications, morphology, and ambiguity
- Morphology studies word structure; it motivates preprocessing and vocabulary management
- Polysemy and synonymy explain why context-aware representations (embeddings) are essential
- Foundation concepts here underpin every later week: preprocessing, tagging, embeddings, transformers, and LLMs

---

← **Previous** · **Index** · **Next** →
What Is Natural Language Processing?

What Is Natural Language Processing?

Intuition First

Every day, billions of people type queries, send messages, and speak to devices. Computers, however, only manipulate bits and floating-point numbers. **Natural Language Processing (NLP)** is the field of artificial intelligence devoted to closing that representational gap — enabling machines to **understand**, **process**, and **interact** using human language as humans do.

NLP sits at the intersection of computer science, artificial intelligence, and linguistics. Its goal is not merely to store text, but to extract meaning and act on it.

---

1. Formal Definition

Natural Language Processing (NLP) is the branch of AI that focuses on enabling computers to:

1. **Read and interpret** text or speech
2. **Extract** structured, useful information from unstructured language
3. **Generate** meaningful responses in natural language

At a high level, NLP bridges **human communication** and **machine computation**.

---

2. Why Human Language Is Hard

Human language carries properties that make naive string processing insufficient:

Property	Meaning	Machine Impact
Unstructured	No rigid schema; free-form sentences	Requires parsing, segmentation, and feature extraction
Ambiguous	Multiple valid interpretations per utterance	Context and world knowledge needed to disambiguate
Context-dependent	Meaning shifts with speaker, domain, prior sentences	Models must track discourse and situational cues

These properties explain why NLP pipelines are multi-stage (tokenise → tag → embed → classify/generate) rather than single-regex solutions.

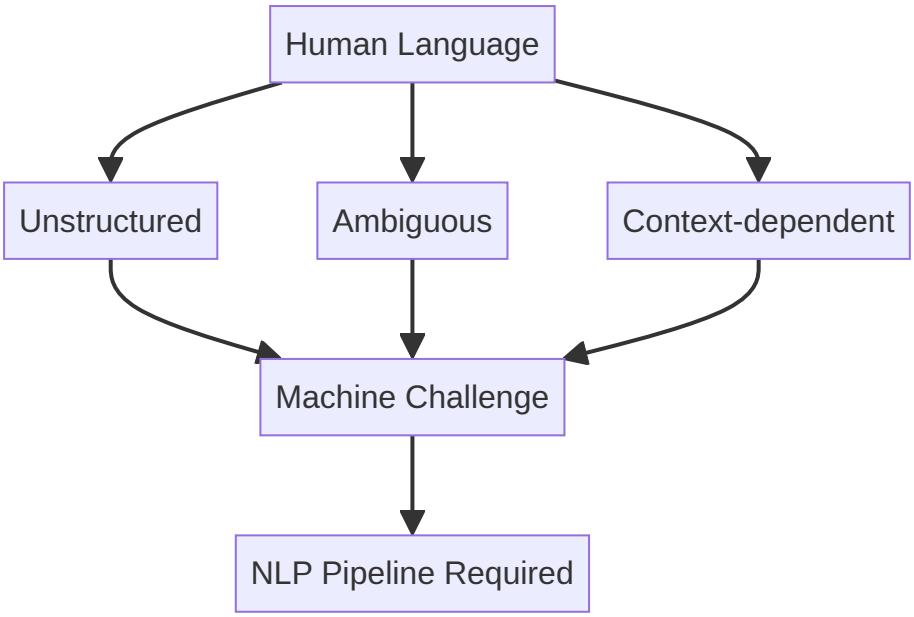

---

3. What NLP Aims to Achieve

At the highest level, NLP systems pursue three capabilities:

- **Read and interpret** — assign structure and meaning to raw text
- **Extract information** — pull entities, intents, sentiments, and facts into machine-usable form
- **Generate responses** — produce human-readable output (summaries, answers, translations)

These map to downstream tasks: search ranking, spam detection, dialogue, summarisation, and machine translation.

---

4. Everyday NLP Applications

Domain	Example System	NLP Task
Search	Google, Bing	Query understanding, document ranking
Voice	Siri, Alexa	Speech recognition + language understanding
Email	Gmail spam filter	Text classification (spam vs not spam)
Mobile keyboards	Next-word prediction	Language modelling
Generative AI	ChatGPT, Gemini	Large-scale language understanding and generation

Each application combines multiple NLP sub-tasks — tokenisation, classification, generation — often in a single production pipeline.

Common Pitfalls / Exam Traps

- Defining NLP as only "text generation" — **understanding and extraction** are equally core
- Ignoring **ambiguity** when designing systems — rule-only approaches fail on real corpora
- Confusing NLP with **speech recognition** — ASR converts audio to text; NLP operates on the resulting (or typed) language
- Assuming one model solves all tasks — production stacks combine specialised components (retrieval, classification, generation)

Quick Revision Summary

- NLP = AI field enabling machines to understand, process, and generate human language
- It spans computer science, AI, and linguistics
- Language is unstructured, ambiguous, and context-dependent — the root difficulty
- High-level goals: read/interpret, extract information, generate responses
- Real examples: search engines, voice assistants, spam filters, autocomplete, LLM chat apps
- NLP bridges human communication and machine computation

← **Previous** · **Index** · **Next** →
NLP, NLU, and NLG: Understanding the Relationship

NLP, NLU, and NLG: Understanding the Relationship

Intuition First

Not every language task asks the same question. Sometimes the machine must **figure out what a human meant**; other times it must **say something useful back**. These two directions — comprehension and production — are the twin engines of practical NLP systems.

Natural Language Processing (NLP) is the umbrella. Under it sit **Natural Language Understanding (NLU)** and **Natural Language Generation (NLG)**. Most production applications use both, often in sequence.

---

1. The Hierarchy

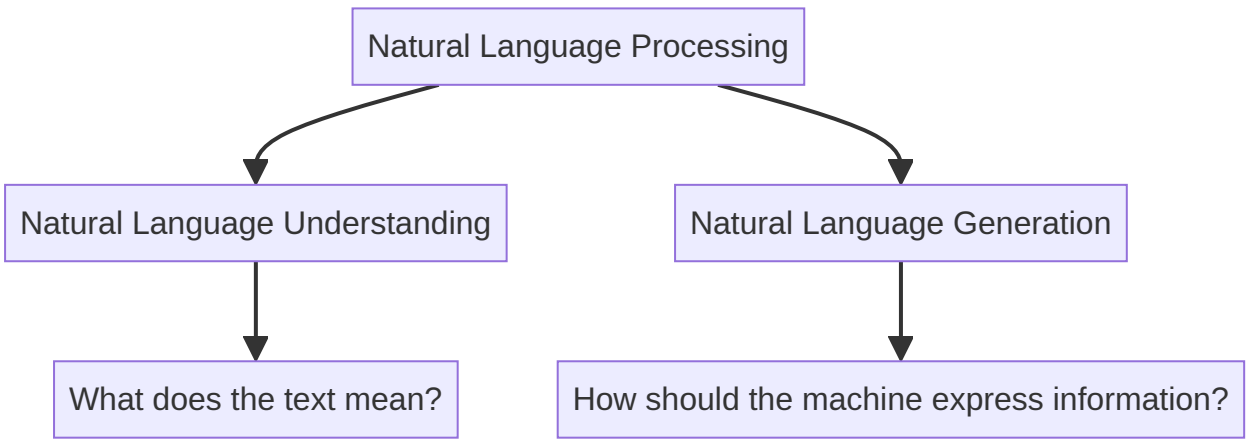

Term	Core Question	Direction
NLP	How do we process human language computationally?	Umbrella — includes both understanding and generation
NLU	What does the input text **mean**?	Input → structured meaning
NLG	How should the machine **express** information?	Structured data → human-readable text

---

2. Natural Language Understanding (NLU)

NLU focuses on **comprehending** text: extracting intent, entities, sentiment, and other semantic signals.

Primary concern: *What does the text mean?*

Typical NLU Tasks

- **Sentiment analysis** — classify opinion polarity (positive, negative, neutral)
- **Named Entity Recognition (NER)** — locate and type entities (person, organisation, date)
- **Text classification** — assign documents or utterances to predefined categories
- **Intent detection** — map user utterances to actions in dialogue systems

NLU transforms unstructured language into **machine-actionable representations** — labels, entity spans, embedding vectors, or structured slots.

---

3. Natural Language Generation (NLG)

NLG focuses on **producing** human-like language from data or internal representations.

Primary concern: *How can a machine communicate information clearly in human language?*

Typical NLG Tasks

- **Text summarisation** — condense long documents into concise summaries
- **Chatbots and question answering** — formulate natural responses to user queries
- **Report generation** — turn analytics or database records into narrative prose
- **Machine translation** — render content in another human language

NLG transforms structured or semi-structured input into **fluent, readable output**.

---

4. How NLU and NLG Work Together

Real-world systems rarely stop at understanding or generation alone.

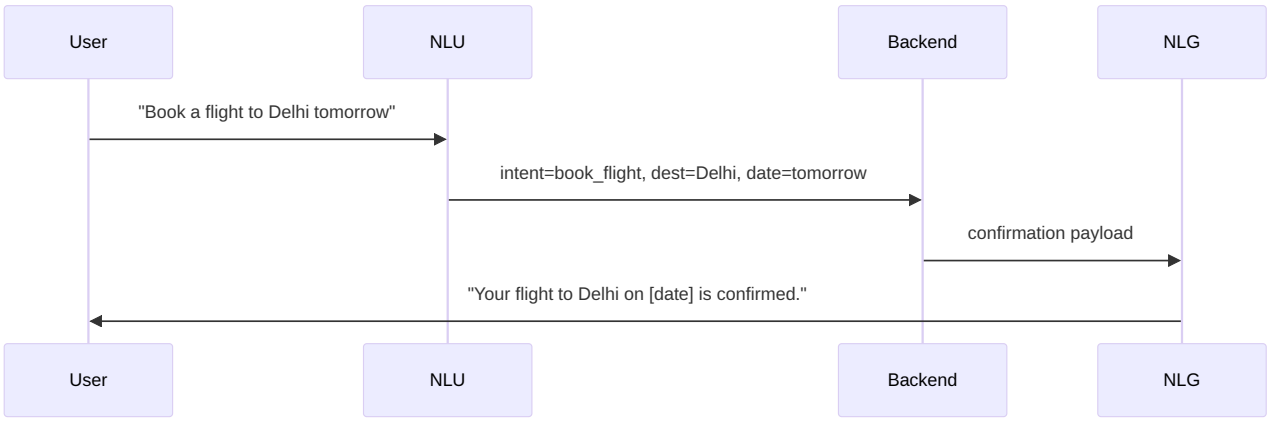

Stage	Component	Role
Input	NLU	Parse intent and slots from user text
Reasoning	Backend / KB	Execute business logic, retrieve data
Output	NLG	Render results as natural language

Examples: customer-support bots (classify ticket → draft reply), search assistants (understand query → summarise results), clinical tools (extract entities from notes → generate discharge summary).

---

5. Comparison Table

Aspect	NLU	NLG
Input	Human language (text/speech)	Structured data, prompts, or internal state
Output	Labels, entities, embeddings, parsed meaning	Sentences, paragraphs, dialogue turns
Evaluation	Accuracy, F1, entity-level metrics	Fluency, coherence, factual correctness
Example stack	BERT classifier, spaCy NER	GPT-style decoder, template + LLM hybrid

---

Common Pitfalls / Exam Traps

- Claiming NLU and NLG are disjoint — **most applications combine both** (understand → act → respond)
- Equating NLG only with chatbots — summarisation, translation, and report generation are equally NLG
- Treating NLP as synonymous with NLU — NLP is the **superset**
- Forgetting that modern LLMs blur the boundary — a single decoder-only model can perform both understanding-style and generation-style tasks, but the conceptual distinction remains useful for system design

---

Quick Revision Summary

- NLP is the umbrella; NLU and NLG are its two major branches
- NLU asks *what does the text mean?* — sentiment, NER, classification, intent
- NLG asks *how should the machine express information?* — summarisation, QA, reports, translation
- NLU = understanding; NLG = generation
- Production pipelines typically chain NLU → backend logic → NLG
- Conceptual split remains valuable even when one model handles multiple tasks

---

← **Previous** · **Index** · **Next** →
Applications of Natural Language Processing by Domain

Applications of Natural Language Processing by Domain

Intuition First

NLP is not a single product — it is a **toolkit** applied wherever unstructured text must be understood or produced at scale. The same underlying techniques (classification, entity extraction, summarisation, retrieval) reappear across industries; only the data, labels, and compliance constraints change.

Mapping applications by **domain** clarifies which NLP tasks matter in practice and how they compose into production systems.

---

1. Application Landscape Overview

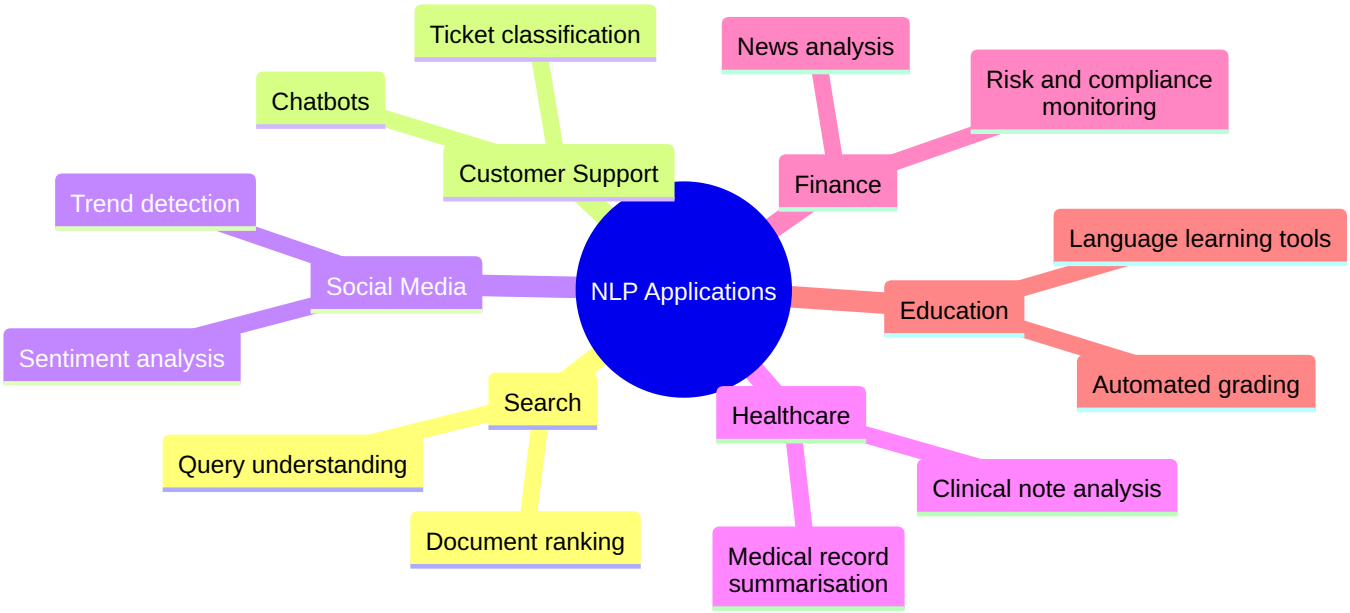

---

2. Search and Information Retrieval

Search engines must interpret vague, misspelled, or conversational queries and return relevant documents from billions of pages.

Application	What It Does	NLP Techniques
Query understanding	Parse intent, expand synonyms, correct spelling	Tokenisation, embeddings, intent classification
Document ranking	Order results by relevance to query	TF-IDF, neural rankers, transformer encoders

Example: A user searches *"cheapest GPU for LLM inference"* — query understanding maps to product-search intent; ranking uses textual similarity and click signals.

---

3. Customer Support

Support organisations process high volumes of repetitive, semi-structured text (tickets, chats, emails).

Application	What It Does	NLP Techniques
Chatbots	Answer FAQs, route complex issues	Intent detection (NLU) + response generation (NLG)
Ticket classification	Auto-label urgency, product area, sentiment	Text classification, NER for product mentions

Example: A SaaS platform classifies incoming tickets as *billing*, *bug*, or *feature request* before human agents see them — reducing mean time to resolution.

---

4. Social Media Analysis

Social platforms generate continuous streams of short, informal, multimodal text.

Application	What It Does	NLP Techniques
Sentiment analysis	Gauge public opinion toward brands, events, policies	Polarity classification, aspect-based sentiment
Trend detection	Surface emerging topics and hashtags	Topic modelling, clustering, time-series aggregation

Example: A brand monitors Twitter/X during a product launch — sentiment dashboards flag negative spikes tied to a specific feature complaint.

---

5. Healthcare

Clinical and administrative text is dense, domain-specific, and regulated.

Application	What It Does	NLP Techniques
Clinical note analysis	Extract diagnoses, medications, procedures	NER, relation extraction, medical ontologies
Medical record summarisation	Condense lengthy charts for handoffs	Abstractive/extractive summarisation

Example: An NLP pipeline scans discharge summaries to flag patients at readmission risk — entities and temporal relations feed downstream predictive models.

---

6. Finance

Financial institutions consume news, filings, earnings calls, and internal compliance documents.

Application	What It Does	NLP Techniques
News analysis	Extract market-moving events from headlines	NER, event extraction, sentiment
Risk and compliance monitoring	Detect policy violations in communications	Classification, anomaly detection on text

Example: A trading desk NLP system tags merger-related headlines in real time to trigger analyst review.

---

7. Education

Educational technology applies NLP to both **assessment** and **personalised learning**.

Application	What It Does	NLP Techniques
Automated grading	Score short answers and essays	Similarity to rubrics, coherence metrics, classifier ensembles
Language learning tools	Correct grammar, suggest vocabulary	POS tagging, error detection, generation of exercises

Example: A MOOC platform auto-grades thousands of short-answer responses using semantic similarity to reference answers — scaling human review.

---

8. Cross-Domain Pattern

Despite surface differences, most domain deployments follow a common architecture:

The **NLU task** (classify, extract, rank) and optional **NLG step** (summarise, reply) vary by domain; preprocessing and evaluation discipline stay constant.

Common Pitfalls / Exam Traps

- Listing applications without linking to **underlying NLP tasks** — exam questions often ask which technique (NER vs classification vs summarisation) fits a scenario
- Assuming one model generalises across domains — **medical NER** requires domain-specific corpora and ontologies
- Ignoring **compliance** in finance and healthcare — accuracy alone is insufficient; auditability and privacy matter
- Confusing **trend detection** with sentiment — trends are topical/clustering problems; sentiment is polarity classification

---

Quick Revision Summary

- NLP applications span search, support, social media, healthcare, finance, and education
- Search: query understanding + document ranking
- Customer support: chatbots + ticket classification
- Social media: sentiment analysis + trend detection
- Healthcare: clinical note analysis + record summarisation
- Finance: news analysis + risk/compliance monitoring
- Education: automated grading + language learning tools
- Most deployments follow preprocess → NLU → business logic → optional NLG

← **Previous** · **Index** · **Next** →
Morphology of Text: Stems, Roots, and Affixes

Morphology of Text: Stems, Roots, and Affixes

Intuition First

Consider the words *running*, *runner*, *runs*, and *ran*. Humans instantly recognise these as variants of the same core idea — *run*. **Morphology** is the linguistic study of how words are built from smaller meaningful parts. For NLP systems, this matters because treating every surface form as a unique token explodes vocabulary size and prevents generalisation.

Morphology directly motivates text preprocessing choices: stemming, lemmatisation, and subword tokenisation all exist to exploit word-internal structure.

---

1. What Is Morphology?

Morphology is the branch of linguistics that studies:

- The **internal structure** of words
- How words are **formed** from morphemes (minimal meaning-bearing units)
- **Relationships** between word variants (inflection, derivation)

In NLP, morphology answers: *Why does the same concept appear as many different strings, and how should a model treat them?*

---

2. Building Blocks of Words

Unit	Definition	Example
Root	Core lexical meaning; often not a standalone word in English	`struct` in *structure*, *construct*, *destruct*
Stem	Reduced form produced by chopping affixes (algorithmic, not always linguistic)	*comput* from *computing*, *computed*, *computer*
Affix	Prefix or suffix attached to a root/stem	*un-* (prefix), *-ing* (suffix) in *unhappiness*
Lemma	Dictionary base form	*run* for *running*, *ran*, *runs*

3. Inflection vs Derivation

Process	Changes	Example	Same part of speech?
Inflection	Grammatical features (tense, number, case)	*walk* → *walked*, *walks*	Yes
Derivation	Creates new word, often new POS	*happy* → *happiness* (adj → noun)	Often no

NLP pipelines must decide whether to collapse inflected forms (usually yes for IR/classification) while being cautious with derivations (*compute* vs *computer* are related but not identical).

4. Why Morphology Matters for NLP

Vocabulary explosion

Without morphological awareness, a corpus containing *run*, *runs*, *running*, and *ran* allocates four separate dimensions in a bag-of-words vector — sparse and redundant.

Generalisation

Models that recognise shared roots generalise better to **unseen inflections**. A spam filter trained on *"click"* should ideally recognise *"clicking"* and *"clicked"*.

Preprocessing design

Morphology motivates two core preprocessing operations:

Operation	Approach	Trade-off
Stemming	Heuristic chop (Porter, Snowball)	Fast; may produce non-words (*studies* → *studi*)
Lemmatisation	Dictionary + POS lookup	Slower; linguistically correct (*studies* → *study*)

Both reduce vocabulary size and align with how humans perceive word families.

5. Morphology and Modern Tokenisation

Even transformer-era models use morphology implicitly:

- **Subword tokenisation** (BPE, WordPiece) splits rare words into frequent morpheme-like pieces — *unhappiness* → `un`, `happi`, `ness`
- This handles **out-of-vocabulary** words without treating every inflection as atomic

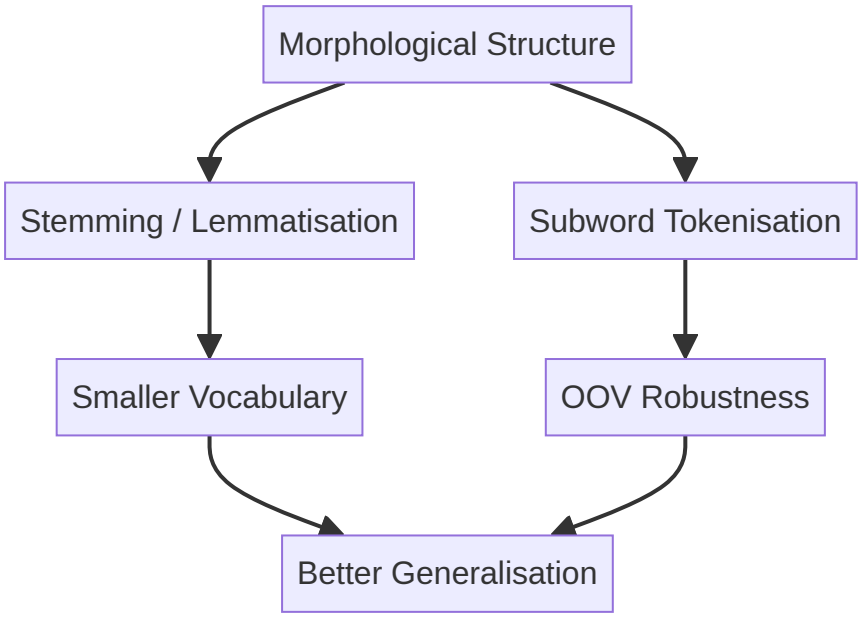

Common Pitfalls / Exam Traps

- Confusing **stem** and **lemma** — stems are algorithmic truncations; lemmas are canonical dictionary forms
- Assuming stemming always improves performance — over-stemming can merge unrelated words (*university* / *universe* with naive Porter)
- Ignoring **POS** in lemmatisation — *running* as verb → *run*; as noun (a running) → *running*
- Believing transformers eliminate morphology concerns — subword splitting is itself a morphological strategy

---

Quick Revision Summary

- Morphology studies word internal structure and word-formation relations
- Roots carry core meaning; affixes modify it; lemmas are dictionary base forms
- Inflection adjusts grammar; derivation often changes word class
- Morphology motivates vocabulary reduction and better generalisation in NLP
- Stemming is fast/heuristic; lemmatisation is slower/linguistically grounded
- Subword tokenisation in modern models handles morphology via learned splits

---

← **Previous** · **Index** · **Next** →
Ambiguity in Language: Polysemy and Synonymy

Ambiguity in Language: Polysemy and Synonymy

Intuition First

Humans resolve ambiguous language effortlessly using context, tone, and world knowledge. Machines start with only character sequences. **Ambiguity** — multiple valid interpretations for the same surface form — is therefore one of the central challenges in NLP.

Two especially important phenomena are **polysemy** (one word, many related meanings) and **synonymy** (many words, similar meanings). Together they explain why dictionary lookups fail and why **context-aware embeddings** became the dominant representation strategy.

---

1. Linguistic Ambiguity

Ambiguity occurs when a word, phrase, or sentence admits more than one valid interpretation.

Humans disambiguate using:

- Surrounding words and sentences
- Domain and situational context
- Pragmatic and cultural knowledge

Machines lack this unless explicitly modelled — making ambiguity a first-class engineering problem.

Classic Example: *"I saw her duck."*

Interpretation	Meaning
Verb reading	I saw her bend down quickly (duck as verb)
Noun reading	I saw her pet bird (duck as animal)

Same tokens, different syntactic parses and semantic frames — the model must choose correctly.

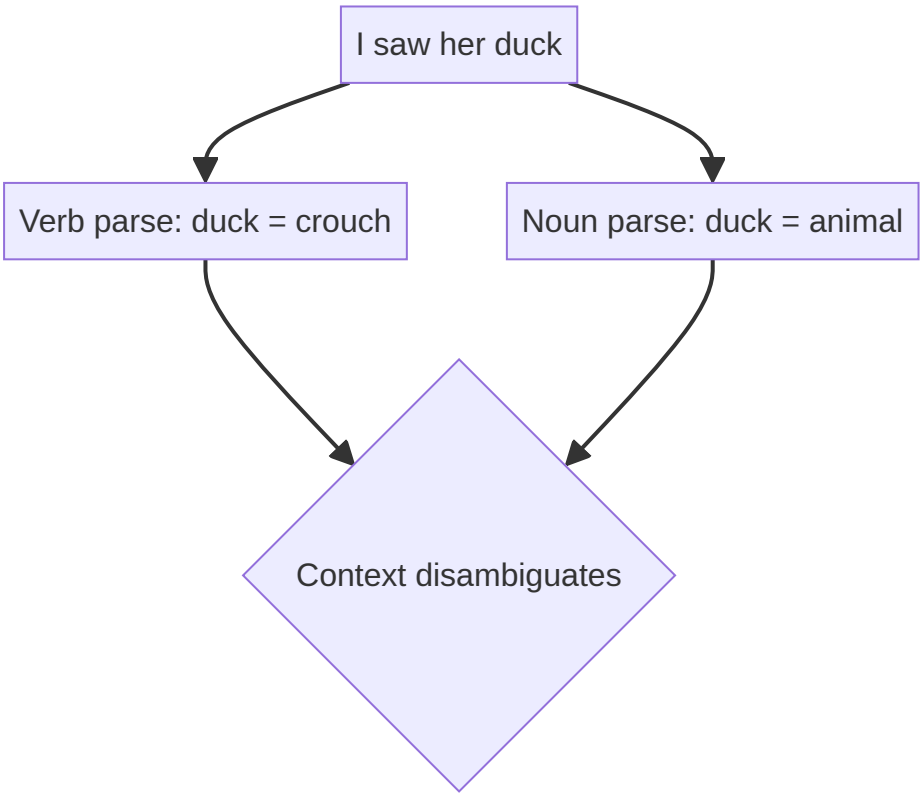

2. Levels of Ambiguity

Ambiguity is not limited to single words:

Level	Example	Challenge
Word	*bank* (financial vs river)	Lexical disambiguation
Phrase	*"telescope man"* (man with telescope vs man visible through telescope)	Attachment ambiguity
Sentence	*"I saw her duck"*	Structural/syntactic ambiguity
Discourse	Pronoun *"it"* referring to multiple antecedents	Coreference resolution

Rule-based systems that map each token to a single dictionary sense struggle at every level.

3. Polysemy

Polysemy is when a **single word** has multiple **related** meanings — senses share an etymological or conceptual connection.

Word	Sense 1	Sense 2	Sense 3
bank	Financial institution	River edge	—
paper	Writing material	Academic article	—
run	Jog physically	Manage (*run a team*)	Execute (*run a program*)

Polysemy is **extremely common** in natural language. The correct sense is almost always determined by **context**, not the word alone.

Implication for NLP

- Static one-hot or dictionary lookups assign one sense per type — insufficient
- Contextual embeddings (ELMo, BERT) produce **different vectors for the same word in different sentences**
- Word Sense Disambiguation (WSD) remains an explicit sub-task in knowledge-heavy pipelines

4. Synonymy

Synonymy is when **different words** express **similar or equivalent** meanings.

Word A	Word B	Relation
big	large	Near-synonyms
buy	purchase	Near-synonyms
happy	joyful	Near-synonyms

Synonyms are rarely perfectly interchangeable — register, collocations, and nuance differ — but they are close enough to confuse naive string matching.

Challenges Synonymy Creates

Challenge	Explanation
Vocabulary sparsity	*buy* and *purchase* occupy separate dimensions in one-hot/BoW — splitting statistical evidence
Matching difficulty	Retrieval systems miss relevant documents that use different but equivalent terms
Training data fragmentation	Supervised labels spread across lexical variants

This is a primary motivation for **embedding-based representations**: words with similar usage contexts map to nearby points in vector space.

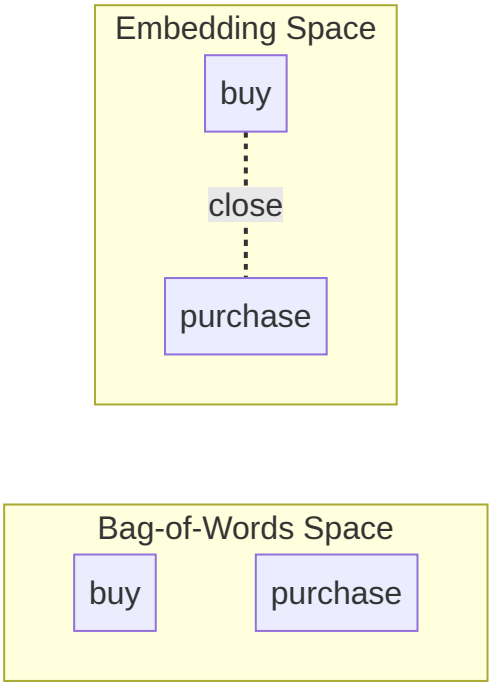

---

5. Why Traditional Rule-Based Systems Struggle

Limitation	Effect
No explicit context	Cannot pick among polysemous senses
Fixed lexicon	Misses synonyms and paraphrases
Domain blindness	*"tablet"* means medicine in healthcare, device in retail
Cultural variation	Idioms and regional usage break hand-crafted rules

Statistical and neural methods learn sense distinctions from **co-occurrence patterns** in large corpora rather than from static rules alone.

---

6. Polysemy vs Synonymy

Aspect	Polysemy	Synonymy
Pattern	One word → many meanings	Many words → similar meaning
Relation among senses/terms	Related senses of same lemma	Distinct words, overlapping semantics
Core problem	Which sense applies?	How to unify equivalent expressions?
Representation fix	Contextual embeddings, WSD	Word2Vec, GloVe, distributional similarity

---

Common Pitfalls / Exam Traps

- Treating polysemy and homonymy as identical — **homonyms** (*bank* river vs bank money) are historically unrelated; polysemous senses are related
- Assuming synonyms are perfectly interchangeable in all contexts — *big* vs *large* have different collocations (*big mistake* vs ?*large mistake*)
- Believing stemming fixes synonymy — *purchase* and *buy* share no common stem
- Ignoring **sentence-level** ambiguity when only studying word senses — exam scenarios often use syntactic ambiguity (*"I saw her duck"*)

Quick Revision Summary

- Ambiguity = multiple valid interpretations; humans use context, machines must learn it
- Ambiguity exists at word, phrase, sentence, and discourse levels
- Polysemy: one word, multiple related meanings (*bank*, *run*, *paper*) — context selects sense
- Synonymy: different words, similar meaning (*big/large*, *buy/purchase*) — causes sparsity and matching issues
- Dictionary lookups alone are insufficient for both phenomena
- Embedding-based and contextual representations address polysemy and synonymy via distributional context
- Rule-based systems fail on implicit context, domain senses, and lexical variation

---

← **Previous** · **Index** · **Next** →

Week 2

Week 2

Text Preprocessing and Cleaning: Foundations for NLP

Text Preprocessing and Cleaning: Foundations for NLP

Why Raw Text Is Not Machine-Ready

Human language is messy: inconsistent spelling, informal grammar, HTML markup in web scrapes, emojis in social posts, and domain-specific jargon. Machine learning models and NLP algorithms operate on structured numerical representations — they cannot directly consume unstructured prose.

Text preprocessing transforms raw text into a form algorithms can process reliably. **Text cleaning** removes non-linguistic noise; **text preprocessing** applies linguistic normalisation. The distinction matters because cleaning without normalisation leaves surface variation intact, while normalisation without cleaning feeds garbage tokens into downstream steps.

Almost every NLP system — from classical TF-IDF classifiers to modern LLM pipelines — depends on some preprocessing stage, even when it is hidden inside a tokenizer API.

Text Cleaning vs Text Preprocessing

Aspect	Text Cleaning	Text Preprocessing
Goal	Remove unwanted elements	Normalise linguistic form
Examples	HTML tags, emojis, extra whitespace	Tokenisation, stemming, lemmatisation
Analogy	Washing vegetables	Chopping and seasoning for a recipe
When skipped	Structured logs with clean fields	Already tokenised subword inputs

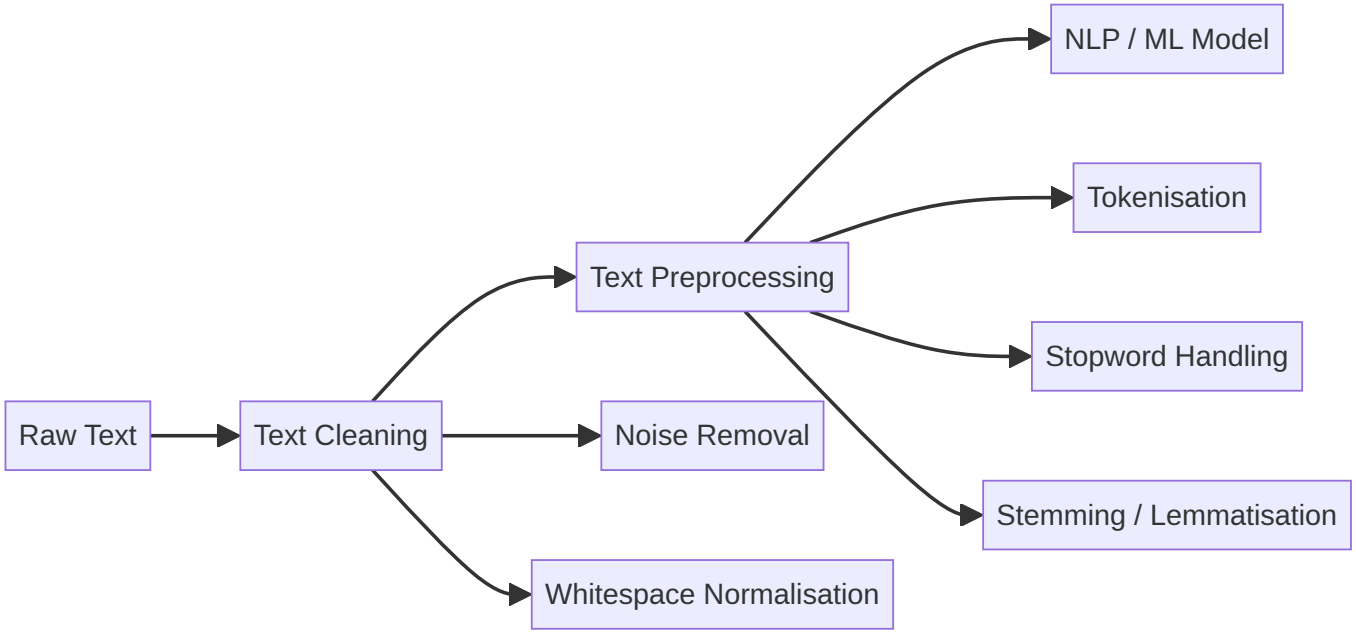

Real-World Context

- **Customer support chatbots** must strip HTML from ticket bodies before intent classification.
- **Search engines** normalise query terms so "running," "runs," and "ran" match the same documents.
- **LLM fine-tuning pipelines** still apply cleaning when ingesting PDFs, web pages, or OCR output — garbage tokens inflate vocabulary and degrade convergence.

Module Scope

This module covers both cleaning and preprocessing as a unified foundation:

1. Tokenisation — breaking text into units
2. Noise removal — regex-based cleaning
3. Stopword removal — filtering high-frequency function words
4. Stemming and lemmatisation — morphological normalisation
5. Pipeline construction — ordering steps for reproducible workflows

Common Pitfalls / Exam Traps

- Treating **cleaning and preprocessing as identical** — exams may ask which step removes HTML vs which reduces "running" to "run"
- Assuming **one universal pipeline fits all tasks** — sentiment analysis keeps negation words; NER needs minimal preprocessing

- **Over-cleaning** — removing punctuation that carries meaning (e.g., "?" in QA systems)
- Believing modern LLMs make preprocessing **obsolete** — raw corpora for training still require cleaning at scale

---

Quick Revision Summary

- Raw human text is unstructured; NLP requires machine-readable form first
- Cleaning removes noise (HTML, symbols, whitespace); preprocessing normalises linguistics
- Both steps underpin classical and modern NLP pipelines
- Task-aware design is essential — no universal cleaning rule
- Module path: tokenisation → noise removal → stopwords → stemming/lemmatisation → pipeline

---

← **Previous** · **Index** · **Next** →
Tokenisation: Splitting Text into Processable Units

Tokenisation: Splitting Text into Processable Units

The Core Idea

Tokenisation is the process of breaking raw text into smaller units called **tokens**. Tokens can represent words, sentences, subwords, or individual characters. Every NLP pipeline — from simple word counting to transformer models — begins here.

If tokenisation is wrong, every downstream step inherits the error: wrong POS tags, missed entities, skewed TF-IDF weights.

Granularity Levels

Level	Unit	Typical Use
Word	Space/punctuation-delimited tokens	Bag-of-words, classical NLP
Sentence	Clause/period boundaries	Summarisation, document segmentation
Subword	BPE/WordPiece fragments	BERT, GPT tokenizers
Character	Individual letters	Spell-checking, very low-resource languages

Why Tokenisation Is Non-Trivial

Natural language resists naive splitting:

- **Punctuation ambiguity** — "U.S.A." vs "USA"; decimals vs sentence endings
- **Hyphenation** — "state-of-the-art" as one token or four
- **Numbers and symbols** — "$1.2B", version strings, URLs
- **Emojis and Unicode** — multi-codepoint sequences

- **Language and domain** — Chinese has no spaces; biomedical text has chemical formulae

There is **no single correct tokenisation** for all applications. The choice depends on language, domain, and downstream task.

Implementation with NLTK

NLTK provides two primary functions after downloading the `punkt` resource:

Word tokenisation — `word_tokenize(text)` splits on word boundaries while preserving punctuation as separate tokens.

Sentence tokenisation — `sent_tokenize(text)` splits on sentence boundaries using a pre-trained Punkt model.

Example flow:

```
import nltk
nltk.download('punkt')
from nltk.tokenize import word_tokenize, sent_tokenize

text = "Natural language processing is fascinating."
word_tokenize(text)
# ['Natural', 'language', 'processing', 'is', 'fascinating', '.']

sent_tokenize("First sentence. Second sentence.")
# ['First sentence.', 'Second sentence.']
```

Real-World Considerations

- **Production search** — query tokenisation must match index tokenisation exactly
- **Multilingual SaaS** — different Punkt models or spaCy pipelines per language
- **Transformer APIs** — subword tokenisation (BPE) differs from NLTK word splits; comparing counts across tokenisers is invalid

---

Common Pitfalls / Exam Traps

- Assuming **split on whitespace** is sufficient — fails on "don't", "U.K.", URLs
- Using **word tokenisation for BERT** — transformer models require their own subword tokenisers
- Forgetting to **download punkt** before calling NLTK tokenisers
- Ignoring that **punctuation as separate tokens** affects vocabulary size and stopword lists

---

Quick Revision Summary

- Tokenisation = breaking text into tokens (words, sentences, subwords, or characters)
- First step in virtually every NLP pipeline; errors propagate downstream
- Complexity arises from punctuation, hyphens, numbers, emojis, and language-specific rules
- NLTK: `word_tokenize()` and `sent_tokenize()` after downloading `punkt`
- No universal tokenisation strategy — match granularity to task and model

---

← **Previous** · **Index** · **Next** →
Noise Removal with Regular Expressions

Noise Removal with Regular Expressions

What Is Text Noise?

In NLP, **noise** refers to text elements that do not contribute semantic meaning but inflate vocabulary and distort statistics. Common sources:

- HTML/XML tags (`<p>` , `<b>` , etc.)
- Special characters, symbols, emojis
- Extra whitespace and newline characters
- Numbers (sometimes — task-dependent)

Noise removal is a **cleaning** step that prepares text for tokenisation and modelling.

Why Regular Expressions?

Regular expressions (regex) provide a flexible, pattern-based way to find and replace text at scale. Key properties:

- **Language-independent** — operates on character patterns, not grammar
- **Industry standard** — used in production ETL and NLP pipelines
- **Composable** — multiple patterns chained in sequence

Common Patterns and Solutions

Noise Type	Example Input	Target Output
HTML tags	`<p>This is <b>important</b></p>`	`This is important`
Special chars	`@NLP!!! is awesome`	`NLP is awesome`
Irregular spacing	`This  sentence   has gaps`	`This sentence has gaps`

Removing HTML Tags

Pattern: `r'<.*?>'` — matches anything between `<` and `>` (non-greedy).

```
import re
re.sub(r'<.*?>', '', html_text)
```

Keeping Only Alphabetic Text

Rather than listing every unwanted symbol, invert the match:

```
re.sub(r'[^a-zA-Z\s]', '', text)  # keep letters and spaces only
```

The `^` inside `[]` negates the character class — replace everything **except** listed characters.

Normalising Whitespace

```
re.sub(r'\s+', ' ', text).strip()  # collapse multiple spaces to one
```

Important: Replace with a single space `' '`, not empty string `''`, or words merge together.

Regex Is Case-Sensitive

By default, `re.sub('important', 'Important', text)` only matches lowercase. Use `re.IGNORECASE` flag or explicit character classes when case-insensitive matching is needed.

Task-Dependent Cleaning

There is **no universal cleaning rule**. Always align with the downstream task:

Element	Remove for topic modelling?	Keep for financial NLP?
Numbers	Often yes	**Must keep**
Emojis	Usually yes	Keep for social media sentiment
Punctuation	Often yes	Keep for QA and sentiment ("?" matters)

Common Pitfalls / Exam Traps

- Replacing `\s+` with **empty string** instead of single space — destroys word boundaries
- Forgetting regex is **case-sensitive** by default
- **Over-cleaning** numbers in financial or medical domains
- Using greedy `.*` instead of `.*?` for HTML — may match across multiple tags incorrectly
- Assuming cleaned text is always **lower-case** — normalisation is a separate step

---

Quick Revision Summary

- Noise = non-meaningful elements that inflate vocabulary and hurt models
- Regex enables scalable, language-independent pattern-based cleaning
- HTML: `r'<.*?>'`; alphabetic-only: `r'[^a-zA-Z\s]'`; whitespace: `r'\s+'` → single space
- Regex is case-sensitive unless flagged otherwise
- Cleaning decisions must match the downstream NLP task — no one-size-fits-all rule

---

← **Previous** · **Index** · **Next** →
Stopword Removal: Filtering Function Words

Stopword Removal: Filtering Function Words

What Are Stopwords?

Stopwords are high-frequency words that carry little standalone semantic meaning — articles, prepositions, conjunctions, and common auxiliaries. Examples: *the*, *is*, *of*, *and*, *a*.

They dominate word frequency distributions, making models overweight function words relative to content-bearing terms like *policy*, *economy*, or *diagnosis*.

Why Remove Stopwords?

Benefit	Explanation
Reduced vocabulary size	Fewer dimensions in bag-of-words / TF-IDF
Improved efficiency	Faster training and smaller feature matrices
Better signal	Models focus on content words that distinguish documents

Example: In topic modelling, *the* and *is* do not reveal themes; *health*, *budget*, and *regulation* do.

Stopword removal is applied **after tokenisation**: tokenise first, then filter tokens against a stopword list.

NLTK Implementation

```
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

nltk.download('stopwords')
nltk.download('punkt')

stop_words = set(stopwords.words('english'))
tokens = word_tokenize("Natural language processing is a fascinating area of computer science.")
filtered = [w for w in tokens if w.lower() not in stop_words]
# ['Natural', 'language', 'processing', 'fascinating', 'area', 'computer', 'science']
```

NLTK ships stopword lists for multiple languages. The English list is generic — customise for domain-specific needs.

When NOT to Remove Stopwords

Stopwords carry grammatical and semantic weight in several tasks:

Task	Why stopwords matter
Sentiment analysis	*not*, *no*, *never* flip polarity — "not good" ≠ "good"
Question answering	Function words define query structure
Language modelling	Next-word prediction needs full grammar
Dependency parsing	Syntactic structure depends on determiners and auxiliaries

Example: "This movie is **not** good" — removing *not* inverts meaning entirely.

Custom Stopword Lists

Domain-specific high-frequency words may add no value:

- **Product reviews** — *product*, *item* appear too often
- **News corpora** — *said*, *reported* are uninformative

Practise: add domain terms to the stopword set, or remove words like *not* from the default list for sentiment tasks.

```
stop_words.add('said')       # news domain
stop_words.discard('not')   # sentiment task
```

Common Pitfalls / Exam Traps

- Treating stopword removal as a **mandatory rule** — it is a design decision
- Removing **negation words** in sentiment pipelines
- Applying stopwords **before tokenisation** — must tokenise first
- Using the default English list for **all domains** without customisation
- Forgetting **case sensitivity** — compare with `.lower()` consistently

---

Quick Revision Summary

- Stopwords = frequent function words with low discriminative power
- Removal reduces vocabulary, improves efficiency, highlights content words
- Apply after tokenisation using NLTK's `stopwords.words('english')`
- Keep stopwords for sentiment, QA, language modelling, and grammar-aware tasks
- Customise lists per domain — add domain noise, preserve negations

---

← **Previous** · **Index** · **Next** →
Stemming: Rule-Based Morphological Reduction

Stemming: Rule-Based Morphological Reduction

What Is Stemming?

Stemming reduces inflected or derived word forms to a common **stem** — a crude base form obtained by stripping suffixes. The goal is vocabulary normalisation: treat related surface forms as one token.

Examples:

- connect, connected, connecting, connection → **connect**
- play, playing, played → **play**

Stemming is **rule-based** — it does not consider word meaning or grammatical context.

Why Stemming Matters

When exact word form is unimportant, stemming:

- Groups morphological variants together
- Reduces vocabulary size
- Speeds preprocessing and retrieval

Use Case	Why stemming helps
Search engines	Query "running" matches documents with "run", "runs"
Information retrieval	Higher recall across word variants
Topic modelling	Fewer sparse dimensions from inflectional forms
Early NLP pipelines	Fast, lightweight normalisation

NLTK Stemmers

NLTK provides two main algorithms:

Stemmer	Basis	Notes
Porter Stemmer	Original Porter rules (1980)	Most widely used; English-focused
Snowball Stemmer	Porter2 / improved rules	Supports multiple languages

```
from nltk.stem import PorterStemmer

ps = PorterStemmer()
ps.stem('running')      # 'run'
ps.stem('studies')      # 'studi'  (over-stemming)
ps.stem('better')       # 'better' (under-stemming)
```

Limitations

Problem	Example
Non-words produced	*studies* → *studi* (not a valid English word)
Over-stemming	*organisation* and *organ* merged despite different meanings
Under-stemming	*better* not reduced to *good*
No context awareness	Same suffix stripped regardless of POS

Because of these trade-offs, stemming sacrifices accuracy for speed.

Stemming vs Lemmatisation (Preview)

Aspect	Stemming	Lemmatisation
Output	May be non-word	Valid dictionary form
Method	Rule-based suffix stripping	Dictionary + POS lookup
Speed	Faster	Slower
Context	Ignored	POS-aware variants exist

Common Pitfalls / Exam Traps

- Expecting stems to be **valid English words** — Porter often produces fragments like *studi*
- Using stemming when **word meaning matters** (e.g., *organ* vs *organisation*)
- Confusing **stem** with **lemma** — stem is heuristic; lemma is dictionary-correct
- Assuming one stemmer works for **all languages** — Porter is English-centric; use Snowball for others

---

Quick Revision Summary

- Stemming = rule-based reduction to a common stem, ignoring context
- Reduces vocabulary, groups inflections, speeds IR and topic modelling
- NLTK: `PorterStemmer()` and `SnowballStemmer()`
- Produces non-words; over/under-stemming are known failure modes
- Prefer stemming when speed matters and slight inaccuracies are acceptable

---

← **Previous** · **Index** · **Next** →
Lemmatisation: Dictionary-Based Morphological Normalisation

Lemmatisation: Dictionary-Based Morphological Normalisation

What Is Lemmatisation?

Lemmatisation reduces a word to its **lemma** — the canonical dictionary form — using vocabulary and morphological rules. Unlike stemming, lemmatisation produces **valid words** and can respect grammatical role.

Examples:

- *running* (verb) → **run**
- *better* → **good**
- *studies* → **study**

Lemmatisation prioritises **semantic correctness** over surface-pattern matching.

Why Lemmatisation Is More Accurate Than Stemming

Stemming blindly strips suffixes. Lemmatisation asks: *What part of speech is this word, and what is its base form?*

Word	Context	Lemma
playing	verb ("They are playing")	run → **play**
playing	noun ("The playing was excellent")	unchanged → **playing**

This POS sensitivity makes lemmatisation valuable for text classification, information extraction, question answering, and semantic analysis.

NLTK: WordNet Lemmatizer

```
import nltk
from nltk.stem import WordNetLemmatizer

nltk.download('wordnet')
nltk.download('omw-1.4')

lemmatizer = WordNetLemmatizer()
lemmatizer.lemmatize('running')   # 'running' (default POS = noun)
lemmatizer.lemmatize('running', pos='v')  # 'run'
lemmatizer.lemmatize('studies')   # 'study'
lemmatizer.lemmatize('better', pos='a')   # 'good'
```

POS Tags for WordNet

Code	Part of Speech
`n`	Noun (default)
`v`	Verb
`a`	Adjective
`r`	Adverb

POS-aware lemmatisation dramatically improves results — always specify POS when available.

---

spaCy Lemmatisation

spaCy performs tokenisation, POS tagging, and lemmatisation in one pipeline:

```
import spacy
nlp = spacy.load('en_core_web_sm')
doc = nlp("The children were playing happily.")
for token in doc:
    print(token.text, token.lemma_)
# The the | children child | were be | playing play | happily happily
```

spaCy does **not offer stemming** — it treats lemmatisation as the production-grade approach.

When to Use Stemming vs Lemmatisation

Choose Stemming	Choose Lemmatisation
Speed is critical	Word meaning matters
Exploratory analysis	Clean, interpretable tokens needed
Slight errors acceptable	Production-grade NLP systems
Search/indexing at scale	Classification, NER, QA

There is no universal best choice — match the method to task constraints.

Common Pitfalls / Exam Traps

- Calling `lemmatize('running')` without POS — defaults to noun, returns *running* not *run*
- Expecting lemmatisation to be **as fast as stemming** — dictionary lookup is slower
- Using stemming in **production NER pipelines** where valid tokens matter
- Forgetting to download **wordnet** and **omw-1.4** for NLTK lemmatizer

---

Quick Revision Summary

- Lemmatisation → valid dictionary lemma using vocabulary + POS
- More accurate than stemming; handles *better* → *good*, verb *playing* → *play*
- NLTK: `WordNetLemmatizer` with POS codes n/v/a/r
- spaCy: integrated pipeline; no separate stemming option
- Use lemmatisation for production and semantic tasks; stemming for speed-critical IR

← **Previous** · **Index** · **Next** →
Building a Text Preprocessing Pipeline

Building a Text Preprocessing Pipeline

Why Pipelines Matter

Individual preprocessing steps — cleaning, tokenisation, stopword removal, lemmatisation — are rarely applied in isolation. Production NLP systems chain them into a **reusable pipeline** that ensures:

- **Consistency** — same transformations on training and inference data
- **Reproducibility** — documented, ordered steps
- **Scalability** — batch processing across datasets and tasks

Incorrect step ordering can destroy information or produce inconsistent tokens.

Recommended Step Order

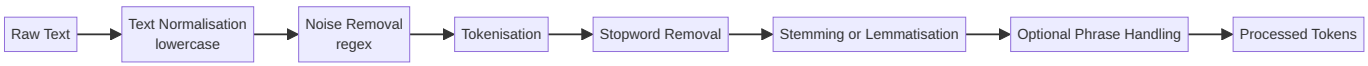

Step	Rationale
Normalise case first	Ensures consistent matching downstream
Clean before tokenise	Prevents HTML fragments becoming tokens
Tokenise before stopwords	Stopword lists operate on tokens
Lemmatise after stopwords	Fewer tokens to lemmatise (efficiency)

Reference Pipeline Implementation

```
import re
import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer

nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')
nltk.download('omw-1.4')

stop_words = set(stopwords.words('english'))
lemmatizer = WordNetLemmatizer()

def clean_text(text):
    text = text.lower()
    text = re.sub(r'[^a-z\s]', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text

def tokenize(text):
    return word_tokenize(text)

def remove_stopwords(tokens):
    return [w for w in tokens if w not in stop_words]

def lemmatize_tokens(tokens):
    return [lemmatizer.lemmatize(w, pos='v') for w in tokens]

def preprocess_text(text):
    text = clean_text(text)
    tokens = tokenize(text)
    tokens = remove_stopwords(tokens)
    tokens = lemmatize_tokens(tokens)
    return tokens
```

Example: `"Natural Language Processing is an exciting field of AI!"` → `['natural', 'language', 'process', 'exciting', 'field', 'ai']`

Task-Aware Pipeline Variants

Task	Pipeline adjustment
Sentiment analysis	Keep negation words (*not*, *never*); minimal cleaning
Topic modelling	Aggressive stopword removal; lemmatisation
Named Entity Recognition	Minimal preprocessing — preserve capitalisation and punctuation
Search engines	Stemming often preferred over lemmatisation for speed

A pipeline must be **task-aware**, **data-aware**, and **flexible** — the template above is a general starting point, not a universal solution.

Common Pitfalls / Exam Traps

- **Tokenising before cleaning** — HTML tags become spurious tokens
- **Removing stopwords before sentiment tasks** — loses negation
- **Same pipeline for NER and topic modelling** — NER needs original casing
- Confusing **function names with variable names** in lab code — test on diverse raw inputs
- Forgetting **punkt** and **wordnet** downloads in fresh environments

---

Quick Revision Summary

- Pipelines chain preprocessing for consistency, reproducibility, and scale
- Order: normalise → clean → tokenise → stopwords → stem/lemmatise
- Task-specific variants: sentiment keeps negations; NER keeps minimal changes
- NLTK pipeline combines regex cleaning, `word_tokenize`, stopword filter, `WordNetLemmatizer`
- Always validate pipeline output on representative raw text before modelling

---

← **Previous** · **Index** · **Next** →

Week 3

Week 3

Structural Text Analysis: POS Tagging and Named Entity Recognition

Structural Text Analysis: POS Tagging and Named Entity Recognition

Module Overview

This module shifts from preprocessing to **structural analysis** — understanding how words function grammatically and which tokens refer to real-world entities.

Two complementary techniques anchor the module:

Technique	Layer	Question Answered
Part-of-Speech (POS) Tagging	Syntactic	What grammatical role does each word play?
Named Entity Recognition (NER)	Semantic	Which spans refer to people, organisations, locations, dates, etc.?

Together, POS and NER form building blocks for dependency parsing, information extraction, knowledge graphs, and question-answering systems.

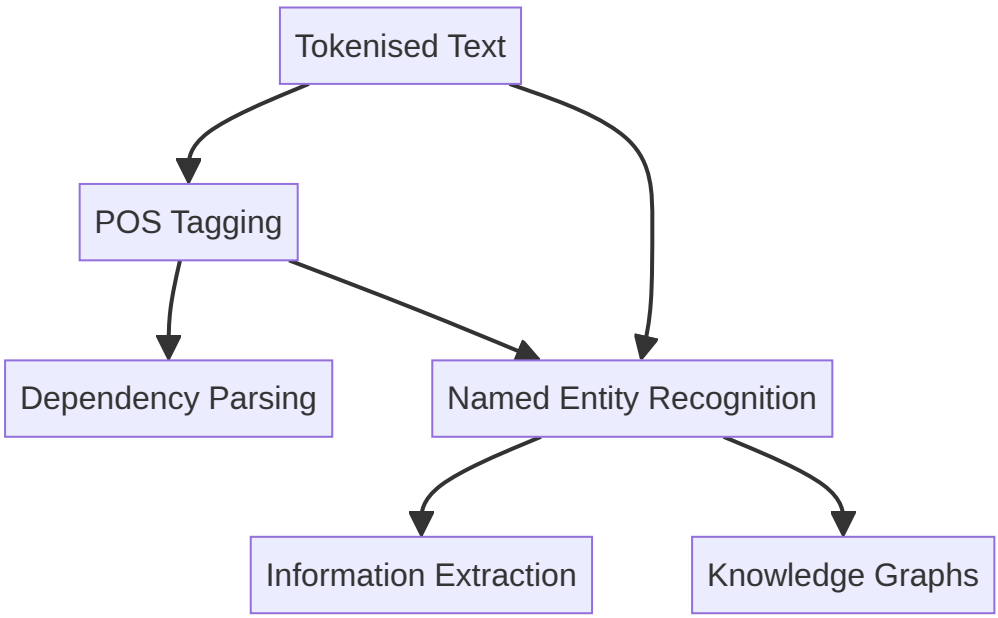

Why Structure Matters Before Semantics

Preprocessing normalises surface form; structural analysis reveals **relationships**:

- POS tags expose syntax — nouns vs verbs, determiners vs auxiliaries
- NER maps text to ontology categories — *Google* as ORG, not just NN
- Modern models learn implicit syntax, but explicit tags remain interpretable and useful for rule-based post-processing

---

Implementation Libraries Covered

Library	Role in Module
spaCy	Production-grade POS and NER with visualisation
NLTK	Traditional modular pipeline; educational baseline
Flair	Neural sequence taggers; accuracy-focused alternative

A comparative framework helps choose the right tool for learning, production, or maximum accuracy.

Common Pitfalls / Exam Traps

- Confusing **POS tagging with NER** — POS assigns grammatical categories; NER identifies real-world entity types
- Assuming POS is a **dictionary lookup** — tags are context-dependent (*book* as verb vs noun)
- Treating all three libraries as **interchangeable** — speed, accuracy, and API complexity differ significantly

---

Quick Revision Summary

- Module focus: syntactic structure (POS) and semantic entities (NER)
- POS → grammatical roles; NER → persons, orgs, locations, dates, money
- Both underpin IE, QA, knowledge graphs, and parsing pipelines
- Implementations: spaCy (production), NLTK (fundamentals), Flair (accuracy)
- Library choice depends on speed vs accuracy vs educational goals

---

← **Previous** · **Index** · **Next** →
Part-of-Speech Tagging: Assigning Grammatical Categories

Part-of-Speech Tagging: Assigning Grammatical Categories

What Is POS Tagging?

Part-of-Speech (POS) tagging assigns a **grammatical category** to each token in a sentence — noun, verb, adjective, adverb, preposition, determiner, pronoun, and finer-grained subtypes.

Example: *"Natural language processing is fascinating."*

Token	POS Tag	Category
Natural	JJ	Adjective
language	NN	Noun
processing	NN	Noun
is	VBZ	Verb (present, 3rd person)
fascinating	JJ	Adjective

POS tagging reveals **syntactic structure** — how words relate within a sentence.

Why POS Tagging Matters

POS tags enable and improve:

- **Dependency parsing** — subject-verb-object relationships
- **Information extraction** — noun phrases as entity candidates
- **Text normalisation** — POS-aware lemmatisation
- **Question answering** — identifying wh-words and focus entities

Even transformer models learn implicit syntactic patterns; explicit POS tags remain interpretable and useful for hybrid pipelines.

Context Sensitivity

POS tagging is **not** dictionary lookup — the same word receives different tags in different contexts:

Sentence	*book* tag	Role
"I will **book** a ticket."	VB (verb)	Action
"Read this **book**."	NN (noun)	Object

Disambiguation requires surrounding context — a core challenge that modern taggers address with statistical or neural models.

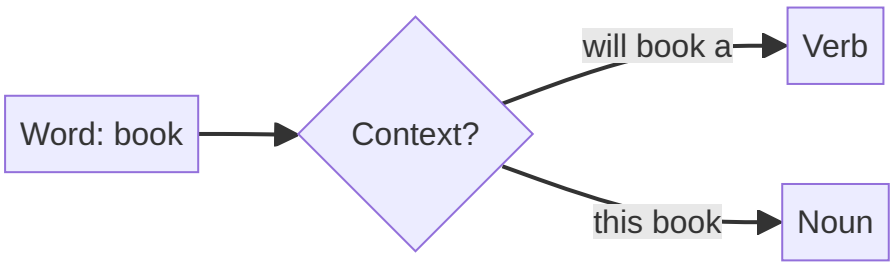

Tag Set Conventions

Different frameworks use different tag sets:

Framework	Tag Set	Example Tags
NLTK / Penn Treebank	Penn POS	NN, NNP, VBZ, DT, JJ
spaCy	Universal POS + fine-grained	NOUN, VERB, DET + `token.tag_`
Flair	Flair / OntoNotes-style	NNP, VBZ, DT

Always consult the tag set documentation when interpreting output.

Common Pitfalls / Exam Traps

- Treating POS as **fixed per word** — context determines tag
- Confusing **POS with NER** — *Apple* as NN (noun) vs ORG (organisation entity)
- Assuming **one tag set** across libraries — Penn tags ≠ Universal POS labels
- Ignoring POS when **lemmatising** — verb vs noun changes the lemma

---

Quick Revision Summary

- POS tagging assigns grammatical categories to each token
- Reveals syntactic structure; supports parsing, IE, QA, and lemmatisation
- Highly context-sensitive — same word, different tags in different sentences
- Penn Treebank tags (NN, VBZ, DT) common in NLTK; spaCy adds Universal POS
- Not a dictionary lookup — requires contextual disambiguation

---

← **Previous** · **Index** · **Next** →
POS Tagging with spaCy

POS Tagging with spaCy

Why spaCy for POS Tagging

spaCy provides an end-to-end NLP pipeline: tokenisation, POS tagging, dependency parsing, and NER in a single `nlp()` call. For production workloads, spaCy balances speed and accuracy with a streamlined API.

Setup and Basic Usage

```
import spacy
nlp = spacy.load('en_core_web_sm')

doc = nlp("I will book a ticket.")
for token in doc:
    print(f"{token.text:12} {token.pos_:5} {token.tag_}")
```

Output:

Token	Universal POS (`pos_`)	Fine-grained (`tag_`)
I	PRON	PRP
will	AUX	MD
book	VERB	VB
a	DET	DT
ticket	NOUN	NN
.	PUNCT	.

- `token.pos_` — Universal POS category (human-readable)
- `token.tag_` — Detailed Penn Treebank-style tag

Context Disambiguation Example

```
doc1 = nlp("I will book a ticket.")   # book → VERB
doc2 = nlp("Read this book.")        # book → NOUN
```

spaCy's transition-based parser uses word embeddings and context to resolve ambiguity automatically — no manual pipeline assembly required.

Key Properties

Property	Detail
Tokenisation	Automatic inside `nlp()`
Model	`en_core_web_sm` (small), `md`, `lg` variants
Speed	Optimised for production batch processing
Output	Token text, POS, tag, lemma, dependency in one pass

Real-World Usage

- **Document intelligence APIs** — extract noun phrases for indexing
- **Chatbot intent pipelines** — verb/noun patterns for slot filling
- **Pre-annotation** — bootstrap NER training data with POS-filtered candidates

---

Common Pitfalls / Exam Traps

- Forgetting to **install and load** the language model (`python -m spacy download en_core_web_sm`)
- Confusing `pos_` (Universal) with `tag_` (Penn fine-grained)
- Expecting spaCy POS tags to **match NLTK Penn tags exactly** — mapping may differ slightly
- Running POS on raw text **without loading model** — `nlp` must be initialised first

---

Quick Revision Summary

- spaCy: `nlp(text)` → tokens with `token.pos_` and `token.tag_`
- Tokenisation included; no separate step needed
- Context resolves ambiguity (*book* verb vs noun)
- Production-ready: fast, integrated pipeline
- Install `en_core_web_sm` before use; practise on varied sentences

← **Previous** · **Index** · **Next** →
Visualising POS Tags and Dependency Trees with spaCy

Visualising POS Tags and Dependency Trees with spaCy

Beyond Text Output

POS tags alone list grammatical categories. **Dependency parsing** adds syntactic relationships — which token modifies which, and what role each plays (subject, object, auxiliary).

spaCy's `displaCy` renders both POS tags and dependency arcs visually — invaluable for debugging pipelines and communicating NLP results.

---

displaCy for POS Visualisation

```
from spacy import displacy

doc = nlp("I will book a ticket.")
displacy.render(doc, style='dep', jupyter=True)
```

The dependency visualisation shows:

- **Tokens** with POS labels beneath each word
- **Directed arcs** linking head tokens to dependents
- **Relation labels** — nsubj (nominal subject), dobj (direct object), aux (auxiliary), det (determiner)

Reading Dependency Relations

Sentence: *"I will book a ticket."*

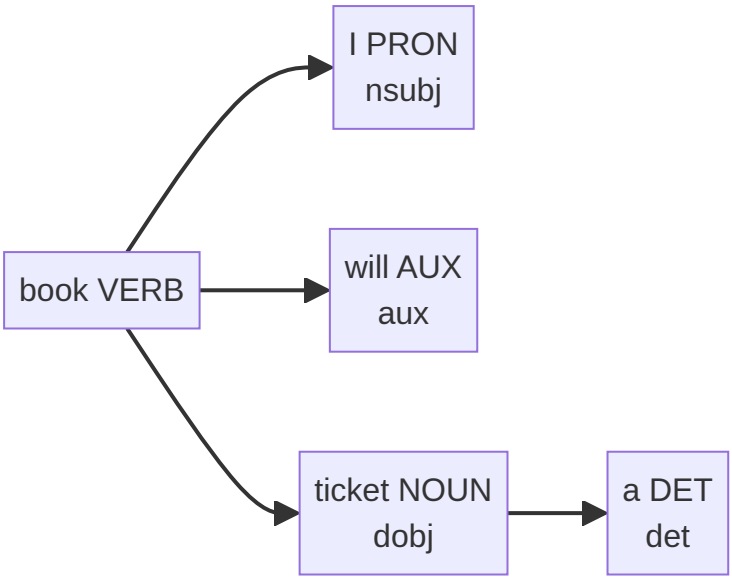

Relation	Meaning	Example
nsubj	Nominal subject	I → book
aux	Auxiliary verb	will → book
dobj	Direct object	ticket → book
det	Determiner	a → ticket

Sentence: *"Read this book."*

- *read* (VERB) — root
- *book* (NOUN) — dobj of *read*
- *this* (DET) — det of *book*

Same word *book* — noun in this sentence, verb in the previous example. Visualisation makes the contrast immediate.

When Visualisation Helps

Scenario	Benefit
Pipeline debugging	Spot wrong attachments quickly
Stakeholder demos	Non-technical audience grasps structure
Annotation QA	Verify parser output before training custom models
Teaching	Illustrate subject-verb-object patterns

Common Pitfalls / Exam Traps

- Using `style='ent'` for POS — use `style='dep'` for dependency/POS trees; `'ent'` is for NER
- Assuming **visualisation replaces evaluation** — pretty graphs don't guarantee correctness
- Ignoring **parse errors** on long or malformed sentences — visual check catches failures
- Confusing **dependency labels** (nsubj, dobj) with **POS tags** (NN, VB)

Quick Revision Summary

- `displaCy.render(doc, style='dep')` visualises POS + dependency arcs
- Dependency relations: nsubj, aux, dobj, det link tokens syntactically
- Same word gets different structure in different sentences (*book* verb vs noun)
- Essential for debugging, demos, and understanding parser behaviour
- NER visualisation uses `style='ent'` — different from dependency view

---

← **Previous** · **Index** · **Next** →
POS Tagging with NLTK

POS Tagging with NLTK

NLTK's Modular Approach

NLTK implements POS tagging as an explicit multi-step pipeline: tokenise first, then tag. This transparency helps learners understand each stage — unlike spaCy's single-call abstraction.

Setup and Resources

Required NLTK downloads:

```
import nltk
from nltk.tokenize import word_tokenize

nltk.download('punkt')
nltk.download('averaged_perceptron_tagger')
nltk.download('tagsets_json')
```

The **averaged perceptron tagger** is a classical ML tagger — a lightweight neural precursor useful for understanding non-transformer tagging.

Implementation

```
text = "NLTK is a powerful library for natural language processing."
tokens = word_tokenize(text)
tagged = nltk.pos_tag(tokens)
# [('NLTK', 'NNP'), ('is', 'VBZ'), ('a', 'DT'), ('powerful', 'JJ'), ...]
```

`nltk.pos_tag()` returns a list of `(token, tag)` tuples using the **Penn Treebank tag set**.

Understanding Penn Treebank Tags

Look up any tag programmatically:

```
nltk.help.upenn_tagset('VBZ')
# verb, present tense, 3rd person singular
```

Common Penn Tags

Tag	Meaning	Example
NN	Noun, singular	language
NNP	Proper noun, singular	NLTK
VBZ	Verb, present 3sg	is
DT	Determiner	a
JJ	Adjective	powerful
IN	Preposition	for

Full reference: `nltk.help.upenn_tagset()` without arguments lists all tags.

NLTK vs spaCy for POS

Aspect	NLTK	spaCy
Pipeline	Manual: tokenise → tag	Automatic in `nlp()`
Tagger	Averaged perceptron	Transition-based NN
Speed	Slower	Faster
Best for	Learning fundamentals	Production

Common Pitfalls / Exam Traps

- Forgetting `averaged_perceptron_tagger` download — tagging fails silently or errors
- Confusing **NN** (common noun) with **NNP** (proper noun)
- Not knowing how to look up tags — `nltk.help.upenn_tagset('TAG')`
- Expecting NLTK output to include **dependency relations** — POS only unless using additional parsers

---

Quick Revision Summary

- NLTK POS: `word_tokenize()` → `nltk.pos_tag()` → Penn Treebank tags
- Download punkt + averaged_perceptron_tagger + tagsets_json
- Common tags: NN, NNP, VBZ, DT, JJ, IN
- Lookup: `nltk.help.upenn_tagset('VBZ')`
- Modular pipeline ideal for learning; spaCy preferred for production speed

---

← **Previous** · **Index** · **Next** →
POS Tagging with Flair

POS Tagging with Flair

What Is Flair?

Flair is a PyTorch-based NLP framework offering **pre-trained neural sequence taggers** for POS tagging, NER, sentiment analysis, and more. Models use contextual embeddings (biLSTM-CRF architectures) to achieve high accuracy at the cost of slower inference compared to spaCy.

Installation and Usage

```
from flair.data import Sentence
from flair.models import SequenceTagger

tagger = SequenceTagger.load('pos')
sentence = Sentence("Flair is a great library for natural language processing.")
tagger.predict(sentence)
print(sentence.to_tagged_string())
```

Output assigns Penn-style tags: *Flair/NNP*, *is/VBZ*, *a/DT*, etc.

How Flair Differs

Aspect	Flair	NLTK	spaCy
Architecture	BiLSTM-CRF + contextual embeddings	Averaged perceptron	Transition-based NN
API	Load model → predict	Tokenise → tag	`nlp(text)`
Preprocessing	Built into `Sentence`	Manual	Built into pipeline
Speed	Moderate	Slowest	Fastest
Accuracy	High on benchmarks	Lower	High; slight trade-off for speed

Flair tag definitions are documented on Hugging Face (Flair POS tag set reference).

Multilingual Support

Flair provides language-specific POS models (e.g., `pos-fast`, German POS taggers). Tag sets may align with Universal Dependencies or language-specific conventions.

When to Prefer Flair for POS

- **Maximum accuracy** on specialised or ambiguous text
- **Research benchmarks** where F1 score matters more than latency
- **Domain text** (legal, medical) where contextual embeddings help disambiguation

For most production POS needs at scale, spaCy remains the default choice.

Common Pitfalls / Exam Traps

- Forgetting **`tagger.predict(sentence)`** — model load alone does not tag
- Passing a **string instead of `Sentence` object** — wrap text in `Sentence(text)`
- Expecting **spaCy-level speed** on millions of documents
- Confusing Flair **POS models** with **NER models** — different `SequenceTagger.load()` checkpoints

Quick Revision Summary

- Flair: PyTorch sequence taggers with contextual embeddings
- `SequenceTagger.load('pos')` → `tagger.predict(sentence)` → `to_tagged_string()`
- Higher accuracy, moderate speed; multilingual models available
- Simple API: create Sentence, load model, predict
- Compare same text across NLTK, spaCy, and Flair to observe tag differences

---

← **Previous** · **Index** · **Next** →
Named Entity Recognition: Identifying Real-World Entities

Named Entity Recognition: Identifying Real-World Entities

What Is NER?

Named Entity Recognition (NER) identifies and classifies **real-world entities** mentioned in text — mapping spans of tokens to predefined categories.

NER moves beyond syntax (POS) to **semantic meaning**: who, what organisation, where, when, and how much.

Common Entity Types

Label	Category	Examples
PERSON	People (real or fictional)	Sundar Pichai, Alice
ORG	Organisations	Google, Apple, UN
GPE / LOC	Geopolitical entities / locations	India, UK, Washington
DATE / TIME	Temporal expressions	December 2026, Monday
MONEY	Monetary amounts	$1 billion, €500

Example: *"Sundar Pichai is the CEO of Google."*

Span	Entity Type
Sundar Pichai	PERSON
Google	ORG

NER Applications

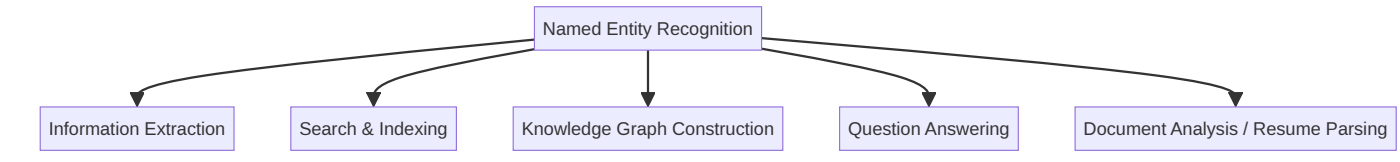

NER enables machines to answer **who did what, where, and when** — essential for enterprise search, compliance monitoring, and automated news aggregation.

POS vs NER: Critical Distinction

Aspect	POS Tagging	NER
Layer	Syntactic / grammatical	Semantic / real-world
Output	Noun, verb, adjective	PERSON, ORG, GPE, DATE
Example: *Apple*	NN (noun)	ORG (organisation)

Sentence: *"Apple released a new product in India."*

- POS: *Apple* → noun
- NER: *Apple* → ORG; *India* → GPE

Same surface form, different annotation layers. NER often **builds on** POS and dependency information internally.

NER Challenges

Challenge	Example
Entity ambiguity	*Washington* — person, location, or organisation?
Boundary detection	*Elon Musk* as one span vs two tokens
Polysemy	*Apple* — fruit (not ORG) vs company (ORG)
Context dependence	Domain-specific entity types

NER is highly **context-dependent** — accuracy requires surrounding text, not isolated token lookup.

Common Pitfalls / Exam Traps

- Equating **NN (noun)** with **named entity** — all entities are typically nouns, but not all nouns are entities
- Assuming **one-word entities only** — multi-token spans (*New York*, *Sundar Pichai*) are common
- Ignoring **entity type hierarchy** — GPE vs LOC vs FAC differ across tag sets
- Treating NER as **optional** in IE pipelines — it is often the primary extraction step

---

Quick Revision Summary

- NER classifies text spans into entity types: PERSON, ORG, GPE, DATE, MONEY
- Semantic layer — distinct from grammatical POS tagging
- Powers IE, search, knowledge graphs, QA, document analysis
- Context-dependent; ambiguity (*Washington*, *Apple*) is a core challenge
- NER builds on syntactic analysis but answers different questions than POS

← **Previous** · **Index** · **Next** →
Named Entity Recognition with spaCy

Named Entity Recognition with spaCy

spaCy as the Industry Standard for NER

spaCy provides production-grade NER integrated into its core pipeline. A single `nlp()` call returns entities accessible via `doc.ents` with labels, spans, and human-readable descriptions.

Basic Implementation

```
import spacy
nlp = spacy.load('en_core_web_sm')

text = "Apple is looking at buying a UK startup for $1 billion. Elon Musk said in December 2026."
doc = nlp(text)

for ent in doc.ents:
    print(ent.text, ent.label_, spacy.explain(ent.label_))
```

Entity	Label	Description
Apple	ORG	Companies, agencies, institutions
UK	GPE	Countries, cities, states
$1 billion	MONEY	Monetary values
Elon Musk	PERSON	People, including fictional
December 2026	DATE	Absolute or relative dates

- `ent.text` — span text
- `ent.label_` — entity type string
- `spacy.explain(label)` — human-readable definition

Visualisation with displaCy

```
from spacy import displacy
displacy.render(doc, style='ent', jupyter=True)
```

Colour-coded entity spans render in Jupyter or browser — each entity type gets a distinct highlight.

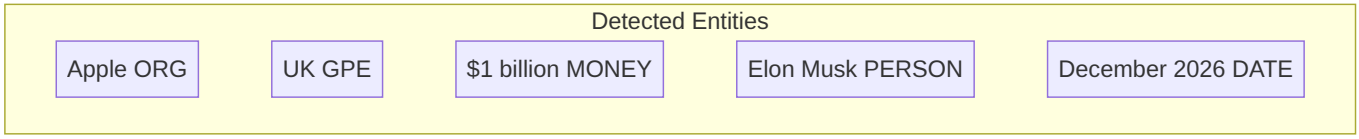

---

spaCy NER Strengths

Strength	Detail
Accuracy	Strong on standard English news and web text
Speed	Fastest among NLTK, spaCy, Flair for large corpora
Integration	Same pipeline as POS and dependency parsing
Customisation	Train `EntityRuler` or fine-tune NER component

Real-World Deployment

- **Financial news monitoring** — extract ORG and MONEY from filings
- **CRM enrichment** — PERSON and ORG from email threads
- **Compliance scanning** — GPE and DATE in contract clauses

Common Pitfalls / Exam Traps

- Accessing entities via **`doc.ents`** not `doc.entities`
- Forgetting **`ent.label_`** (string) vs `ent.label` (hash ID)
- Using **`style='dep'`** instead of **`style='ent'`** for NER visualisation
- Expecting perfect NER on **domain jargon** without custom training

---

Quick Revision Summary

- spaCy NER: `doc = nlp(text)` → iterate `doc.ents`
- Labels: ORG, GPE, MONEY, PERSON, DATE + `spacy.explain()`
- Visualise: `displacy.render(doc, style='ent')`
- Industry default for production NER — fast and accurate
- Practise on varied text; consider custom models for specialised domains

---

← **Previous** · **Index** · **Next** →
Named Entity Recognition with NLTK

Named Entity Recognition with NLTK

NLTK's Chunk-Based NER Pipeline

NLTK implements NER through a **multi-step classical pipeline** — fundamentally different from spaCy's integrated neural approach:

1. Tokenise text
2. Apply POS tagging
3. Chunk named entities with `ne_chunk()`

This modular design exposes each stage for learning, but produces lower accuracy than modern neural taggers.

Setup

```
import nltk
from nltk.tokenize import word_tokenize
from nltk.tag import pos_tag
from nltk.chunk import ne_chunk

nltk.download('punkt')
nltk.download('averaged_perceptron_tagger')
nltk.download('maxent_ne_chunker')
nltk.download('words')
```

Implementation

```
text = "Apple is looking at buying a UK startup for $1 billion. Elon Musk said in December 2026."
tokens = word_tokenize(text)
tagged = pos_tag(tokens)
ne_tree = ne_chunk(tagged)
```

`ne_chunk()` returns a nested **Tree** structure where entity subtrees are labelled (e.g., `GPE`, `PERSON`, `ORG`).

Extract named entities programmatically:

```
for chunk in ne_tree:
    if hasattr(chunk, 'label'):
        print(chunk.label(), chunk.leaves())
```

NLTK Entity Types

NLTK's `maxent_ne_chunker` supports a **limited set** of entity types:

Tag	Meaning
PERSON	People
ORG	Organisations
GPE	Geopolitical entities
LOCATION	Non-GPE locations
FACILITY	Buildings, airports, highways

Accuracy Limitations

On the same benchmark text, NLTK typically detects **fewer entities** than spaCy:

Span	spaCy	NLTK (typical)
Apple	ORG	GPE (misclassified)
UK	GPE	Often missed
$1 billion	MONEY	Missed
Elon Musk	PERSON	Partial (*Musk* only)
December 2026	DATE	Missed

NLTK merges or misses multi-token entities and lacks MONEY/DATE categories in the default chunker.

When NLTK NER Still Matters

- **Educational contexts** — understanding chunking and POS-dependent NER
- **Legacy systems** — maintaining older rule-based pipelines
- **Fallback** — when neural models unavailable (resource-constrained environments)

For production NER, spaCy or Flair is strongly preferred.

Common Pitfalls / Exam Traps

- Forgetting `maxent_ne_chunker` and `words` downloads
- Expecting **MONEY and DATE** entities — NLTK default chunker lacks these
- Not checking `hasattr(chunk, 'label')` when iterating — non-entity chunks are plain tuples
- Assuming NLTK NER **matches spaCy output** — significant accuracy gap

---

Quick Revision Summary

- NLTK NER: tokenise → `pos_tag()` → `ne_chunk()`
- Download maxent_ne_chunker + words + tagger resources
- Limited entity types: PERSON, ORG, GPE, LOCATION, FACILITY
- Lower accuracy than spaCy — misses dates, money, multi-token spans
- Useful for learning pipeline mechanics; not recommended for production

---

← **Previous** · **Index** · **Next** →
Named Entity Recognition with Flair

Named Entity Recognition with Flair

Flair for High-Accuracy NER

Flair applies **biLSTM-CRF** neural architectures with contextual string embeddings to NER. The same `SequenceTagger` framework used for POS tagging loads a dedicated NER checkpoint.

Trade-off: higher accuracy potential vs slower inference than spaCy.

Implementation

```
from flair.data import Sentence
from flair.models import SequenceTagger

tagger = SequenceTagger.load('ner')
sentence = Sentence("Apple is looking at buying a UK startup for $1 billion. Elon Musk said in December 2026.")
tagger.predict(sentence)  # Required — predictions not automatic

for entity in sentence.get_spans('ner'):
    print(entity.text, entity.tag)
```

Output spans with labels such as ORG, LOC, MONEY, PER, MISC — depending on the model's tag set.

Critical Step: predict()

Unlike model loading alone, `tagger.predict(sentence)` **must be called** before accessing entities. Omitting this step returns empty results — a common implementation error.

Flair vs spaCy vs NLTK on Same Text

Entity	spaCy	Flair	NLTK
Apple	ORG	ORG	GPE (partial)
UK	GPE	LOC	Missed
$1 billion	MONEY	Detected	Missed
Elon Musk	PERSON	Detected	Partial
December 2026	DATE	Detected	Missed

Flair generally outperforms NLTK; compared to spaCy, results are competitive though occasional misclassifications occur (e.g., *Elon Musk* as ORG in some runs). Benchmark on your specific domain.

Architecture (Conceptual)

Contextual embeddings capture surrounding words — critical for disambiguating *Apple* (company vs fruit).

When to Choose Flair for NER

- Specialised domains (medical, legal) where maximum F1 matters
- Research requiring state-of-the-art sequence labelling
- Batch offline processing where latency is acceptable

For real-time tweet streams or high-volume document ingestion, spaCy's speed advantage dominates.

Common Pitfalls / Exam Traps

- **Skipping** `predict()` — most common Flair NER bug
- Using `load('pos')` instead of `load('ner')` — wrong model checkpoint
- Expecting **instant inference** on large corpora without GPU
- Not comparing **tag sets** — Flair PER vs spaCy PERSON naming differs

---

Quick Revision Summary

- Flair NER: `SequenceTagger.load('ner')` → `Sentence(text)` → `predict()` → `get_spans('ner')`
- BiLSTM-CRF + contextual embeddings for high accuracy
- Must call `predict()` before reading entities
- Strong alternative to spaCy when accuracy > speed
- Compare outputs on identical text across all three libraries

---

← **Previous** · **Index** · **Next** →
Comparing NLTK, spaCy, and Flair for NLP Tasks

Comparing NLTK, spaCy, and Flair for NLP Tasks

Choosing the Right Tool

For POS tagging and NER, library selection depends on the project's **primary goal**:

Goal	Recommended Library
Learn NLP fundamentals	NLTK
Deploy fast production systems	spaCy
Maximise accuracy (compute available)	Flair

Always match tool to task — no single library dominates all dimensions.

Side-by-Side Comparison

Dimension	NLTK	spaCy	Flair
Primary use	Education, research POC	Production, industry	Research, specialised domains
Speed	Slowest	**Fastest**	Moderate (slowest neural option)
Accuracy	Lower (classical ML)	High; ~1% below heaviest models	**Highest F1** on benchmarks
Architecture	Rule-based + classical ML	Transition-based shallow NN	BiLSTM-CRF + contextual embeddings
Embeddings	None (taggers use features)	Bloom embeddings (context-aware)	Flair embeddings (contextual strings)
API complexity	Manual multi-step pipeline	Streamlined single call	Load model → predict
Visualisation	Limited (tree plots)	**displaCy** built-in	Text output primarily
Ease of use	Manual steps	User-friendly	Simplest API (few lines)

NLTK: The Educator

Methodology: Modular pipeline — tokenise → POS tag → chunk for NER.

Strengths:

- Transparent step-by-step mechanics
- Excellent for teaching and small-scale experiments
- Rich corpora and linguistic resources

Weaknesses:

- Slower and less accurate on modern benchmarks
- NER limited entity types and span detection
- Not ideal for production microservices at scale

Best for: Coursework, POCs, understanding building blocks.

---

spaCy: The Production Workhorse

Methodology: Efficient transition-based parser powered by shallow feed-forward neural networks and Bloom embeddings that capture word context.

Strengths:

- Fastest processing — critical for millions of documents or real-time streams
- Integrated POS, NER, dependency parsing in one pass
- displaCy for visual debugging and demos
- Industry-standard deployment path

Weaknesses:

- Slightly lower accuracy than heaviest deep models (often <1% F1 difference)

Best for: Production APIs, real-time tweet processing, document pipelines at scale.

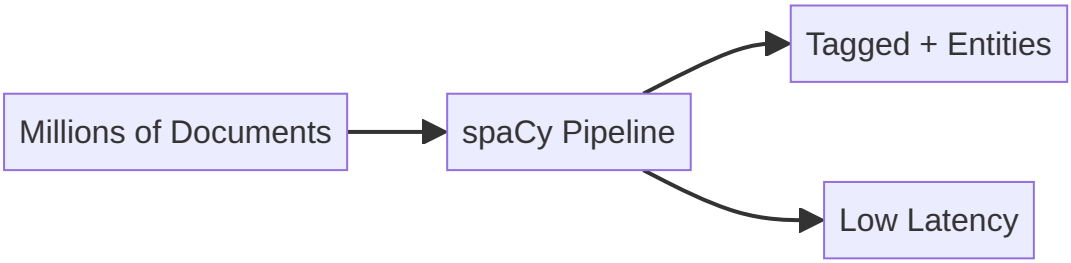

Flair: The Perfectionist

Methodology: BiLSTM-CRF with contextual string embeddings — deep understanding of word meaning from surrounding text. Handles polysemy well.

Strengths:

- State-of-the-art accuracy on standard benchmarks
- Strong on ambiguous and domain-specific text
- Simple high-level API

Weaknesses:

- Heavier model — slower inference
- Less ecosystem tooling than spaCy (visualisation, custom pipeline components)

Best for: Medical/legal NLP, research papers, offline batch where F1 is paramount.

Decision Framework

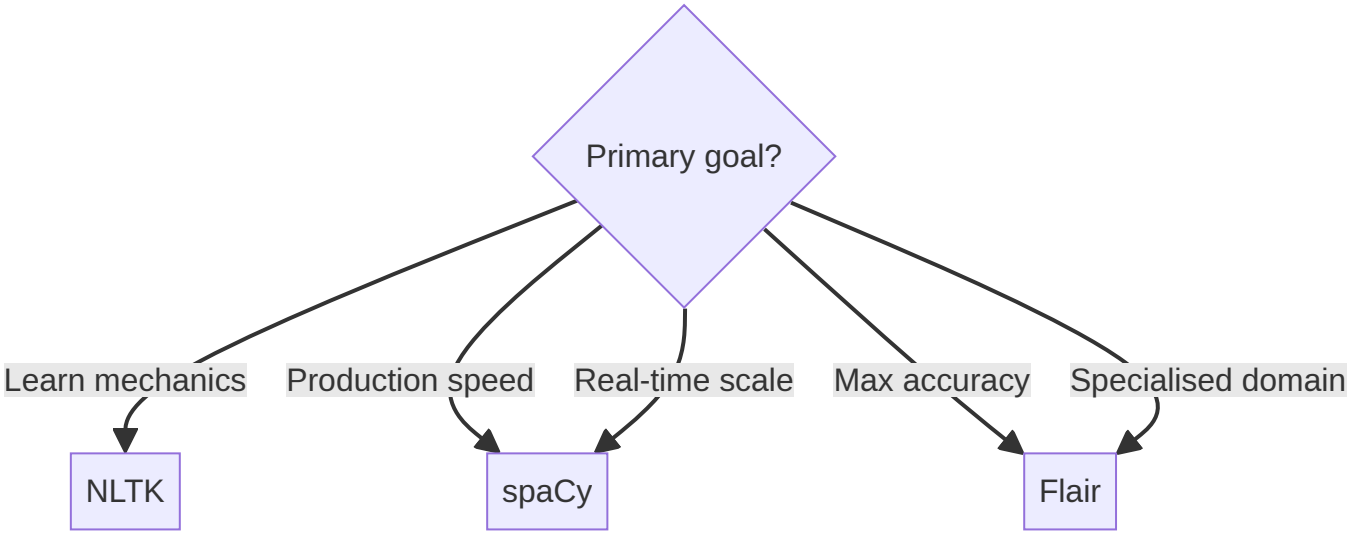

Example scenarios:

- Process **millions of tweets in real time** → spaCy
- Squeeze **maximum accuracy** from critical medical records → Flair
- Build a **teaching demo** of NER pipeline steps → NLTK
- NLTK as **backup** when other libraries fail in constrained environments

---

Common Pitfalls / Exam Traps

- Claiming **NLTK is obsolete** — still valuable for education and linguistics resources
- Assuming **Flair is always better** — speed penalty matters at scale
- Stating spaCy uses **deep transformers** for POS/NER — it uses transition-based shallow NNs (distinct from BERT)
- Confusing **Flair speed ranking** — Flair is slower than spaCy, not faster than NLTK in all cases (NLTK slowest overall)

Quick Revision Summary

- NLTK: modular, educational, slowest, lowest accuracy — learn fundamentals
- spaCy: fastest, production-ready, displaCy, industry standard
- Flair: highest accuracy, BiLSTM-CRF, slower, best for specialised/research use
- Choose by goal: learn (NLTK), deploy (spaCy), accuracy (Flair)
- Real-time scale → spaCy; critical F1 → Flair; pipeline transparency → NLTK

← Previous · Index · Next →

Week 4

Week 4

Text Corpora and Corpus Linguistics: Module Overview

Text Corpora and Corpus Linguistics: Module Overview

Why Corpora Matter for NLP

Every NLP model learns from text data. The quality, composition, and bias of that data directly shape model behaviour — from word embeddings to large language models.

This module examines:

1. **What text corpora are** and how they differ from documents and datasets
2. **Types of corpora** — general-purpose, domain-specific, raw, annotated
3. **Bias and representation** — how human choices in data collection propagate into models
4. **Hands-on analysis** — exploring the Gutenberg and Brown corpora to discover statistical properties of language

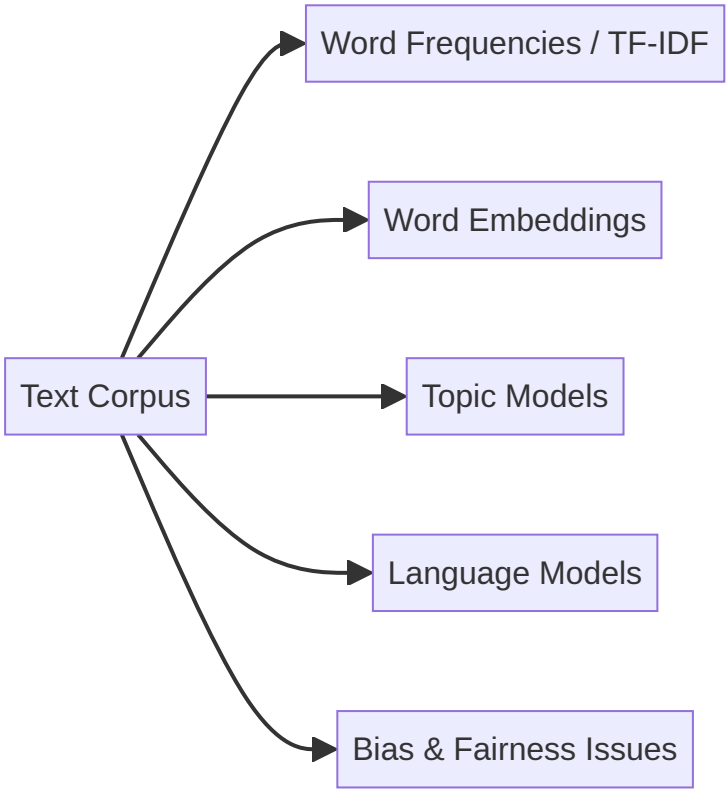

Key Insight: Garbage In, Garbage Out

Models do not understand language like humans — they learn **statistical patterns** from corpora. Change the corpus, change the model:

\text{Model behaviour} = f(\text{corpus statistics})
Understanding your data is as important as understanding your architecture.

Module Lab Focus

Corpus	Character	Analysis Goals
Gutenberg	Literary classics	Vocabulary richness, word frequency, preprocessing effects
Brown	Balanced modern American English across genres	Cross-domain comparison, genre-specific language patterns

Analysis reveals laws governing natural language — Zipfian word distributions, type-token ratios, and domain signatures in high-frequency terms.

Common Pitfalls / Exam Traps

- Using **corpus**, **document**, and **dataset** interchangeably — each has a distinct definition
- Assuming **large corpus = good corpus** — representation and bias matter more than size alone
- Skipping **corpus exploration** before modelling — leads to silent failures in production

---

Quick Revision Summary

- Module covers corpora foundations, types, bias, and hands-on Gutenberg/Brown analysis
- Models learn statistics from corpora — corpus quality determines model quality
- Gutenberg: literary; Brown: genre-balanced modern English
- Explore corpora before training — avoid treating text as a black box
- Bias in data propagates to bias in predictions

← **Previous** · **Index** · **Next** →
What Is a Text Corpus and Why It Matters

What Is a Text Corpus and Why It Matters

Defining a Text Corpus

A **text corpus** (plural: corpora) is a large, structured collection of text assembled to study, analyse, and model language. It is more than a single document or an unorganised folder of files — it is **carefully collected** to represent language usage for a purpose.

Corpus vs Document vs Dataset

Term	Definition	Example
Document	Single unit of text	One book, article, or review
Corpus	Collection of documents	500 books; 1 million reviews
Dataset	Corpus organised for a specific task, often with labels	IMDb sentiment dataset; SQuAD QA pairs

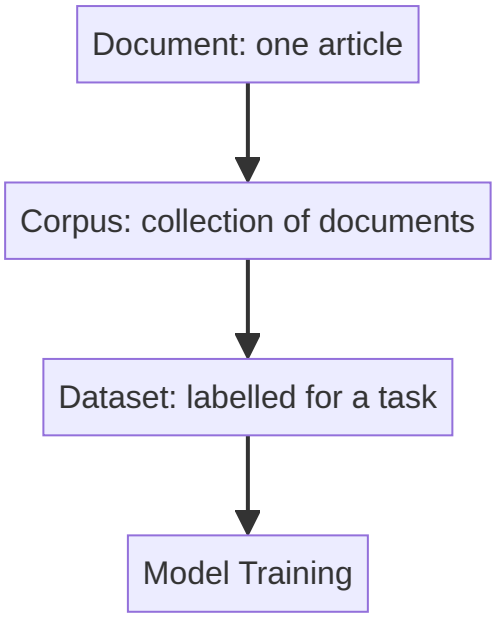

Key distinction: Corpora provide **language**; datasets provide **supervision** (labels for training/evaluation).

Examples of Text Corpora

Corpus Type	Example	Typical Use
Literary collection	Gutenberg corpus	Language modelling, literary NLP
News across categories	Brown corpus	Genre analysis, POS research
Labelled reviews	Movie sentiment corpus	Sentiment classification
QA pairs	SQuAD	Question answering

When we say "train an NLP model," we mean train on a **text corpus** (or labelled dataset derived from one).

Why Corpora Drive Everything in NLP

Machines process zeros and ones — they learn **patterns, statistics, and distributions** from text:

NLP Technique	Corpus Dependency
TF-IDF	Word frequencies across documents
Word embeddings	Word co-occurrence patterns
Topic models (LDA)	Document-word distributions
Language models	Sequential token patterns

$$\text{TF-IDF}(t, d) \propto \text{count}(t \text{ in } d) \times \log\frac{N}{\text{df}(t)}$$
All statistics originate from the corpus. **Any change in corpus changes model behaviour.**

Language Is Statistical, Not Rule-Based

Traditional linguistics emphasises grammatical rules. Modern NLP emphasises **usage patterns**:

- *Bank* disambiguates via co-occurrence (*river bank* vs *investment bank*), not explicit rules
- Frequency and context from the corpus drive meaning — not a hand-crafted lexicon

Principle: Garbage corpus → garbage model (GIGO).

Common Pitfalls / Exam Traps

- Calling a **single PDF** a corpus — a corpus requires multiple documents
- Confusing **corpus** (raw language) with **dataset** (task-specific labels)
- Assuming models **understand** language — they reproduce corpus statistics
- Ignoring that **TF-IDF, embeddings, and LMs** all derive statistics from the same underlying text

---

Quick Revision Summary

- Corpus = large, structured text collection for language study and modelling
- Document < Corpus < Dataset (with labels for tasks)
- Corpora provide language; datasets provide supervision
- TF-IDF, embeddings, topic models, LMs all depend on corpus statistics
- Language learning in NLP is statistical — GIGO applies

← **Previous** · **Index** · **Next** →
Types of Text Corpora

Types of Text Corpora

Why Classify Corpora?

Not all text corpora serve the same purpose. Genre, annotation level, and domain determine which corpus fits which NLP task. Choosing wrongly leads to **poor generalisation** and misinterpreted model outputs.

News language differs from conversational text; literary prose differs from product reviews; raw text differs from labelled annotations.

Taxonomy Overview

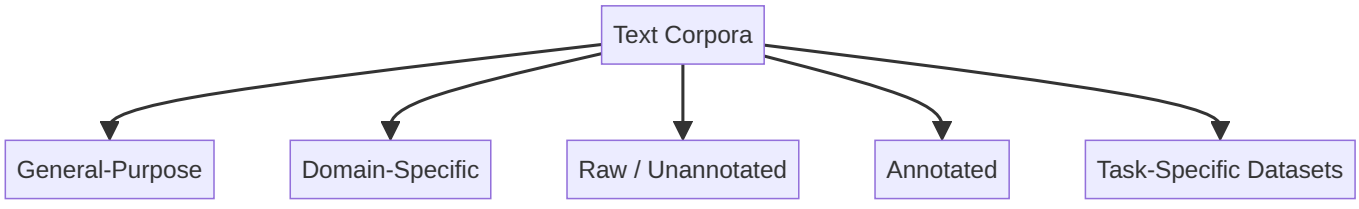

---

1. General-Purpose Corpora

Designed to represent **broad language usage** across genres and styles.

Corpus	Composition	Strengths
Brown	Balanced across news, fiction, reviews, etc.	Diverse vocabulary; multiple writing styles
Gutenberg	Literary classics and books	Rich vocabulary; long-form narrative

Limitation: Poor performance on specialised domains (finance, healthcare, legal) without fine-tuning or domain corpora.

---

2. Domain-Specific Corpora

Focused on a **particular field or application**.

Domain	Corpus Example	Task
Sentiment	Movie reviews (IMDb)	Sentiment classification
Clinical	Medical reports	Clinical NLP
Legal	Contracts, case law	Contract analysis

Properties: Specialised terminology, narrow context, high task relevance.

Critical rule: A sentiment model trained on movie reviews **will not** perform well on medical text.

---

3. Raw (Unannotated) Corpora

Plain text with **no linguistic or task labels**.

Examples	Use Cases
Gutenberg books	Language modelling
Wikipedia dumps	Word embeddings, topic modelling

Suited for **unsupervised** methods that learn from co-occurrence and distribution alone.

---

4. Annotated Corpora

Enriched with **linguistic or task-specific labels**:

- POS tags
- Named entities
- Sentiment labels (positive / negative / neutral)
- Syntactic dependencies

Examples	Use
POS-tagged Brown corpus	Supervised tagging evaluation
CoNLL NER datasets	NER training and benchmarking

Used for **supervised learning** and **evaluation**.

5. Task-Specific Datasets

Corpora created for a **defined input-output task**:

Task	Dataset
Sentiment analysis	IMDb reviews
Question answering	SQuAD
Text classification	AG News, etc.

Properties: Clear input-output pairs, high-quality labels, ideal for training and comparison.

Limitation: Narrow task scope; limited linguistic diversity beyond the task.

Comparison Table

Type	Labels	Diversity	Best For
General-purpose	Usually none	High	Foundational NLP, exploration
Domain-specific	Varies	Low (within domain)	Specialised applications
Raw	None	Varies	LM, embeddings, topic modelling
Annotated	Linguistic	Medium	Supervised tagging, NER
Task-specific	Task labels	Low	Model training, benchmarks

Common Pitfalls / Exam Traps

- Training a **medical NER model** on Gutenberg literary text
- Using **task-specific datasets** for general language modelling — too narrow
- Assuming **annotated corpora** are always large — many are small but high-quality
- Confusing **Gutenberg as general-purpose** — it skews heavily literary/archaic

---

Quick Revision Summary

- Five corpus types: general-purpose, domain-specific, raw, annotated, task-specific
- Brown = balanced genres; Gutenberg = literary classics
- Domain corpora essential for finance, health, legal NLP
- Raw corpora for unsupervised methods; annotated for supervised tagging/NER
- Match corpus type to task — wrong corpus → poor generalisation

---

← **Previous** · **Index** · **Next** →
Bias and Representation in Text Corpora

Bias and Representation in Text Corpora

Corpora Are Not Neutral

A text corpus is written by humans, selected by humans, and collected from specific sources. It naturally reflects:

- Social norms of the source community
- Cultural assumptions
- Historical inequalities

NLP models learn **directly** from corpora — they do not invent bias; they **absorb** it from training data.

\text{Model bias} \leftarrow \text{Corpus bias}
Bias is a **data problem before it is a model problem.**

Types of Corpus Bias

Bias Type	Description	Example
Source bias	Each source has its own lens	News vs social media vs academic text
Representation bias	Who appears, how often	Under-represented demographics; absent dialects
Historical bias	Old texts reflect past norms	Outdated terminology; stereotypical role language
Annotation bias	Human labelers inject subjectivity	Ambiguous cases simplified; cultural framing in labels

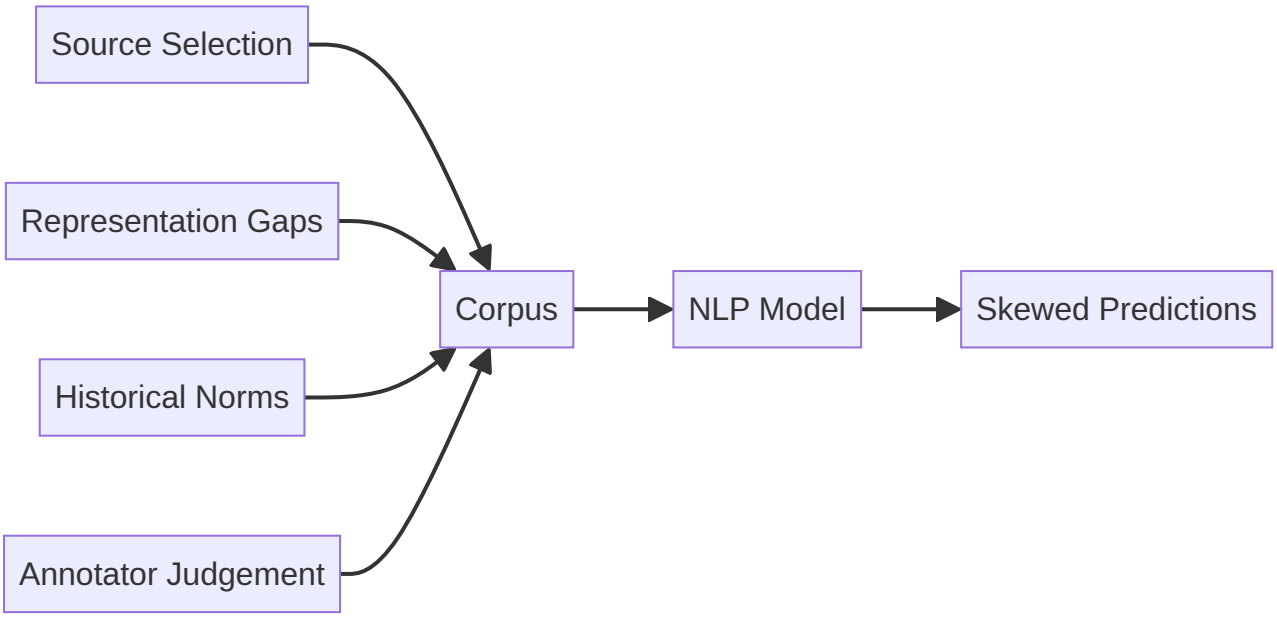

How Bias Manifests in Models

Manifestation	Example
Skewed predictions	Model performs worse on under-represented groups
Stereotypes in embeddings	*doctor* closer to *man*; *nurse* closer to *woman*
Unequal performance	Higher error rates on dialects absent from training data

Word embeddings capture **historical co-occurrence patterns** from text — models optimise statistical regularities, not fairness.

This is not a bug in the algorithm alone; it reflects **patterns in the corpus**.

Pre-Deployment Checklist

Before training or deploying any NLP model, ask:

1. **Where does this corpus come from?** (Source, time period, collection method)
2. **Who does it represent well?** (Demographics, domains, languages)
3. **Who might it miss?** (Minority dialects, non-Western contexts, rare entities)

Understanding data is as important as understanding model architecture.

Mitigation Strategies (Conceptual)

Strategy	Action
Audit	Explore corpus demographics and term frequencies
Diversify	Combine multiple sources to balance representation
Evaluate per group	Measure performance across subpopulations
De-bias embeddings	Post-processing or adversarial training (advanced)

Common Pitfalls / Exam Traps

- Claiming **"the model is biased"** without examining **training data** first
- Assuming **larger corpora eliminate bias** — scale can amplify dominant viewpoints
- Ignoring **annotation bias** in supervised datasets — labels carry human judgement
- Treating **historical text** as representative of current social norms

---

Quick Revision Summary

- Corpora reflect human choices — not neutral ground truth
- Bias types: source, representation, historical, annotation
- Models absorb corpus bias → skewed predictions and stereotyped embeddings
- Ask: where from, who represented, who missed — before deployment
- Bias is a data problem first; understanding corpus > blaming model alone

---

← Previous · Index · Next →

Corpus Exploration Setup with NLTK

Corpus Exploration Setup with NLTK

Why Explore a Corpus Before Modelling?

Corpus analysis reveals what data is actually available — structure, size, tokenisation, and genre composition. Skipping exploration treats text as a **black box**, hiding vocabulary gaps, preprocessing needs, and domain mismatch until models fail silently in production.

Exploration answers:

- How many documents and tokens exist?
- Is text pre-tokenised or raw?
- What genres or authors are represented?

NLTK Setup for Classic Corpora

```
import nltk
from nltk.corpus import gutenberg, brown

nltk.download('gutenberg')
nltk.download('brown')
```

Gutenberg Corpus Overview

The Gutenberg corpus contains **literary classics** — novels, poetry, and historical texts.

```
gutenberg.fileids()
# ['austen-emma.txt', 'austen-persuasion.txt', 'blake-poems.txt',
#  'bryant-stories.txt', 'shakespeare-hamlet.txt', ...]
```

Each file ID corresponds to one classic work.

Inspecting a Single Book

```
words = gutenberg.words('carroll-alice.txt')
sents = gutenberg.sents('carroll-alice.txt')

len(words)   # e.g., 34,110 tokens
len(sents)   # e.g., 1,703 sentences
```

NLTK provides **pre-tokenised** words and sentences — ready for frequency analysis without manual tokenisation.

Brown Corpus Preview

The Brown corpus represents **modern American English** balanced across multiple genres (news, fiction, religion, etc.). Deep analysis follows in a dedicated note; here we confirm availability:

```
brown.categories()
# ['adventure', 'belles_lettres', 'editorial', 'fiction', 'government',
#  'hobbies', 'humor', 'learned', 'lore', 'mystery', 'news', ...]
```

Exploration Workflow

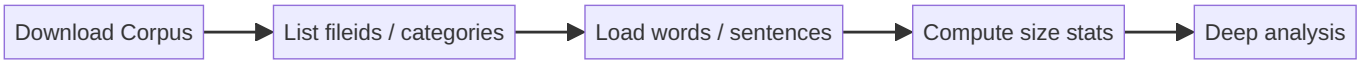

Step	Function	Output
List contents	`gutenberg.fileids()`	Available works
Word tokens	`gutenberg.words(fileid)`	Token list
Sentences	`gutenberg.sents(fileid)`	Sentence list
Categories	`brown.categories()`	Genre labels

Key Observation

Pre-tokenised NLTK corpora include **punctuation and capitalisation** from original texts. Raw counts differ from cleaned counts — always document whether statistics apply to raw or preprocessed text.

Common Pitfalls / Exam Traps

- Forgetting `nltk.download()` before accessing corpora
- Confusing **word count** with **unique vocabulary size**
- Assuming pre-tokenised text is **clean** — punctuation and caps remain
- Using **Gutenberg statistics** to characterise modern web language

---

Quick Revision Summary

- Explore corpora before modelling — avoid black-box assumptions
- NLTK: `gutenberg` and `brown` with `nltk.download()`
- Gutenberg: literary classics; `words()` and `sents()` pre-tokenised
- Brown: genre-balanced modern American English
- Raw corpus tokens include punctuation and capitalisation

← **Previous** · **Index** · **Next** →
Gutenberg Corpus Analysis: Vocabulary and Word Frequency

Gutenberg Corpus Analysis: Vocabulary and Word Frequency

Analysis Objectives

Using a single literary text (*Alice's Adventures in Wonderland*), this analysis explores:

- Corpus size (tokens and sentences)
- Effect of preprocessing on counts
- Vocabulary richness (type-token ratio)
- Word frequency distributions
- Content words after stopword removal

Loading and Raw Inspection

```
from nltk.corpus import gutenberg

fileid = 'carroll-alice.txt'
words = gutenberg.words(fileid)
sents = gutenberg.sents(fileid)

len(words)  # 34,110 total tokens
len(sents)  # 1,703 sentences
```

Raw tokens preserve **original capitalisation and punctuation** — reflecting the source text, not model-ready form.

Preprocessing for Analysis

Normalise for statistical analysis (distinct from modelling preprocessing):

```
clean_words = [w.lower() for w in words if w.isalpha()]
len(clean_words)  # 27,333 (down from 34,110)
```

Effect	Impact
Lowercasing	Reduces vocabulary inflation (*The* vs *the*)
`isalpha()` filter	Removes punctuation tokens and numbers

Vocabulary Richness

```
vocab = set(clean_words)
len(vocab)  # 2,569 unique words (types)
```

Type-Token Ratio (TTR):

$$\text{TTR} = \frac{|\text{types}|}{|\text{tokens}|} = \frac{2569}{27333} \approx 0.094$$

TTR measures **lexical diversity** — higher values suggest richer vocabulary relative to text length. Literary texts often show higher TTR than repetitive genres.

TTR provides **intuition**, not absolute judgement — compare within similar text lengths and genres.

Word Frequency Distribution

```
from collections import Counter
freq = Counter(clean_words)
freq.most_common(10)
# [('the', 1642), ('to', 872), ('and', 702), ...]
```

Top tokens are **function words** (stopwords) — low semantic content, high frequency.

After stopword removal:

```
from nltk.corpus import stopwords
stops = set(stopwords.words('english'))
content_words = [w for w in clean_words if w not in stops]
Counter(content_words).most_common(10)
# [('said', ...), ('alice', ...), ('queen', ...), ...]
```

Content words reveal **thematic signals** — character names, plot vocabulary — useful for topic modelling, summarisation, and sentiment analysis.

Zipfian Distribution Pattern

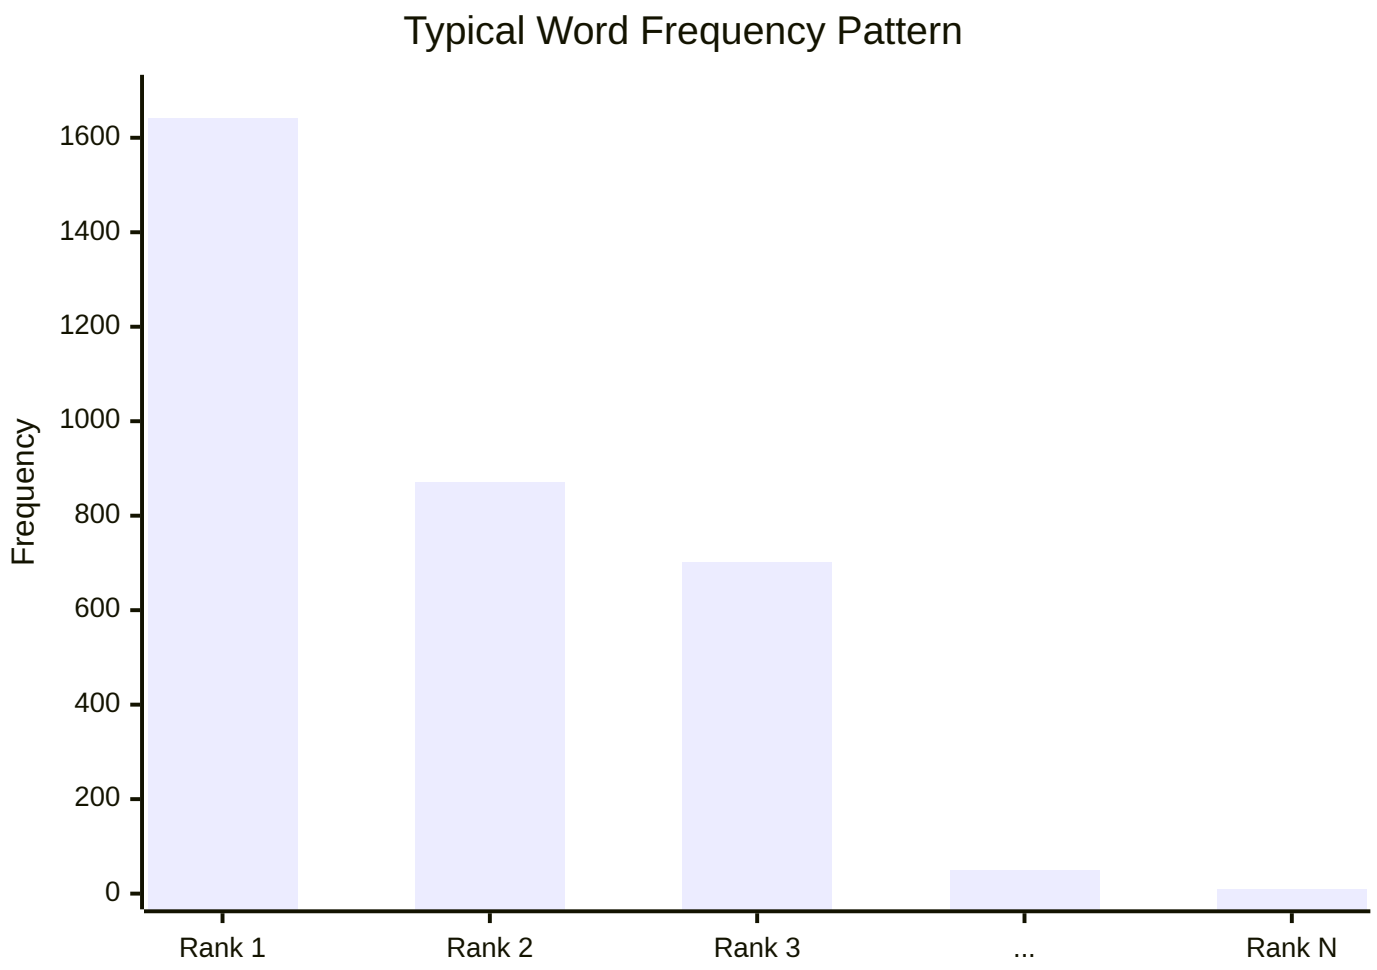

Few words dominate (head); a long tail of rare words follows. This **power-law** pattern appears across virtually all natural language corpora.

Visualisation

Matplotlib bar charts of `freq.most_common(20)` for raw vs stopword-filtered distributions make the shift from function words to content words visually clear.

Common Pitfalls / Exam Traps

- Reporting **raw token count** without noting punctuation inclusion
- Confusing **types** (unique words) with **tokens** (total count)
- Using TTR to compare **texts of very different lengths** without normalisation
- Assuming stopword removal is appropriate for **all downstream tasks**

Quick Revision Summary

- Gutenberg *Alice*: ~34K raw tokens → ~27K alphabetic lowercase tokens
- Vocabulary: 2,569 unique types; TTR ≈ 0.094 (lexical diversity measure)
- Raw freq dominated by *the*, *to*, *and*; content words emerge after stopword removal
- Few words dominate; long tail of rare words (Zipfian pattern)
- Preprocessing for analysis ≠ preprocessing for modelling — document which applies

---

← **Previous** · **Index** · **Next** →
Brown Corpus Analysis: Cross-Genre Language Comparison

Brown Corpus Analysis: Cross-Genre Language Comparison

Why the Brown Corpus?

The Brown corpus differs fundamentally from Gutenberg:

Aspect	Gutenberg	Brown
Era	Classic literature	Modern American English
Composition	Single-author books	Balanced across genres
Purpose	Literary language study	Representative language usage
Design	Author/style-specific	Multi-genre sampling

Brown enables **quantitative comparison** of how language varies across domains.

Exploring Categories

```
from nltk.corpus import brown

brown.categories()
# ['adventure', 'belles_lettres', 'editorial', 'fiction', 'government',
#  'hobbies', 'humor', 'learned', 'lore', 'mystery', 'news', 'religion',
#  'reviews', 'romance', 'science_fiction', ...]

len(brown.words())  # ~1,161,192 total tokens (entire corpus)
```

Each category represents a **genre** of English — news, fiction, government prose, etc.

Comparing Genre Sizes

```
news_words = brown.words(categories='news')
fiction_words = brown.words(categories='fiction')

len(news_words)    # e.g., ~46,000+ tokens (varies by preprocessing)
len(fiction_words) # e.g., ~68,000+ tokens
```

Token count alone does not explain lexical richness — compare vocabulary metrics next.

Vocabulary Metrics by Genre

Preprocess: keep alphabetic tokens only.

```
def preprocess(words):
    return [w.lower() for w in words if w.isalpha()]

news_clean = preprocess(news_words)
fiction_clean = preprocess(fiction_words)

news_vocab = set(news_clean)
fiction_vocab = set(fiction_clean)

news_ttr = len(news_vocab) / len(news_clean)
fiction_ttr = len(fiction_vocab) / len(fiction_clean)
```

Example observations:

Genre	Tokens (clean)	Unique types	TTR
News	~11,151	~1,488	~0.133
Fiction	~8,141	~1,161	~0.143

Fiction often shows **higher lexical diversity** (TTR) despite fewer total tokens. News language is more **standardised and repetitive** — institutional terms recur (*said*, *county*, *jury*).

Frequency Analysis with Stopwords

```
from collections import Counter
from nltk.corpus import stopwords

stops = set(stopwords.words('english'))

def top_content_words(words, n=10):
    filtered = [w for w in words if w not in stops]
    return Counter(filtered).most_common(n)
```

Genre	Top content words (examples)
News	*fulton*, *county*, *grand*, *jury*, *investigation*, *said*
Fiction	*would*, *said*, *one*, *could*, *like*, *man*, *back*

News exhibits **institutional/formal** vocabulary; fiction uses **narrative/descriptive** terms — clear **domain signatures**.

Custom Stopwords

Domain-specific high-frequency words may lack semantic value:

```
stops.add('said')  # common in news but uninformative for topic analysis
```

Custom stopword lists improve task-specific analysis — *said* disappears from top-N lists after addition.

Visualisation

Matplotlib bar charts comparing `most_common(20)` for news vs fiction (after stopword removal) with rotated x-labels (`plt.xticks(rotation=45)`) reveal genre-specific frequency profiles.

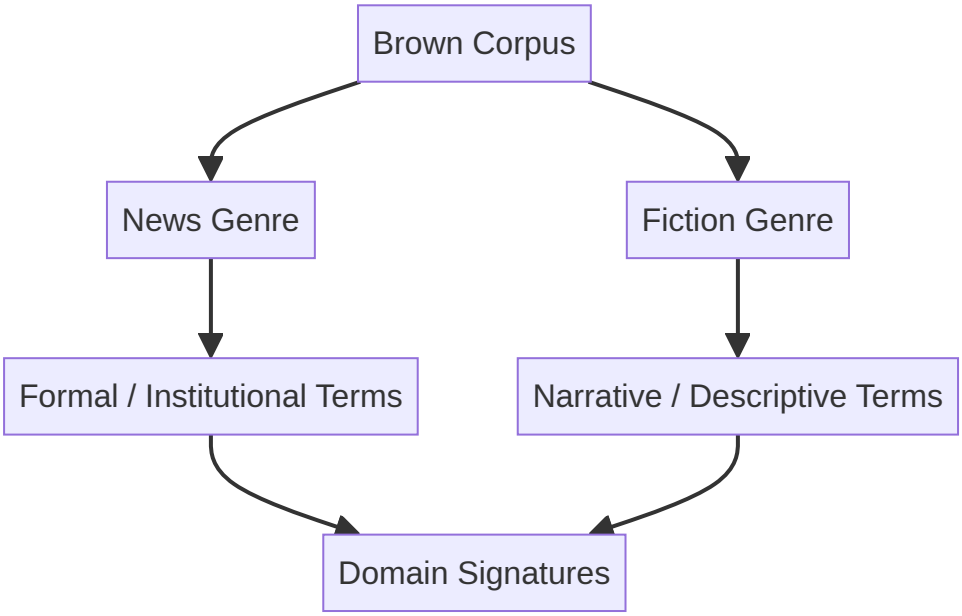

Key Observations

1. **Same language, different statistics** — English words behave differently by genre
2. **High-frequency words are domain-specific** after stopword removal
3. **TTR and token counts** must be interpreted together — not in isolation
4. **Custom stopwords** refine analysis for specific use cases

---

Common Pitfalls / Exam Traps

- Comparing **raw token counts** without alphabetic preprocessing
- Assuming **higher token count = richer vocabulary** — TTR may disagree
- Using **default stopwords only** — domain terms like *said* in news may need custom removal
- Generalising **Brown (American English, 1960s)** to all modern English web text

Quick Revision Summary

- Brown: ~1.16M tokens, balanced across 15+ genres of modern American English
- News vs fiction: different token counts, TTR, and content-word profiles
- News → institutional terms; fiction → narrative vocabulary
- TTR: fiction often higher lexical diversity; news more repetitive
- Custom stopwords refine genre analysis; visualise with frequency bar charts

---

← **Previous** · **Index** · **Next** →

Week 5

Week 5

Traditional Word Embeddings: Module Overview

Traditional Word Embeddings: Module Overview

Why Machines Need Numbers, Not Text

Natural language processing models are mathematical engines. They operate on tensors of floating-point values — not on characters, sentences, or meaning in the human sense. Before any classifier, clustering algorithm, or neural network can process text, that text must be converted into a numerical representation. This conversion is called **vectorization** or **word embedding**.

Consider a cloud-based spam filter at scale: millions of emails arrive per hour. The pipeline cannot "read" each message; it must transform subject lines and body text into feature vectors that a logistic regression model or gradient-boosted tree can score in milliseconds.

The Evolution of Text Representation

Text processing in NLP has progressed through distinct eras:

Era	Approach	Example Techniques
Rule-based	Hand-crafted linguistic rules	Regex patterns, stemming, lemmatization
Statistical	Count-based numerical features	One-hot encoding, BoW, TF-IDF, n-grams
Dense semantic	Learned continuous vectors	Word2Vec, GloVe, contextual embeddings

Earlier modules covered **preprocessing** — cleaning and normalizing text. This module enters **statistical NLP**, where words become vectors and mathematical operations (dot products, cosine similarity, matrix multiplication) become meaningful.

Core Techniques Covered

This module introduces four foundational vectorization methods:

- **One-hot encoding (OHE)** — each word is an isolated category; presence/absence only

- **Bag of Words (BoW)** — word frequencies within a document; order discarded
- **TF-IDF** — frequency weighted by rarity across the corpus; downweights common words
- **N-grams** — contiguous token sequences that partially recover word order

Each technique trades off between **sparsity** (memory efficiency vs. wasted zeros) and **semantic retention** (whether related words appear mathematically similar).

The Central Trade-off

Dimension	Sparse methods (OHE, BoW, TF-IDF)	Dense methods (Word2Vec, GloVe)
Vector size	Vocabulary size (10,000+)	Fixed small dimension (50–300)
Semantic similarity	None — all words orthogonal	Captured — similar words cluster
Interpretability	High — each dimension is a word	Low — dimensions are latent features
Compute cost	Low for small vocab	Higher training cost

Understanding these trade-offs is essential before moving to dense embeddings in the next module.

Common Pitfalls / Exam Traps

- **Confusing vectorization with preprocessing** — stemming and stopword removal clean text; vectorization converts cleaned text to numbers. They are sequential steps, not the same operation.
- **Assuming "embedding" always means dense vectors** — one-hot encoding and BoW are also called word embeddings in the broad sense, but they are sparse and non-semantic.
- **Ignoring the vocabulary problem** — every vectorization method requires building a vocabulary from the training corpus. Words unseen at test time are typically dropped or mapped to an unknown token.
- **Treating all vectorization methods as interchangeable** — TF-IDF suits document classification; one-hot suits small categorical vocabularies; neither captures that "car" and "automobile" are synonyms.

Quick Revision Summary

- NLP models require numerical input; vectorization bridges human language and machine computation.
- Statistical NLP converts words to vectors so mathematical operations can be applied.
- One-hot encoding, BoW, TF-IDF, and n-grams are the foundational sparse techniques.
- The core trade-off is sparsity (memory, dimensionality) versus semantic retention.
- Rule-based preprocessing precedes vectorization; they solve different problems.
- Dense embeddings (next module) address the semantic gap that sparse methods leave open.
- Real-world pipelines (spam filters, search ranking, document classifiers) depend on choosing the right representation for the task.

← **Previous** · **Index** · **Next** →
One-Hot Encoding for Text

One-Hot Encoding for Text

Intuition: Words as Categories

One-hot encoding (OHE) is the simplest way to represent text numerically. It treats every unique word in a **vocabulary** as an isolated category — much like encoding a color field with values `red`, `blue`, `green` in a tabular dataset.

The key idea: a word is either present in a document (1) or absent (0). No partial credit, no frequency count, no semantic relationship between words.

Building a Vocabulary

Given a corpus of documents, extract every **unique** word across all sentences. This collection is the **vocabulary**.

Example corpus:

Sentence
I love NLP
I love AI
NLP is the future

Vocabulary (unique words): `AI`, `I`, `NLP`, `future`, `is`, `love` — six words total.

Constructing the One-Hot Matrix

Each sentence becomes a row; each vocabulary word becomes a column. Cell values indicate whether that word appears in that sentence (1 = present, 0 = absent).

Sentence	AI	I	NLP	future	is	love
I love NLP	0	1	1	0	0	1
I love AI	1	1	0	0	0	1
NLP is the future	0	0	1	1	1	0

Notice:

- "I love NLP" and "I love AI" share three active dimensions (`I`, `love`, and either `NLP` or `AI`).
- Repeated words within a sentence still register as 1 — OHE does not count frequency.

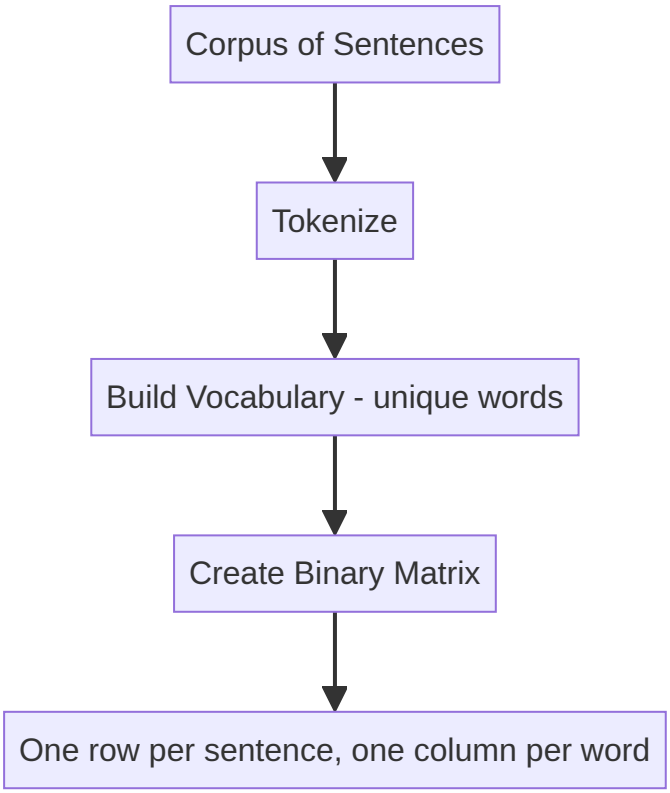

Mathematical Form

For vocabulary size |V| and document d, the one-hot vector \mathbf{x}_d \in \{0, 1\}^{|V|} satisfies:

x_i =

$$\begin{cases} 1 & \text{if word } w_i \text{ appears in } d \\ 0 & \text{otherwise} \end{cases}$$

At most |V| entries can be 1 per sentence (fewer if the sentence is shorter than the vocabulary).

When OHE Makes Sense

OHE works well when:

- The vocabulary is **small** (dozens to low hundreds of words)
- Only **presence** matters, not frequency or order
- Interpretability is valued — each dimension has a clear label

In cloud ML contexts, OHE appears in feature stores for categorical fields (e.g., product category tags in a recommendation pipeline) and as a baseline text representation before investing in embedding infrastructure.

Common Pitfalls / Exam Traps

- **Confusing OHE with BoW** — OHE marks presence (0/1); BoW counts frequency (0, 1, 2, ...). "I love love" would be `[0,1,0,0,2]` in BoW but `[0,1,0,0,1]` in OHE.
- **Forgetting vocabulary ordering** — columns must stay in a consistent order across all documents; shuffling vocabulary indices breaks the representation.
- **Assuming OHE captures word importance** — common words and rare words receive equal weight if both appear once.

- **Exam trap: "OHE preserves word order"** — it does not. "dog bites man" and "man bites dog" produce identical vectors.

Quick Revision Summary

- One-hot encoding represents each word as a binary indicator in a fixed-size vocabulary vector.
- Vocabulary = set of all unique words in the corpus.
- Each sentence becomes a row of 0s and 1s — 1 where the word is present, 0 otherwise.
- OHE is the simplest vectorization method but ignores frequency, order, and semantics.
- Works best for small vocabularies and presence-only tasks.
- Distinct from BoW (which counts) and TF-IDF (which weights by rarity).

← **Previous** · **Index** · **Next** →
Implementing One-Hot Encoding in Python

Implementing One-Hot Encoding in Python

Intuition: From Concept to Code

Understanding the matrix on paper is one step; implementing it reveals the pipeline every NLP system follows: **tokenize** → **build vocabulary** → **encode**. Manual implementation clarifies what libraries like scikit-learn's `CountVectorizer` automate under the hood.

Pipeline Overview

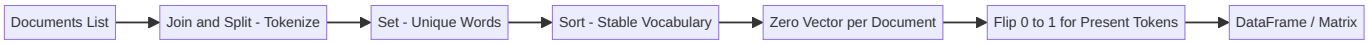

Step 1: Prepare the Corpus

```
documents = [
    "I love NLP",
    "I love AI",
    "NLP is the future"
]
```

Step 2: Build the Vocabulary

Join all sentences, split on whitespace, deduplicate with a set, then sort for reproducibility:

```
vocab = sorted(set(" ".join(documents).split()))
# ['AI', 'I', 'NLP', 'future', 'is', 'love']
```

Operation	Purpose
`" ".join(documents)`	Flatten corpus into one string
`.split()`	Tokenize on whitespace
`set(...)`	Remove duplicates
`sorted(...)`	Deterministic column order

Step 3: Encode Each Document

For each document:

1. Tokenize the sentence
2. Create a zero vector of length `len(vocab)`
3. For each token present in the vocabulary, set the corresponding index to 1
4. Append the vector to the result list

```
one_hot_vectors = []
for doc in documents:
    tokens = doc.split()
    vector = [0] * len(vocab)
    for token in tokens:
        if token in vocab:
            vector[vocab.index(token)] = 1
    one_hot_vectors.append(vector)
```

Result:

Sentence	Vector
I love NLP	`[0, 1, 1, 0, 0, 1]`
I love AI	`[1, 1, 0, 0, 0, 1]`
NLP is the future	`[0, 0, 1, 1, 1, 0]`

Step 4: Display as a DataFrame

```
import pandas as pd

df = pd.DataFrame(one_hot_vectors, columns=vocab, index=documents)
```

The DataFrame makes the matrix human-readable: rows are sentences, columns are vocabulary words, values are 0 or 1.

Production Alternatives

Tool	Class	Notes
scikit-learn	`CountVectorizer(binary=True)`	`binary=True` enforces 0/1 instead of counts
pandas	`pd.get_dummies()`	Works for token-level, not document-level
Keras	`keras.utils.to_categorical()`	Common for label encoding, adaptable for tokens

For large-scale pipelines (e.g., AWS SageMaker feature engineering), pre-built vectorizers handle out-of-vocabulary words, n-grams, and sparse matrix storage automatically.

Common Pitfalls / Exam Traps

- **Case sensitivity** — `"NLP"` and `"nlp"` are different tokens unless lowercased during preprocessing.
- **Using `list.index()` in a loop** — works for demos but is O(n) per lookup; production code uses a `dict` mapping word → index.
- **Forgetting to sort the vocabulary** — without `sorted()`, set iteration order is non-deterministic across Python runs.
- **Exam trap: "join then split equals original tokens"** — only true for space-separated tokens; punctuation attached to words (`"NLP."`) becomes a separate token.

Quick Revision Summary

- OHE implementation: tokenize corpus → build sorted vocabulary → create zero vectors → set 1 for present tokens.
- `set()` removes duplicates; `sorted()` ensures reproducible column order.
- Output is a list of lists (or sparse matrix) with one row per document.
- Pandas DataFrame aids visualization with sentences as index and words as columns.
- Production systems use `CountVectorizer(binary=True)` or equivalent for efficiency.
- Case handling and tokenization strategy must be decided before vocabulary construction.

← **Previous** · **Index** · **Next** →
Advantages and Limitations of One-Hot Encoding

Advantages and Limitations of One-Hot Encoding

Intuition: Know When to Use a Tool

One-hot encoding is the "Hello World" of text vectorization. Its strengths are transparency and simplicity; its weaknesses become critical at scale. Choosing OHE without understanding these trade-offs leads to wasted memory, poor model performance, and false confidence in semantic understanding.

Advantages

Advantage	Explanation
Simplicity	Easy to understand, implement, and debug — no training required
No information loss on presence	If a word appears, it is recorded exactly; nothing is averaged or compressed away
Interpretability	Each dimension maps to a known word; feature importance is directly readable
No training data dependency	Vocabulary is built from any corpus instantly — no GPU, no epochs

In practice, OHE suits **small, controlled vocabularies**: intent labels in a chatbot router, product SKU categories in an e-commerce classifier, or keyword flags in a rule-augmented ML pipeline.

Limitations

1. Extreme Sparsity

For a vocabulary of 10,000 words, each sentence vector has at most a handful of 1s and **9,995+ zeros**. This wastes memory and slows computation.

\text{Sparsity ratio} \approx 1 - \frac{\text{words per sentence}}{|V|}
A 10-word sentence in a 10,000-word vocabulary is 99.9% zeros.

2. No Semantic Meaning

Related words are **orthogonal** — their dot product is zero:

\text{dog} \cdot \text{puppy} = 0
The model has no signal that these words are related. In a customer-support ticket classifier, "refund" and "reimbursement" are treated as completely unrelated features.

3. Variable-Length Sentences

Sentences of different lengths produce vectors of the **same fixed size** (vocabulary length), but the number of active (non-zero) dimensions varies. Without padding or aggregation strategies, comparing raw sentence vectors across documents of very different lengths can mislead distance-based algorithms.

4. No Frequency Information

"I love love love NLP" and "I love NLP" produce identical vectors — OHE cannot distinguish emphasis or repetition.

5. Vocabulary Explosion

Every new unique word adds a dimension. Open-vocabulary domains (social media, medical notes) make the matrix impractically wide.

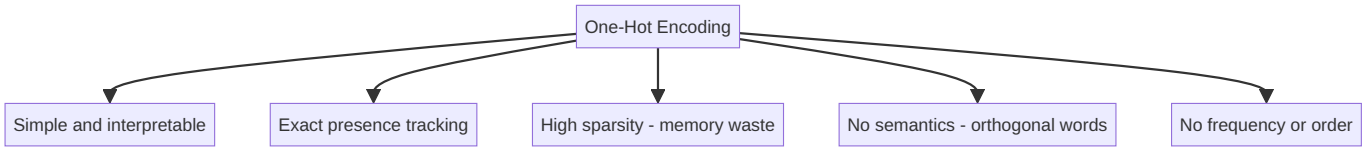

Comparison with Alternatives

Property	One-Hot	Bag of Words	TF-IDF
Values	0 or 1	Integer counts	Float weights
Frequency	Ignored	Captured	Captured + weighted
Semantics	None	None	None
Sparsity	Very high	Very high	Very high

All three sparse methods share the semantic blindness — that gap is addressed by dense embeddings (Word2Vec, GloVe, BERT).

Common Pitfalls / Exam Traps

- **"OHE is lossless"** — true for presence, false overall. Frequency, order, and meaning are all discarded.
- **Using OHE for large vocabularies** — exam questions may ask when OHE fails; answer: high dimensionality and no semantic similarity.

- **Confusing orthogonality with independence** — orthogonal vectors are mathematically unrelated; linguistically related words should not be orthogonal, but OHE makes them so.
- **Assuming sparsity is always bad** — sparse matrices with efficient storage (CSR format) mitigate memory; the deeper issue is lack of semantics.

Quick Revision Summary

- OHE advantages: simple, interpretable, no training, preserves word presence exactly.
- OHE limitations: extreme sparsity, no semantic similarity, no frequency, no word order.
- "Dog" and "puppy" are orthogonal under OHE — the model cannot learn they are related.
- A 10,000-word vocabulary means ~9,999 zeros per sentence vector.
- OHE suits small vocabularies and presence-only tasks; not large-scale semantic NLP.
- BoW and TF-IDF share OHE's semantic blindness but add frequency and weighting respectively.

← **Previous** · **Index** · **Next** →
Bag of Words (BoW)

Bag of Words (BoW)

Intuition: A Multisets of Words

Bag of Words treats a document as a **multiset** (bag) of its words — like shaking a Scrabble tile bag and counting what fell out. Grammar, syntax, and word order are discarded entirely. What remains is **how many times** each word appears.

The name captures the metaphor: tip a document into a bag, shake it, and count the tiles. "The cat sat on the mat" and "mat the on sat cat the" produce identical bags.

What BoW Answers

BoW answers two questions for each word in the vocabulary:

1. **Does this word appear in the document?**
2. **How many times?**

This is a step beyond one-hot encoding, which only answers the first question.

Worked Example

Corpus:

Document
The cat sat on the mat
The dog sat on the log
Cats and dogs are enemies

Vocabulary (sample): `the`, `cat`, `sat`, `on`, `mat`, `dog`, `log`, `cats`, `and`, `dogs`, `are`, `enemies`

BoW vector for "The cat sat on the mat":

Word	Count
the	2
cat	1
sat	1
on	1
mat	1
all others	0

The word `the` appears twice — BoW records `2`, whereas one-hot encoding would record `1`.

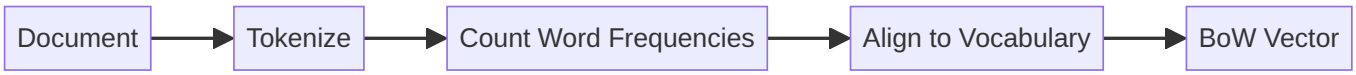

BoW vs One-Hot Encoding

Aspect	One-Hot	Bag of Words
Values	0 or 1	Non-negative integers
Repeated words	Counted once	Counted each time
Information captured	Presence	Frequency
Word order	Lost	Lost
Semantics	None	None

Why BoW Persists in Production

Despite its simplicity, BoW remains useful in cloud ML pipelines for:

- **Spam detection** — keywords like "OTP", "lottery", "urgent" appear with characteristic frequencies
- **Topic routing** — support tickets classified by dominant term counts
- **Baseline models** — fast to build, easy to benchmark before investing in embeddings

BoW is a strong baseline: if a complex model cannot beat BoW on a keyword-driven task, the added complexity is not justified.

Common Pitfalls / Exam Traps

- **"Bag of words preserves multiplicity but not order"** — a classic exam phrasing; multiplicity means repeat counts are kept.
- **Confusing BoW with TF-IDF** — BoW uses raw counts; TF-IDF scales counts by inverse document frequency.
- **Case and stemming matter** — `"Cat"` and `"cats"` are different dimensions unless normalized.
- **Identical BoW for different sentences** — "the man bit the dog" and "the dog bit the man" produce the same vector.

Quick Revision Summary

- Bag of Words represents a document as word frequency counts over a fixed vocabulary.

- Word order and grammar are completely ignored; only multiplicity (count) is preserved.
- BoW extends one-hot encoding by recording frequency, not just presence.
- Effective for keyword-driven tasks like spam detection and simple classification.
- Does not capture semantics — related words remain unrelated in vector space.
- Serves as a fast, interpretable baseline in production ML pipelines.

← **Previous** · **Index** · **Next** →
Implementing Bag of Words with scikit-learn

Implementing Bag of Words with scikit-learn

Intuition: Automating the Counting Pipeline

Manual BoW implementation mirrors one-hot encoding but increments counters instead of flipping bits. In production, `CountVectorizer` from scikit-learn handles tokenization, vocabulary building, counting, and sparse matrix storage in a single `fit_transform` call.

Pipeline Overview

Step 1: Import and Prepare Corpus

```
from sklearn.feature_extraction.text import CountVectorizer

corpus = [
    "The cat sat on the mat",
    "The dog sat on the log",
    "Cats and dogs are enemies"
]
```

Step 2: Fit and Transform

```
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(corpus)
```

`fit_transform` performs two operations:

- **fit** — learns the vocabulary from the corpus
- **transform** — converts each document to a count vector

Output shape: `(3, 12)` — 3 documents, 12 unique tokens in vocabulary.

The result is a **compressed sparse row (CSR) matrix** — efficient storage when most values are zero.

Step 3: Inspect the Matrix

```
import numpy as np
X.toarray()
```

BoW vector for "The cat sat on the mat":

Word	Count
cat	1
mat	1
on	1
sat	1
the	2

The word `the` appears twice — the distinguishing feature of BoW over one-hot encoding.

```
import pandas as pd

df = pd.DataFrame(
    X.toarray(),
    columns=vectorizer.get_feature_names_out(),
    index=corpus
)
```

Key CountVectorizer Parameters

Parameter	Default	Effect
`binary`	`False`	Set `True` for one-hot-style presence only
`lowercase`	`True`	Converts all tokens to lowercase
`stop_words`	`None`	Remove common words (`'english'` available)
`ngram_range`	`(1, 1)`	Use `(1, 2)` for unigrams + bigrams
`max_features`	`None`	Limit vocabulary to top-N frequent terms
`min_df` / `max_df`	`1` / `1.0`	Filter rare or overly common terms

For a document classification API on Azure ML, `max_features=10000` and `min_df=2` are common starting points to control dimensionality.

Sparse vs Dense Storage

Format	When to Use
CSR sparse matrix	Default; efficient for high-dimensional text
Dense array (`.toarray()`)	Debugging, small vocabularies, visualization
DataFrame	Human-readable inspection

A `(10000, 50000)$ sparse matrix may store only millions of non-zero entries instead of 500 million floats.

Common Pitfalls / Exam Traps

- **Forgetting `fit` before `transform` on new data** — vocabulary must be learned on training data only; applying a test corpus to `fit_transform` leaks future vocabulary.
- **Not lowercasing consistently** — `CountVectorizer` lowercases by default; custom preprocessing that preserves case creates mismatches.
- **Confusing `binary=True` with BoW** — `binary=True` produces one-hot-style vectors, not frequency counts.
- **Exam trap: output dtype** — BoW produces integers; TF-IDF produces floats.

Quick Revision Summary

- `CountVectorizer` from scikit-learn implements BoW via `fit_transform`.
- Output is a sparse matrix: rows = documents, columns = vocabulary, values = word counts.
- `the` appearing twice in a sentence yields count 2, not 1.
- Convert to DataFrame with `get_feature_names_out()` for readable inspection.
- Key parameters: `binary`, `stop_words`, `ngram_range`, `max_features`, `min_df`.
- Always fit on training data only; transform test data with the learned vocabulary.

← **Previous** · **Index** · **Next** →
Advantages and Limitations of Bag of Words

Advantages and Limitations of Bag of Words

Intuition: A Workhorse with Blind Spots

Bag of Words is deceptively powerful for keyword-driven tasks and deceptively weak for anything requiring structure or meaning. Knowing both sides prevents over-engineering (using transformers for spam detection) and under-engineering (using BoW for machine translation).

Advantages

Advantage	Explanation
Simple to implement	One `CountVectorizer` call in scikit-learn
Computationally efficient	No training; linear in corpus size
Surprisingly effective for keyword tasks	Spam, topic detection, intent routing
Interpretable features	Each dimension is a word; feature weights are directly explainable
Frequency signal	Unlike OHE, repeated keywords increase the count

Real-World Example: Spam Detection

Spam emails share characteristic vocabulary patterns:

- High counts of: `OTP`, `lottery`, `winner`, `urgent`, `click`, `free`
- Low counts of: normal conversational words

A BoW + logistic regression pipeline on AWS Lambda can score incoming emails in sub-millisecond latency without GPU infrastructure. Grammar is irrelevant — the keywords tell the story.

Limitations

1. Complete Loss of Context and Order

These two sentences produce **identical** BoW vectors:

- "The man bit the dog"
- "The dog bit the man"

The meaning is opposite; the representation is the same. Any task where word order determines semantics (negation, question vs statement, subject-verb-object) fails with BoW.

2. High Dimensionality and Sparsity

Like OHE, BoW vectors have dimension equal to vocabulary size. A 50,000-word vocabulary means 50,000-dimensional vectors that are ~99% zeros. Memory and compute scale poorly without sparse matrix formats.

3. No Semantic Understanding

"Excellent" and "outstanding" are unrelated dimensions. The model must learn their equivalence purely from co-occurrence patterns in labeled data — with no structural hint.

4. Vocabulary Growth

Every new unique token (typos, hashtags, product names) expands the matrix. Open-domain text (Twitter, customer reviews) creates unstable, high-dimensional features.

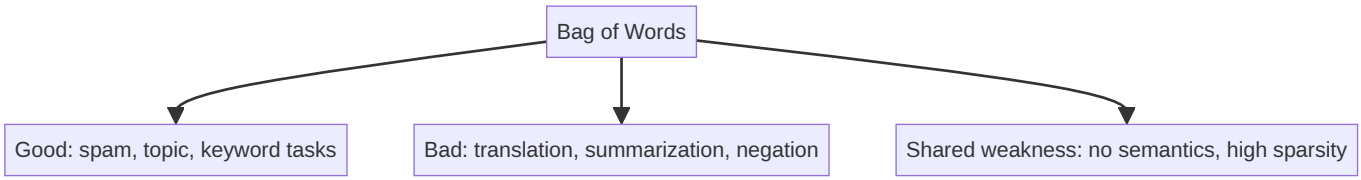

BoW vs TF-IDF vs Dense Embeddings

Criterion	BoW	TF-IDF	Word2Vec / BERT
Captures frequency	Yes	Yes (weighted)	N/A
Downweights common words	No	Yes	Learned
Captures word order	No	No	Partially / Yes
Captures semantics	No	No	Yes
Training required	No	No	Yes

Common Pitfalls / Exam Traps

- **"BoW is better than OHE because it counts"** — true for frequency, but both lose order and semantics.
- **Using BoW for sentiment with negation** — "not good" and "good" share the word `good`; BoW cannot distinguish them without n-grams.
- **Exam classic: identical vectors** — always cite "the man bit the dog" vs "the dog bit the man".
- **Assuming BoW fails on all tasks** — it remains competitive for keyword-driven classification at scale.

Quick Revision Summary

- BoW advantages: simple, fast, interpretable, effective for keyword-driven tasks like spam detection.
- BoW limitations: no word order, no semantics, high sparsity, vocabulary explosion.
- "The man bit the dog" and "the dog bit the man" are identical under BoW.
- Frequency counts distinguish BoW from one-hot encoding.
- TF-IDF improves BoW by downweighting common words; dense embeddings address semantics.
- Choose BoW as a baseline; upgrade when order, negation, or meaning matter.

← **Previous** · **Index** · **Next** →
TF-IDF: Term Frequency–Inverse Document Frequency

TF-IDF: Term Frequency–Inverse Document Frequency

Intuition: Not All Words Are Equal

Bag of Words treats every word equally — `the` and `mitochondria` both contribute to the vector. But common words carry little discriminative power, while rare, domain-specific terms are highly informative.

TF-IDF corrects this imbalance by asking two questions:

1. **How important is this word in this document?** (Term Frequency)
2. **How rare is this word across the entire corpus?** (Inverse Document Frequency)

The result: common words are downweighted; distinctive words are highlighted.

The Two Components

Term Frequency (TF)

How often word x appears in a single document d:

\text{TF}(x, d) = \text{count of } x \text{ in } d
A word appearing 5 times in a document is more locally important than one appearing once.

Inverse Document Frequency (IDF)

How rare word x is across all N documents in the corpus:

\text{IDF}(x) = \log\frac{N}{\text{DF}(x)}
where \text{DF}(x) = number of documents containing word x.

- Word in **every** document → \text{DF} = N → \log(N/N) = \log(1) = 0 → no weight
- Word in **one** document → \text{DF} = 1 → \log(N/1) → maximum weight

---

Combined Formula

\text{TF-IDF}(x, d) = \text{TF}(x, d) \times \log\frac{N}{\text{DF}(x)}
The \log dampens the effect so that very rare words do not dominate excessively.

Worked Example

Corpus:

Document
The car is fast
The race car is fast
The car is parked

Intuition for the word `car`:

- Appears in all 3 documents → lower IDF than `race` or `parked`
- Appears twice in document 2 → higher TF in that document
- TF-IDF score for `car` varies per document based on local frequency and global rarity

Intuition for the word `the`:

- Appears in every document → IDF approaches 0 → heavily downweighted

Intuition for `mitochondria` in a biology corpus:

- Appears in few documents → high IDF → strongly weighted when present

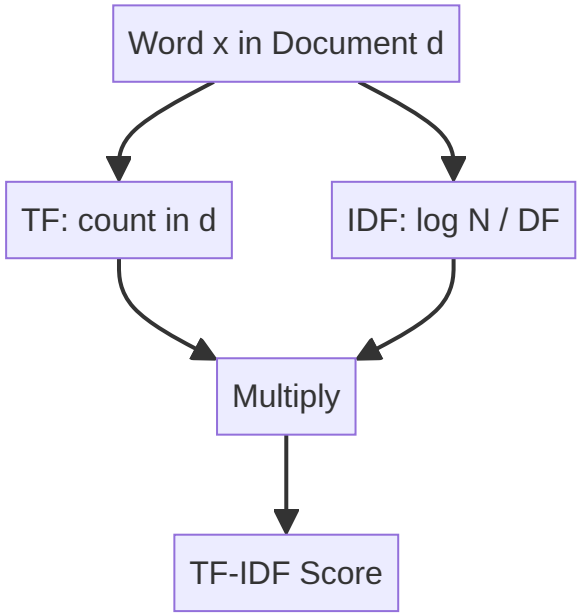

TF-IDF vs Bag of Words

Aspect	BoW	TF-IDF
Values	Raw integer counts	Float weights
Common words	High counts, high weight	Downweighted automatically
Rare domain terms	Low counts, low weight	Boosted by high IDF
Stopword removal	Often needed manually	Partially automatic

When TF-IDF Excels

- **Document classification** — news categorization, legal document routing
- **Information retrieval** — search engines rank documents by query term TF-IDF overlap
- **Similarity matching** — cosine similarity on TF-IDF vectors for FAQ deduplication in cloud support systems

---

Common Pitfalls / Exam Traps

- **Confusing TF and IDF** — TF is per-document frequency; IDF is corpus-wide rarity. Exam questions often swap them.
- **Forgetting the log in IDF** — without log, rare words dominate too aggressively.
- **"TF-IDF eliminates need for stopwords"** — mostly true, but not perfectly; explicit stopword removal can still help.
- **Assuming TF-IDF captures semantics** — it does not. "Car" and "automobile" remain unrelated.

Quick Revision Summary

- TF-IDF = Term Frequency × Inverse Document Frequency.
- TF measures how often a word appears in one document.
- IDF = \log(N / \text{DF}) — rare words get higher weight; universal words get near-zero weight.
- Automatically downweights common words like `the`, `is`, `and`.
- Boosts domain-specific rare terms like `mitochondria` in specialized corpora.
- Improves on BoW for classification and retrieval but still lacks semantic understanding.

← **Previous** · **Index** · **Next** →
Implementing TF-IDF in Python

Implementing TF-IDF in Python

Intuition: From Formula to Sparse Matrix

Implementing TF-IDF by hand requires computing term counts, document frequencies, and log ratios for every word-document pair. scikit-learn's `TfidfVectorizer` encapsulates this entire pipeline — the same `fit_transform` pattern as `CountVectorizer`, but output values are floats, not integers.

Pipeline Overview

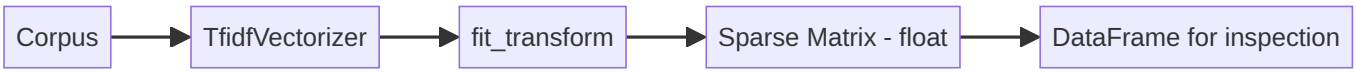

Step 1: Import and Prepare Corpus

```
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd

corpus = [
    "The car is fast",
    "The race car is fast",
    "The car is parked"
]
```

Step 2: Fit and Transform

```
tfidf = TfidfVectorizer()
X = tfidf.fit_transform(corpus)
```

Output: CSR sparse matrix, `dtype=float`, shape `(3, 6)` — 3 documents, 6 vocabulary terms.

Unlike BoW (integer counts), TF-IDF produces **floating-point weights** reflecting both local frequency and global rarity.

Step 3: Inspect Results

```
df = pd.DataFrame(
    X.toarray(),
    columns=tfidf.get_feature_names_out(),
    index=corpus
)
```

Key observation — same word, different scores:

Document	`car` score
The car is fast	0.46
The race car is fast	0.36
The car is parked	0.41

The word `car` receives a **different TF-IDF weight in each document** because:

- Local term frequency differs (`race car` has `car` alongside another distinctive term)
- The relative importance of `car` vs other words in each document differs

Words appearing in all documents (like `the`, `is`) receive lower weights due to low IDF.

TfidfVectorizer vs CountVectorizer

Property	CountVectorizer	TfidfVectorizer
Output type	Integer counts	Float weights
Common word handling	No adjustment	Automatic downweighting
Underlying steps	Tokenize + count	Tokenize + count + TF-IDF scaling
Typical use	BoW baseline	Document classification, retrieval

Internally, `TfidfVectorizer` applies TF-IDF normalization to the count matrix produced by the same tokenization logic.

Key Parameters

Parameter	Effect
`sublinear_tf`	Use 1 + \log(\text{TF}) instead of raw TF — dampens high frequencies
`smooth_idf`	Adds 1 to DF to avoid division by zero for unseen terms
`norm`	`'l2'` (default) — normalizes vectors for cosine similarity
`use_idf`	Set `False` to get raw counts (equivalent to BoW)
`ngram_range`	Include bigrams/trigrams for partial order capture

Production Usage

In a document search microservice:

1. `fit` TF-IDF on the indexed document corpus
2. Transform incoming queries with the same vectorizer
3. Compute cosine similarity between query vector and document vectors
4. Return top-k matches

This pattern powers FAQ matching in Azure Bot Service and similar retrieval systems without neural models.

Common Pitfalls / Exam Traps

- **Expecting integer output** — TF-IDF values are floats; BoW values are integers.
- **Fitting on test data** — same leakage risk as BoW; fit only on training corpus.
- **Same word, same score misconception** — TF-IDF scores are per-document; `car` in sentence 1 ≠ `car` in sentence 2.
- **Exam trap: dtype** — if asked what changes from BoW to TF-IDF, answer: integer counts → float weights with IDF scaling.

Quick Revision Summary

- `TfidfVectorizer` from scikit-learn implements TF-IDF via `fit_transform`.
- Output is a float sparse matrix, not integer counts.
- Same word gets different TF-IDF scores in different documents.
- Common words (`the` , `is`) receive low weights; distinctive words receive high weights.
- Fit on training data only; transform new documents with learned vocabulary and IDF.
- Default L2 normalization makes vectors ready for cosine similarity comparisons.

← **Previous** · **Index** · **Next** →
Advantages and Limitations of TF-IDF

Advantages and Limitations of TF-IDF

Intuition: The Best Sparse Representation

TF-IDF is often the strongest sparse vectorization method — it improves on BoW without requiring neural network training. But it inherits fundamental limitations of all count-based methods. Understanding where TF-IDF stops is as important as knowing where it shines.

Advantages

1. Automatic Stopword Downweighting

Common words (`the` , `is` , `and` , `of`) appear in nearly every document, so their IDF approaches zero. TF-IDF reduces their influence without an explicit stopword removal step.

Word	Typical IDF Behavior
`the`	Near 0 — appears everywhere
`mitochondria`	High — appears in few biology documents
`refund`	Moderate-high — domain-specific to support tickets

2. Domain-Specific Term Emphasis

Rare but meaningful terms receive high weights. In a biology paper corpus, `mitochondria` scores far higher than `cell` (common even within the domain) or `the` (common everywhere).

3. Strong Baseline for Classification and Retrieval

TF-IDF + linear SVM or logistic regression remains competitive for:

- News categorization
- Email routing in enterprise helpdesk systems
- Patent or legal document classification on cloud ML platforms

4. Efficient and Interpretable

No GPU training, no epochs. Each feature is a word with a float weight — explainable via top-weighted terms per document.

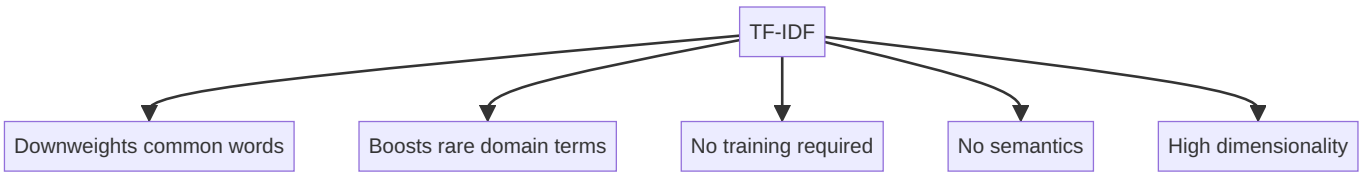

Limitations

1. No Semantic Meaning

"Car" and "automobile" remain orthogonal features. "Happy" and "joyful" have no structural relationship. The model must discover synonymy from label co-occurrence alone.

2. High Dimensionality

Vocabulary size equals feature count. A 100,000-term vocabulary produces 100,000-dimensional vectors. Sparse storage mitigates memory, but compute for dense operations remains costly.

3. No Word Order

Like BoW, TF-IDF discards sequence. "Not good" and "good" share the feature `good` unless n-grams are added.

4. Cold Start for New Vocabulary

Words not seen during `fit` are silently dropped at `transform` time. Rapidly evolving domains (social media, product catalogs) require periodic vocabulary refresh.

5. Sensitive to Corpus Composition

IDF values depend entirely on the training corpus. A term rare in one corpus may be common in another — TF-IDF weights are not transferable across domains without refitting.

Comparison Table: All Sparse Methods

Property	OHE	BoW	TF-IDF
Values	0/1	Integers	Floats
Frequency	No	Yes	Yes (weighted)
Stopword handling	Manual	Manual	Automatic (partial)
Semantics	No	No	No
Best use case	Small categorical vocab	Keyword tasks	Classification, retrieval

---

When to Move Beyond TF-IDF

Upgrade to dense or contextual embeddings when:

- Synonymy and paraphrase matter (semantic search)
- Word order and negation are critical (sentiment, NLI)
- Transfer learning from large pretrained models is available (BERT, sentence-transformers)

Common Pitfalls / Exam Traps

- **"TF-IDF captures semantics"** — false. It captures statistical importance, not meaning.
- **"Stopword removal is unnecessary with TF-IDF"** — mostly true but not absolute; explicit removal can still improve results.
- **Confusing TF-IDF with Word2Vec** — TF-IDF is sparse and count-based; Word2Vec is dense and learned.
- **Exam trap: TF-IDF still has high dimensionality** — yes, vocabulary-sized vectors remain.

Quick Revision Summary

- TF-IDF advantages: auto-downweights common words, boosts domain terms, no training, strong classification baseline.
- TF-IDF limitations: no semantics, high dimensionality, no word order, corpus-dependent IDF.
- `mitochondria` scores high in biology corpora; `the` scores near zero.
- TF-IDF is the strongest sparse method but still treats words as independent features.
- N-grams can partially recover word order within TF-IDF pipelines.
- Dense embeddings (Word2Vec, BERT) address the semantic gap TF-IDF cannot close.

← **Previous** · **Index** · **Next** →

Week 6

Week 6

Dense Word Embeddings: Module Overview

Intuition: From Symbols to Meaning

Sparse methods (one-hot, BoW, TF-IDF) treat every word as an isolated symbol. The vectors for `car` and `automobile` share no dimensions and have zero similarity — even though humans know they mean nearly the same thing. This **semantic gap** limits every downstream task that depends on word relatedness.

Dense word embeddings solve this by mapping each word into a continuous vector space of 50–300 real numbers, where **geometric proximity reflects semantic similarity**.

The Problem with Sparse Representations

Issue	Example
Orthogonality	`dog` and `puppy` have dot product = 0
Dimensionality	10,000-word vocabulary → 10,000-dim vectors, 99.9% zeros
No compositionality	Cannot compute "king − man + woman ≈ queen"

In a cloud search system, sparse methods cannot match a query for "automobile insurance" to documents about "car coverage" without exact keyword overlap.

The Dense Embedding Solution

Instead of a vector of 10,000 zeros and one 1, represent each word as:

$$\mathbf{v}_{\text{king}} = [0.2,\ -0.5,\ 0.8,\ \ldots] \in \mathbb{R}^{d}$$
where d is typically 50–300.

Core insight (distributional hypothesis): words that appear in similar contexts will have similar vectors.

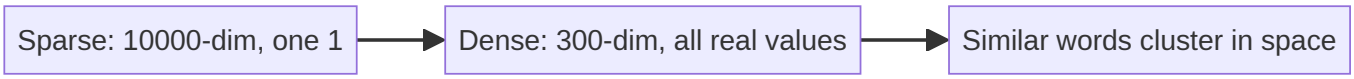

The Famous Analogy: Vector Arithmetic

Word embeddings capture relational structure through vector arithmetic:

$$\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}} \approx \mathbf{v}_{\text{queen}}$$
Intuition: subtract "maleness" from "king", add "femaleness" → arrive near "queen".

Similarly: $\mathbf{v}_{\text{Paris}} - \mathbf{v}_{\text{France}} + \mathbf{v}_{\text{Germany}} \approx \mathbf{v}_{\text{Berlin}}$

This works because embeddings encode latent attributes (gender, royalty, geography) as directions in vector space.

Techniques in This Module

Method	Approach	Key Idea
Word2Vec	Shallow neural network	Predict context from word (or vice versa)
GloVe	Matrix factorization	Global co-occurrence statistics
Visualization	PCA / t-SNE	Project high-dim vectors to 2D
Static vs Contextual	Comparison	One vector per word vs context-dependent

---

Module Transition

Previous Module	This Module
Frequency counting	Semantic understanding
Sparse, high-dimensional	Dense, low-dimensional
Words as atomic symbols	Words as points in continuous space
No similarity between synonyms	Cosine similarity reflects meaning

---

Common Pitfalls / Exam Traps

- **Confusing "embedding" with only dense vectors** — TF-IDF is also called a word embedding in the broad sense.
- **Assuming all embeddings are contextual** — Word2Vec and GloVe produce one fixed vector per word (static).
- **"Dense = better always"** — TF-IDF can outperform Word2Vec on small datasets or keyword tasks.
- **Exam trap: dimensionality** — dense embeddings use 50–300 dims; sparse use vocabulary size.

Quick Revision Summary

- Dense embeddings map words to low-dimensional continuous vectors (50–300 dims).
- Similar contexts → similar vectors (distributional hypothesis).

- Solves the semantic gap: `car` ≈ `automobile` in vector space.
- Vector arithmetic captures relationships: king − man + woman ≈ queen.
- Word2Vec (neural) and GloVe (co-occurrence matrix) are the two main static methods.
- This module transitions from frequency counting to semantic representation.

← **Previous** · **Index** · **Next** →
Word2Vec: Learning Word Embeddings

Word2Vec: Learning Word Embeddings

Intuition: Predict to Compress

Word2Vec learns word meanings by training a shallow neural network on a simple task: **predict a missing word from its neighbors** (or predict neighbors from a word). To solve this prediction efficiently, the network compresses word meanings into dense vectors in its hidden layer. These hidden-layer weights become the embeddings.

The training objective forces words in similar contexts to develop similar representations.

Architecture Overview

Word2Vec is not a single architecture — it comes in two flavors:

Architecture	Input	Output	Task
CBOW (Continuous Bag of Words)	Context words	Target word	Fill in the blank
Skip-gram	Target word	Context words	Predict surroundings

Both are shallow neural networks (one hidden layer) trained on large corpora.

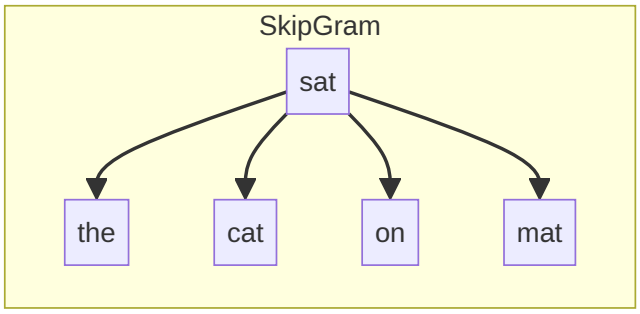

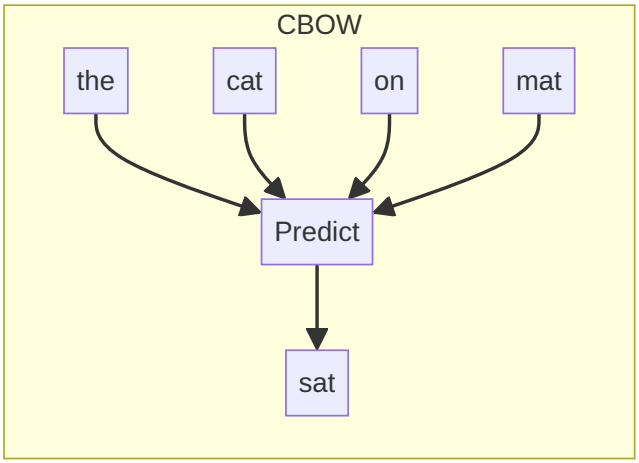

CBOW: Context → Target

Continuous Bag of Words averages context word vectors and predicts the center word.

Example sentence: "The cat sat on the mat"

Training instance:

- Input (context): `the`, `cat`, `on`, `mat`
- Output (target): `sat`

Like a fill-in-the-blanks game: given surrounding words, predict the missing center word.

CBOW is **faster** to train because it predicts one word from many context inputs averaged together.

Skip-gram: Target → Context

Skip-gram takes the center word and predicts each surrounding context word.

Training instance:

- Input (target): `sat`
- Output (context): `the`, `cat`, `on`, `mat`

Skip-gram is **slower** but often performs better on rare words because it generates more training examples per word.

How Embeddings Emerge

During training:

1. Each word is initially assigned a random vector
2. The network adjusts vectors to minimize prediction error
3. Words that share contexts (e.g., `cat` and `dog` both follow `the`) move closer in vector space
4. After training, the hidden-layer weight matrix rows are the word embeddings

\mathbf{v}_w \in \mathbb{R}^{d}, \quad d \in [50, 300]

CBOW vs Skip-gram

Criterion	CBOW	Skip-gram
Direction	Context → target	Target → context
Speed	Faster	Slower
Rare words	Less effective	Better
Training examples per word	Fewer	More
Best for	Large, clean corpora	Smaller corpora, rare terms

Real-World Usage

Pretrained Word2Vec models power:

- Semantic similarity in recommendation systems (similar product descriptions)
- Feature initialization for downstream NLP models
- Vocabulary expansion in search (find related terms)

In Gensim, pretrained models like `word2vec-google-news-300` provide 3 million words in 300 dimensions — ready for cosine similarity queries without training from scratch.

Common Pitfalls / Exam Traps

- **Confusing CBOW and Skip-gram direction** — CBOW: many → one; Skip-gram: one → many.
- **"Word2Vec is a deep network"** — it is a shallow (single hidden layer) network.
- **Assuming Word2Vec is contextual** — each word gets one fixed vector regardless of sentence context.
- **Exam trap: training objective** — Word2Vec learns by prediction, not by counting co-occurrences (that's GloVe).

Quick Revision Summary

- Word2Vec learns dense embeddings via a shallow neural network trained on word prediction.
- CBOW: context words → predict target word (fill in the blank).
- Skip-gram: target word → predict context words (opposite direction).
- Similar contexts produce similar vectors (distributional hypothesis).
- CBOW is faster; Skip-gram handles rare words better.
- Embeddings live in the hidden layer weights; typical dimension: 50–300.

← **Previous** · **Index** · **Next** →
Implementing Word2Vec with Gensim

Implementing Word2Vec with Gensim

Intuition: Training Embeddings on Your Own Corpus

Conceptual understanding of CBOW and Skip-gram is necessary; training a model on a real corpus reveals how hyperparameters control context window, dimensionality, and vocabulary filtering. Gensim provides a production-ready `Word2Vec` implementation.

Pipeline Overview

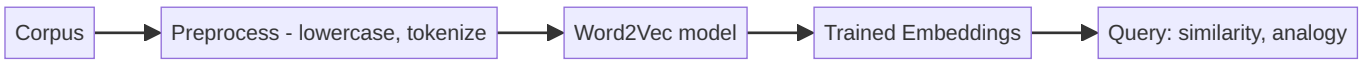

Step 1: Install and Import

```
# pip install gensim
from gensim.models import Word2Vec
```

Step 2: Prepare Corpus

```
corpus = [
    "I love machine learning",
    "Machine learning is fascinating",
    "I love dogs",
    # ... more sentences
]

# Preprocess: lowercase + tokenize
tokenized_corpus = [sentence.lower().split() for sentence in corpus]
```

Gensim expects a **list of lists** — each inner list is a tokenized sentence.

Step 3: Train the Model

```
model = Word2Vec(
    sentences=tokenized_corpus,
    vector_size=50,
    window=3,
    min_count=1,
    workers=4
)
```

Hyperparameters

Parameter	Value	Meaning
`sentences`	tokenized corpus	Training data
`vector_size`	50	Embedding dimensionality (d)
`window`	3	Context window: 3 words before and after the target
`min_count`	1	Ignore words appearing fewer than this many times
`workers`	4	Number of CPU threads for parallel training

Window example: for the word `learning` in "I am learning machine learning love":

- Window = 3 looks at 3 words before and 3 words after
- Context includes: `am`, `machine`, `love`, etc.

Step 4: Query Embeddings

Get a word vector

```
model.wv['machine']
# array([0.02, -0.15, 0.33, ...])  — 50-dimensional
```

Find similar words

```
model.wv.most_similar('machine', topn=3)
# [('learning', 0.85), ('dog', 0.42), ('love', 0.38)]
```

Similarity quality depends on:

- **Corpus size** — small corpora produce noisy similarities
- **Hyperparameters** — window size, vector dimension, min_count
- **Domain relevance** — a general corpus won't capture specialized terminology

On a small demo corpus, `machine` may be similar to `learning` (co-occurrence) but also spuriously to unrelated words — this is expected with insufficient data.

Production Considerations

Approach	When
Train from scratch	Domain-specific corpus (medical, legal, internal docs)
Load pretrained	General English (Google News 300, fastText)
Fine-tune	Start pretrained, continue training on domain data

For cloud deployments, pretrained embeddings reduce training cost; custom training is justified when domain vocabulary diverges significantly from general English.

Common Pitfalls / Exam Traps

- **Not lowercasing** — `Machine` and `machine` become separate vocabulary entries.
- **Passing raw strings instead of tokenized lists** — Gensim requires `List[List[str]]`.
- **`min_count` too high on small corpora** — rare but important words get dropped.
- **Overinterpreting similarities from tiny corpora** — exam/demo models are not production quality.
- **Exam trap:** **`vector_size`** — this is embedding dimensionality, not context window.

Quick Revision Summary

- Gensim `Word2Vec` trains embeddings from a tokenized corpus (list of lists).
- Key hyperparameters: `vector_size` (dims), `window` (context radius), `min_count`, `workers`.
- `model.wv['word']` returns the embedding vector; `most_similar()` finds nearest neighbors.
- Preprocess with lowercase and whitespace tokenization before training.
- Similarity quality scales with corpus size and domain coherence.
- Pretrained models are preferred for general English; train custom for specialized domains.

← **Previous** · **Index** · **Next** →
GloVe: Global Vectors for Word Representation

GloVe: Global Vectors for Word Representation

Intuition: Global Statistics, Not Local Windows

Word2Vec learns from local context windows — it sees a few neighboring words at a time. **GloVe** (Global Vectors) takes a different approach: build a **global co-occurrence matrix** counting how often every word pair appears together across the entire corpus, then factorize this matrix to produce word vectors.

GloVe combines the strengths of count-based methods (global statistics) and prediction-based methods (dense vectors).

How GloVe Differs from Word2Vec

Aspect	Word2Vec	GloVe
Training signal	Local context windows	Global co-occurrence matrix
Architecture	Shallow neural network	Matrix factorization
Scope	Neighborhood around each word	Entire corpus statistics
Pretrained availability	Google News, etc.	Wikipedia, Twitter, Common Crawl

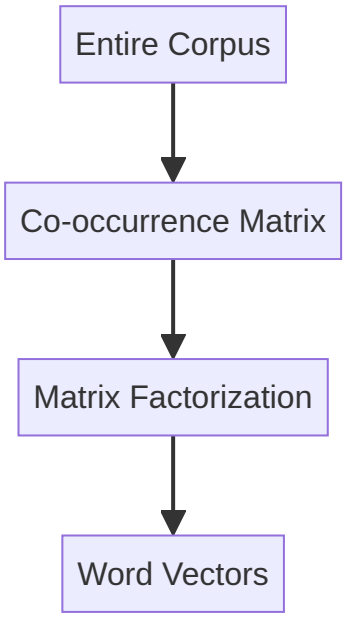

The Co-occurrence Matrix

For vocabulary of size |V|, construct a |V| \times |V| matrix X where:

X_{ij} = \text{number of times word } j \text{ appears in the context of word } i
This matrix captures **global** word-word relationships — not just immediate neighbors.

Matrix factorization decomposes X into lower-dimensional word vectors \mathbf{w}_i and \mathbf{w}_j such that:

\mathbf{w}_i^T \mathbf{w}_j \approx \log X_{ij}

The Ratio Intuition

GloVe's key insight: the **ratio of co-occurrence probabilities** encodes meaning.

Consider words `ice`, `steam`, and probe words `solid`, `gas`, `water`:

Probe word	P(word | ice)	P(word | steam)	Ratio
solid	High	Very low	Large
gas	Very low	High	Small
water	High	High	≈ 1

The ratio P(\text{solid} \mid \text{ice}) / P(\text{solid} \mid \text{steam}) is large — revealing that `ice` relates to `solid`. Similarly, ratios near 1 indicate words common to both (`water`), and small ratios reveal `steam` → `gas` associations.

This ratio-based encoding captures abstract concepts (states of matter, thermodynamics) from raw co-occurrence statistics.

GloVe vs Word2Vec in Practice

Criterion	GloVe	Word2Vec
Training data efficiency	Uses global counts (efficient)	Requires many window passes
Rare words	Good (global counts help)	Skip-gram better for rare words
Pretrained models	Widely available (Gensim, spaCy)	Widely available
Vector arithmetic	Strong (king − man + woman)	Strong
Training from scratch	Slower (matrix construction)	Faster with Gensim

---

Common Pretrained Models

Model	Dimensions	Corpus
`glove-wiki-gigaword-100`	100	Wikipedia + Gigaword
`glove-twitter-25`	25	Twitter data
`glove-840B-300d`	300	Common Crawl (840B tokens)

Pretrained GloVe models are ideal for quick prototyping — download and query without training.

Common Pitfalls / Exam Traps

- **"GloVe is a neural network"** — it uses matrix factorization, not backpropagation through a prediction task.
- **Confusing local vs global** — Word2Vec = local windows; GloVe = global co-occurrence.
- **Assuming GloVe is contextual** — like Word2Vec, one fixed vector per word.
- **Exam trap: what GloVe adds over Word2Vec** — global co-occurrence statistics via matrix factorization.

Quick Revision Summary

- GloVe = Global Vectors for Word Representation.
- Builds a global word co-occurrence matrix, then factorizes it into dense vectors.
- Ratio of co-occurrence probabilities encodes semantic relationships.
- Combines count-based global statistics with dense vector output.
- Pretrained models available (Wikipedia, Twitter, Common Crawl).
- Static embeddings — one vector per word, no context sensitivity.

← **Previous** · **Index** · **Next** →
Implementing GloVe with Gensim

Implementing GloVe with Gensim

Intuition: Pretrained Embeddings Ready to Use

Training GloVe from scratch requires building a co-occurrence matrix over a large corpus — computationally expensive. Gensim provides pretrained GloVe models via its downloader API, enabling immediate similarity and analogy queries without training.

Pipeline Overview

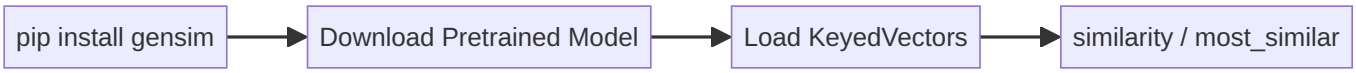

Step 1: Install and Download

```
# pip install gensim
import gensim.downloader as api

# Download pretrained model (one-time)
glove_model = api.load("glove-twitter-25")
```

Available Models

Model Name	Dimensions	Training Data
`glove-twitter-25`	25	Twitter
`glove-twitter-100`	100	Twitter
`glove-wiki-gigaword-50`	50	Wikipedia + Gigaword
`glove-wiki-gigaword-100`	100	Wikipedia + Gigaword
`glove-840B-300d`	300	Common Crawl (840B tokens)

For production semantic search, `glove-wiki-gigaword-100` or `glove-840B-300d` offer better quality than the small Twitter model.

Step 2: Word Similarity

```
glove_model.similarity("king", "queen")
# ~0.92 — high similarity (royalty, gender-related concepts)
```

```
glove_model.similarity("king", "car")
# ~0.67 — lower similarity (unrelated domains)
```

Cosine similarity ranges from −1 to 1; values near 1 indicate strong semantic relatedness.

Step 3: Word Analogies

Gensim supports analogy queries via `most_similar`:

```
glove_model.most_similar(
    positive=["woman", "king"],
    negative=["man"],
    topn=1
)
# [('queen', 0.85)]
```

This implements the vector arithmetic:

$$\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}} \approx \mathbf{v}_{\text{queen}}$$

How it works:

- `positive` words are **added** to the query vector
- `negative` words are **subtracted**
- The model returns words whose vectors are closest to the resulting vector

Another Example

```
glove_model.most_similar(
    positive=["woman", "man"],
    negative=["king"],
    topn=1
)
# Returns a word in the direction of "commoner" or similar
```

GloVe vs Training Word2Vec

Approach	Code	Training Time	Quality
Pretrained GloVe	`api.load("glove-wiki-gigaword-100")`	None (download only)	High for general English
Train Word2Vec	`Word2Vec(sentences=...)`	Minutes to hours	Depends on corpus
Train GloVe from scratch	External tools (Stanford GloVe)	Hours to days	Depends on corpus

For most cloud ML prototypes, pretrained GloVe is the fastest path to semantic features.

Common Pitfalls / Exam Traps

- **Using a tiny model for serious tasks** — `glove-twitter-25` (25 dims) is for demos; production needs 100–300 dims.
- **Words not in vocabulary** — `KeyError` if the word was not in training data; handle with try/except or `model.key_to_index` check.
- **Confusing `positive` / `negative` order** — order within lists does not matter; addition/subtraction direction does.
- **Exam trap: GloVe in Gensim is pretrained only** — Gensim loads pretrained GloVe; training GloVe requires Stanford's original implementation.

Quick Revision Summary

- Gensim `api.load()` downloads and loads pretrained GloVe models.
- `similarity(w1, w2)` returns cosine similarity between two word vectors.
- `most_similar(positive=[...], negative=[...])` solves analogy tasks via vector arithmetic.
- king − man + woman ≈ queen is the canonical analogy example.

- Pretrained models avoid training cost; choose dimensionality based on task requirements.
- Handle out-of-vocabulary words explicitly in production code.

← **Previous** · **Index** · **Next** →
Visualizing Word Embeddings

Visualizing Word Embeddings

Intuition: Seeing Semantic Space

Word embeddings live in 50–300 dimensions — impossible to visualize directly. Dimensionality reduction projects these high-dimensional vectors onto a 2D plane, revealing clusters of semantically related words. Seeing `king`, `queen`, `man`, `woman` grouped together makes the abstract concept of embedding space concrete.

Why Visualization Matters

Without visualization	With visualization
Vectors are abstract number arrays	Clusters reveal semantic groups
Similarity is a cosine score	Proximity on a plot is intuitive
Hard to debug bad embeddings	Outliers and noise become visible

In ML experimentation (e.g., comparing custom-trained embeddings vs pretrained), 2D plots provide a quick sanity check before deploying to production.

Pipeline Overview

Step 1: Load Pretrained Model and Select Words

```
import gensim.downloader as api

glove_model = api.load("glove-twitter-25")

words = [
    "king", "queen", "man", "woman",
    "apple", "orange", "fruit",
    "car", "bus"
]
```

Step 2: Extract Embedding Vectors

```
vectors = [glove_model[word] for word in words]
```

Each vector has dimension equal to the model's `vector_size` (e.g., 25 for `glove-twitter-25`).

Step 3: Dimensionality Reduction with PCA

```
from sklearn.decomposition import PCA

pca = PCA(n_components=2)
reduced = pca.fit_transform(vectors)
```

PCA (Principal Component Analysis) finds the two directions of maximum variance in the high-dimensional data and projects onto them.

Method	Speed	Preserves	Best for
PCA	Fast	Global variance	Quick exploration
t-SNE	Slow	Local neighborhoods	Cluster visualization
UMAP	Moderate	Both local and global	Publication-quality plots

PCA is preferred for quick inspection; t-SNE better reveals tight clusters at the cost of global structure distortion.

Step 4: Scatter Plot

```
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 8))
plt.scatter(reduced[:, 0], reduced[:, 1])

for i, word in enumerate(words):
    plt.annotate(word, (reduced[i, 0], reduced[i, 1]))

plt.title("Word Embeddings in 2D (PCA)")
plt.xlabel("PC1")
plt.ylabel("PC2")
plt.show()
```

Expected Cluster Patterns

Cluster	Words	Why close
Royalty/gender	king, queen, man, woman	Shared gender and authority attributes
Food	apple, orange, fruit	Culinary co-occurrence
Transport	car, bus	Vehicle domain

Words in the same semantic domain cluster together; unrelated domains occupy different regions of the plot.

Interpreting the Plot

- **Distance ≈ dissimilarity** — but only approximately; PCA loses information
- **Cluster tightness** — reflects consistency of co-occurrence patterns in training data
- **Outliers** — words with unusual usage patterns or polysemy may appear misplaced

---

Common Pitfalls / Exam Traps

- **Treating 2D distances as exact** — PCA/t-SNE distort high-dimensional relationships; cosine similarity in original space is authoritative.
- **Too few words** — clusters are meaningless with 3–4 words; use 20+ for meaningful patterns.
- **Mixing models** — vectors from different models/dimensions cannot be plotted together.
- **Exam trap: PCA purpose** — reduces d-dimensional vectors to 2D for visualization, not for production inference.

Quick Revision Summary

- High-dimensional embeddings (50–300D) require dimensionality reduction for visualization.
- PCA projects vectors to 2D by preserving maximum variance directions.
- Semantic clusters emerge: royalty words together, food words together, transport words together.
- Pipeline: load model → extract vectors → PCA → scatter plot with labels.
- t-SNE is an alternative that better preserves local neighborhoods.
- 2D plots are for exploration; use cosine similarity in original space for production.

← **Previous** · **Index** · **Next** →
Static vs Contextual Word Embeddings

Static vs Contextual Word Embeddings

Why Word Embeddings Exist

Human language is understood through meaning; machines process numbers. **Word embeddings** bridge this gap by converting each word into a mathematical vector — a list of real numbers that encodes semantic properties.

$$\text{king} \rightarrow [0.2,\ -0.5,\ 0.8,\ \ldots]$$
The numbers are not arbitrary: they capture the essence of the word so that mathematical operations reveal relationships humans recognize intuitively.

How Embeddings Are Created: The Property Encoding Model

Imagine encoding words by listing observable properties:

Property	king	queen
Has authority	1	1
Gender (0=male, 1=female)	0	1
Is rich	1	1

This yields:

- $\mathbf{v}_{\text{king}} = [1,\ 0,\ 1]$
- $\mathbf{v}_{\text{queen}} = [1,\ 1,\ 1]$

Add more properties (has moustache, war experience, trained in cooking) and the vectors grow in dimensionality. Different people might choose different properties — the scheme is arbitrary, but the principle is universal: **encode meaning as numbers along meaningful dimensions**.

Real models (Word2Vec, GloVe, BERT) learn hundreds of latent properties automatically — we do not know what each dimension represents, but we know the vectors work.

Embedding Dimensionality

Model	Typical Dimensions
Custom toy model	3–8
GloVe	100–300
Word2Vec	300–500
BERT-base	768
BERT-large	1024
OpenAI text-embedding-small	1536

More dimensions capture finer-grained meaning but increase compute and storage cost.

Vector Arithmetic and Relationships

The King − Man + Woman = Queen Identity

$$\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}} \approx \mathbf{v}_{\text{queen}}$$

Intuition: subtract "maleness" from "king" (leaving royalty/crown), add "femaleness" → arrive at "queen".

Verification with property vectors:

$$\begin{bmatrix} 1 \\ 0 \\ 1 \end{bmatrix} - \begin{bmatrix} 0.1 \\ 0 \\ 0.2 \end{bmatrix} + \begin{bmatrix} 0.1 \\ 1 \\ 0.2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \approx \mathbf{v}_{\text{queen}}$$

Geographic Analogy

$$\mathbf{v}_{\text{Delhi}} - \mathbf{v}_{\text{India}} + \mathbf{v}_{\text{Russia}} \approx \mathbf{v}_{\text{Moscow}}$$

Remove "India" (leaving "capital-of"), apply to "Russia" → "Moscow".

Relationship Vectors

The difference between two word vectors encodes a **relationship**:

$$\mathbf{r}_{\text{gender}} = \mathbf{v}_{\text{woman}} - \mathbf{v}_{\text{man}}$$
Transposing this relationship to royalty:

$$\mathbf{v}_{\text{king}} + \mathbf{r}_{\text{gender}} \approx \mathbf{v}_{\text{queen}}$$
Similarly, $\mathbf{v}_{\text{Delhi}} - \mathbf{v}_{\text{India}}$ encodes the "capital city" relationship, which transposes to Moscow − Russia.

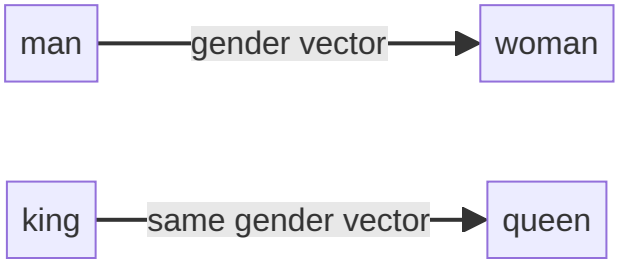

Static Word Embeddings

Definition: Each word has exactly **one fixed vector** regardless of context.

Property	Static (Word2Vec, GloVe)
Vectors per word	1
Context sensitivity	None
Polysemy handling	Poor
Training	Corpus-wide co-occurrence or prediction
Speed	Fast lookup

The Polysemy Problem

The word `bank` has one vector whether the sentence discusses a river bank or a financial bank. The model cannot distinguish meanings without surrounding context.

Contextual Word Embeddings

Definition: The vector for a word **changes based on its surrounding context** in the sentence.

Property	Contextual (BERT, GPT)
Vectors per word	One per occurrence
Context sensitivity	Full
Polysemy handling	Excellent
Training	Massive neural networks
Speed	Slower (requires full sentence)

Resolving "dish" Through Context

Consider: "I ate the **dish** last night."

Context level	Embedding shifts toward
`dish` alone	Ambiguous (food or utensil)
`rice dish`	Rice-based foods (biryani, risotto, kheer)
`Indian rice dish`	Indian rice foods (biryani, kheer)
`Indian rice dish with spices and meat`	**Biryani**

Each additional context word shifts the embedding through **relationship vectors** in space:

The final contextual embedding for `dish` in this sentence is near biryani, not risotto or pizza.

Context from Surrounding Text

Example: Two people with different food preferences:

- Utkarsh loves chole bhature, mutton biryani, kachori (Indian cuisine)
- Don Corleone loves risotto, pasta, calzone (Italian cuisine)
- Don Corleone invites Utkarsh for dinner at his home

Prediction: The blank "_____ cuisine" is most likely **Italian** (highest probability), not Indian or Brazilian — because contextual cues (Italian name, Italian food preferences, host's home) dominate.

Static embeddings would struggle; contextual models capture these cross-sentence relationships.

Static vs Contextual: Comparison Table

Criterion	Static	Contextual
Vectors per word	1 (fixed)	Context-dependent
Polysemy	Cannot resolve	Resolves via context
Model examples	Word2Vec, GloVe	BERT, GPT, ELMo
Dimensions	100–500	768–1536
Compute	Low (lookup)	High (forward pass)
Analogy arithmetic	Strong	Less interpretable
Production use	Similarity, clustering	Classification, QA, generation

---

How Contextual Embeddings Are Built

Modern models (BERT, GPT) use massive neural networks with millions of parameters across many layers. These layers learn to:

1. Read the entire sentence (or document)
2. Compute attention-weighted combinations of all words

3. Produce a context-specific vector for each token

The components of a contextual embedding decompose into contributions from surrounding words — analogous to how biryani's embedding decomposes into dish + rice + Indian + spices + meat.

Common Pitfalls / Exam Traps

- **"All embeddings are contextual"** — Word2Vec and GloVe are static; only one vector per word.
- **Confusing dimensions with properties** — we know dimension counts (768 for BERT) but not what each dimension means.
- **Assuming vector arithmetic works equally on contextual embeddings** — king − man + woman is demonstrated on static embeddings; contextual vectors change per sentence.
- **Exam trap: polysemy** — "bank" (river) vs "bank" (financial) is the classic example of static embedding failure.
- **"Higher dimensions always better"** — more dimensions help but increase cost; task-dependent trade-off.

Quick Revision Summary

- Word embeddings convert words to numerical vectors that encode meaning.
- Embeddings are created by encoding properties as dimensions; real models learn latent properties automatically.
- Vector arithmetic captures relationships: king − man + woman ≈ queen.
- Relationship vectors (gender, capital-city) transpose across word pairs.
- Static embeddings: one fixed vector per word (Word2Vec, GloVe); fail on polysemy.
- Contextual embeddings: vectors change with sentence context (BERT, GPT); resolve ambiguity.
- "dish" → "Indian rice dish with spices and meat" → biryani demonstrates contextual resolution.
- BERT-base: 768 dims; GPT embeddings: up to 1536 dims.

← **Previous** · **Index** · **Next** →

Week 7

Week 7

Sequential Modelling: Module Overview

Sequential Modelling: Module Overview

Intuition: Language Is a Sequence, Not a Bag

Previous vectorization methods — TF-IDF, feedforward neural networks — treat text as a fixed-size feature vector. Two fundamental problems make this inadequate for real language understanding:

1. **Fixed input size** — vocabulary-sized vectors cannot naturally handle a 3-word sentence and a 1000-word paragraph differently
2. **Loss of word order** — "the dog bit the man" and "the man bit the dog" are identical to BoW/TF-IDF

Language is inherently **sequential**: the meaning of each word depends heavily on what came before it.

The Two Core Problems

Problem 1: Fixed Input Size

Method	Input constraint
BoW / TF-IDF	Vector length = vocabulary size (fixed)
Feedforward NN	Input layer has fixed number of neurons
One-hot per sentence	Same dimension regardless of sentence length

Variable-length text requires padding, truncation, or aggregation — all of which lose information.

Problem 2: Loss of Order

Sentence A	Sentence B
The dog bit the man	The man bit the man

BoW representation: **identical** (same words, same counts).

But meaning: **opposite** (who is the biter?).

Negation is equally broken: "not good" vs "good" share the feature `good`.

The Sequential Modelling Solution

Sequential models address both problems:

Capability	How
Variable length	Process one token at a time; no fixed input size
Temporal order	Hidden state carries information from previous tokens

Key insight: "not" coming before "good" changes meaning entirely — "not good" is bad. Order matters.

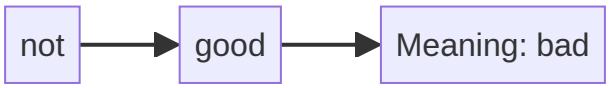

What This Module Covers

Topic	Purpose
Seq2Seq modelling	Encoder-decoder for sequence-to-sequence tasks
RNNs	Recurrent architecture with memory
RNN applications	Many-to-one, one-to-many, many-to-many
Vanishing gradients	Why RNNs forget long-range context
LSTM	Gated memory for long sequences
GRU	Simplified gated alternative

Why Sequential Models Matter in Production

Application	Why sequence matters
Machine translation	Word order differs across languages
Speech recognition	Audio is a temporal signal
Text summarization	Key facts at the start must influence the summary
Chatbots	Conversation history provides context

Cloud services like Google Translate, AWS Transcribe, and Azure Speech all depend on sequential architectures (historically RNN/LSTM; now largely Transformer-based).

Common Pitfalls / Exam Traps

- **"TF-IDF handles variable length"** — the vector size is fixed (vocabulary length); only the number of non-zero entries varies.
- **Confusing sequential models with n-grams** — n-grams capture local order in a fixed window; RNNs capture order across the full sequence.
- **Assuming BoW is never useful** — it remains valid for keyword tasks; sequential models are needed when order matters.
- **Exam trap: two problems** — fixed input size and loss of order are the two motivations for sequential modelling.

Quick Revision Summary

- Two problems with prior methods: fixed input size and loss of word order.
- BoW treats "the dog bit the man" and "the man bit the dog" as identical.
- Sequential models process variable-length text token by token.
- Temporal order is preserved via hidden state memory.
- "not good" ≠ "good" — negation requires sequence awareness.
- This module covers Seq2Seq, RNN, LSTM, GRU, and vanishing gradients.

← **Previous** · **Index** · **Next** →
Sequence-to-Sequence (Seq2Seq) Modelling

Sequence-to-Sequence (Seq2Seq) Modelling

Intuition: One Sequence In, Another Sequence Out

Many NLP tasks transform an input sequence into an output sequence of different length and content:

- English sentence → French sentence (translation)
- Long article → short summary (summarization)
- Audio waveform → text transcript (speech recognition)

Sequence-to-sequence (Seq2Seq) modelling trains systems to convert sequences from one domain into sequences in another. The architecture that powers this is the **encoder-decoder** framework.

The Encoder-Decoder Architecture

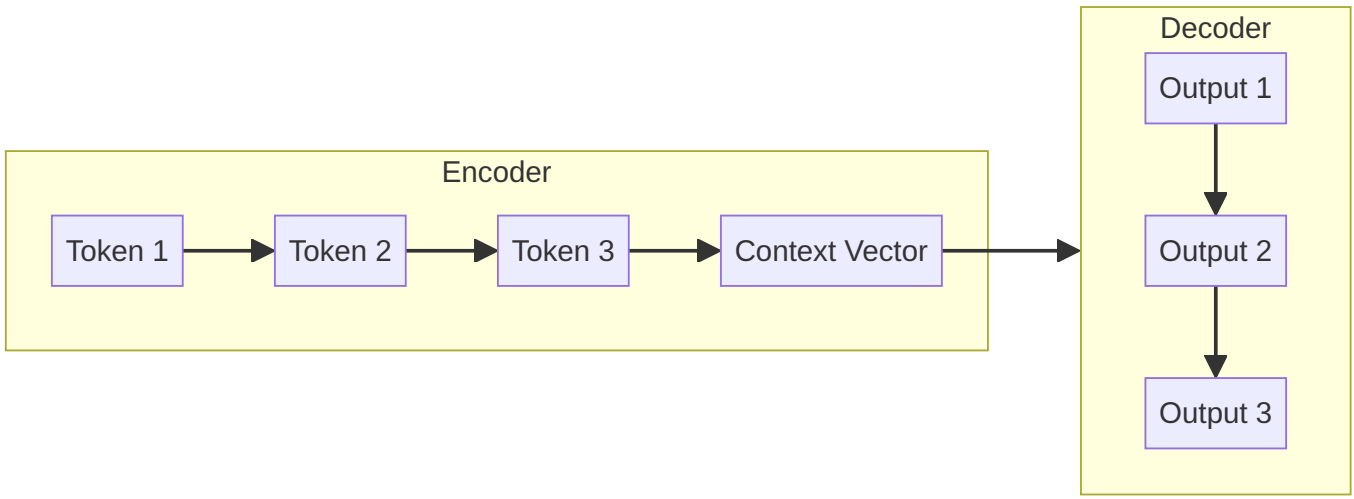

Encoder

- Consumes the input sequence **one token at a time**
- Progressively updates its internal state
- Compresses the entire input meaning into a single fixed-size vector: the **context vector** (or final hidden state)
- Analogy: reading an entire book and forming one mental summary

Decoder

- Receives the context vector as its starting point
- Generates the output sequence **one token at a time**
- Each generated token feeds back as input for the next step
- Analogy: taking that mental summary and writing it in a different language

---

Worked Example: English → French

Input: "I am fine"

Step	Encoder reads	Hidden state updates
1	`I`	h_1
2	`am`	h_2
3	`fine`	h_3 (context vector)

Decoder generates:

Step	Decoder outputs
1	`Je`
2	`vais`
3	`bien`

Output: "Je vais bien" — French for "I am fine".

Input and output lengths differ (3 → 3 here, but can vary arbitrarily).

Real-World Applications

Application	Input Sequence	Output Sequence
Machine translation	Source language sentence	Target language sentence
Text summarization	Long article (1000+ words)	Short summary (50 words)
Speech recognition	Audio signal frames	Text transcript
Chatbot response	User message + history	System reply

All share the pattern: **sequence in → sequence out**.

The Information Bottleneck

The encoder must compress **everything** about a long input — grammar, tone, facts, entities — into a single fixed-size context vector.

Input length	Challenge
Short (5 words)	Context vector is sufficient
Long (500 words)	Model forgets the beginning by the time it reaches the end
Very long (5000 words)	Severe information loss

This bottleneck is the fundamental limitation of basic Seq2Seq and motivates **attention mechanisms** (covered in later modules), which allow the decoder to look back at all encoder states rather than relying on one compressed vector.

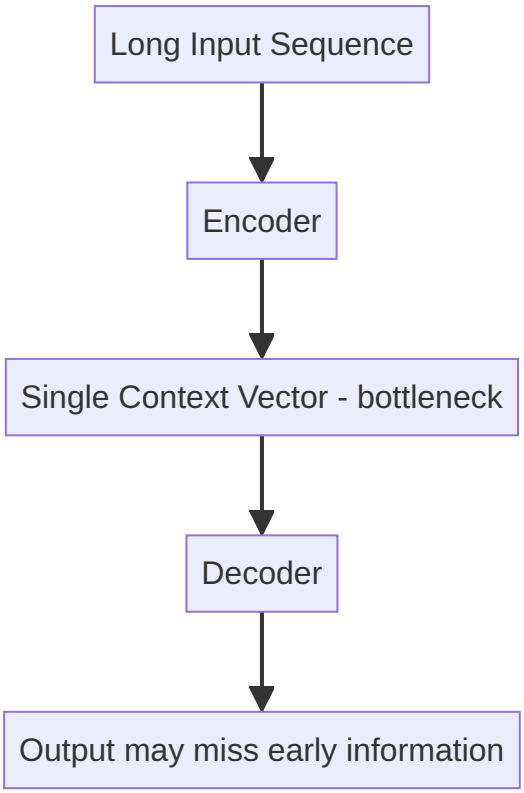

Seq2Seq vs Other Architectures

Architecture	Input	Output	Example
Many-to-one	Sequence	Single label	Sentiment analysis
One-to-many	Single input	Sequence	Image captioning
Many-to-many (synced)	Sequence	Sequence (aligned)	POS tagging
Seq2Seq (many-to-many)	Sequence	Sequence (unaligned)	Translation

Seq2Seq is the general many-to-many case where input and output lengths need not match.

Common Pitfalls / Exam Traps

- **Confusing encoder and decoder roles** — encoder reads input; decoder generates output.
- **"Context vector captures everything perfectly"** — false for long sequences; this is the information bottleneck.
- **Assuming equal input/output length** — Seq2Seq handles different lengths (article → summary).
- **Exam trap: what solves the bottleneck** — attention mechanisms (not covered in basic Seq2Seq).

Quick Revision Summary

- Seq2Seq converts input sequences to output sequences in a different domain.
- Encoder-decoder architecture: encoder compresses input → context vector → decoder generates output.
- Applications: translation, summarization, speech recognition.
- Encoder processes one token at a time; decoder generates one token at a time.
- Information bottleneck: single context vector cannot capture very long inputs.

- Attention mechanisms (future topic) address the bottleneck by allowing selective lookback.

← **Previous** · **Index** · **Next** →
Recurrent Neural Networks (RNNs)

Recurrent Neural Networks (RNNs)

Intuition: Memory Through Feedback

A standard feedforward neural network maps input X directly to output Y — each input is processed independently with no memory of previous inputs. Language requires memory: the meaning of "bank" depends on whether the previous word was "river" or "financial".

RNNs introduce a **feedback loop**: the hidden state from the previous time step feeds into the current time step, creating persistent memory across the sequence.

Feedforward vs Recurrent

Feedforward NN	RNN
X \rightarrow H \rightarrow Y	X_t + H_{t-1} \rightarrow H_t \rightarrow Y_t
No memory between inputs	Hidden state carries past context
Fixed input size	Variable sequence length

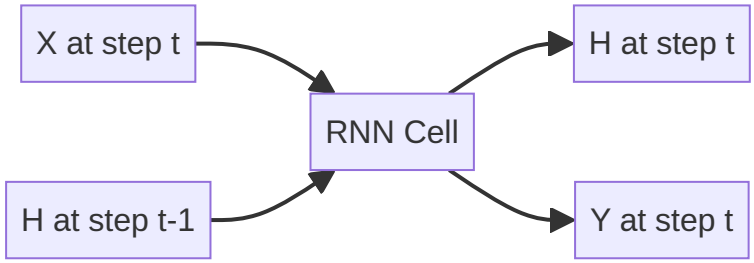

The RNN Computation at Time Step t

At each time step t:

- **Input:** X_t (current token) and H_{t-1} (previous hidden state / memory)
- **Output:** Y_t (current prediction) and H_t (updated hidden state)

H_t = \tanh\!\left(W_{hh} \cdot H_{t-1} + W_{xh} \cdot X_t + b_h\right)
Y_t = W_{hy} \cdot H_t + b_y

Symbol	Meaning
X_t	Input at current time step
H_{t-1}	Previous memory (hidden state)
H_t	Updated memory
Y_t	Output at current time step
W_{hh}, W_{xh}, W_{hy}	Weight matrices (shared across all time steps)
\tanh	Activation function (bounds values to [−1, 1])

The Hidden State as Memory

The hidden state H_t is the **memory bucket** of the network:

- As the RNN reads word by word, it updates this bucket
- New information is added; old information may fade
- At step t, the network decides based on:
  - Current input X_t
  - Current context H_t (being computed)
  - Past context H_{t-1}

Unrolling Over Time

A 5-word sentence produces a chain of 5 identical RNN cells:

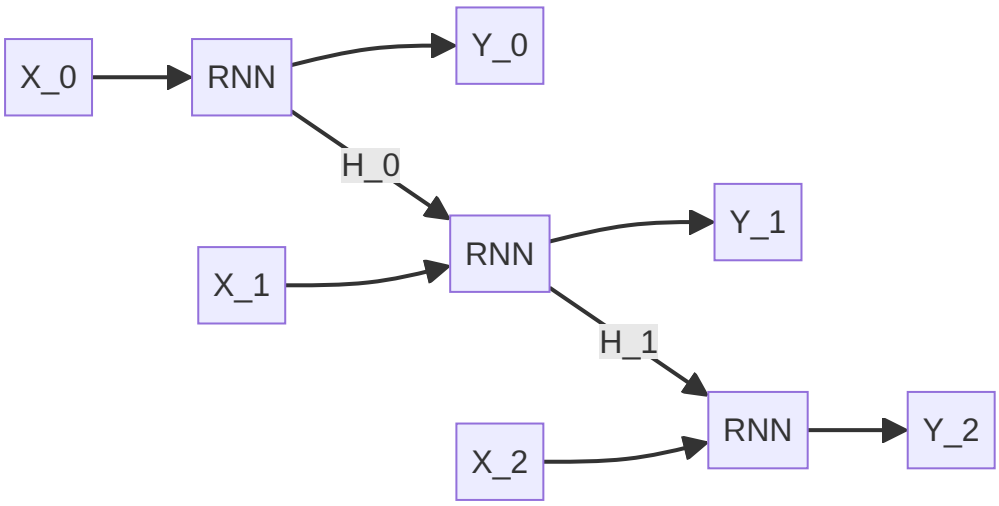

Each cell uses the **same weight matrices** (W_{xh}, W_{hh}, W_{hy}) — the model learns one general rule for processing language regardless of position in the sequence.

Why Shared Weights Matter

Property	Implication
Same weights at every step	Position-invariant processing rules
Parameter efficiency	Few weights regardless of sequence length
Generalization	Rules learned at position 3 apply at position 300

This is analogous to applying the same reading comprehension strategy to every word in a sentence.

RNNs and the Information Bottleneck

Basic Seq2Seq encoders are RNNs. The final hidden state H_T serves as the context vector for the decoder. For short sequences, this works well. For long sequences, early information in H_0, H_1, \ldots gets overwritten by later inputs — the memory bucket overflows.

This limitation leads to LSTM and GRU architectures (covered next).

Common Pitfalls / Exam Traps

- **"RNN has different weights at each step"** — false; weights are shared (tied) across time steps.
- **Confusing H_t and Y_t** — H_t is internal memory; Y_t is the output prediction.
- **Assuming RNN solves long-term memory** — vanilla RNNs suffer from vanishing gradients; LSTM/GRU are needed.
- **Exam trap: activation function** — RNN hidden state uses \tanh, not sigmoid.

Quick Revision Summary

- RNNs add a feedback loop: previous hidden state H_{t-1} feeds into current step.
- H_t = \tanh(W_{hh} H_{t-1} + W_{xh} X_t + b_h)
- Hidden state H_t is the network's memory bucket.
- Unrolled RNN = chain of identical cells sharing weights.
- Same weight matrices at every time step — position-invariant processing.
- Vanilla RNNs handle short context but fail on long-range dependencies (vanishing gradients).

← **Previous** · **Index** · **Next** →
RNN Architecture Patterns: Many-to-One, One-to-Many, Many-to-Many

RNN Architecture Patterns: Many-to-One, One-to-Many, Many-to-Many

Intuition: Same RNN, Different Shapes

A single RNN cell can be wired differently depending on the task. The naming convention describes how many input time steps and how many output time steps the architecture uses:

\text{Input tokens} \rightarrow \text{RNN} \rightarrow \text{Output tokens}
Four patterns cover most sequence modelling tasks in production NLP and speech systems.

The Four Architecture Types

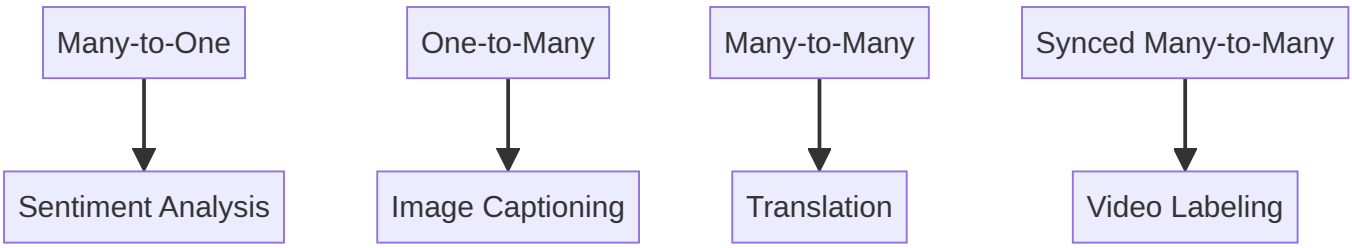

---

1. Many-to-One

Input: sequence of tokens → **Output:** single label

Component	Description
Input	Multiple words (sequence)
RNN	Processes all tokens, final hidden state summarizes
Output	One prediction (class label, score)

Example — Sentiment Analysis:

- Input: "The movie is terrible"
- Output: `negative`

Like reading an entire review to decide thumbs up or down.

Production uses: product review scoring, ticket urgency classification, social media brand monitoring.

---

2. One-to-Many

Input: single token/vector → **Output:** sequence of tokens

Component	Description
Input	One vector (image features, seed token)
RNN	Generates tokens sequentially from initial state
Output	Sequence of words

Example — Image Captioning:

- Input: image feature vector (from CNN)
- Output: "A cat sat on a sofa"

The RNN looks at the photo representation and describes it word by word.

Production uses: automated alt-text generation, video thumbnail descriptions.

---

3. Many-to-Many (Seq2Seq)

Input: sequence → **Output:** sequence (lengths may differ)

Component	Description
Input	Source language sentence
Encoder-decoder RNN	Compresses input, generates output
Output	Target language sentence

Example — Machine Translation:

- Input: English sentence (5 words)
- Output: French sentence (7 words)

Input and output lengths are **not required to match**.

Production uses: Google Translate (historically), text summarization, chatbot response generation.

---

4. Synced Many-to-Many

Input: sequence → **Output:** sequence (aligned one-to-one)

Component	Description
Input	One token per time step
RNN	Produces one output per input step
Output	One label per input step

Example — Video Frame Labeling / Subtitle Generation:

- Input: video frames [F_1, F_2, F_3, F_4]
- Output: labels [L_1, L_2, L_3, L_4] (running, jumping, running, sitting)

Each input time step maps to exactly one output time step.

Production uses: frame-level action recognition, phoneme-level speech alignment.

Comparison Table

Architecture	Input	Output	Input:Output ratio	Example
Many-to-one	Sequence	1 label	N:1	Sentiment analysis
One-to-many	1 vector	Sequence	1:N	Image captioning
Many-to-many	Sequence	Sequence	N:M (variable)	Translation
Synced many-to-many	Sequence	Sequence	N:N (aligned)	POS tagging, frame labels

---

Choosing the Right Architecture

Task	Architecture	Why
Classify a document	Many-to-one	Whole document → one category
Generate text from image	One-to-many	Single image → word sequence
Translate or summarize	Many-to-many (Seq2Seq)	Variable-length in and out
Tag each word	Synced many-to-many	One label per token

Common Pitfalls / Exam Traps

- **Confusing many-to-many with synced many-to-many** — Seq2Seq allows different lengths; synced requires 1:1 alignment.
- **"Sentiment analysis is one-to-many"** — false; it is many-to-one (many words → one label).
- **Image captioning as many-to-one** — false; it is one-to-many (one image → many words).
- **Exam trap: translation architecture** — many-to-many (Seq2Seq), not synced.

Quick Revision Summary

- Four RNN patterns: many-to-one, one-to-many, many-to-many, synced many-to-many.
- Many-to-one: sequence input → single output (sentiment analysis).
- One-to-many: single input → sequence output (image captioning).
- Many-to-many: sequence → different-length sequence (translation, summarization).
- Synced many-to-many: one output per input step (POS tagging, frame labeling).
- Architecture choice depends on input/output cardinality, not the RNN cell itself.

← **Previous** · **Index** · **Next** →
The Vanishing and Exploding Gradient Problem

The Vanishing and Exploding Gradient Problem

Intuition: The Whisper Game

Humans maintain context over long passages effortlessly. Consider a paragraph that begins "I grew up in France" and ends with "I speak fluent ___" — the answer is obviously "French". A vanilla RNN struggles with this: by the time it reaches the last word, it has forgotten "France" from the beginning.

The root cause is how gradients flow backward through time during training.

Why RNNs Fail at Long-Term Context

In theory, the hidden state H_{t-1} provides memory across the full sequence. In practice, vanilla RNNs:

- Handle **short-term context** well: "The clouds are in the ___" → `sky`
- Fail at **long-term context**: connecting "France" (word 1) to "fluent" (word 100)

This is RNN **amnesia** — only the last few words influence predictions.

Backpropagation Through Time (BPTT)

Training an RNN requires propagating the error backward from the last time step to the first:

To update weights affecting an early word, the gradient must travel through **every intermediate time step**.

The Chain Rule Trap

By the chain rule, the gradient at time step 0 involves multiplying gradients from all subsequent steps:

\frac{\partial \mathcal{L}}{\partial H_0} = \frac{\partial \mathcal{L}}{\partial H_T} \prod_{t=1}^{T} \frac{\partial H_t}{\partial H_{t-1}}
Each \frac{\partial H_t}{\partial H_{t-1}} involves the weight matrix W_{hh}.

Vanishing Gradients (weights < 1)

If |W_{hh}| < 1 (e.g., 0.9):

0.9 \times 0.9 \times \ldots \text{(100 times)} \approx 0.00003 \approx 0
The gradient **shrinks to zero** → weights at the beginning of the sequence stop updating → the network cannot learn from the distant past.

Exploding Gradients (weights > 1)

If |W_{hh}| > 1 (e.g., 1.1):

1.1^{100} \approx 13{,}780
The gradient **grows exponentially** → weights become NaN → training crashes.

The Whisper Analogy

Problem	Analogy	Effect
Vanishing	Whisper gets quieter through 10 people	Early words have no influence
Exploding	Whisper becomes a scream	Training diverges (NaN)

Solutions

Problem	Solution	Mechanism
Exploding gradients	**Gradient clipping**	Cap gradient magnitude at a threshold (e.g., 5.0)
Vanishing gradients	**LSTM / GRU**	Gated cell state preserves gradient flow
Vanishing gradients	**Attention**	Direct connections to all positions (Transformers)

Gradient Clipping

$$\text{if } \|\mathbf{g}\| > \text{threshold}: \quad \mathbf{g} \leftarrow \frac{\text{threshold}}{\|\mathbf{g}\|} \cdot \mathbf{g}$$

Simple and effective for exploding gradients — does not help vanishing.

LSTM: The Walkie-Talkie Solution

Instead of whispering through a chain, LSTM provides a **cell state highway** where information flows with minimal distortion — like giving each person a walkie-talkie instead of whispering.

The France → French Example Revisited

Model	Can connect "France" to "fluent ___"?
Vanilla RNN	No — gradient vanishes over ~100 steps
LSTM / GRU	Yes — cell state preserves long-range info
Transformer	Yes — attention connects any two positions directly

---

Common Pitfalls / Exam Traps

- **Confusing vanishing and exploding** — vanishing: weights → 0, no learning; exploding: weights → NaN, crash.
- **"Gradient clipping fixes vanishing gradients"** — false; clipping only caps large gradients (exploding).
- **Attributing amnesia to architecture alone** — it is specifically the gradient multiplication in BPTT.
- **Exam trap: solution for vanishing** — LSTM/GRU (gated cell state), not gradient clipping.

Quick Revision Summary

- Vanilla RNNs forget long-range context due to vanishing gradients in BPTT.
- Chain rule multiplies gradients across all time steps.
- Weights < 1 → vanishing (gradient → 0); weights > 1 → exploding (gradient → NaN).
- Gradient clipping fixes exploding gradients only.
- LSTM/GRU solve vanishing via gated cell state (information highway).
- "I grew up in France … I speak fluent ___" is the canonical long-context test.

← **Previous** · **Index** · **Next** →
Long Short-Term Memory (LSTM) Networks

Long Short-Term Memory (LSTM) Networks

Intuition: A Computer Chip with Gates

Vanilla RNNs suffer from short-term memory — gradients vanish over long sequences. **LSTM** (Long Short-Term Memory) solves this by separating **short-term processing** (hidden state) from **long-term storage** (cell state) and using **gates** to control what information flows, stays, or gets discarded.

Think of LSTM not as a simple neuron but as a logic chip with selective memory management.

Two Memory Stores

Store	Role	Analogy
Hidden state H_t	Short-term working memory	Current thought
Cell state C_t	Long-term storage highway	Persistent notebook

The cell state runs through the entire chain with minimal transformation — gradients flow backward through it largely unchanged, solving the vanishing gradient problem.

The Four Components

1. Cell State (C_t) — The Highway

Information flows straight down the cell state path with few interactions. This "superhighway" allows gradients to propagate backward without vanishing.

2. Forget Gate — What to Discard

Decides what information to **throw away** from the cell state.

- Uses **sigmoid** activation → output in [0, 1]
- 0 = forget completely; 1 = keep everything

Example: When the subject changes from singular to plural, forget the old singular subject information.

$$f_t = \sigma(W_f \cdot [H_{t-1}, X_t] + b_f)$$

3. Input Gate — What to Store

Decides what **new information** to write into the cell state.

- Sigmoid: what values to update
- Tanh: candidate new values

Example: Store the new gender of a plural subject.

$$i_t = \sigma(W_i \cdot [H_{t-1}, X_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [H_{t-1}, X_t] + b_C)$$

4. Output Gate — What to Reveal

Decides what to **output** based on the current cell state.

Example: Output a verb conjugated for the new plural subject.

$$o_t = \sigma(W_o \cdot [H_{t-1}, X_t] + b_o)$$
$$H_t = o_t \odot \tanh(C_t)$$

Simplified LSTM Flow

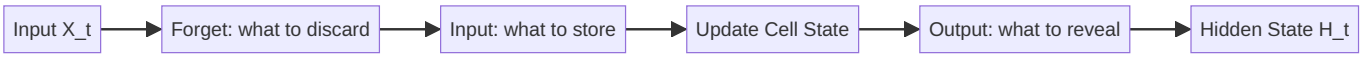

1. Receive input X_t and previous states H_{t-1}, C_{t-1}
2. **Forget gate** removes irrelevant old information
3. **Input gate** adds relevant new information
4. **Cell state** is updated
5. **Output gate** filters cell state to produce H_t

LSTM vs Vanilla RNN

Criterion	Vanilla RNN	LSTM
Long-term memory	Poor (vanishing gradients)	Strong (cell state highway)
Gates	None	Forget, input, output
Parameters	Fewer	More (4 gate weight matrices)
Compute cost	Lower	Higher
Best for	Short sequences	Long paragraphs, translation

Practical Note

Understanding LSTM gate equations in full detail is valuable for research but **not required** to use LSTMs effectively in practice. Frameworks like PyTorch (`nn.LSTM`) and TensorFlow (`tf.keras.layers.LSTM`) handle the internals; practitioners focus on architecture design, hyperparameters, and data.

Production Applications

Task	Why LSTM
Machine translation	Long source sentences require persistent memory
Speech recognition	Audio sequences span thousands of frames
Time-series forecasting	Long historical context (stock prices, sensor data)
Text generation	Maintain thematic coherence over paragraphs

Common Pitfalls / Exam Traps

- **Confusing hidden state and cell state** — C_t is long-term storage; H_t is short-term output.
- **"LSTM has no vanishing gradient problem"** — the cell state mitigates it; very extreme cases can still struggle.
- **Forget gate uses tanh** — false; forget gate uses sigmoid (0 to 1); tanh is for candidate values.
- **Exam trap: three gates** — forget, input, output (cell state is not a gate, it is the highway).

Quick Revision Summary

- LSTM separates short-term (hidden state) and long-term (cell state) memory.
- Three gates: forget (discard), input (store), output (reveal).
- Cell state is a gradient highway — solves vanishing gradients.
- Forget/input gates use sigmoid; candidate values use tanh.
- More parameters than vanilla RNN but handles long sequences effectively.
- Frameworks abstract gate internals; focus on architecture and data for applied use.

← **Previous** · **Index** · **Next** →
Gated Recurrent Unit (GRU)

Gated Recurrent Unit (GRU)

Intuition: LSTM's Leaner Sibling

LSTM solved the long-term memory problem but at a cost: four gate weight matrices, separate cell and hidden states, and significant computational overhead. **GRU** (Gated Recurrent Unit) achieves comparable performance with a simplified architecture — merging the cell and hidden states and reducing three LSTM gates to two.

GRU vs LSTM: Architectural Comparison

Component	LSTM	GRU
Memory stores	Cell state + hidden state	Hidden state only (merged)
Gates	3 (forget, input, output)	2 (update, reset)
Parameters	More	Fewer
Training speed	Slower	Faster
Long sequence performance	Slightly better	Comparable

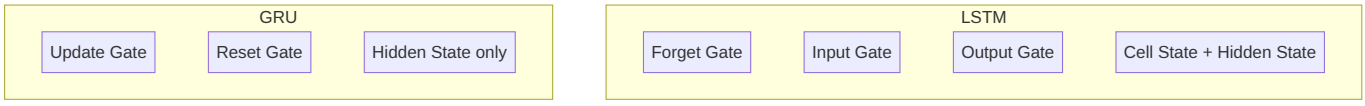

The Two GRU Gates

1. Update Gate — Merge of Forget + Input

Combines LSTM's forget and input gates into one decision:

> "How much of the past should I keep, and how much new information should I let in?"

Value	Behavior
Near 0	Discard past, accept new information
Near 1	Keep past, ignore new information

Analogy: "Should I update my memory now?"

$$z_t = \sigma(W_z \cdot [H_{t-1}, X_t])$$

2. Reset Gate — Selective Context Ignoring

Decides how much of the **past information to ignore** when computing the new candidate state.

Analogy: "Drop the previous context for a moment — does this new word make sense on its own?"

$$r_t = \sigma(W_r \cdot [H_{t-1}, X_t])$$

The reset gate is useful when a new topic begins and prior context would mislead the prediction.

GRU Computation (Simplified)

1. Compute update gate z_t and reset gate r_t from H_{t-1} and X_t
2. Compute candidate hidden state $\tilde{H}_t$ (using reset gate to optionally ignore past)
3. Interpolate between old and new state using update gate:

$$H_t = (1 - z_t) \odot H_{t-1} + z_t \odot \tilde{H}_t$$

The update gate controls the blend: mostly old memory or mostly new information.

When to Choose GRU vs LSTM

Scenario	Recommendation	Reason
Small dataset (tweets, reviews)	**GRU**	Fewer parameters → less overfitting
Fast prototyping	**GRU**	Faster training
Very long sequences (translation, long text gen)	**LSTM**	Slightly better long-range memory
Limited compute budget	**GRU**	Fewer FLOPs per step
GRU underperforms on validation	**Switch to LSTM**	More capacity for complex patterns

Rule of thumb: Start with GRU for speed. If validation performance is insufficient, upgrade to LSTM.

Real-World Usage

Application	Typical Choice
Tweet sentiment analysis	GRU (short sequences, small data)
Product review classification	GRU
Machine translation (production)	LSTM or Transformer (long sequences)
Time-series anomaly detection	GRU (speed matters for real-time)

In cloud ML pipelines, GRU-based models deploy faster and cost less per inference — a meaningful advantage at scale when LSTM's marginal accuracy gain does not justify the overhead.

GRU vs LSTM vs Transformer

Model	Memory mechanism	Long-range	Speed	Era
Vanilla RNN	Hidden state only	Poor	Fast	Legacy
GRU	Gated hidden state	Good	Fast	2014+
LSTM	Gated cell + hidden	Better	Moderate	1997+/dominant 2014–2017
Transformer	Self-attention	Best	Parallelizable	2017+/current

Modern production NLP (BERT, GPT) has largely replaced RNN/LSTM/GRU, but understanding gated architectures remains essential for legacy systems, resource-constrained deployments, and exam contexts.

Common Pitfalls / Exam Traps

- **"GRU has a separate cell state"** — false; GRU merges cell and hidden state.
- **"GRU is always worse than LSTM"** — false; often comparable, especially on smaller datasets.
- **Confusing update and reset gates** — update: how much past to keep; reset: how much past to ignore for new candidate.
- **Exam trap: GRU gate count** — 2 gates (update + reset), not 3.

Quick Revision Summary

- GRU is a simplified LSTM: 2 gates instead of 3, no separate cell state.
- Update gate: blends past and new information (replaces forget + input gates).
- Reset gate: optionally ignores past when computing new candidate state.
- Faster to train, fewer parameters, often matches LSTM on smaller datasets.
- Rule of thumb: start with GRU; switch to LSTM if underperforming on long/complex sequences.
- Transformers have largely superseded both in modern NLP, but GRU/LSTM remain relevant for constrained environments.

← **Previous** · **Index** · **Next** →

Week 8

Week 8

From Sequential Models to Transformers: The Motivation

Why RNNs and LSTMs Hit a Wall

Recurrent Neural Networks (RNNs) and their improved variant, Long Short-Term Memory (LSTM) networks, were designed to handle sequences by maintaining hidden state across time steps. LSTMs address the vanishing gradient problem better than plain RNNs and can retain information over longer spans — a genuine advance over earlier sequence models.

However, both architectures share a structural constraint: **sequential processing**. To compute the representation of the 100th token, the model must first process tokens 1 through 99. Each step depends on the previous one, so the computation graph unfolds strictly in order.

The GPU Parallelism Problem

Modern GPUs excel at **massive parallel matrix operations** — the same kind of work that powers training on cloud ML platforms (AWS SageMaker, Google Vertex AI, Azure ML). Sequential models underutilize this hardware:

- Only one time step runs at a time per sequence
- Batch parallelism helps across sequences, but not within a single long document
- Training on large corpora (millions of sentences) becomes prohibitively slow

For a 100-word sentence, the bottleneck is manageable. For a document with 1,000–10,000 words, sequential dependency makes both **training time** and **inference latency** unacceptable at production scale.

The Long-Distance Dependency Problem

Even with LSTM gating, information from early tokens tends to **fade** by the time the model reaches token 100. Real NLP tasks — coreference resolution ("it" referring to "animal"), document-level classification, long-form summarization — require linking distant words. LSTMs partially mitigate this but do not eliminate it.

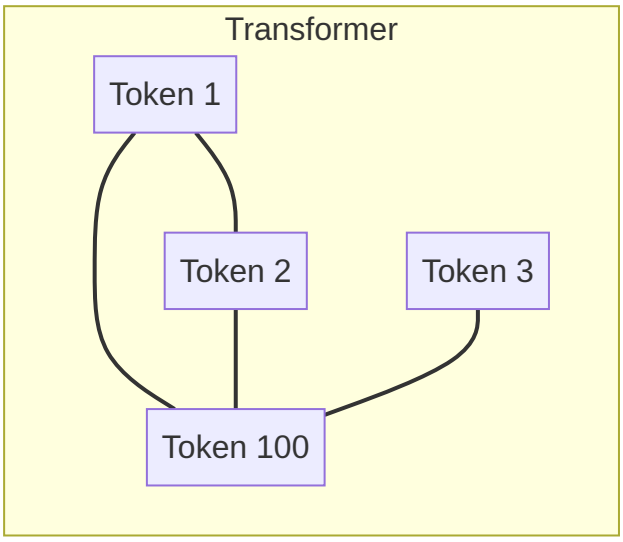

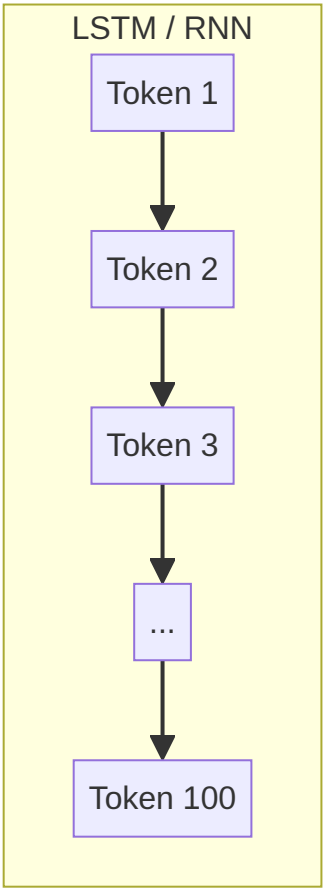

The Transformer Idea

Transformers abandon recurrence entirely. Instead, they rely on a mechanism that lets **any token attend to any other token in a single parallel pass**:

- Process the entire sentence (or chunk) at once
- Each word can directly "look at" every other word and compute a relevance score
- No waiting for prior tokens to finish processing

This design unlocks full GPU parallelism and direct long-range connections — the foundation of BERT, GPT, Gemini, Claude, and virtually every modern NLP system deployed in production.

Limitation	LSTM / RNN	Transformer
Processing order	Sequential (token by token)	Parallel (all tokens at once)
GPU utilization	Low within a sequence	High — matrix ops across all positions
Long-range links	Indirect, decay over distance	Direct via attention weights
Training speed on large data	Slow	Much faster at scale

Common Pitfalls / Exam Traps

- **Trap:** "LSTMs solve all long-sequence problems." LSTMs improve memory but still process sequentially and still struggle with very long dependencies compared to attention.
- **Trap:** Confusing "parallel batching" (processing many sequences at once) with "parallel within-sequence processing" — only Transformers achieve the latter.
- **Trap:** Assuming Transformers eliminate the need for positional information — they process tokens in parallel but still need positional encoding to preserve word order.

Quick Revision Summary

- LSTMs fix vanishing gradients but retain sequential processing as a core bottleneck.
- Sequential models cannot efficiently use GPU parallelism within a single long sequence.
- Long documents (1,000+ words) expose both speed and long-distance dependency failures in RNN/LSTM architectures.
- Transformers replace recurrence with attention: every token can relate to every other token in one step.
- This parallel design is the prerequisite for training on internet-scale data and building modern LLMs.
- The shift from LSTM-era NLP to Transformer-era NLP is driven by hardware efficiency and direct long-range modeling, not just incremental accuracy gains.

← **Previous** · **Index** · **Next** →
Attention Is All You Need: Origins of the Transformer

Attention Is All You Need: Origins of the Transformer

Historical Context

Before 2017, state-of-the-art sequence modeling relied heavily on **Recurrent Neural Networks (RNNs)** and **Convolutional Neural Networks (CNNs)** for tasks like machine translation. These architectures processed text step-by-step or through local convolution windows — workable, but slow and limited in capturing global context efficiently.

In 2017, researchers at Google Brain published **"Attention Is All You Need"** (Vaswani et al.). The paper's central proposal was radical in its simplicity:

> Discard RNNs and CNNs for sequence transduction. Use **self-attention** as the sole mechanism for relating tokens.

The Core Idea: Self-Attention

Self-attention allows each position in a sequence to compute a weighted combination of representations from **all other positions**. No hidden state needs to propagate sequentially from left to right. The model learns which tokens matter for interpreting each token — in parallel.

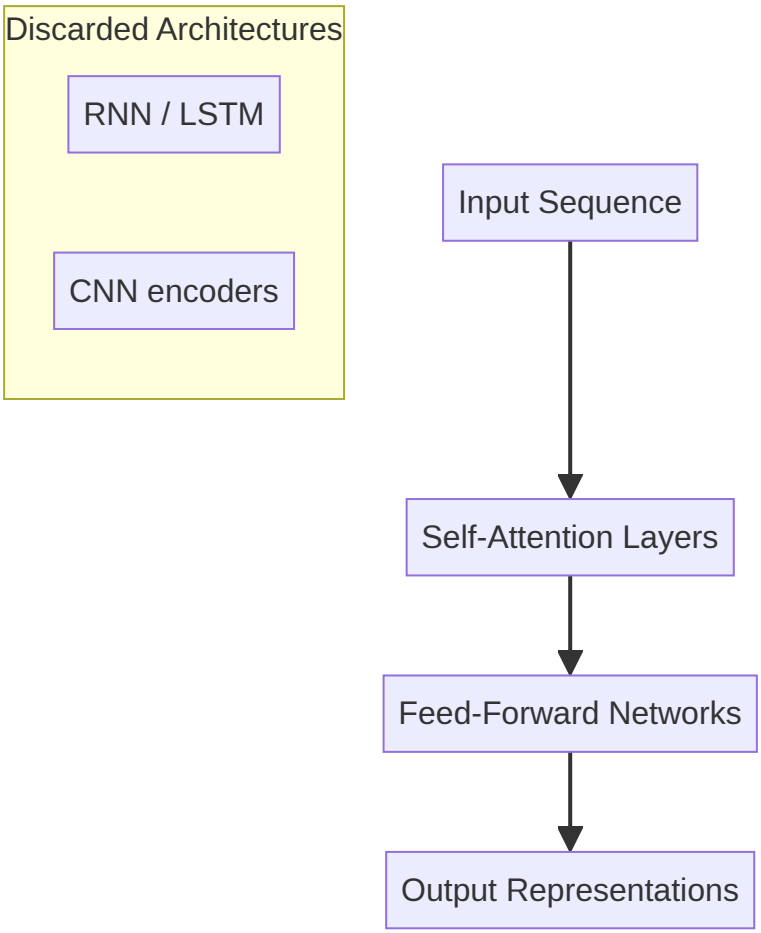

Impact on Scale and the LLM Era

The Transformer architecture enabled training on **massive corpora** that were impractical under sequential models:

- Full Wikipedia dumps
- Reddit comment archives
- Stack Overflow Q&A
- Broad web-scale crawls

Training became **orders of magnitude faster** because attention layers map cleanly to parallel GPU kernels. The result was not just faster experiments — it unlocked models large enough to exhibit emergent language understanding.

What Followed

The paper directly enabled the modern large language model ecosystem:

Model family	Role	Transformer variant
BERT	Language understanding	Encoder-only
GPT-3 / GPT-4	Text generation	Decoder-only
Gemini, Claude	General-purpose LLMs	Decoder-only (scaled)
T5, BART	Translation, summarization	Encoder-decoder

Every major cloud NLP API — OpenAI, Google AI, Anthropic, Azure OpenAI — rests on Transformer derivatives trained on web-scale data.

Why the Paper Mattered for Industry

- **Transfer learning at scale:** Pre-train once on huge data, fine-tune for downstream tasks (sentiment, NER, QA).
- **Infrastructure alignment:** Transformers match the hardware trend toward tensor-core GPUs and TPU pods.
- **Open ecosystem:** Architectures and weights published on Hugging Face accelerated adoption across startups and enterprises.

Real-world example: A fintech company routing support tickets previously used bag-of-words + logistic regression. After BERT-based fine-tuning (enabled by the Transformer paper's architecture), accuracy on intent classification jumped while the same cloud GPU budget supported larger batch sizes during training.

Common Pitfalls / Exam Traps

- **Trap:** Crediting LSTMs as the foundation of GPT — GPT is **decoder-only Transformer**, not LSTM-based.
- **Trap:** Thinking "Attention Is All You Need" introduced attention for the first time — attention existed in seq2seq models before; the novelty was using **only** self-attention, removing recurrence entirely.
- **Trap:** Assuming the original Transformer is identical to GPT — the 2017 model had **both encoder and decoder**; GPT uses only the decoder stack.

Quick Revision Summary

- Published in 2017 by Google Brain; title: "Attention Is All You Need."
- Proposed replacing RNNs and CNNs with pure self-attention for sequence modeling.
- Self-attention computes relationships between all token pairs in parallel.
- Enabled training on web-scale datasets (Wikipedia, Reddit, Stack Overflow) much faster than before.
- Directly birthed BERT, GPT, Gemini, Claude, and the modern LLM industry.
- The original architecture combined encoder and decoder; later models often use only one side.
- Cloud-scale NLP today is fundamentally a consequence of this architectural shift.

← **Previous** · **Index** · **Next** →
Transformer Architecture: Self-Attention, Math, and Positional Encoding

Transformer Architecture: Self-Attention, Math, and Positional Encoding

Self-Attention: Intuition First

Consider the sentence:

> The animal did not cross the street because **it** was too tired.

Who is "it"? A human reader resolves this instantly: **it** refers to **animal**, not **street**. Resolution requires comparing **it** against every other word and judging relevance.

Self-attention automates this process:

1. Each word generates a representation
2. Every word "looks at" every other word
3. A **relevance score** is computed for each pair
4. Higher scores mean stronger influence when building the final representation

For **it**, the relevance score with **animal** exceeds the score with **street**, so the model's representation of **it** is dominated by **animal**.

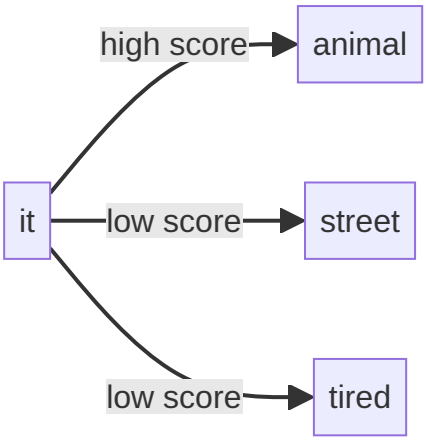

This bidirectional comparison happens for **every token simultaneously** — the defining property of self-attention.

Full Transformer Architecture Overview

The original Transformer is an **encoder-decoder** model for sequence-to-sequence tasks (e.g., translation).

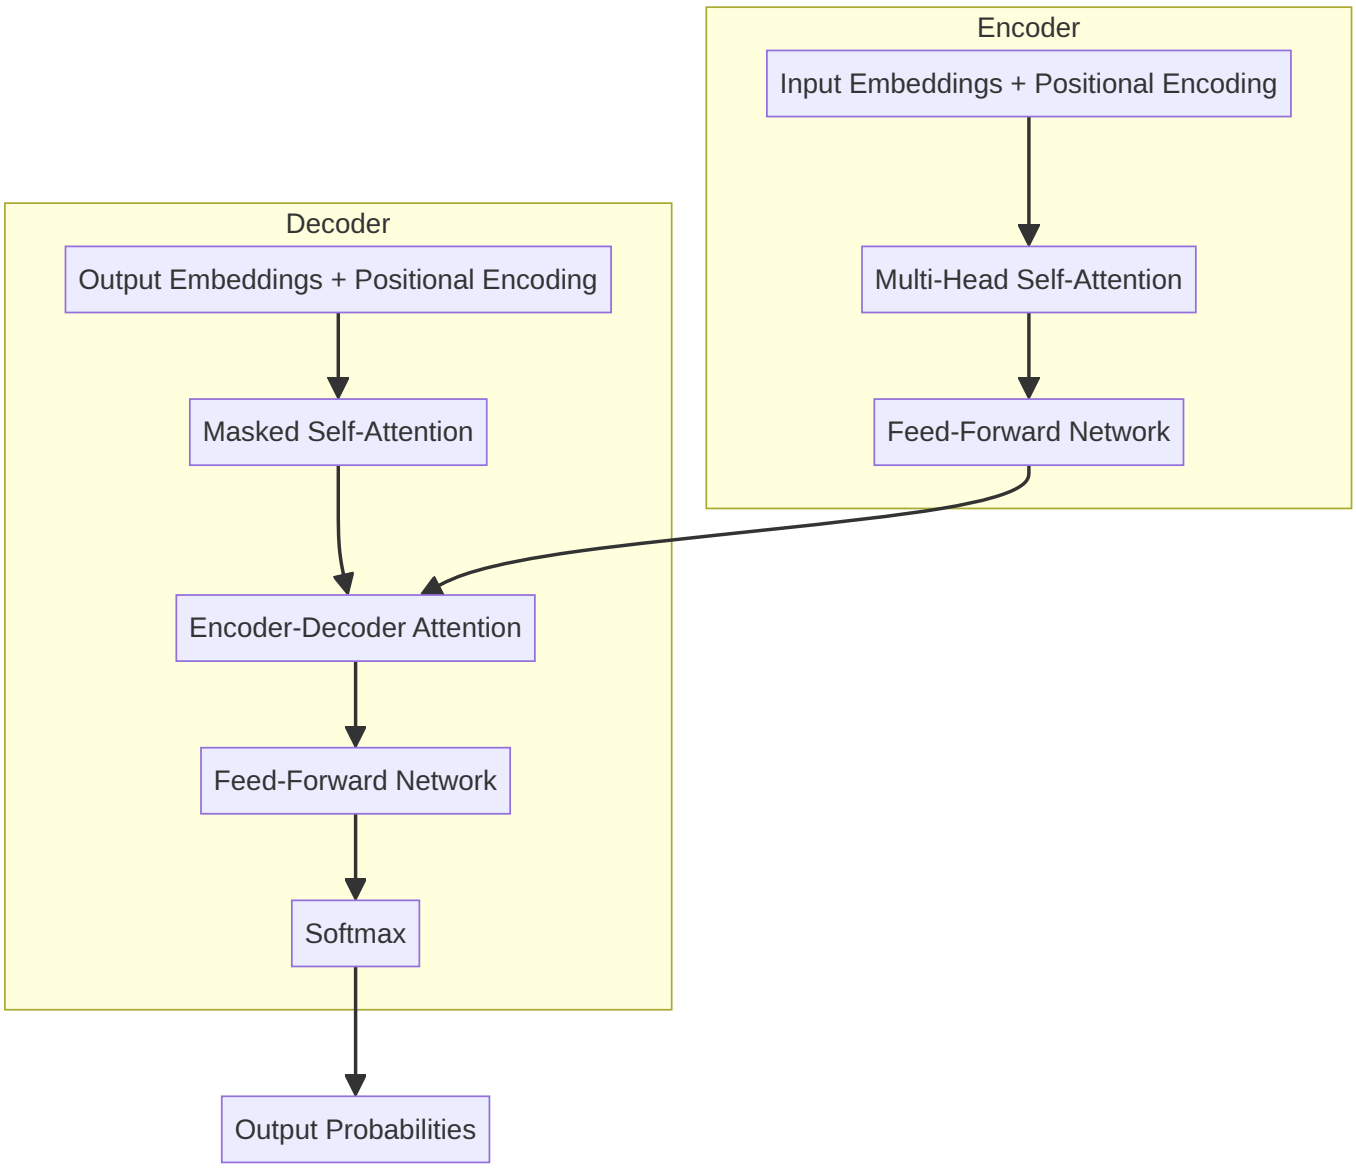

Pipeline stages

Stage	Purpose
Input embeddings	Convert token IDs to dense vectors
Positional encoding	Inject order information (see below)
Encoder self-attention	Tokens attend to all tokens in the source
Feed-forward network (FFN)	Non-linear transformation per position
Decoder masked attention	Tokens attend only to prior decoder positions
Encoder-decoder attention	Decoder queries attend to encoder outputs
Softmax	Convert logits to probability distribution over vocabulary

"Shifted right" in decoder inputs means the decoder receives previously generated tokens as input during training (teacher forcing), shifted one position so it predicts the next token at each step.

The Attention Equation

Scaled dot-product attention is defined as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$$

Symbol	Meaning	Intuition
Q (Query)	"What am I looking for?"	The token asking for context
K (Key)	"What do I contain?"	Labels/tags used for matching
V (Value)	"What information do I pass?"	Actual content aggregated after matching
d_k	Key dimension	Scaling factor prevents dot products from growing too large

Softmax converts raw scores into a probability distribution (values between 0 and 1 that sum to 1), determining how much each value vector contributes.

The scaling by $\sqrt{d_k}$ stabilizes gradients when dimensionality is high — a detail often tested in exams.

Multi-Head Attention

One attention head captures one type of relationship. **Multi-head attention** runs several heads **in parallel**, each learning different patterns:

- Head 1: syntactic structure (subject–verb agreement)
- Head 2: semantic similarity (synonyms, related concepts)
- Head 3: coreference (pronouns → antecedents)

Outputs from all heads are **concatenated** and projected. This replaces sequential processing with parallel specialized attention — directly addressing the GPU efficiency problem.

Positional Encoding: Why Order Matters

Because all tokens are processed simultaneously, the raw embedding of "The dog bit the man" is **identical** to "The man bit the dog" without additional information. Word order is lost.

Solution: Add a positional signal to each embedding using **sine and cosine waves** of different frequencies. Each position receives a unique encoding vector.

Example for two sentences with the same words in different order:

Word	Sentence 1 position	Sentence 2 position
The	1	1 (different role)
dog	2	5
bit	3	3
the	4	4
man	5	2

The model can now distinguish "dog bit man" from "man bit dog" because positional encodings differ even when word embeddings repeat.

Common Pitfalls / Exam Traps

- **Trap:** Saying self-attention is unidirectional in the encoder — encoder self-attention is **fully bidirectional**; only the decoder uses masking for autoregressive generation.
- **Trap:** Forgetting \sqrt{d_k} in the attention formula — omitting it is a common exam error.
- **Trap:** Confusing Query/Key/Value with input/output embeddings — they are derived from the same input via learned linear projections.
- **Trap:** Assuming Transformers inherently know word order — positional encoding (or learned position embeddings in BERT/GPT) is **mandatory**.

Quick Revision Summary

- Self-attention: each token scores its relationship with every other token; high scores drive the final representation.
- Coreference ("it" → "animal") is a canonical example of attention-based disambiguation.
- Encoder: bidirectional self-attention + FFN; Decoder: masked self-attention + cross-attention + FFN + softmax.
- Attention formula: \text{softmax}(QK^T / \sqrt{d_k}) \cdot V with Query, Key, Value roles.
- Multi-head attention runs parallel specialized heads (grammar, semantics, coreference) and concatenates results.
- Positional encoding (sin/cos waves) injects order because parallel processing discards sequence information by default.
- Scaling by \sqrt{d_k} prevents softmax saturation in high dimensions.

← **Previous** · **Index** · **Next** →
Popular Transformer Model Families: Encoders, Decoders, and Seq2Seq

Popular Transformer Model Families: Encoders, Decoders, and Seq2Seq

From One Architecture to Three Ecosystems

The original 2017 Transformer combined an **encoder** (reads and understands input) and a **decoder** (generates output). Modern production models rarely use both stacks unchanged. Instead, practitioners pick the half that matches the task — or use the full encoder-decoder for transformation tasks.

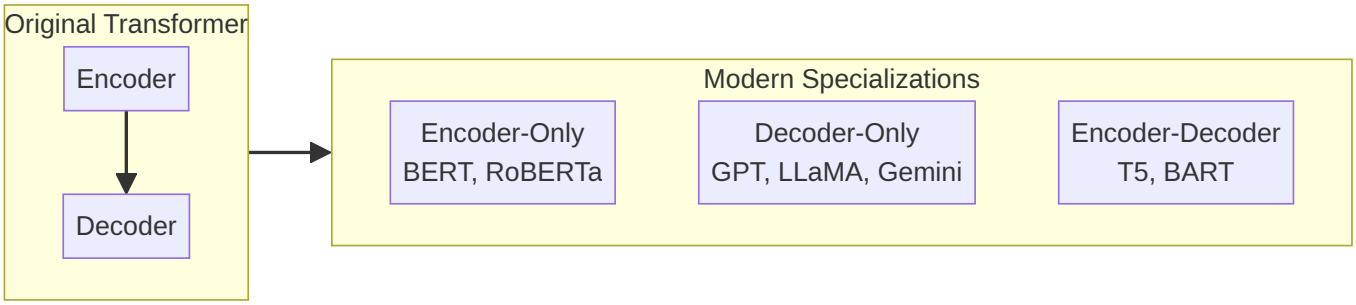

Encoder-Only Models: The Readers

Encoder-only models process the **entire input bidirectionally** — every token sees every other token. They excel at **understanding**, not generation.

Model	Notes
BERT	Bidirectional Encoder Representations from Transformers
RoBERTa	Robustly optimized BERT (more data, no NSP)
DistilBERT	Smaller, faster distilled BERT

Strengths:

- Text classification (spam, intent, topic)
- Sentiment analysis
- Named Entity Recognition (NER)
- Question answering (with fine-tuning)
- Semantic similarity

Real-world use: A cloud support platform fine-tunes DistilBERT on ticket text to route "billing" vs "technical" vs "cancellation" requests — low latency on CPU-friendly deployments.

Key property: Sees context from **both left and right** simultaneously (bidirectional).

Decoder-Only Models: The Writers

Decoder-only models generate text **autoregressively** — each new token is predicted from all **prior** tokens only (left context). They cannot see future tokens during generation.

Model	Notes
GPT (Generative Pre-trained Transformer)	The **T** stands for Transformer
LLaMA	Meta's open-weight decoder family
Gemini	Google's multimodal decoder stack

Strengths:

- Open-ended text generation
- Code completion
- Chat and dialogue
- Creative writing, summarization (via prompting)

Real-world use: GitHub Copilot and ChatGPT-style assistants are decoder-only models scaled to billions of parameters and aligned with human feedback.

Key property: Causal / unidirectional attention — only past tokens inform the next prediction.

Encoder-Decoder Models: The Translators

These retain both stacks: the encoder builds a rich representation of the source; the decoder generates the target sequence while attending to the encoder.

Model	Notes
T5 (Text-to-Text Transfer Transformer)	Frames all tasks as text-to-text
BART	Denoising autoencoder pre-training

Strengths:

- Machine translation
- Abstractive summarization
- Data-to-text generation
- Any task mapping one sequence to another

Real-world use: Google Translate pipelines and enterprise document summarization APIs often use encoder-decoder Transformers fine-tuned per language pair.

Comparison Table

Aspect	Encoder-Only	Decoder-Only	Encoder-Decoder
Direction	Bidirectional	Unidirectional (left-to-right)	Encoder: bidirectional; Decoder: causal
Primary task	Understanding / classification	Generation	Sequence transformation
Examples	BERT, RoBERTa, DistilBERT	GPT, LLaMA, Gemini	T5, BART
Sees full input at once?	Yes	No (during generation)	Encoder yes; decoder incrementally
Typical deployment	Classification APIs, search	Chatbots, copilots	Translation, summarization

Common Pitfalls / Exam Traps

- **Trap:** Using GPT (decoder-only) for bidirectional NER without adaptation — GPT sees only left context unless fine-tuned with special prompting; BERT is the natural choice for token-level understanding.
- **Trap:** Forgetting that **T in GPT = Transformer**, not "trained" or "text" alone.
- **Trap:** Assuming BERT can generate long coherent paragraphs out of the box — BERT is not autoregressive; it needs a classification head or separate generation stack.
- **Trap:** Treating T5 and BART as encoder-only — they require **both** halves for seq2seq.

Quick Revision Summary

- Original Transformer = encoder + decoder; modern models often use one side or both depending on task.
- Encoder-only (BERT family): bidirectional readers — classification, sentiment, NER, QA.
- Decoder-only (GPT family): left-to-right writers — generation, chat, code; T = Transformer.
- Encoder-decoder (T5, BART): translators — translation, summarization, seq2seq.
- Choose architecture by task: understand → encoder; generate → decoder; transform sequence → encoder-decoder.
- DistilBERT and similar variants trade minimal accuracy for speed — common in production edge deployments.

← **Previous** · **Index** · **Next** →
Hugging Face: The Open ML Hub for Models, Data, and Apps

Hugging Face: The Open ML Hub for Models, Data, and Apps

What Hugging Face Is

Hugging Face (`huggingface.co`) is a collaborative platform where the machine learning community publishes and consumes **models**, **datasets**, and **applications**. It began as an aggregator for open-source Transformer weights and evolved into a full ecosystem for AI developers — comparable to GitHub for code, but centered on ML artifacts.

As of recent counts, the hub hosts **over 2 million models**, spanning NLP, vision, audio, and multimodal systems. Major contributors include Meta (Facebook), Mistral, Baidu, Cohere, and thousands of independent researchers.

Platform Layout

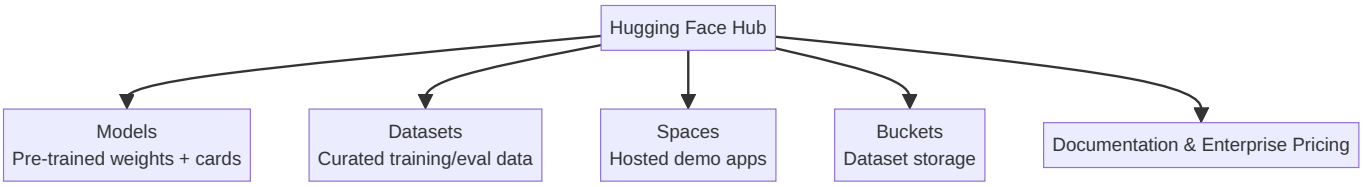

Models Hub

The core value proposition: **download and run state-of-the-art models in a few lines of Python**.

Exploring a model page (e.g., `bert-base-uncased`)

Each model page typically includes:

- **Model card:** Description, training data, intended use, limitations
- **Variations:** Related checkpoints (uncased, cased, fine-tuned variants)
- **Usage snippet:** Copy-paste Python with `transformers` library
- **Interactive widget:** Test inference in the browser without local setup

Fill-in-the-blank example: Input "Paris is the ___" → model predicts **capital** with ~99% confidence. Alternatives (heart, centre, city) rank lower — demonstrating masked language modeling behavior.

Custom input example: "The cat ___ on the mat" → top prediction **sat** (16.8% among plausible verbs: lay, landed, collapsed). The highest-probability token reflects training data statistics, not universal truth.

Integration paths

- **Colab / Kaggle:** One-click notebook links from model pages
- **Local Python:** Copy `pipeline` or `AutoModel` snippets into scripts
- **Enterprise:** Hosted inference endpoints with pricing tiers

Real-world use: A startup building a document QA product prototypes with `bert-base-uncased` from Hugging Face, then swaps to a domain fine-tune (`legal-bert`) without changing pipeline code — only the model ID changes.

Datasets Hub

Hugging Face also hosts high-value datasets with documentation and loading scripts:

- Reasoning traces synthesized by frontier models
- Multilingual corpora
- Domain-specific medical, legal, and financial text

Each dataset page describes schema, license, size, and a `load_dataset()` code example — enabling reproducible training pipelines on cloud VMs.

Spaces: Hosted Applications

Spaces let developers deploy Gradio or Streamlit apps powered by Hugging Face models. Anyone can try the app in a browser.

Example workflow:

1. Build a speech-to-text app using a multilingual ASR model
2. Deploy to Spaces
3. Users upload audio and receive transcripts without installing dependencies

This lowers the barrier from "research checkpoint" to "demoable product" — useful for portfolio projects and stakeholder demos.

Buckets: Dataset Storage

A newer offering: **managed storage** for private or team datasets, analogous to cloud object storage integrated with the hub workflow. Teams training custom models on proprietary data can keep artifacts in one ecosystem.

Getting Started

1. Create a free account (email + password)
2. Browse trending models on the homepage or search by name (e.g., "BERT", "Mistral")
3. Read the model card for license and limitations
4. Copy the usage snippet into Colab or local environment
5. Optionally follow organizations for updates on new releases

Common Pitfalls / Exam Traps

- **Trap:** Assuming all Hugging Face models are free for commercial use — always check the **license** on the model card (Apache 2.0, MIT, CC, research-only, etc.).
- **Trap:** Confusing the **website widget** results with guaranteed production accuracy — browser demos use default checkpoints and may differ from fine-tuned deployments.
- **Trap:** Thinking Hugging Face only hosts NLP models — the hub includes vision (OCR), audio, and multimodal checkpoints.
- **Trap:** Ignoring **model limitations** sections — bias, language coverage, and domain mismatch are documented for a reason.

Quick Revision Summary

- Hugging Face is the primary open hub for sharing and running ML models, datasets, and demo apps.
- Models hub: 2M+ checkpoints with cards, usage code, and browser testing widgets.
- Datasets hub: documented corpora loadable via `datasets` library.
- Spaces: host interactive Gradio/Streamlit apps for public demos.
- Buckets: storage for custom datasets within the platform.
- Typical workflow: search model → read card → copy Python snippet → run in Colab or cloud VM.
- Always verify license, limitations, and domain fit before production deployment.

← **Previous** · **Index** · **Next** →
Hugging Face Tokens and Model Implementation in Colab

Hugging Face Tokens and Model Implementation in Colab

Why You Need an Access Token

Many models on Hugging Face are public and download without authentication. However, **gated models**, private checkpoints, and higher rate limits require an **access token**. Best practice: configure the token once per notebook session so any model download succeeds without embedding secrets in source code.

Step 1: Generate an Access Token

1. Log in at `huggingface.co`
2. Open profile menu → **Access Tokens** (`huggingface.co/settings/tokens`)
3. Click **Create new token**
4. Assign a descriptive name (e.g., `colab-research`)
5. Click **Create token** at the bottom
6. **Copy the token immediately** — it may not be shown again in full

Token management options after creation:

- Edit permissions (read vs write)
- Refresh or revoke

- Delete compromised tokens

Security rule: Never commit tokens to Git, share them in screenshots, or hardcode them in notebook cells visible to others.

Step 2: Store the Token Securely in Google Colab

Colab **Secrets** keep credentials out of notebook source:

1. Click the **key icon** (Secrets) in the left sidebar
2. Add new secret:
   - **Name:** `HF_TOKEN` (exact spelling matters for lookup)
   - **Value:** paste your token
3. Toggle **Notebook access** to grant the runtime permission

Retrieve the secret in code:

```
from google.colab import userdata
hf_token = userdata.get('HF_TOKEN')
```

Store in a variable — do not print the token after loading.

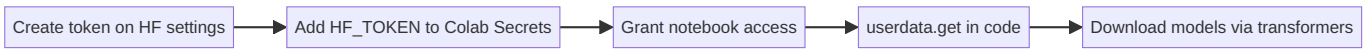

Step 3: Load and Run a Model (Fill-in-the-Blank / MLM)

Copy the usage snippet from a model page (e.g., `bert-base-uncased`):

```
from transformers import pipeline

fill_mask = pipeline("fill-mask", model="bert-base-uncased")
fill_mask("Hello I'm a [MASK] model.")
# Top prediction: "fashion" (among role, new, super, fine, etc.)

fill_mask("The cat sat on the [MASK].")
# Top predictions: floor, bed, couch, sofa, ground
```

Checkpoint: The model identifier string (e.g., `bert-base-uncased`) points to weights + config on the hub. `from_pretrained(checkpoint)` downloads and caches locally.

Not every model requires a token, but loading `HF_TOKEN` upfront avoids failures on gated checkpoints.

Step 4: Sequence-to-Sequence Translation Example

For translation (encoder-decoder), the pattern uses **AutoTokenizer** and **AutoModelForSeq2SeqLM**:

```
from transformers import AutoTokenizer, AutoModelForSeq2SeqLM

checkpoint = "surya7/en-ta-translator"  # example English→Tamil
tokenizer = AutoTokenizer.from_pretrained(checkpoint)
model = AutoModelForSeq2SeqLM.from_pretrained(checkpoint)

def translate(text):
    inputs = tokenizer(text, return_tensors="pt")
    outputs = model.generate(**inputs)
    return tokenizer.decode(outputs[0], skip_special_tokens=True)

translate("Hard work never fails.")
translate("The cat sat on the mat.")
```

Component	Role
`checkpoint`	Hub model ID — selects weights and vocabulary
`AutoTokenizer`	Text → token IDs (subword segmentation)
`AutoModelForSeq2SeqLM`	Encoder-decoder LM for input sequence → output sequence
`model.generate()`	Autoregressive decoding to produce target text

Seq2Seq = sequence-to-sequence: one text sequence in, one text sequence out (translation, summarization).

Tokenizer vs Model

- **Tokenizer:** Splits text into subword tokens, maps to IDs, handles padding/truncation
- **Model:** Neural network forward pass producing logits or generated sequences

Both must come from the **same checkpoint** — mixing tokenizer and model from different families causes garbage output.

Common Pitfalls / Exam Traps

- **Trap:** Hardcoding `hf_...` tokens in notebook cells — use Colab Secrets or environment variables.
- **Trap:** Mismatching tokenizer and model checkpoints — always pair them from the same repo.
- **Trap:** Assuming all models work with `pipeline("fill-mask")` — task type must match architecture (MLM vs seq2seq vs classification).
- **Trap:** Forgetting `truncation=True` on long inputs — BERT-family models typically cap at 512 tokens; longer text is silently truncated or errors without the flag.
- **Trap:** Printing secrets after `userdata.get()` — exposes credentials in notebook output logs.

Quick Revision Summary

- Generate HF access tokens at Settings → Access Tokens; copy once at creation.
- Store as `HF_TOKEN` in Colab Secrets; load via `userdata.get('HF_TOKEN')`.
- Never expose tokens in source code, commits, or printed output.
- `pipeline()` offers high-level task APIs (fill-mask, sentiment, translation).
- Seq2Seq: `AutoTokenizer` + `AutoModelForSeq2SeqLM.from_pretrained(checkpoint)` + `generate()`.
- Checkpoint = hub model ID linking tokenizer weights and architecture.
- Pre-load token even for public models — gated models will fail without it.

← **Previous** · **Index** · **Next** →

Week 9

Week 9

BERT and Contextual Embeddings: Beyond Static Word Vectors

BERT and Contextual Embeddings: Beyond Static Word Vectors

The Problem with Static Embeddings

Before BERT, dominant word representations — Word2Vec, GloVe, fastText — assigned **one fixed vector per word type**. The word "bank" received the same embedding whether the sentence discussed a **riverbank** or **Bank of America**.

This is the **polysemy problem**: one surface form, multiple meanings. Static embeddings average or collapse senses into a single point in vector space, limiting downstream task accuracy.

Approach	"bank" (river)	"bank" (financial)
Word2Vec / GloVe	Same vector	Same vector
BERT	Context-dependent vector	Different vector

What BERT Is

BERT = **B**idirectional **E**ncoder **R**epresentations from **T**ransformers

BERT is an **encoder-only Transformer stack** pre-trained on large text corpora. It revolutionized Natural Language Understanding (NLU) by producing **contextual embeddings** — the vector for each token depends on **all surrounding tokens**.

Contextual Embeddings: Intuition

When BERT processes "I deposited money at the bank near the river," the representation of **bank** incorporates **deposited**, **money**, and **river**. The same word form in a financial headline gets a different internal representation because its neighbors differ.

This directly addresses limitations of Word2Vec-era pipelines used in early cloud search and recommendation systems.

Transfer Learning in the BERT Era

Training BERT from scratch requires enormous compute (Wikipedia-scale data, many GPU-days). **Transfer learning** reuses pre-trained weights:

1. **Pre-training:** Learn general language structure (self-supervised on unlabeled text)
2. **Fine-tuning:** Add a task-specific head; train on labeled data for classification, QA, NER, etc.
3. **Distillation:** Compress a large teacher model into a smaller student (e.g., DistilBERT) for mobile or real-time inference

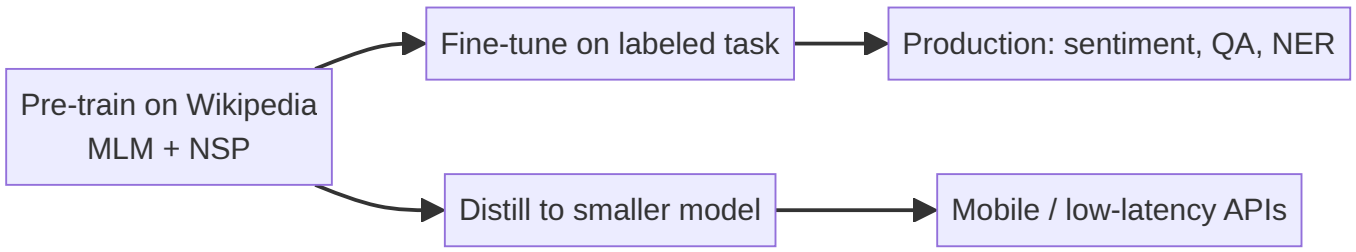

Real-world example: An e-commerce platform fine-tunes BERT on 50k product reviews instead of training a classifier on bag-of-words from scratch — achieving higher F1 on aspect sentiment with less labeled data.

Why BERT Matters for Downstream Tasks

- **Classification:** `[CLS]` token embedding → softmax over labels
- **Question answering:** Predict start/end spans in passage
- **NER:** Per-token labels from contextual representations

All benefit from bidirectional context unavailable to left-to-right models at pre-training time.

Common Pitfalls / Exam Traps

- **Trap:** Calling BERT embeddings "static" — they are **contextual**; only the pre-training objective is fixed, not the runtime vectors.
- **Trap:** Confusing BERT with GPT — BERT is **encoder-only** and **bidirectional**; GPT is **decoder-only** and **unidirectional**.
- **Trap:** Assuming fine-tuning replaces pre-training entirely — fine-tuning **starts from** pre-trained weights; training from random init rarely matches performance.
- **Trap:** Using "BERT" and "Transformers" interchangeably — BERT is one specific encoder-only instantiation.

Quick Revision Summary

- Static embeddings (Word2Vec, GloVe): one vector per word regardless of context.
- BERT: Bidirectional Encoder Representations from Transformers — encoder-only stack.
- Contextual embeddings change per occurrence based on surrounding tokens.
- Polysemy ("bank" river vs financial) motivated the shift from static to contextual models.
- Transfer learning: pre-train once, fine-tune for classification, QA, NER; distill for speed.
- BERT is the foundation model for modern NLU pipelines before the GPT chat era dominated generation tasks.

← **Previous** · **Index** · **Next** →
BERT Architecture and Pre-Training Objectives

BERT Architecture and Pre-Training Objectives

BERT as Encoder-Only Transformer

BERT is a **stack of Transformer encoder blocks** — the encoder half of the original "Attention Is All You Need" architecture, without the decoder. GPT occupies the decoder side; BERT occupies the encoder side.

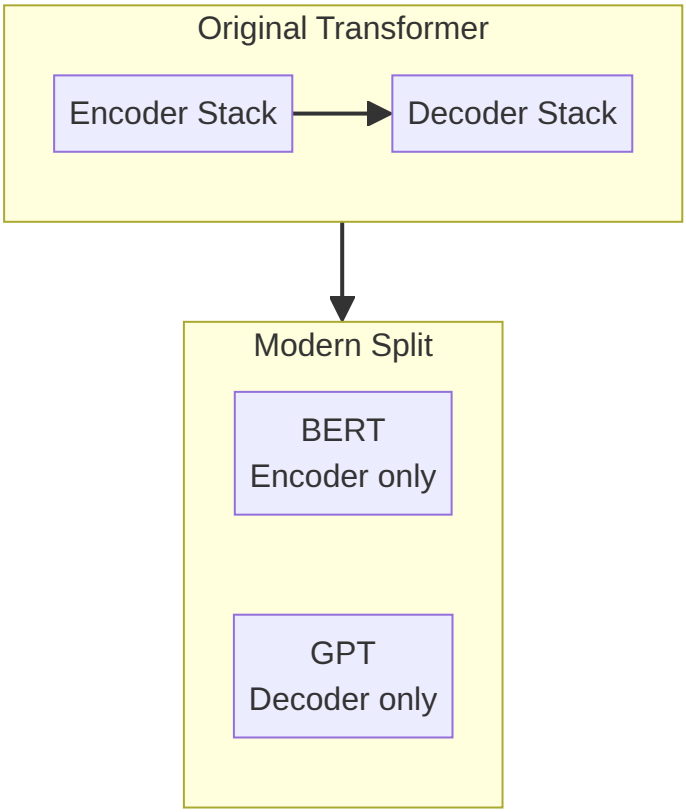

Model sizes

Variant	Encoder layers	Typical use
BERT-base	12	Default for fine-tuning experiments
BERT-large	24	Higher capacity; more compute

BERT was **not** trained for translation or autoregressive generation. It was pre-trained to **understand** language via two self-supervised tasks on Wikipedia (and BooksCorpus).

Pre-Training Task 1: Masked Language Modeling (MLM)

Goal: Force the model to use **bidirectional context** to predict hidden words.

Procedure:

1. Take a sentence: "The cat sat on the mat"
2. Randomly mask ~15% of tokens (e.g., hide **sat** → "The cat [MASK] on the mat")
3. Model predicts the original token using **both left and right** context
4. Repeat across millions of sentences with different masks each epoch

Why "bidirectional": Unlike GPT (left-to-right only), BERT sees **sat**'s neighbors on both sides when filling the blank. The label for **sat** depends on **cat**, **on**, and **mat** simultaneously.

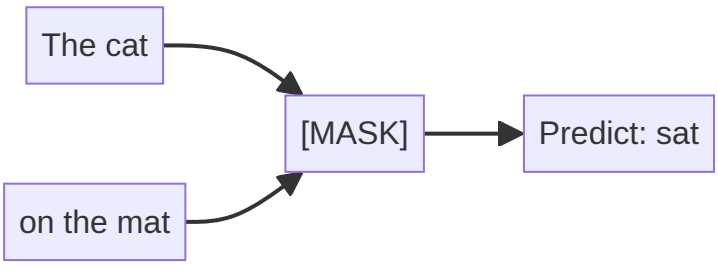

Training detail (often tested): Of selected tokens, 80% are replaced with `[MASK]`, 10% with random words, 10% unchanged — reducing train/serve mismatch.

Pre-Training Task 2: Next Sentence Prediction (NSP)

Goal: Teach **inter-sentence relationships** — critical for QA, natural language inference, and summarization.

Procedure:

1. Split text into sentence pairs:
   - Sentence A: "The man went to the store"
   - Sentence B: "He bought milk"
2. Training examples include:
   - **IsNext:** B truly follows A in the original document
   - **NotNext:** B is a random sentence from elsewhere
3. Model predicts whether B is the actual next sentence

Why it matters: Question-passage pairs, premise-hypothesis pairs, and multi-sentence inputs require knowing whether two spans are related — not just whether individual words co-occur.

Task	What the model learns
MLM	Token-level meaning from full context
NSP	Discourse / sentence-pair coherence

Input Format

BERT inputs concatenate segments with special tokens:

`[CLS] sentence_A [SEP] sentence_B [SEP]`

- `[CLS]` — classification aggregate used for sentence-level tasks
- `[SEP]` — separator between segments

Segment embeddings + token embeddings + positional embeddings are summed before entering the encoder stack.

Common Pitfalls / Exam Traps

- **Trap:** Saying BERT is trained left-to-right — MLM is **bidirectional**; only the decoder in GPT is strictly causal.
- **Trap:** Confusing MLM with autoregressive LM — MLM predicts **masked** positions with full context; GPT predicts **next** token with left context only.
- **Trap:** Assuming NSP is universally kept — RoBERTa **removed NSP** after ablations showed limited benefit; exam questions may contrast BERT vs RoBERTa on this point.
- **Trap:** Stating BERT-base has 24 layers — **12 layers**; 24 is BERT-large.

Quick Revision Summary

- BERT = stack of Transformer **encoders only** (no decoder).
- BERT-base: 12 layers; BERT-large: 24 layers.
- Pre-trained on Wikipedia-scale text, not for translation/generation.
- **MLM:** mask tokens; predict using bidirectional context → contextual embeddings.
- **NSP:** predict if sentence B follows sentence A → sentence-pair understanding.
- Input uses `[CLS]`, `[SEP]`, token, segment, and positional embeddings.
- RoBERTa later dropped NSP; BERT original recipe includes both MLM and NSP.

← **Previous** · **Index** · **Next** →
Applications of BERT in Production NLP

Applications of BERT in Production NLP

Why BERT Excels at Understanding Tasks

BERT produces **deep contextual representations** for every token plus a sentence-level `[CLS]` embedding. After pre-training, a lightweight task-specific layer is added and fine-tuned on labeled data — far less data and compute than training from scratch.

Because BERT sees **full bidirectional context**, it outperforms bag-of-words, TF-IDF, and static embedding classifiers on nuanced text.

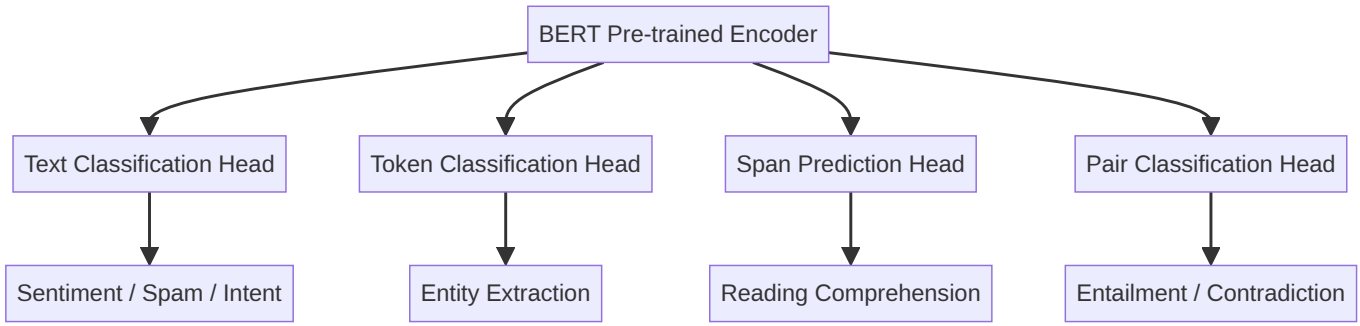

Text Classification

Assign a **single label** (or multi-label set) to an entire document or sentence.

Application	Label examples
Sentiment analysis	positive, negative, neutral
Spam detection	spam, not spam
Intent recognition	book_flight, cancel_reservation, track_order
Topic labeling	sports, finance, politics
Ticket routing	billing, technical, account

Mechanism: Fine-tune BERT with `[CLS]` output → linear layer → softmax.

Real-world use: Gmail-style spam filtering and Zendesk-style support ticket categorization on cloud NLP APIs often use fine-tuned BERT or DistilBERT backends.

Named Entity Recognition (NER)

Extract **structured entities** from unstructured text — persons, organizations, locations, dates, medical terms.

Example: "Apple announced a new office in Austin" → `ORG: Apple`, `LOC: Austin`

Mechanism: Fine-tune with a **token classification head** — each subword token receives a BIO tag (Begin, Inside, Outside).

Real-world use: Healthcare pipelines extract drug names and dosages from clinical notes; legal tech extracts party names and jurisdictions from contracts.

Question Answering (QA)

Given a **question** and a **context passage**, predict the **start and end span** of the answer in the passage.

Example:

- Context: "The Transformer was published in 2017 by Google Brain."
- Question: "Who published the Transformer?"
- Answer span: "Google Brain"

Mechanism: Two logits per token (start probability, end probability); highest-scoring valid span wins.

Real-world use: Enterprise search over internal wikis; customer support bots retrieving policy clauses.

Natural Language Inference (NLI)

Determine the logical relationship between a **premise** and a **hypothesis**:

- **Entailment:** hypothesis must be true if premise is true
- **Contradiction:** hypothesis cannot be true if premise is true
- **Neutral:** neither entailed nor contradicted

Example:

- Premise: "It rained all day."
- Hypothesis: "The ground is wet." → Entailment

Mechanism: Sentence pair input `[CLS] premise [SEP] hypothesis [SEP]` → classification head.

Real-world use: Fact-checking pipelines, contract compliance checking (does clause A imply obligation B?).

Application Comparison

Task	Output granularity	Typical head
Text classification	One label per document	`[CLS]` + linear
NER	Label per token	Token-wise linear
QA	Span (start, end)	Two span logits
NLI	Label per pair	`[CLS]` + linear

Common Pitfalls / Exam Traps

- **Trap:** Using BERT for open-ended story generation without a generative decoder — BERT is not autoregressive; use GPT/T5 for free-form generation.
- **Trap:** Confusing QA **extractive** (span from passage) with **generative** QA (model writes answer in own words) — classic BERT QA is extractive.
- **Trap:** Assuming NER tags apply to whole words — BERT uses **subword tokens**; post-processing merges WordPiece fragments.
- **Trap:** Listing only sentiment as a BERT application — NER, QA, and NLI are equally core exam topics.

Quick Revision Summary

- BERT fine-tuning adds a small head on pre-trained encoders for downstream tasks.
- **Text classification:** sentiment, spam, intent, topic — `[CLS]` embedding → label.
- **NER:** token-level BIO tags for persons, orgs, locations, domain entities.
- **QA:** predict answer start/end spans in context passages.
- **NLI:** classify premise-hypothesis as entailment, contradiction, or neutral.
- All applications leverage bidirectional contextual representations from MLM pre-training.
- Production systems (search, support, healthcare, legal) deploy fine-tuned BERT variants via Hugging Face pipelines or custom serving.

← **Previous** · **Index** · **Next** →
BERT Variants: RoBERTa, ALBERT, DistilBERT, and Domain Fine-Tunes

BERT Variants: RoBERTa, ALBERT, DistilBERT, and Domain Fine-Tunes

Why So Many "BERTs"?

The original BERT checkpoint is a starting point, not the final word. Researchers and practitioners have created **architectural variants** (different pre-training recipes), **efficiency variants** (smaller/faster), and **domain fine-tunes** (specialized vocabulary and corpora). Hugging Face lists hundreds of checkpoints when searching "BERT."

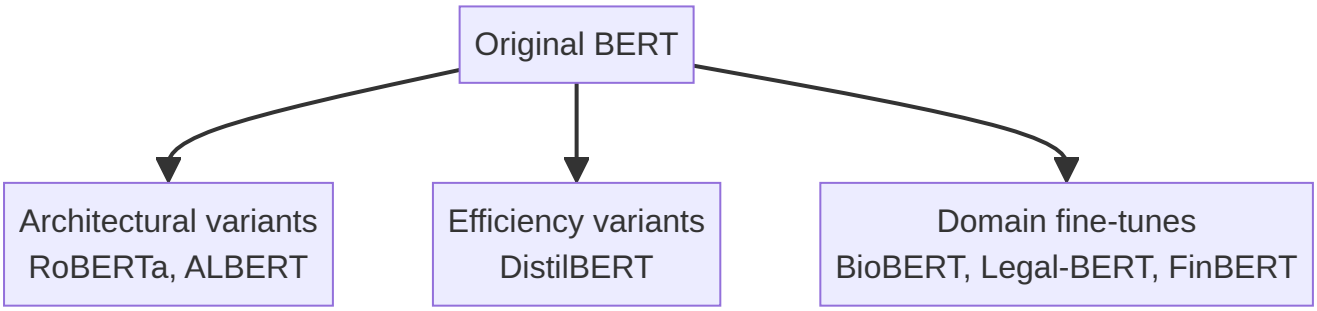

Architectural and Training Variants

RoBERTa (Robustly Optimized BERT)

- Developed by Facebook AI
- Same Transformer encoder architecture as BERT
- Trained **longer** on **more data** (including CC-News, OpenWebText, Stories)
- **Removed Next Sentence Prediction (NSP)** — ablations showed NSP added little
- **Case-sensitive** (cased checkpoints common)
- Typically **outperforms** original BERT on GLUE and related benchmarks

Feature	BERT	RoBERTa
NSP task	Yes	No
Training duration	Shorter	Longer
Data volume	Wikipedia + BooksCorpus	Larger mix
Case handling	uncased & cased variants	primarily cased

ALBERT (A Lite BERT)

- Reduces memory via **parameter sharing** across layers (same weights reused)
- **Case-sensitive**
- Fewer unique parameters → lower GPU memory for same depth
- Trade-off: sometimes slightly lower accuracy vs full BERT-large at equal compute

Use case: Training large-depth models on memory-constrained cloud instances.

DistilBERT

- **Knowledge distillation** from full BERT teacher to smaller student
- Student mimics teacher's **output probability distributions**, not just hard labels
- Results: **~40% smaller**, **~60% faster**, retains **~97%** of BERT-base performance
- Ideal for **mobile**, **edge**, and **real-time** APIs (sentiment on live chat streams)

Domain-Specific Fine-Tunes

General BERT knows broad English; domain models adapt vocabulary and semantics to specialized text.

Model	Domain	Language / notes
FinBERT	Financial literature	Market reports, earnings calls
PubMedBERT	PubMed biomedical abstracts	Clinical NLP, literature mining
BioBERT / SciBERT	Biomedical & scientific text	Entity extraction in papers
Legal-BERT	Legal documents	Contracts, case law
German BERT variants	Medical/legal German	e.g., GermEval-BERT

Real-world use: A hospital NLP pipeline uses PubMedBERT for extracting diagnosis codes from notes — general `bert-base-uncased` underperforms on medical abbreviations and drug names.

Choosing a Variant

Requirement	Recommended direction
Maximum accuracy, GPU budget available	RoBERTa-large or BERT-large
Latency-sensitive / CPU inference	DistilBERT
Memory-constrained training	ALBERT
Financial text	FinBERT
Biomedical NER	PubMedBERT / BioBERT
Legal clause classification	Legal-BERT

Common Pitfalls / Exam Traps

- **Trap:** RoBERTa "changes the Transformer architecture entirely" — it mainly changes **training procedure** (no NSP, more data).
- **Trap:** DistilBERT numbers — remember **40% smaller, 60% faster, ~97% performance** (approximate; common exam fodder).
- **Trap:** Assuming all BERT searches on Hugging Face return interchangeable models — **cased vs uncased** and **domain** matter for tokenization and vocabulary.
- **Trap:** Confusing **fine-tuning** (task labels) with **distillation** (compressing teacher → student) — DistilBERT uses distillation during pre-training of the student.
- **Trap:** ALBERT being "uncased by default" — ALBERT is **case-sensitive**.

Quick Revision Summary

- BERT spawned architectural, efficiency, and domain-specific variants on Hugging Face.
- **RoBERTa:** longer training, more data, **no NSP**, case-sensitive, usually beats BERT.
- **ALBERT:** parameter sharing → lower memory footprint; case-sensitive.
- **DistilBERT:** knowledge distillation → 40% smaller, 60% faster, ~97% of BERT performance.
- Domain models: FinBERT (finance), PubMedBERT (medicine), Legal-BERT (law), language-specific medical BERTs.
- Pick variant by accuracy budget, latency, memory, and domain vocabulary — not all "BERT" checkpoints are equivalent.

← Previous · Index · Next →

Week 10

Week 10

Text Classification: Motivation and Real-World Applications

Text Classification: Motivation and Real-World Applications

The Scale of Unstructured Text

Modern organizations generate enormous volumes of unstructured text every minute:

- Email and messaging
- Social media posts and product reviews
- Support tickets and chat logs
- Video titles, descriptions, and transcripts
- Internal documents and knowledge bases

Manual reading and labeling cannot keep pace. **Text classification** automates the assignment of predefined **categories or labels** to text documents using NLP and machine learning.

Manual vs Automated Approaches

Approach	How it works	Limitation
Manual	Human reads each document and assigns a label	Impossible at scale; slow, expensive, inconsistent
Automated	Model predicts label from text features	Requires training data and model maintenance; scales horizontally on cloud infrastructure

Example — Gmail spam filtering: Every incoming email must be classified as **spam** or **not spam** in milliseconds. A human-in-the-loop for each message is infeasible; a classifier runs on Google's infrastructure for every user simultaneously.

What Text Classification Is

Text classification is the process of mapping an input text document x to a label y from a fixed set:

y \in \{c_1, c_2, \ldots, c_k\}
Variants include **multi-class** (exactly one label) and **multi-label** (zero or more labels per document).

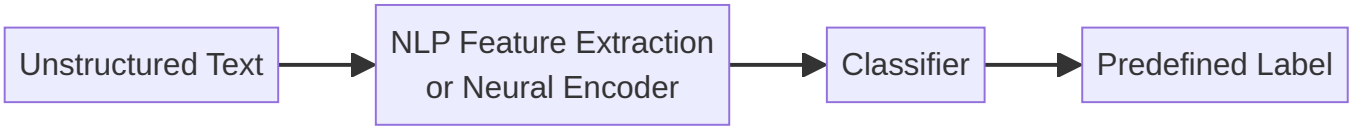

Major Application Types

Spam detection

Binary or multi-class filtering of unwanted email, comments, or messages.

Language identification

Detect document language to trigger translation UI ("Translate this page?") — used by browsers and global content platforms.

Topic labeling (thematic classification)

Assign themes such as **sports**, **politics**, **finance** to articles for news aggregators and recommendation engines.

Intent recognition

Classify **user goal** from utterances in chatbots and voice assistants:

- "Book a flight to Delhi" → `book_flight`
- "Cancel my reservation" → `cancel_booking`

Critical for dialog systems on AWS Lex, Google Dialogflow, and Azure Bot Service.

Sentiment analysis

Classify emotional tone (positive, negative, neutral) in reviews, tweets, and feedback — covered in depth in subsequent notes.

Business Value

Text classification is among the **highest-ROI NLP applications** because it:

- Automates repetitive triage workflows

- Enables real-time routing (support queues, moderation pipelines)
- Powers analytics dashboards (topic trends, sentiment over time)
- Reduces operational cost vs human-only review

Cloud ML context: Models deploy as REST endpoints (SageMaker, Vertex AI) or stream processors (Kafka → classifier → alert) for continuous ingestion.

Common Pitfalls / Exam Traps

- **Trap:** Treating sentiment analysis as separate from text classification — it **is** a type of text classification with emotion labels.
- **Trap:** Assuming text classification only means topic labeling — intent, spam, and language ID are equally valid exam categories.
- **Trap:** Ignoring **class imbalance** (99% not-spam, 1% spam) — accuracy alone is misleading; precision/recall on minority class matters.
- **Trap:** Confusing **topic modeling** (unsupervised discovery) with **topic labeling** (supervised assignment to predefined categories).

Quick Revision Summary

- Unstructured text grows faster than humans can label it — automation is mandatory at scale.
- Text classification assigns predefined categories using NLP + ML.
- Core applications: spam detection, language ID, topic labeling, intent recognition, sentiment analysis.
- Gmail spam filtering is the canonical scale example.
- Intent recognition powers chatbots and virtual assistants in cloud dialog platforms.
- Text classification delivers direct business value through workflow automation and routing.
- Sentiment analysis is a specialized subset of text classification focused on emotional tone.

← **Previous** · **Index** · **Next** →
Sentiment Analysis: Levels, Examples, and Challenges

Sentiment Analysis: Levels, Examples, and Challenges

Definition

Sentiment analysis is a specialized form of text classification whose goal is to determine the **emotional tone** or **polarity** expressed in text. Standard labels:

- **Positive** — approval, satisfaction, enthusiasm
- **Negative** — dissatisfaction, criticism, anger
- **Neutral** — factual or mixed without clear polarity

Example review:

> "I absolutely love the screen resolution, but the battery life is terrible."

This sentence mixes positive and negative clauses — resolving the **overall** label requires aggregation strategy (often negative or neutral depending on method and granularity).

Granularity Levels

1. Document-level

One label for the **entire** document.

- Entire movie review → positive / negative / neutral
- Single tweet → polarity of whole post

Use case: App store rating trend dashboards aggregating thousands of reviews per day.

2. Sentence-level

Each **sentence** receives its own label within a longer document.

- Multi-sentence product review: sentence 1 positive, sentence 2 negative

Use case: Fine-grained monitoring of long-form feedback where different sentences express different opinions.

3. Aspect-Based Sentiment Analysis (ABSA)

The most advanced level: identify **what** the sentiment targets and assign polarity **per aspect**.

Example:

> "The food was great, but the service was slow."

Aspect	Sentiment
food	positive
service	negative

Use case: Restaurant chains and e-commerce platforms track sentiment per feature (battery, camera, delivery speed) rather than one score per review.

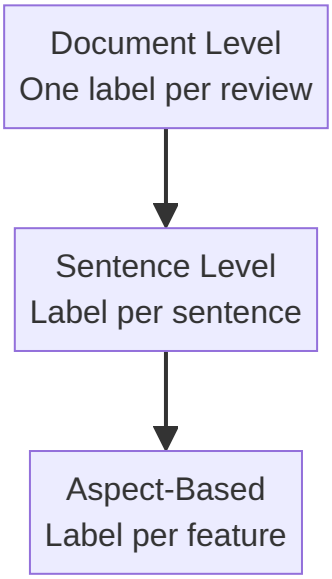

Inherent Challenges

Human language nuance breaks naive lexicon-based systems. Key challenge categories:

Sarcasm

Surface words may be positive while intended meaning is negative.

> "Oh great! My flight is delayed again."

Words like "great" are positive in isolation; context flips polarity to negative.

Negation and double negatives

Negators invert polarity; stacked negatives invert again.

> "The movie was **not** bad." → effectively positive

Both "not" and "bad" carry negative valence individually; together they yield positive sentiment.

Domain-dependent context

The same adjective carries opposite sentiment by domain.

Text	Domain	Sentiment
"The plot was unpredictable."	Thriller movie	Positive (engaging)
"The plot was unpredictable."	Driving manual	Negative (confusing)

Polysemy (same word, different senses)

> "The battery life is **long**." → positive (duration is good) "The wait time was **long**." → negative (delay is bad)

Long is not inherently positive or negative — context determines valence.

Why These Challenges Favor Contextual Models

Rule-based lexicons (VADER) encode heuristics for negation, intensifiers, and contrast conjunctions but struggle with sarcasm and domain polysemy. **Contextual models (BERT, Flair)** encode full-sentence meaning and generally handle complex cases better — at higher compute cost.

Challenge	Lexicon difficulty	Contextual model advantage
Sarcasm	High	Moderate (still not perfect)
Negation	Moderate (rules help)	Strong
Domain context	High	Strong with domain fine-tuning
Polysemy	High	Strong

Common Pitfalls / Exam Traps

- **Trap:** Labeling mixed reviews as purely positive or negative without specifying **granularity level** — document-level vs ABSA yields different correct answers.
- **Trap:** Treating "not bad" as negative because it contains "bad" — double negation flips to positive.
- **Trap:** Assuming "long" is always positive or always negative — **polysemy** makes valence context-dependent.
- **Trap:** Confusing **subjectivity** (opinion vs fact) with **polarity** (positive vs negative) — TextBlob reports both separately.
- **Trap:** Expecting 100% accuracy on sarcastic social media — even BERT fails on subtle irony without domain-specific training.

Quick Revision Summary

- Sentiment analysis classifies emotional tone: positive, negative, or neutral.
- Three granularity levels: document, sentence, aspect-based (ABSA).
- ABSA assigns polarity per feature ("food" positive, "service" negative).
- Sarcasm inverts surface word polarity — lexicons often fail here.
- Negation and double negatives require compositional handling ("not bad" → positive).
- Domain and polysemy make the same word positive or negative depending on context.
- Mixed reviews expose limitations of single document-level labels.
- Contextual Transformer models generally outperform rule-based lexicons on hard cases.

← **Previous** · **Index** · **Next** →
Sentiment Analysis with NLTK VADER: Rule-Based Lexicon Approach

Sentiment Analysis with NLTK VADER: Rule-Based Lexicon Approach

Two Approaches in Contrast

Sentiment analysis can be implemented via:

1. **Traditional rule-based / lexicon methods** — fast, interpretable, lightweight
2. **Modern deep learning (Transformer) methods** — context-aware, higher accuracy, heavier compute

This note covers the first approach using **NLTK's VADER** (Valence Aware Dictionary and sEntiment Reasoner).

VADER Overview

VADER is a **lexicon-based**, **rule-based** model tuned for **social media text**. It ships with:

- A dictionary of words with **valence scores** (positive/negative intensity)
- Hand-crafted rules for capitalization, punctuation, degree modifiers, negation, and contrast conjunctions

Word	Example valence
good	+1.9 (approximate scale in lexicon)
horrible	-2.5

Unlike static embeddings, VADER scores are designed for sentiment — not general semantics.

Setup and Lexicon Download

```
import nltk
from nltk.sentiment import SentimentIntensityAnalyzer

nltk.download('vader_lexicon')
sia = SentimentIntensityAnalyzer()
```

The **VADER lexicon** is a specialized dictionary mapping words (and some phrases) to sentiment intensities.

Scoring Output

For each input sentence, `sia.polarity_scores(text)` returns:

Key	Meaning
`neg`	Proportion of negative tone
`neu`	Proportion of neutral tone
`pos`	Proportion of positive tone
`compound`	Normalized aggregate score in [-1, 1]

Classification thresholds (customizable)

Typical defaults used in coursework:

- `compound > 0.05` → **positive**
- `compound < -0.05` → **negative**
- otherwise → **neutral**

```
scores = sia.polarity_scores(sentence)
compound = scores['compound']
if compound > 0.05:
    label = 'positive'
elif compound < -0.05:
    label = 'negative'
else:
    label = 'neutral'
```

Worked Examples

Sentence	compound (approx.)	Label
"I love this product."	+0.67	positive
"This is the worst experience ever."	strongly negative	negative
"The movie was OK. Nothing special."	moderately negative	negative
"I usually hate waiting, but this was worth it."	~0.0	neutral
"The food was good, but the service was terrible."	negative	negative

Mixed sentence note: VADER applies a **contrast rule** — words after **"but"** receive higher weight, partially explaining why the food/service example trends negative overall.

VADER Rule Features

VADER goes beyond summing word scores:

- **Capitalization:** "GOOD" is more intense than "good"
- **Punctuation:** "good!!!" > "good"
- **Degree modifiers:** "extremely good" > "good"
- **Negation:** "not good" flips valence
- **Conjunctions:** "good but bad" shifts weight toward the clause after **but**

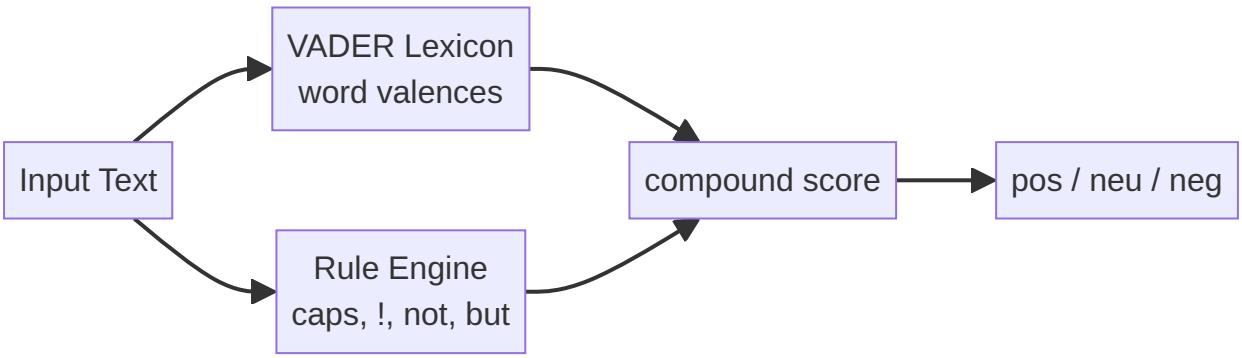

When to Use VADER

Strengths:

- No GPU required — runs on CPU edge devices
- **Interpretable** — inspect lexicon to explain scores
- Fast enough for high-throughput social media streams
- Reasonable on informal short text (tweets, reviews)

Weaknesses:

- Struggles with sarcasm, deep context, domain polysemy
- "I usually hate waiting, but this was worth it." → **neutral** (compound ~0) while humans and BERT often label **positive**

Real-world use: Brand monitoring dashboards polling Twitter firehoses often start with VADER for speed, escalating uncertain cases to Transformer models.

Common Pitfalls / Exam Traps

- **Trap:** Using `pos` alone instead of **`compound`** for final classification — compound is the standard aggregated score.
- **Trap:** Forgetting thresholds are **configurable** — 0.05 / -0.05 are conventions, not universal constants.
- **Trap:** Expecting VADER to match human judgment on **contrast sentences** — "but" weighting can yield neutral when overall intent feels positive.
- **Trap:** Confusing VADER with **TextBlob** — different libraries, different scoring scales.
- **Trap:** Skipping `nltk.download('vader_lexicon')` — analyzer fails without lexicon file.

Quick Revision Summary

- VADER = Valence Aware Dictionary and sEntiment Reasoner (NLTK); rule + lexicon based.
- Returns `neg`, `neu`, `pos`, and **`compound`** — classify using compound thresholds (±0.05).
- Lexicon assigns valence scores; rules handle caps, punctuation, intensifiers, negation, "but".
- Fast, CPU-friendly, interpretable — ideal for social media monitoring at scale.
- Weak on sarcasm and nuanced contrast; mixed sentences may score neutral incorrectly.
- Contrasts with BERT: VADER = traditional; BERT = contextual deep learning (next notes).
- Always download `vader_lexicon` before initializing `SentimentIntensityAnalyzer`.

← **Previous** · **Index** · **Next** →
Sentiment Analysis with BERT via Hugging Face Pipeline

Sentiment Analysis with BERT via Hugging Face Pipeline

Why BERT for Sentiment?

VADER scores individual words and applies rules — it never builds a deep representation of full sentence meaning.

BERT applies **bidirectional self-attention** over the entire sentence. Every token's representation incorporates context from all neighbors, enabling better handling of:

- Sarcasm and irony (partially)
- Complex negation ("not uninteresting")
- Contrast structures ("hate waiting, but worth it")
- Polysemy ("long battery" vs "long wait")

BERT for sentiment uses a **fine-tuned classification head** on top of pre-trained encoders — weights already adapted on millions of labeled reviews available on Hugging Face.

Setup: Hugging Face Token

1. Generate token at `huggingface.co/settings/tokens`
2. Store as `HF_TOKEN` in Colab Secrets
3. Grant notebook access

The token enables downloading fine-tuned sentiment checkpoints from the hub.

Implementation with `pipeline`

```
from transformers import pipeline

sentiment_pipeline = pipeline(
    "sentiment-analysis",
    truncation=True  # required for texts > 512 tokens
)

sentences = [
    "I love this product.",
    "This is the worst experience ever.",
    "The movie was OK. Nothing special.",
    "I usually hate waiting, but this was worth it.",
    "The food was good, but the service was terrible.",
]

for text in sentences:
    result = sentiment_pipeline(text)[0]
    label = result['label']      # POSITIVE or NEGATIVE (model-dependent casing)
    score = result['score']      # confidence in [0, 1]
    print(text, label, score)
```

Key parameter: `truncation=True`

BERT-family models typically accept **512 tokens** maximum. Longer inputs are truncated (usually from the end or with a strategy defined by the tokenizer). Without truncation, pipeline calls may error on long documents.

Expected Results on Standard Test Sentences

Sentence	BERT label	Confidence (typical)
"I love this product."	positive	~0.999+
"This is the worst experience ever."	negative	~0.99+
"The movie was OK. Nothing special."	negative	~0.99 (slightly lower)
"I usually hate waiting, but this was worth it."	positive	~0.999+
"The food was good, but the service was terrible."	negative	high (~0.9+)

Contrast with VADER: The contrast sentence ("hate waiting, but worth it") is labeled **neutral** by VADER (compound ≈ 0) but **positive** by BERT — illustrating contextual understanding.

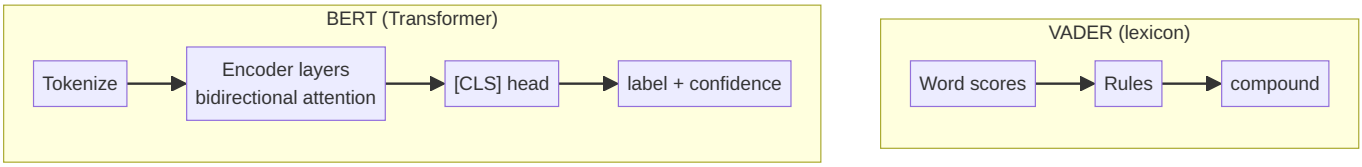

Architecture Recap

1. Tokenizer converts text → subword IDs
2. BERT encoder produces contextual embeddings
3. `[CLS]` token embedding passes through a linear classification layer
4. Softmax yields label probabilities; highest wins

The Hugging Face `sentiment-analysis` pipeline wraps tokenizer, model, and post-processing — no manual forward pass required for inference.

Practical Notes

- Default checkpoint (e.g., `distilbert-base-uncased-finetuned-sst-2-english`) is English review sentiment — other languages need multilingual models.
- Labels may appear as `POSITIVE` / `NEGATIVE` or `LABEL_0` / `LABEL_1` depending on checkpoint — check model card.
- Confidence reflects model probability, not calibrated human uncertainty.

Real-world use: Automated support ticket routing with high accuracy requirements often uses fine-tuned BERT served on GPU endpoints — acceptable latency trade-off for reduced mis-routing cost.

Common Pitfalls / Exam Traps

- **Trap:** Claiming BERT "looks at individual words like VADER" — BERT uses **full-sentence contextual** representations.
- **Trap:** Omitting `truncation=True` on long texts — silent truncation failures or errors on >512 tokens.
- **Trap:** Assuming three-way positive/neutral/negative — many HF sentiment models are **binary** (pos/neg only); neutral must come from low confidence or a 3-class fine-tune.
- **Trap:** Hardcoding API keys in notebook cells instead of Colab Secrets.
- **Trap:** Expecting identical labels across all BERT checkpoints — SST-2 fine-tunes differ from 3-star review models.

Quick Revision Summary

- BERT sentiment uses bidirectional context — superior to lexicon methods on negation and contrast.
- Hugging Face `pipeline("sentiment-analysis")` downloads a pre-fine-tuned review classifier.
- Set `truncation=True` for inputs longer than 512 subword tokens.
- Output: `label` + `score` (confidence); typically binary positive/negative.
- BERT correctly labels "hate waiting, but worth it" as positive; VADER often neutral.
- Requires HF token for gated models; heavier GPU compute than VADER.
- Best for high-accuracy production routing (support tickets, compliance triage).

← **Previous** · **Index** · **Next** →
Comparing VADER and BERT for Sentiment Analysis

Comparing VADER and BERT for Sentiment Analysis

Two Paradigms

After implementing sentiment analysis with **NLTK VADER** (rule-based lexicon) and **BERT** (fine-tuned Transformer), the natural question is: **when should each be used?**

They differ across speed, hardware, interpretability, accuracy, and deployment context.

Side-by-Side Comparison

Dimension	VADER (NLTK)	BERT (Hugging Face)
Approach	Lexicon + hand rules	Deep learning, contextual embeddings
Context	Word-level + rule composition	Full-sentence bidirectional attention
Speed	Very fast	Slower
Hardware	CPU sufficient	GPU recommended for batch inference
Interpretability	High — inspect lexicon & rules	Low — "black box" neural weights
Sarcasm / complex negation	Weak	Stronger (not perfect)
Mixed/contrast sentences	Often neutral or wrong	Generally better ("but" clauses)
Setup complexity	`pip install nltk` + lexicon download	HF token, large model download
Training required	None (pre-built)	Uses pre-fine-tuned checkpoint

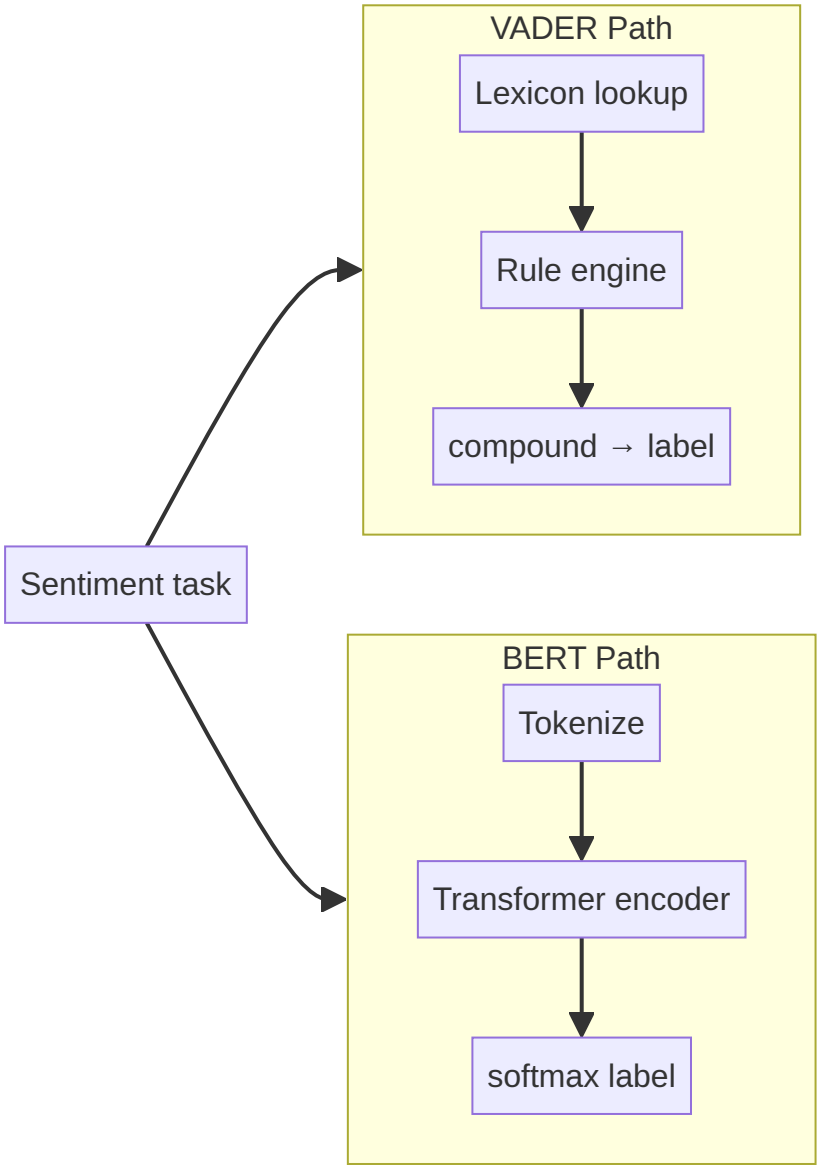

Strengths and Best Use Cases

VADER — choose when:

- **Social media monitoring** at high throughput (Twitter/X firehose, Reddit streams)
- **Edge / mobile** deployment with no GPU
- **Explainability** is required (regulatory, audit trails — "word X scored -2.5")
- **Prototyping** before investing in GPU infrastructure
- Latency budget is **milliseconds** on cheap CPU instances

Example: Real-time brand mention dashboard classifying 10k tweets/minute on a single cloud VM without GPU.

BERT — choose when:

- **High accuracy** is critical — misclassification has real cost
- **Automated support ticket routing** (wrong queue = delayed resolution)
- **Complex language** — negation, contrast, informal mixed sentiment
- GPU budget available (cloud inference endpoints, batch nightly processing)
- Willing to sacrifice interpretability for F1 score

Example: Enterprise helpdesk auto-assigning P1 vs P3 priority from customer email bodies.

Observed Behavioral Differences (Same Test Set)

Sentence	VADER	BERT
"I love this product."	positive	positive
"This is the worst experience ever."	negative	negative
"The movie was OK. Nothing special."	negative	negative
"I usually hate waiting, but this was worth it."	**neutral**	**positive**
"The food was good, but the service was terrible."	negative	negative

The contrast sentence exposes the core trade-off: VADER's **but** rule and compound aggregation miss overall positive intent; BERT's contextual encoding captures that the second clause dominates speaker intent.

Hybrid Production Pattern

Many systems combine both:

1. Run VADER on all incoming text (fast filter)
2. Escalate low-confidence or high-stakes items to BERT
3. Human review for remaining edge cases

This optimizes **cost vs accuracy** on cloud billing — GPU time only where needed.

Common Pitfalls / Exam Traps

- **Trap:** "BERT is always better" — VADER wins on **speed, cost, interpretability** for social media at scale.
- **Trap:** "VADER is not ML" — it is ML-adjacent rule-based NLP; exams contrast it with **deep learning**, not with "non-computational."
- **Trap:** Ignoring **black box** as a deployment constraint — regulated industries may require lexicon explainability despite lower accuracy.
- **Trap:** Using BERT on a CPU-only Lambda with strict latency SLA — may timeout; VADER or DistilBERT better fit.
- **Trap:** Forgetting VADER is **tuned for social media** — formal legal prose may score poorly.

Quick Revision Summary

- VADER: fast, CPU, interpretable lexicon+rules — best for social media monitoring.

- BERT: slower, GPU-heavy, black box — best for high-accuracy ticket routing and complex sentences.
- BERT handles contrast ("hate X but worth it") better; VADER often returns neutral.
- VADER explainability = inspect word valences; BERT = no simple dictionary explanation.
- Choose by accuracy requirements, hardware budget, latency, and explainability needs.
- Hybrid pipelines: VADER first pass, BERT for uncertain/high-value cases.
- Neither solves sarcasm perfectly; BERT is materially stronger on negation and context.

← **Previous** · **Index** · **Next** →
Sentiment Analysis with Flair

Sentiment Analysis with Flair

What Flair Is

Flair is a PyTorch-based NLP library supporting sequence labeling tasks including **POS tagging**, **NER**, and **sentiment analysis**. Like BERT-based pipelines, Flair uses **deep learning** and contextual embeddings — not hand-written lexicon rules.

Flair models often download weights from Hugging Face (may require `HF_TOKEN` for PyTorch checkpoints).

Installation and Imports

```
# Install: pip install flair

from flair.data import Sentence
from flair.models import TextClassifier
```

Package size is **large** — first install and model download take noticeable time compared to NLTK VADER.

Single-Sentence Workflow

```
text = "I love Delhi"
sentence = Sentence(text)

classifier = TextClassifier.load('sentiment')
classifier.predict(sentence)

print(sentence)  # shows labels and confidence
# Example: Sentence[...] - positive (0.9912)
```

Steps:

1. Wrap raw string in `Sentence` object
2. Load pre-trained `TextClassifier` with `'sentiment'` tag
3. Call `classifier.predict(sentence)`
4. Read predicted label and probability from sentence annotations

Batch Processing on Test Sentences

```
sentences = [
    "I love this product.",
    "This is the worst experience ever.",
    "The movie was OK. Nothing special.",
    "I usually hate waiting, but this was worth it.",
    "The food was good, but the service was terrible.",
]

classifier = TextClassifier.load('sentiment')

for text in sentences:
    sent = Sentence(text)
    classifier.predict(sent)
    print(text, sent.labels)
```

Results on Standard Benchmark Sentences

Sentence	Flair prediction	Confidence
"I love this product."	positive	~99%
"This is the worst experience ever."	negative	~100%
"The movie was OK. Nothing special."	negative	~99.9%
"I usually hate waiting, but this was worth it."	positive	~99%
"The food was good, but the service was terrible."	negative	~99%

Flair aligns with **BERT** (not VADER) on the contrast sentence — labeling "worth it" dominance as **positive**.

Flair vs Other Libraries

Library	Backend	Speed	Contextual?
VADER	Lexicon + rules	Fastest	No
TextBlob	Pattern + TextBlob lexicon	Fast	Limited
Flair	PyTorch neural tagger	Slow (large download)	Yes
BERT pipeline	Transformers	Moderate–slow	Yes

When to Use Flair

Strengths:

- Unified API for NER, POS, and sentiment (same library as earlier course modules)
- Strong accuracy on contextual sentences
- Pre-trained models readily available

Weaknesses:

- Heavy dependencies (PyTorch + model weights)
- Slow cold start (download + load)
- Less ecosystem breadth than Hugging Face `transformers` for arbitrary checkpoints

Real-world use: Research prototypes and coursework pipelines where Flair is already used for NER/POS and sentiment is added with one consistent API — less common in large-scale production than raw Hugging Face serving.

Common Pitfalls / Exam Traps

- **Trap:** Forgetting to wrap strings in `Sentence` before `predict()` — plain strings won't work.
- **Trap:** Expecting instant startup — first `load('sentiment')` downloads weights; plan for cache or pre-baked Docker images in production.
- **Trap:** Assuming Flair = rule-based because it's "another alternative to BERT" — Flair sentiment is **neural**, like BERT.
- **Trap:** Not granting HF token access when Flair pulls gated PyTorch checkpoints from the hub.
- **Trap:** Confusing `TextClassifier` with sequence labeling taggers used for NER — different loaded model names.

Quick Revision Summary

- Flair: PyTorch NLP library with pre-trained sentiment `TextClassifier`.
- Pipeline: `Sentence(text)` → `TextClassifier.load('sentiment')` → `predict()` → label + confidence.
- Large package; slow download but strong contextual accuracy.
- Matches BERT on contrast sentences (positive for "hate waiting, but worth it").
- Same ecosystem as Flair NER/POS from earlier modules — consistent API.
- Requires HF token when downloading certain hub-hosted weights.
- Trade-off: heavier than VADER/TextBlob; more unified than ad-hoc multi-library stacks.

← **Previous** · **Index** · **Next** →
Sentiment Analysis with spaCy and TextBlob

Sentiment Analysis with spaCy and TextBlob

Overview

spaCy is an industrial-strength NLP library focused on efficient pipelines (tokenization, tagging, parsing, NER). Sentiment analysis is **not native** to core spaCy — it is added via the **`spacytextblob`** extension, which integrates **TextBlob** polarity and subjectivity into spaCy `Doc` objects.

This provides a **third lightweight alternative** alongside VADER (NLTK) and Transformer models (BERT, Flair).

Installation

```
# Quiet install of English model + extension
# pip install spacy spacytextblob
# python -m spacy download en_core_web_sm

import spacy
from spacytextblob.spacytextblob import SpacyTextBlob

nlp = spacy.load('en_core_web_sm')
nlp.add_pipe('spacytextblob')
```

Requires:

- `en_core_web_sm` (or larger English pipeline)
- `spacytextblob` pipe registered on the `nlp` object

Polarity and Subjectivity

TextBlob exposes **two distinct scores** per document:

Polarity

Emotional tone from negative to positive.

- Range: **-1 to +1**
- -1 = very negative
- +1 = very positive
- 0 = neutral

Subjectivity

Opinion vs factual content.

- Range: **0 to 1**
- 0 = objective / factual
- 1 = highly subjective / opinion-based

Score	Measures	Range
Polarity	Sentiment direction	[-1, 1]
Subjectivity	Opinion vs fact	[0, 1]

Important: Subjectivity is **independent** of polarity — a factual negative news headline can have negative polarity but low subjectivity.

Single-Sentence Example

```
text = "I am feeling very good."
doc = nlp(text)

print(doc._.polarity)       # e.g., 0.90 (positive)
print(doc._.subjectivity)   # e.g., 0.78 (subjective opinion)
```

Access via `doc._.` extension attributes after processing through the pipeline.

Batch Implementation on Test Sentences

```
sentences = [
    "I love this product.",
    "This is the worst experience ever.",
    "The movie was OK. Nothing special.",
    "I usually hate waiting, but this was worth it.",
    "The food was good, but the service was terrible.",
]

for text in sentences:
    doc = nlp(text)
    print(f"{text}")
    print(f"  Polarity: {doc._.polarity:.2f}")
    print(f"  Subjectivity: {doc._.subjectivity:.2f}")
```

Observed Results vs Other Libraries

Sentence	spaCy/TextBlob polarity	Other libraries
"I love this product."	~0.62 (moderate positive)	VADER/BERT/Flair: strong positive
"This is the worst experience ever."	~-1.0 (very negative)	All agree: negative
"The movie was OK. Nothing special."	~-0.42 (neutral–negative)	BERT/Flair: negative
"I usually hate waiting, but this was worth it."	~**-0.21** (negative)	BERT/Flair: **positive**; VADER: neutral
"The food was good, but the service was terrible."	~-0.15 (negative)	All tend negative

Key divergence: TextBlob/spaCy misclassifies the contrast sentence as **negative** while BERT and Flair label it **positive** — demonstrating that lightweight Pattern-based TextBlob lacks sophisticated contrast handling.

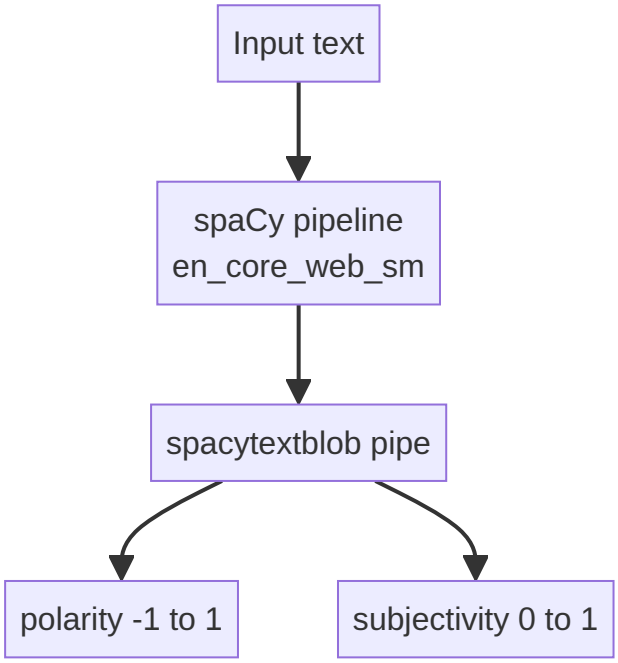

When to Use spaCy + TextBlob

Strengths:

- Integrates sentiment into existing **spaCy pipelines** (same `Doc` as entities, POS tags)
- Returns **subjectivity** — useful for filtering opinion reviews vs factual reports
- Lighter than BERT/Flair

Weaknesses:

- Polarity often **underconfident** on clearly positive short praise
- Poor on **contrast clauses** compared to contextual models
- TextBlob backend is older Pattern-based logic — not Transformer-grade

Real-world use: Content moderation pipeline using spaCy NER + TextBlob polarity to flag subjective negative reviews mentioning product defects — subjectivity score helps separate opinion from spec-sheet comparisons.

Comparing All Four Implementations

Sentence	VADER	BERT/Flair	spaCy/TextBlob
"I love this product."	positive	positive	weak positive (0.62)
"hate waiting, but worth it"	neutral	positive	**negative**
Subjectivity score	No	No	**Yes**

Common Pitfalls / Exam Traps

- **Trap:** Treating polarity and subjectivity as the same — **subjectivity ≠ sentiment**; factual text can be objective with zero polarity.
- **Trap:** Forgetting `nlp.add_pipe('spacytextblob')` — without it, `doc._.polarity` doesn't exist.
- **Trap:** Expecting spaCy core to include sentiment — must install **spacytextblob** extension explicitly.
- **Trap:** Assuming all libraries agree on mixed sentences — TextBlob often disagrees with BERT on "but" constructions.
- **Trap:** Polarity ranges: TextBlob/spaCy uses **-1 to 1**; VADER compound also [-1,1] but VADER also outputs separate pos/neu/neg ratios — scales differ conceptually.

Quick Revision Summary

- spaCy sentiment via `spacytextblob` extension integrating TextBlob into `Doc` objects.
- Install `en_core_web_sm` + register `spacytextblob` pipe on `nlp`.
- **Polarity** ([-1,1]): negative ↔ positive emotional tone.
- **Subjectivity** ([0,1]): factual ↔ opinion-based — unique among methods covered.
- Access scores: `doc._.polarity`, `doc._.subjectivity`.
- TextBlob backend: fast, integrated with spaCy NER pipeline, but weaker on contrast sentences.
- "hate waiting, but worth it" → TextBlob negative; BERT/Flair positive; VADER neutral — classic exam comparison.
- Compare all libraries on same test set when choosing a production approach.

← **Previous** · **Index** · **Next** →

Week 11

Week 11

Introduction to Topic Modelling

What Is Topic Modelling?

Topic modelling is an **unsupervised** machine learning technique that discovers hidden themes in a collection of documents without requiring human labels. Instead of reading thousands of texts manually, the algorithm finds groups of words that tend to co-occur and treats each group as a latent "topic."

A topic is not a single word — it is a **probability distribution over words**. Words like *software*, *data*, *system*, *network*, and *algorithm* might cluster into a technology topic, each with a different weight within that theme.

Intuition: One Document, Multiple Themes

Consider a travel review that mentions monuments, hiking, restaurants, transport, and words like *great* and *bad*. A single document can simultaneously express:

Topic	Representative Words
Tourism	monuments, adventure, hiking, scenic
Facilities	restaurants, transport, hotels
Feedback	great, good, bad, awesome

Topic modelling recovers these overlapping themes automatically.

Why Topic Modelling Matters

Modern systems generate text continuously — news feeds, research papers, customer reviews, social media posts, support tickets. Manual categorisation does not scale.

Topic modelling helps organisations:

- **Summarise** large document collections at a glance
- **Identify** recurring themes without predefined labels
- **Group** similar documents by latent theme
- **Support** downstream tasks such as search, recommendation, and trend monitoring

The central question it answers: **What are the main themes present in this corpus?**

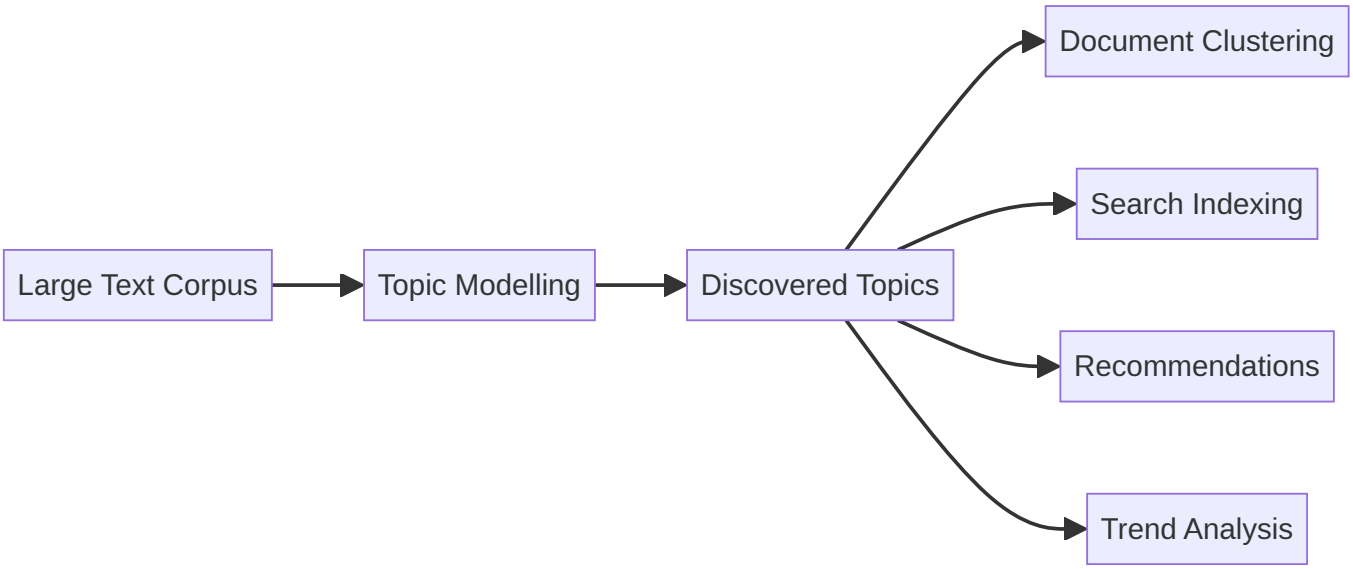

Topic Modelling vs Text Classification

Both work with text, but they solve different problems.

Aspect	Text Classification	Topic Modelling
Learning type	Supervised	Unsupervised
Labels required	Yes — predefined classes	No
Output	Assigns documents to **known** categories	Discovers **unknown** structure
Example	Spam vs non-spam email	Themes in news articles

Classification answers: "Which of these predefined buckets does this document belong to?"

Topic modelling answers: "What hidden themes exist, and how is each document composed of them?"

Common Use Cases

- **Document clustering** — group similar documents by recurrent themes (e.g., clustering AWS support tickets by issue type)
- **Exploratory data analysis** — quickly understand an unfamiliar corpus before building a classifier
- **Trend analysis** — track how topic prevalence shifts over time (e.g., COVID-related themes in news from 2020–2022)
- **Information retrieval** — improve search with topic-based indexing instead of keyword-only matching
- **Content recommendation** — suggest articles sharing latent topics with what a user has already read

---

Common Pitfalls / Exam Traps

- **Confusing topic modelling with classification** — topic modelling does not need labels and does not use predefined classes; exam questions often swap these properties.
- **Assuming one topic per document** — many algorithms (including LDA) model documents as *mixtures* of topics, not single assignments.
- **Expecting semantic understanding from all topic models** — classical methods like LDA use bag-of-words and ignore word order; they find co-occurrence patterns, not deep meaning.
- **Skipping preprocessing** — stop-word removal, tokenisation, and normalisation strongly affect topic quality for traditional models.

Quick Revision Summary

- Topic modelling is unsupervised discovery of latent themes in text corpora.
- A topic is a distribution over words, not a single keyword.
- Documents can express multiple topics in different proportions.
- Unlike classification, no labelled data or predefined categories are required.
- Industrial use cases include clustering, EDA, trend tracking, search, and recommendations.
- LDA is the most widely used classical algorithm; embedding-based alternatives address its limitations.

← **Previous** · **Index** · **Next** →
Latent Dirichlet Allocation (LDA)

Latent Dirichlet Allocation (LDA)

Why LDA Dominates Classical Topic Modelling

Latent Dirichlet Allocation (LDA) is the most widely used algorithm for topic modelling because it is:

- **Probabilistic** — outputs interpretable probability distributions, not hard labels
- **Interpretable** — topics are described by their top-weighted words
- **Scalable** — handles large document collections efficiently

LDA answers two questions about a corpus:

1. What topics exist?
2. How is each document composed of those topics?

Two Core Ideas

LDA rests on two simple but powerful assumptions:

1. Documents Are Mixtures of Topics

A single document rarely discusses only one subject. A tech-health article might be 70% technology and 30% healthcare. LDA assigns **topic proportions** (not a single label) to each document.

2. Topics Are Distributions Over Words

Each topic is defined by a set of words with associated probabilities. The technology topic might weight *software*, *system*, *data*, and *algorithm* highly; the healthcare topic might weight *patient*, *treatment*, and *hospital*.

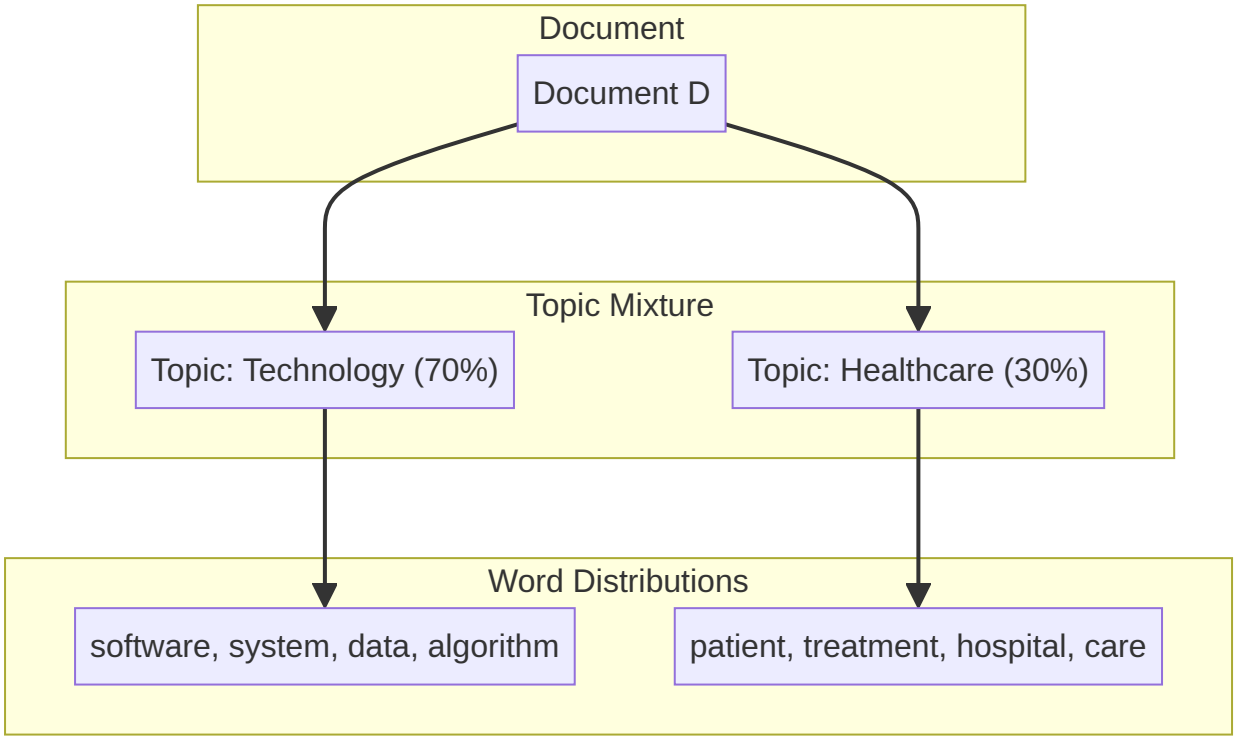

Worked Example

Entity	Composition
Document	70% technology + 30% healthcare
Technology topic	software, system, data, algorithm, ...
Healthcare topic	patient, treatment, hospital, care, ...

Critical distinction: LDA does **not** assign one topic per document. It outputs a vector of proportions such as (0.70, 0.30) for a two-topic model.

How to Read LDA Output

For a document with topic proportions (0.07, 0.86, 0.06) across three topics:

- The document is **primarily** about topic 1 (86% weight)
- Topics 0 and 2 contribute minimally
- Topic labels (e.g., "finance", "sports") must be inferred by inspecting the top words per topic — LDA does not name topics automatically

---

Common Pitfalls / Exam Traps

- **"LDA assigns one topic per document"** — false; it assigns proportions. The highest-proportion topic is often used as a proxy label, but that is a post-hoc simplification.
- **Confusing LDA with LSA** — both are topic-modelling families, but LDA is generative and probabilistic; LSA uses matrix factorisation (SVD).
- **Expecting topic count to be automatic** — the number of topics K must be specified manually and tuned.
- **Ignoring that topics are unlabelled** — exam answers listing "topic 0, topic 1" without interpreting top words miss the interpretability requirement.

Quick Revision Summary

- LDA is probabilistic, interpretable, and scales to large corpora.
- Documents = mixtures of topics; topics = distributions over words.
- Output is topic **proportions** per document, not a single hard assignment.
- Topic names are inferred manually from top-weighted words per topic.
- K (number of topics) is a hyperparameter set by the practitioner.
- LDA is the baseline classical algorithm before embedding-based alternatives like BERTopic.

← **Previous** · **Index** · **Next** →
Key Assumptions of LDA

Key Assumptions of LDA

Why Assumptions Matter

Every machine learning model simplifies reality to make inference tractable. LDA's assumptions explain both its strengths and its well-known failure modes. Understanding them is essential for preprocessing decisions and for knowing when to switch to modern alternatives.

The Four Core Assumptions

1. Word Order Does Not Matter (Bag of Words)

LDA treats a document as an unordered collection of words — identical to the bag-of-words (BoW) representation. *"The dog bit the man"* and *"The man bit the dog"* produce the same input.

Consequence: No syntactic or sequential context is preserved.

2. Topics Are Shared Across Documents

The same set of latent topics appears throughout the entire corpus. A "finance" topic discovered in one document is the same "finance" topic used to explain another.

Consequence: LDA builds a global topic vocabulary; it does not create document-specific topic sets.

3. Documents Are Mixtures of Topics

Each document can contain multiple topics in varying proportions. A news article might be 60% politics and 40% economics.

Consequence: LDA is well-suited to long, multi-theme documents but poorly suited to very short texts that express a single idea.

4. Words Are Generated Independently Given a Topic

Once a topic is chosen, each word is drawn independently from that topic's word distribution. Co-occurrence is modelled at the topic level, not at the word-sequence level.

Consequence: Phrases and collocations (*"New York"*, *"machine learning"*) are not natively captured.

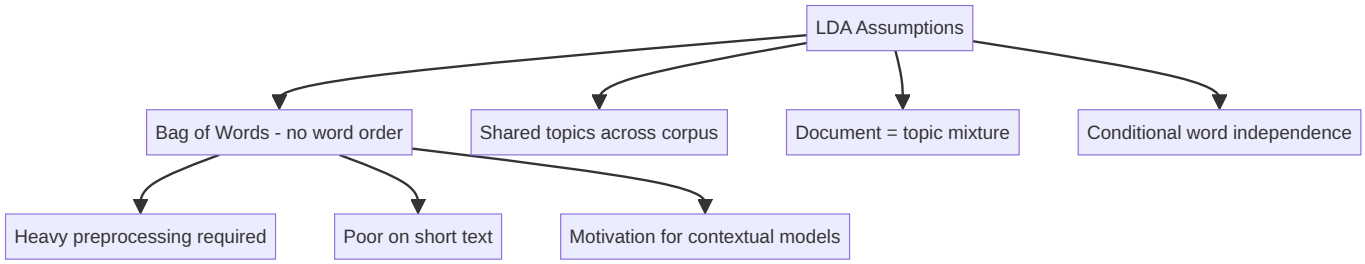

Practical Implications

Assumption	Practical Impact
Bag of words	Tokenisation, stop-word removal, and normalisation are critical
Shared topics	Topic count K must be chosen carefully for the whole corpus
Topic mixtures	Works on articles and reports; struggles on tweets and captions
Word independence	Synonyms and paraphrases are not recognised unless they co-occur

Why Preprocessing Matters

Because LDA ignores context, the signal comes entirely from **which words appear and how often**. Noisy tokens (stop words, punctuation, casing variants) dilute topic structure. A clean BoW representation directly improves topic coherence.

Why LDA Struggles with Short Text

A tweet with 15 words provides too few co-occurrence statistics for LDA to infer a meaningful topic mixture. Models like GSDMM (which assume one topic per short document) address this gap.

Why Context-Aware Models Emerged

LDA's inability to capture semantics, word order, and synonymy motivated embedding-based approaches (Word2Vec, BERT) and modern topic models like BERTopic that cluster documents in semantic embedding space.

Common Pitfalls / Exam Traps

- **Claiming LDA preserves context** — it explicitly does not; word order is discarded.
- **Using LDA on tweets without preprocessing or alternative models** — short text violates the mixture assumption in practice and yields incoherent topics.
- **Skipping preprocessing because "LDA is unsupervised"** — unsupervised does not mean preprocessing-free; BoW quality directly determines output quality.
- **Confusing "topics are shared" with "one topic per document"** — topics are global, but each document draws from multiple of them.

Quick Revision Summary

- LDA assumes bag of words: no word order or syntactic context.
- Topics are shared globally across all documents in the corpus.
- Each document is a mixture of topics, not a single-topic assignment.
- Words are conditionally independent given the assigned topic.

- These assumptions make preprocessing essential and explain LDA's weakness on short text.
- Context-aware and embedding-based models were developed to overcome LDA's semantic blindness.

← **Previous** · **Index** · **Next** →
Implementing Topic Modelling with LDA

Implementing Topic Modelling with LDA

Overview

This lab implements classical topic modelling using **Gensim's LDA** on a small custom corpus. The pipeline covers preprocessing, bag-of-words vectorisation, model training, and topic interpretation — the standard workflow used in production analytics pipelines before deploying transformer-based alternatives.

Pipeline Architecture

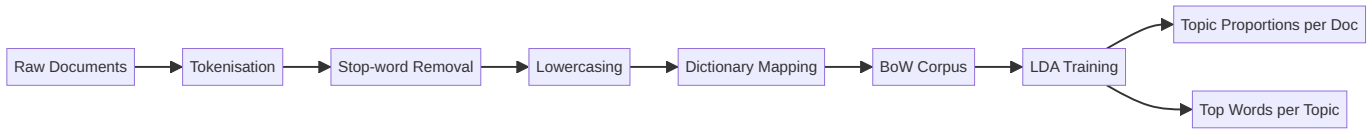

Step 1: Environment Setup

Required libraries:

Library	Role
Gensim	`Dictionary`, `LdaModel` for BoW and LDA
NLTK	Stop words, tokenisation

```
from gensim import corpora
from gensim.models import LdaModel
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
```

Step 2: Corpus Preparation

A sample corpus might contain sentences about stock markets, machine learning, and football:

- *"Stock markets reacted to inflation news."* → themes: finance, economics
- *"Deep learning models require large datasets."* → themes: ML, data
- *"The football team is preparing for the World Cup."* → themes: sports

Before running the model, manually hypothesise 2–3 themes per sentence — then compare against LDA output to build intuition.

Step 3: Text Preprocessing

Standard steps applied to each document:

1. **Tokenise** into words
2. **Remove stop words** (the, is, to, ...)
3. **Lowercase** all tokens (normalisation)

Output: a list of token lists, one per document.

Step 4: Bag-of-Words Vectorisation

```
dictionary = corpora.Dictionary(processed_docs)
corpus = [dictionary.doc2bow(doc) for doc in processed_docs]
```

- `Dictionary` maps each unique word to an integer ID
- `doc2bow` converts each document into a sparse vector of `(word_id, count)` pairs

This is the input representation LDA expects.

Step 5: Train the LDA Model

```
lda_model = LdaModel(
    corpus=corpus,
    id2word=dictionary,
    num_topics=3,  # K — must be set manually
    random_state=42
)
```

`num_topics=3` is a hyperparameter. There is no automatic way to choose K; practitioners experiment or use coherence metrics.

Step 6: Interpret Document-Topic Assignments

```
for doc_bow in corpus:
    print(lda_model.get_document_topics(doc_bow))
```

Example output for document 0: `[(0, 0.07), (1, 0.86), (2, 0.06)]`

Topic Index	Confidence	Interpretation
0	7%	Minor contribution
1	86%	**Dominant topic**
2	6%	Minor contribution

The dominant topic (highest probability) is often used as a proxy label, but the full distribution is the true output.

Step 7: Inspect and Label Topics

```
lda_model.print_topics(num_words=5)
# or
lda_model.show_topics(num_words=10)
```

Example discovered topics from a mixed corpus:

Topic	Top Words	Inferred Label
0	learning, datasets, model, deep, large	Machine Learning
1	markets, reacted, inflation, news, stock	Finance / Economics
2	football, team, World, Cup, preparing	Sports

Topic labels are manual — inspect top words and assign meaningful names. This is standard practice in industry dashboards.

Self-Assessment Questions

After running the model, evaluate quality:

- Do topics look **clean** and semantically distinct?
- Are the same words **repeated across multiple topics** (topic overlap)?
- Do the inferred labels match your manual hypotheses from Step 2?
- Does changing `num_topics` improve or degrade coherence?

Extension Exercise

Replace the sample corpus with your own dataset (e.g., product reviews, news headlines) and rerun the full pipeline. Compare topic quality across domains.

Common Pitfalls / Exam Traps

- **Forgetting to set `num_topics`** — LDA requires K upfront; there is no default optimal value.
- **Treating confidence scores as classification probabilities** — they are topic proportions within a generative model, not calibrated class probabilities.
- **Expecting automatic topic names** — `print_topics` shows words; human interpretation is required.
- **Skipping stop-word removal** — function words dominate BoW counts and produce meaningless topics.
- **Using raw text without lowercasing** — *"Stock"* and *"stock"* become separate dictionary entries.

Quick Revision Summary

- LDA pipeline: preprocess → dictionary → BoW corpus → train → interpret.
- Gensim `LdaModel` takes `corpus`, `id2word`, and `num_topics`.
- `get_document_topics()` returns proportional weights, not hard labels.
- `print_topics()` / `show_topics()` reveal top words for manual labelling.
- Preprocessing (tokenise, stop words, lowercase) is mandatory for clean topics.
- Always validate topic quality and experiment with K on your own corpus.

← Previous · Index · Next →

Why Alternatives to LDA Are Needed

Why Alternatives to LDA Are Needed

LDA's Strengths and Ceiling

LDA remains widely deployed in industry for exploratory analysis on large document collections. Its probabilistic framework, interpretability, and scalability made it the default topic model for over a decade. However, as NLP advanced and text sources diversified, LDA's simplifying assumptions became limiting rather than enabling.

Core Limitations of LDA

Limitation	What It Means	Real-World Impact
Bag of words	No context retained	*"bank"* (river) and *"bank"* (finance) treated identically unless co-occurrence differs
No word order	Sentence structure ignored	Negation and phrasing lost (*"not good"* ≈ *"good"*)
No semantic similarity	Synonyms not linked unless co-occurring	*"automobile"* and *"car"* may split across topics
Manual topic count	K must be tuned by hand	Wrong K produces merged or fragmented topics
Short text failure	Too few words for mixture inference	Tweets, captions, product reviews yield incoherent topics
General-purpose design	Not optimised for domain-specific language	Legal, medical, or social-media corpora need specialised handling

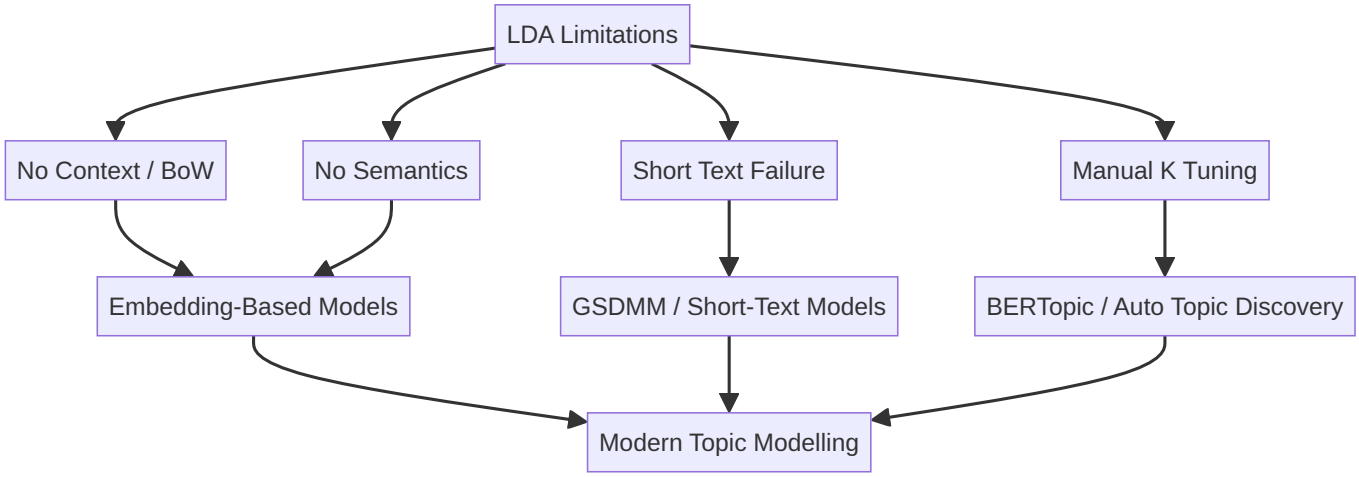

Why Embedding-Based Methods Emerged

As NLP evolved from statistical co-occurrence to distributed representations, a new class of topic models emerged:

- **Sentence/document embeddings** (from BERT and similar transformers) capture semantic meaning
- **Clustering in embedding space** groups semantically similar documents without bag-of-words assumptions
- **Classical TF-IDF on clusters** still provides interpretable topic words

This hybrid approach — deep embeddings for grouping, classical NLP for labelling — powers models like **BERTopic**.

Choosing the Right Tool

Scenario	Recommended Approach
Long articles, reports, research papers	LDA (with careful preprocessing)
Short social media posts, reviews	GSDMM or BERTopic
Need semantic grouping of customer feedback	BERTopic
Large corpus, need fast baseline	LDA
Domain with rich contextual language	Embedding-based models

Common Pitfalls / Exam Traps

- **Assuming LDA is always the best topic model** — it is a strong baseline for long text, not a universal solution.
- **Applying LDA to tweets without acknowledging limitations** — exam scenarios involving short text should trigger GSDMM or BERTopic as answers.
- **Confusing "unsupervised" with "no preprocessing needed"** — LDA's limitations make preprocessing even more critical.
- **Listing only one alternative** — know at least GSDMM (short text) and BERTopic (semantic embeddings).

Quick Revision Summary

- LDA ignores context, word order, and semantic similarity (bag-of-words assumption).
- Topic count K must be manually tuned — no automatic selection.
- LDA performs poorly on short text (< 50 words): tweets, captions, reviews.
- Embedding-based methods address LDA's semantic blindness.
- GSDMM targets short-text, single-topic documents; BERTopic uses transformer embeddings.
- Model choice depends on text length, domain, and whether semantic grouping is required.

← **Previous** · **Index** · **Next** →
BERTopic: Semantic Topic Modelling

BERTopic: Semantic Topic Modelling

The Problem BERTopic Solves

Classical LDA models word co-occurrence in a bag-of-words space. It cannot recognise that *"cloud computing"* and *"AWS infrastructure"* belong to the same theme unless those exact words co-occur frequently. BERTopic addresses this by operating in **semantic embedding space** rather than raw word-count space.

What Is BERTopic?

BERTopic is a modern topic modelling technique that:

1. Uses a **transformer model** (typically BERT-based) to create dense sentence/document embeddings
2. **Clusters** documents by semantic similarity in embedding space
3. **Extracts** interpretable topic words from each cluster using classical NLP techniques

The name reflects its foundation: **BERT** embeddings + **topic** discovery.

Key Idea

Instead of modelling word distributions directly (as LDA does), BERTopic:

- Groups **semantically similar documents** together
- Derives topic labels from the words most characteristic of each cluster

How BERTopic Differs from LDA

Aspect	LDA	BERTopic
Representation	Bag of words	Dense embeddings
Semantics	Co-occurrence only	Contextual meaning
Word order	Ignored	Captured in embeddings
Short text	Poor	Better (embedding quality helps)
Topic words	Global word distributions	Per-cluster TF-IDF
Topic count	Manual K	Clustering algorithm determines groups

Why It Works in Practice

Consider a corpus of ML platform reviews mentioning *"SageMaker"*, *"model deployment"*, *"inference latency"*, and *"endpoint scaling"*. LDA might split these across topics based on word frequency alone. BERTopic groups the reviews semantically because the embeddings capture that these phrases describe the same operational concern — model serving infrastructure.

When to Use BERTopic

- Customer feedback analysis where paraphrasing is common
- Research paper clustering with varied terminology
- Any corpus where **meaning** matters more than **exact word overlap**
- Situations where LDA produces overlapping or incoherent topics

---

Common Pitfalls / Exam Traps

- **Saying BERTopic is "just LDA with BERT"** — the architecture is fundamentally different: embed → reduce → cluster → c-TF-IDF.
- **Expecting BERTopic to need no classical NLP** — it still uses TF-IDF (c-TF-IDF) for interpretable topic words.
- **Forgetting the embedding model dependency** — topic quality depends on the chosen sentence transformer.
- **Confusing BERTopic with BERT fine-tuning** — BERTopic uses pretrained embeddings for clustering; it does not fine-tune BERT for classification.

Quick Revision Summary

- BERTopic uses transformer embeddings to capture semantic document similarity.
- Documents are clustered in embedding space, not bag-of-words space.

- Topic words are extracted per cluster using a TF-IDF variant (c-TF-IDF).
- It addresses LDA's lack of context, semantics, and poor short-text handling.
- The pipeline: embed → dimensionality reduction → cluster → topic representation.
- Best for corpora where meaning and paraphrase variation matter.

← **Previous** · **Index** · **Next** →
BERTopic Architecture: End-to-End Pipeline

BERTopic Architecture: End-to-End Pipeline

High-Level Flow

BERTopic transforms raw documents into named, interpretable topics through five sequential stages. Each stage solves a specific problem: semantic representation, computational tractability, grouping, labelling, and diversity optimisation.

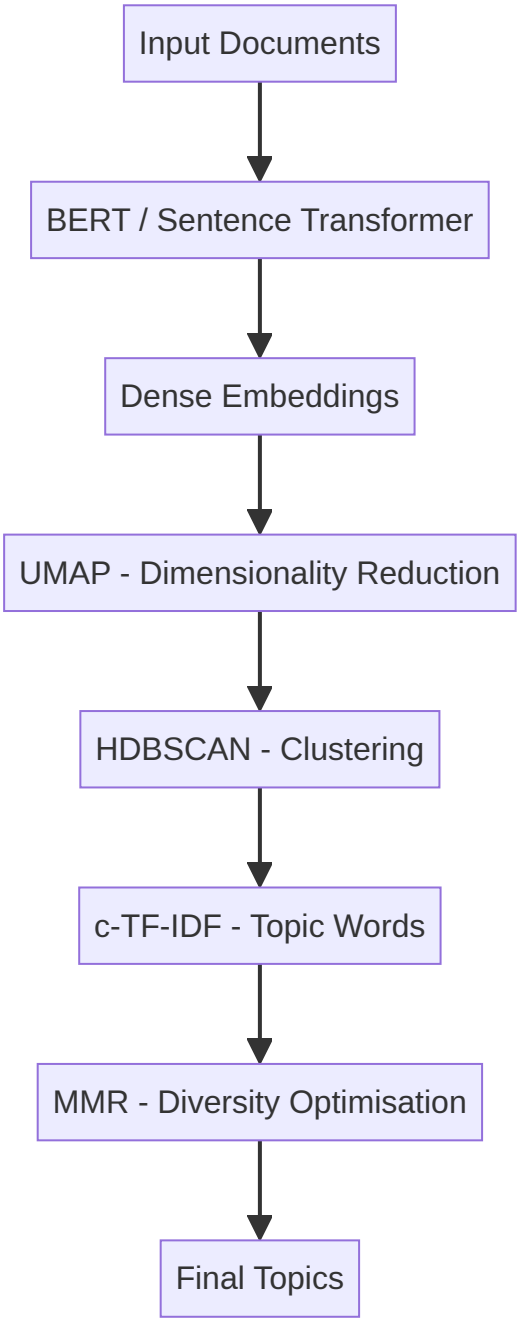

Stage 1: Embedding Generation

Component: Pretrained transformer (e.g., BERT, Sentence-BERT)

Purpose: Convert each document into a dense vector that encodes semantic meaning.

- Captures context, synonyms, and paraphrases
- Output: high-dimensional embedding per document (e.g., 384 or 768 dimensions)

Intuition: Documents about "Kubernetes scaling" and "container orchestration autoscaling" land near each other in vector space even if they share few exact words.

Stage 2: Dimensionality Reduction (UMAP)

Component: UMAP (Uniform Manifold Approximation and Projection)

Purpose: Compress high-dimensional embeddings into a lower-dimensional space suitable for clustering.

Property	UMAP	PCA
Preserves local structure	Yes (emphasised)	Global variance
Non-linear	Yes	No
Clustering-friendly	Yes	Moderate

UMAP keeps semantically similar documents close together while reducing computational cost for the clustering step.

Stage 3: Clustering (HDBSCAN)

Component: HDBSCAN (Hierarchical Density-Based Spatial Clustering)

Purpose: Group documents by semantic similarity using density-based clustering.

- Documents in the same cluster discuss similar themes
- **No fixed K required** — the number of topics emerges from data density
- Uses **cosine similarity** on reduced embeddings
- Outlier documents may be assigned to no cluster (noise points)

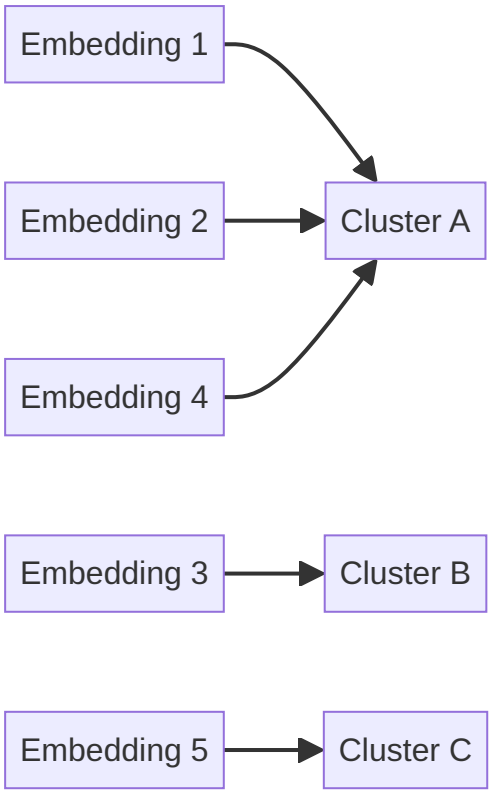

Stage 4: Topic Representation (c-TF-IDF)

Component: Class-based TF-IDF (c-TF-IDF)

Purpose: Extract the most representative words for each cluster.

Standard TF-IDF scores words across all documents. **c-TF-IDF** computes term frequency and inverse document frequency **per cluster**:

- Words frequent in one cluster but rare in others become top topic terms
- Produces human-readable topic descriptions (e.g., *markets, inflation, stock, reacted*)

Stage 5: MMR (Optional)

Component: Maximal Marginal Relevance

Purpose: Optimise topic word lists for **diversity** — reduce redundant or near-duplicate terms in the top-N words.

- Not present in all BERTopic variants
- Enhances readability of topic labels in dashboards and reports

---

End-to-End Data Flow Summary

Step	Input	Output	Key Algorithm
1	Raw text	Dense vectors	Sentence Transformer
2	Dense vectors	Reduced vectors	UMAP
3	Reduced vectors	Cluster assignments	HDBSCAN
4	Clusters + documents	Top words per topic	c-TF-IDF
5	Top words	Diversified word list	MMR (optional)

Common Pitfalls / Exam Traps

- **Listing only "BERT + clustering"** — the full pipeline includes UMAP, HDBSCAN, c-TF-IDF, and optionally MMR.
- **Confusing UMAP with PCA** — UMAP is non-linear and preserves local neighbourhood structure; PCA is linear.
- **Assuming HDBSCAN requires preset K** — unlike LDA, topic count emerges from density; this is a key exam differentiator.
- **Calling it "TF-IDF" without the "c-"** — c-TF-IDF computes IDF per cluster, not globally.
- **Treating MMR as mandatory** — it is optional; smaller BERTopic configurations may omit it.

Quick Revision Summary

- BERTopic pipeline: Embed → UMAP → HDBSCAN → c-TF-IDF → MMR (optional).
- BERT/sentence transformers produce semantic document embeddings.
- UMAP reduces dimensionality while preserving local structure for clustering.
- HDBSCAN clusters by density using cosine similarity; topic count is data-driven.
- c-TF-IDF extracts representative words per cluster for interpretability.
- MMR optionally diversifies top topic words to reduce redundancy.

← **Previous** · **Index** · **Next** →
GSDMM: Topic Modelling for Short Text

GSDMM: Topic Modelling for Short Text

The Short-Text Problem

LDA assumes each document is a **mixture of topics**. Long articles naturally satisfy this — a 2,000-word report might discuss technology, regulation, and market impact simultaneously. Short texts violate this assumption in practice:

- Tweets (≤ 280 characters)
- Product reviews (1–3 sentences)
- Image captions on social media

These snippets typically express **one theme**. With fewer than 50 words, LDA lacks sufficient co-occurrence signal to infer meaningful topic mixtures.

GSDMM (Gibbs Sampling Dirichlet Multinomial Mixture) was designed specifically for this scenario.

LDA vs GSDMM: Core Assumption Difference

Property	LDA	GSDMM
Document-topic model	Mixture of topics	**One topic per document**
Best for	Long documents, articles	Short text, tweets, reviews
Text length	Hundreds+ words	< 50 words typical
Algorithm	Variational inference	Gibbs sampling

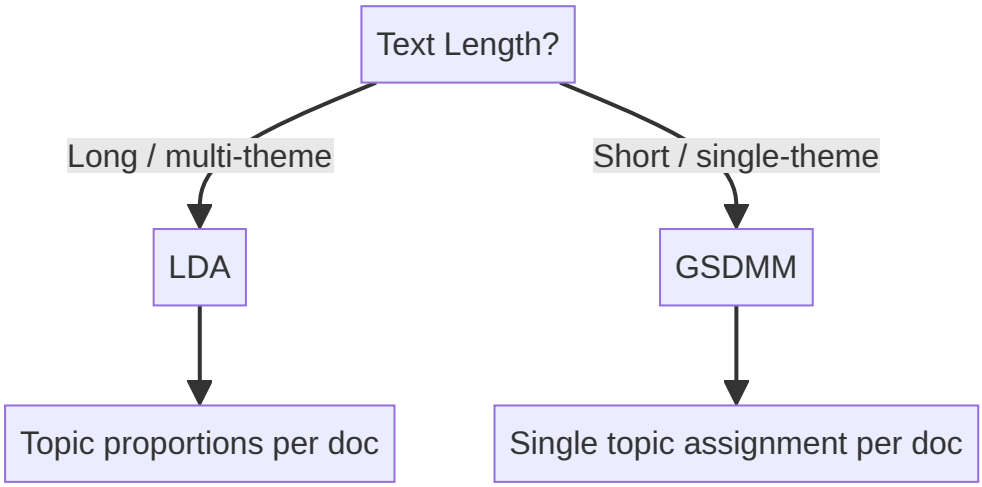

GSDMM Overview

- **Full name:** Gibbs Sampling Dirichlet Multinomial Mixture
- **Key class:** `MovieGroupProcess` (from the GSDMM Python library)
- **Assumption:** Each short document belongs to exactly one topic
- **Method:** Gibbs sampling over document-cluster assignments

This one-topic-per-document assumption aligns with how people write tweets and reviews — one thought, one theme.

Implementation Pipeline

1. Data Source

Short-text corpus example: tweets from a public archive (JSON format with a `text` column containing tweet content and metadata columns).

2. Setup

```
git clone <gsdmm-repo>  # provides MovieGroupProcess class
```

Dependencies: `pandas`, `numpy`, `json`, `nltk`

3. Preprocessing

Step	Purpose
Drop missing values	Ensure clean input
Unicode normalisation	Handle encoding variants (UTF-8, ASCII)
Lowercasing	Normalise case
Regex cleaning	Remove non-alphanumeric characters
Stop-word removal	Filter function words
Stemming (Porter)	Reduce *elections* → *elect*, *remembered* → *rememb*

4. Model Training

```
from gsdmm import MovieGroupProcess

mgp = MovieGroupProcess(
    K=15,           # number of topics (manual)
    alpha=0.1,      # cluster-document hyperparameter
    beta=1.0,       # cluster-word hyperparameter
    n_iters=30      # Gibbs sampling iterations
)
mgp.fit(docs, vocab_size)
```

Hyperparameters:

Parameter	Role	Tuning
`K`	Number of topics	Set based on domain knowledge
`alpha`	Document-cluster density	Experiment per dataset
`beta`	Word-cluster density	Experiment per dataset
`n_iters`	Sampling iterations	Increase until convergence

5. Interpreting Clusters

Rank clusters by **document count** (most populated = most important theme):

```
Cluster index 4: 630 documents → Topic: Elections/Voting
Cluster index 10: 233 documents → Topic: Vaccination/Health
Cluster index 8: 45 documents → Topic: National Security
```

Inspect top words per cluster via `mgp.cluster_word_distribution()`:

- Cluster 4: *elect, vote, state, Georgia, people* → Elections
- Cluster 10: *great, thank, news, vaccine, get* → Vaccination
- Cluster 8: *nation, section, defence, security* → National Security

6. Packaging Results

Build a DataFrame with columns: `text`, `topic`, `stems`

```
def create_topics_dataframe(data_text, mgp, threshold=0.3, topic_dict, stem_texts):
    # Assign topic if probability >= threshold, else "Other"
    ...
```

- `threshold=0.3`: documents below 30% confidence are grouped into an "Other" category
- `topic_dict`: maps cluster indices to human-readable names

Example Output

Topic Name	Document Count
Elections	654
Vaccination	233
National Security	29
Other	(below threshold)

Common Pitfalls / Exam Traps

- **Using LDA for tweets in an exam** — the correct short-text answer is GSDMM (or BERTopic).
- **Confusing GSDMM's one-topic assumption with LDA's mixture assumption** — this is the primary differentiator.
- **Skipping stemming for short text** — with few words per document, normalising word forms improves cluster coherence.
- **Ignoring hyperparameter tuning** — `alpha`, `beta`, and `n_iters` require experimentation; defaults may not converge.
- **Misreading cluster importance** — rank by document count, not cluster index order.

Quick Revision Summary

- GSDMM = Gibbs Sampling Dirichlet Multinomial Mixture, designed for short text.
- Assumes one topic per document (opposite of LDA's mixture assumption).
- Best for tweets, captions, and product reviews (< 50 words).
- Implemented via `MovieGroupProcess` class from the GSDMM library.
- Pipeline: load JSON → clean → stem → stop words → fit → inspect clusters → package DataFrame.
- Rank clusters by document count; label topics from top word distributions manually.
- Use probability threshold to filter low-confidence assignments into "Other".

← Previous · Index · Next →

Week 12

Week 12

Why Decoder-Only Models Dominate NLG

The Generation Problem

Natural Language Generation (NLG) requires producing coherent, fluent text one token at a time. Not all transformer architectures are equally suited to this task. **Decoder-only** models — the architecture behind GPT, Gemini, and Llama — dominate NLG because their training objective directly matches how text is produced.

Autoregressive Next-Token Prediction

Decoder-only transformers model the probability of the **next token** given all preceding tokens:

$$P(t_n \mid t_1, t_2, \ldots, t_{n-1})$$

Generation is **sequential and autoregressive**:

1. Given context, predict token t_1
2. Append t_1 to context; predict t_2
3. Append t_2; predict t_3
4. Continue until an end-of-sequence token or length limit

Each generated token becomes part of the context for the next prediction. The model learns: *given this sequence so far, what is the most likely next token?*

Why This Architecture Wins for NLG

Property	Decoder-Only	Encoder-Only (BERT)	Encoder-Decoder (T5)
Training objective	Next-token prediction	Masked token prediction	Seq2seq (input → output)
Natural for text generation	Yes — same as inference	No — bidirectional, not generative	Yes, but more complex
Scales with data and parameters	Extremely well	Good for understanding	Good for translation/summarisation
Dominant in chat/completion APIs	Yes (GPT, Gemini, Llama)	No	Moderate (specialised tasks)

The simplicity of next-token prediction scales remarkably well with model size and training data — a key reason decoder-only models power modern chatbots, code assistants, and content generators.

Contrast with Encoder-Only Models

BERT-style encoder-only models use **masked language modelling** — predicting missing tokens using bidirectional context. This excels at understanding (classification, NER, embeddings) but does not naturally generate text left-to-right. Decoder-only models are trained exactly as they are used at inference time.

Common Pitfalls / Exam Traps

- **Claiming BERT is used for open-ended text generation** — BERT is encoder-only; GPT-style decoder-only models dominate NLG.
- **Confusing autoregressive with parallel generation** — tokens are generated one at a time; parallelism happens only across batch items, not within a single sequence.
- **Assuming the model "plans" the full answer first** — each token is chosen based only on prior tokens; there is no lookahead.
- **Mixing up encoder-decoder and decoder-only** — T5/BART use both components; GPT/Gemini/Llama use decoder-only stacks.

Quick Revision Summary

- Decoder-only transformers dominate NLG because they model next-token probability.
- Generation is autoregressive: each token is appended to context for the next prediction.

- Training objective (next-token prediction) matches inference behaviour exactly.
- This simple objective scales well with parameters and data.
- Encoder-only models (BERT) excel at understanding, not open-ended generation.
- GPT, Gemini, and Llama are decoder-only architectures.

← **Previous** · **Index** · **Next** →
What Is a Large Language Model (LLM)?

What Is a Large Language Model (LLM)?

Defining "Large"

A model earns the label **Large Language Model** based on three scaling dimensions — not a fixed numerical threshold. "Large" is relative and evolves as hardware and research advance.

Dimension	What Scales	Typical Scale
Parameters	Model weights	Millions to billions (10^9+)
Training data	Text corpus size	Internet-scale: Wikipedia, Stack Overflow, Reddit, books, web crawl
Sequence length	Context window for next-token prediction	Thousands of tokens (modern models: 128K+)

What matters is that **scaling produces emergent capabilities** — behaviours not present in smaller models, such as multi-step reasoning, code generation, and instruction following.

The Pre-Training Objective

LLMs are pre-trained with one deceptively simple objective:

> Given a sequence of tokens (context), predict the next token.

Examples

Context	Possible Next Tokens
"I love natural language processing. NLP is very"	powerful, interesting, useful, ...
"The cat sat on the"	mat, table, floor, desk, ...

The chosen token depends on **training data statistics** and **context**. The model learns grammar, factual associations, writing style, and discourse structure — all implicitly, without task-specific labels.

No Labels Required

Pre-training is **self-supervised**: the training signal comes from the text itself (predict the next word). No human annotation of "correct answers" is needed. This is why LLMs can be trained on trillions of tokens from the open web.

After pre-training, models may undergo **fine-tuning** or **RLHF** for instruction following — but the foundation is always next-token prediction at scale.

Emergent Capabilities

As parameters, data, and compute increase, LLMs exhibit capabilities that smaller models lack:

- Multi-paragraph coherent writing
- Code generation and debugging
- Translation and summarisation without task-specific training
- Following complex natural-language instructions (with alignment tuning)

These emerge from scale — they are not explicitly programmed.

Common Pitfalls / Exam Traps

- **Citing a fixed parameter count as the LLM threshold** — there is no strict cutoff; "large" is relative across eras.
- **Confusing pre-training with fine-tuning** — pre-training uses next-token prediction on raw text; fine-tuning uses labelled or preference data.
- **Assuming LLMs need labelled data for pre-training** — pre-training is self-supervised.
- **Treating next-token prediction as trivial** — its simplicity is precisely why it scales; the complexity emerges from data and parameters.

Quick Revision Summary

- LLMs are defined by scale: parameters, training data, and sequence length.
- No strict quantitative threshold — "large" is relative.
- Pre-training objective: given context, predict the next token.
- Pre-training is self-supervised — no task labels required.
- Scaling leads to emergent NLG capabilities.
- Training data (Wikipedia, web, code) determines factual associations and style.

← **Previous** · **Index** · **Next** →
Why LLMs Excel at Natural Language Generation

Why LLMs Excel at Natural Language Generation

The NLG Leap

Before transformer-based LLMs, generating multi-paragraph coherent text was difficult. RNNs and LSTMs degraded over long sequences. LLMs changed this through scale, attention, and a training objective aligned with generation.

Three Key Strengths

1. Long-Range Dependencies

RNNs suffer from vanishing gradients — after several sentences, early context is effectively forgotten. Transformer self-attention lets every token attend to every other token, preserving context across hundreds or thousands of tokens.

Example: In a 10-paragraph essay prompt, an LLM can maintain thematic consistency from paragraph 1 to paragraph 10.

2. Coherent Multi-Sentence Output

LLMs generate essays, explanations, and multi-paragraph responses that maintain logical flow — a capability largely absent in pre-transformer NLG systems.

3. Instruction Adaptation via Prompts

The same model produces different outputs depending on the **prompt** (input instructions). Change the prompt from "summarise this" to "explain this to a 10-year-old" and the output style, depth, and vocabulary shift accordingly.

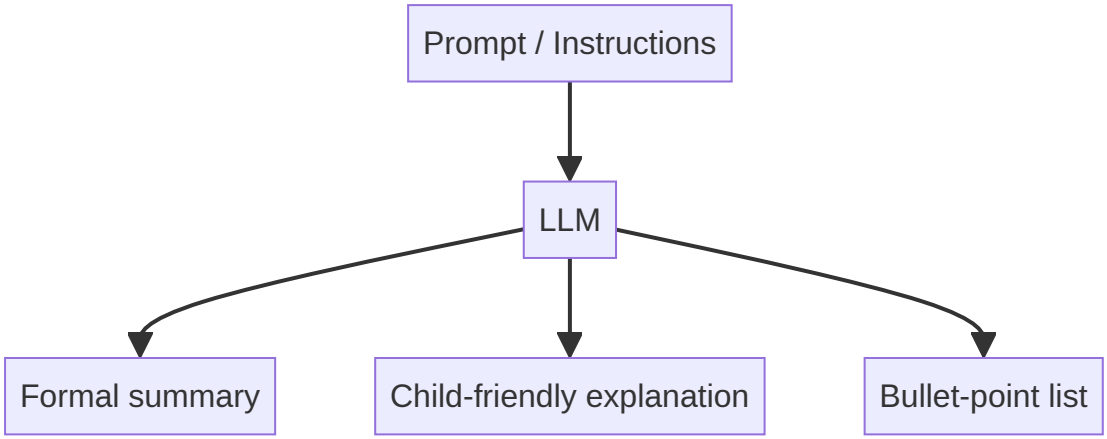

The Critical Reality Check

LLMs appear intelligent and reasoning-capable, but they are fundamentally **next-token prediction engines**.

What It Seems Like	What Actually Happens
Reasoning and thinking	Statistical pattern matching over training data
Understanding the question	Conditioning on token sequences
Choosing the "right" answer	Selecting the highest-probability next token

Example: "The cat sat on the ___"

The model might output *mat*, but *table*, *floor*, and *desk* were also candidates. *Mat* wins because its conditional probability $P(\text{mat} \mid \text{context})$ was highest — not because the model "understood" furniture semantics.

Prompt Quality Determines Output Quality

Because generation is probabilistic and prompt-conditioned:

- Vague prompts produce vague outputs
- Specific format instructions produce structured outputs
- The same model can be brilliant or useless depending on prompt design

This directly motivates **prompt engineering** as a core LLM skill.

Common Pitfalls / Exam Traps

- **Attributing genuine reasoning to LLMs** — they predict tokens probabilistically; chain-of-thought is pattern mimicry, not symbolic reasoning.
- **Ignoring prompt impact** — exam questions about output quality should reference prompting, not just model size.
- **Comparing LLMs to retrieval systems** — LLMs generate text; they do not look up verified facts unless augmented with RAG/search.
- **Assuming deterministic output** — sampling strategies (temperature, top-p) introduce variability.

Quick Revision Summary

- LLMs capture long-range dependencies via transformer attention (unlike RNNs).
- They generate coherent multi-sentence and multi-paragraph text.
- Output adapts to prompt instructions without retraining.
- LLMs are next-token prediction engines, not reasoning systems.
- Token selection is probabilistic — highest-probability token wins.
- Output quality depends heavily on prompt design.

← **Previous** · **Index** · **Next** →
Natural Language Generation (NLG)

Natural Language Generation (NLG)

What Is NLG?

Natural Language Generation (NLG) is the task of producing **human-readable text** from some form of structured or unstructured input. In modern systems, that input is typically a **prompt** — a natural-language instruction that tells the model what to generate.

NLG sits on the output side of the NLP pipeline: where NLU extracts meaning from text, NLG creates text from meaning, data, or instructions.

The Paradigm Shift

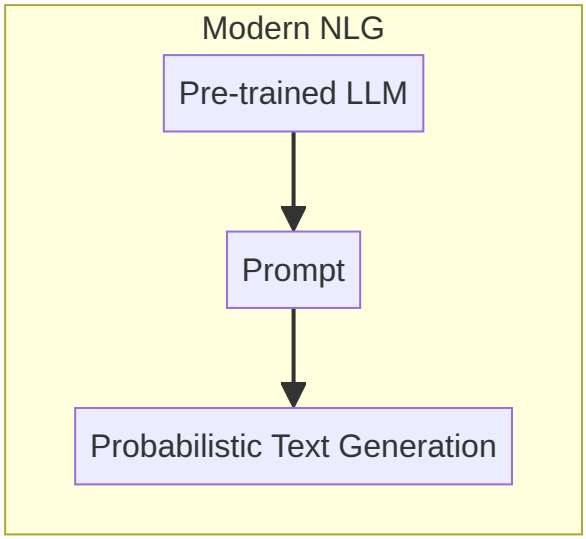

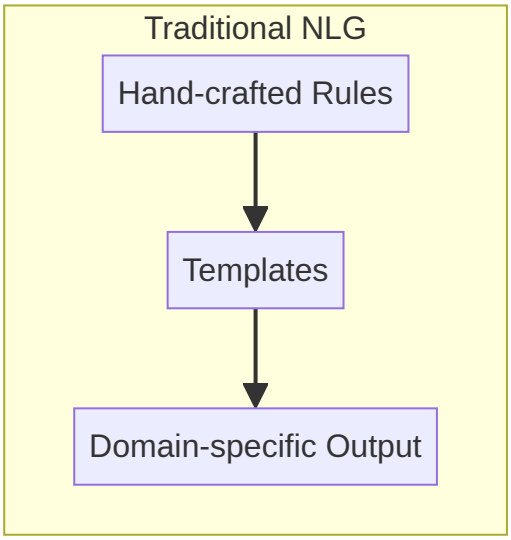

Traditional NLG Systems

- **Rule-based** — explicit if-then logic for sentence construction
- **Template-driven** — fill-in-the-blank patterns (*"Your order {id} will arrive on {date}"*)
- **Highly task-specific** — an e-commerce NLG system could not generalise to healthcare without rebuilding from scratch
- **Brittle** — edge cases and language variation broke rules easily
- **Poor generalisation** — each new domain required new engineering

LLM-Based NLG

- **Single pre-trained model** serves multiple domains
- **Implicit knowledge** learned from training data — no explicit rules
- **Prompt-driven adaptation** — change the task by changing the prompt, not the code
- **Fluent and coherent** — probabilistic generation produces natural-sounding prose

Comparison Table

Aspect	Traditional NLG	LLM-Based NLG
Knowledge source	Explicit rules and templates	Implicit patterns from training data
Domain adaptation	Rebuild system per domain	Change the prompt
Fluency	Often robotic	Human-like
Robustness	Fragile to input variation	More tolerant of varied inputs
Generalisation	Poor across tasks	Strong across tasks via prompting
Control	Precise but rigid	Flexible but probabilistic

The Core Shift

> From **explicit rules** to **implicit knowledge** learned from data.

A traditional weather-report generator needed templates for every phrasing variant. An LLM generates varied, natural weather descriptions from a single prompt: *"Write a morning weather briefing for Mumbai in a conversational tone."*

Common NLG Tasks with LLMs

- **Summarisation** — condense long text into a short summary
- **Text expansion** — expand bullet points into full explanations
- **Paraphrasing** — rewrite while preserving meaning
- **Explanation** — generate educational content on any topic
- **Creative writing** — stories, poetry, marketing copy

Common Pitfalls / Exam Traps

- **Treating NLG as deterministic** — LLM output is probabilistic; traditional NLG with fixed templates is deterministic.
- **Claiming LLMs use explicit rules** — they use learned statistical patterns.
- **Ignoring the prompt's role** — in modern NLG, the prompt is the primary control mechanism.
- **Assuming one LLM prompt works for all domains** — prompting must be adapted per task, even with a single model.

Quick Revision Summary

- NLG produces human-readable text from input (typically a prompt).
- Traditional NLG: rule-based, template-driven, domain-specific, brittle.
- LLM-based NLG: single model, prompt-adapted, fluent, generalises across tasks.
- Key shift: explicit rules → implicit knowledge from training data.
- Common tasks: summarisation, expansion, paraphrasing, explanation.
- Modern NLG quality depends on both model capability and prompt design.

← **Previous** · **Index** · **Next** →
How LLMs Generate a Response

How LLMs Generate a Response

Single-Turn vs Multi-Turn Generation

This course focuses on **single-turn** text generation: one prompt, one response, no memory of prior interactions.

Mode	Behaviour	Example
Single-turn	Each prompt is independent	"Summarise this paragraph"
Multi-turn	Prior Q&A is retained as context	Q1: "Who is the US president?" → Q2: "What is his age?" (resolves "his")

Single-Turn Limitation Demonstrated

```
Prompt 1: "Who is the president of USA?"
Response: "Donald Trump."

Prompt 2: "What is his age?"   (no shared context in single-turn)
Response: "Whose age do you want?" or "Who is he?"
```

In multi-turn mode, the model retains conversation history and correctly resolves pronouns.

Single-turn use cases: summarisation, concept explanation, text rewriting, one-shot classification.

What LLMs Actually Do

LLMs do **not**:

- Retrieve answers from a database
- Perform symbolic reasoning
- Verify facts against ground truth

LLMs **only** generate the most probable next token, repeated autoregressively until completion.

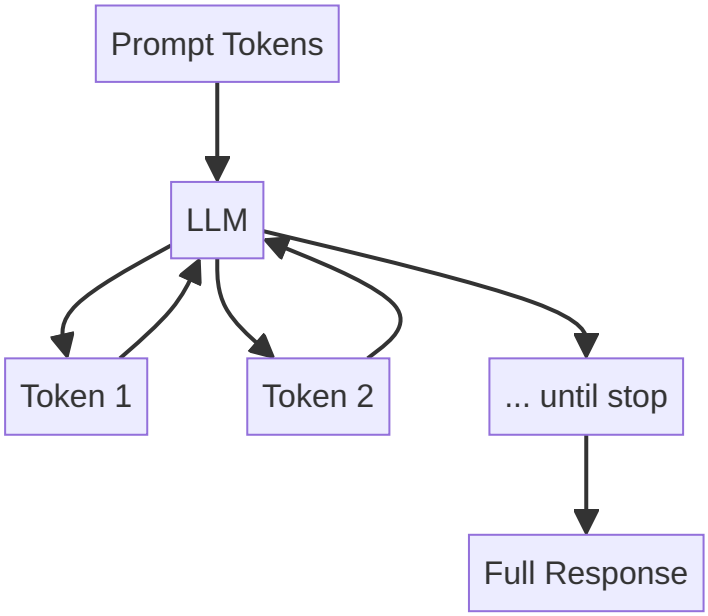

Implications for Production Systems

The Confident Liar Problem

LLM output may sound authoritative but be **factually incorrect**. The model optimises for plausible text, not truth.

- If training data states Bill Clinton was president during the data cutoff, the model may answer "Bill Clinton" even when the real-world answer has changed.
- **Knowledge cutoff date** limits factual accuracy for events after training.

Probability, Not Truth

Output reflects $P(\text{token} \mid \text{context})$ from training statistics — not verified real-world facts. This is why critical applications need:

- Retrieval-Augmented Generation (RAG)
- Human review
- Grounding with search APIs

Common NLG Tasks

Task	Description	Example Prompt
Summarisation	Condense long text	"Summarise this article in 3 sentences"
Text expansion	Expand brief input	"Expand these bullet points into a paragraph"
Paraphrasing	Rewrite preserving meaning	"Rephrase this in simpler language"
Explanation	Teach a concept	"Explain quantum entanglement to a beginner"

Common Pitfalls / Exam Traps

- **Treating LLM responses as retrieved facts** — they are generated, not looked up.
- **Expecting pronoun resolution in single-turn mode** — no conversation memory exists.
- **Ignoring knowledge cutoff** — models cannot know events after their training data date.
- **Trusting confident tone as accuracy** — fluency does not imply correctness.
- **Confusing single-turn with stateless API calls** — even multi-turn requires explicit history passing in API design.

Quick Revision Summary

- Single-turn: one prompt, one response, no prior context retained.
- Multi-turn: conversation history enables pronoun resolution and follow-ups.
- LLMs generate probable tokens — they do not retrieve or reason symbolically.
- Output may be confident but incorrect (hallucination).
- Factual accuracy is limited by training data and knowledge cutoff.
- Common NLG tasks: summarisation, expansion, paraphrasing, explanation.
- Critical interpretation and better prompting mitigate risks.

← **Previous** · **Index** · **Next** →
Controlling LLM Output: Decoding Strategies

Controlling LLM Output: Decoding Strategies

Why the Same Prompt Gives Different Answers

LLMs are largely black boxes — internal weights cannot be directly controlled at inference time. However, **decoding strategies** and **hyperparameters** shape how the model samples from its output probability distribution, giving practitioners control over creativity, determinism, and diversity.

Top-K Sampling

Top-K restricts candidate tokens to the **K most probable** options, renormalises their probabilities, and samples from that subset.

K Value	Behaviour	Use Case
K = 1	Greedy decoding — always picks highest-probability token	Maximum determinism
Small K (e.g., 5)	Focused, conservative output	Factual Q&A, data extraction
Large K (e.g., 50)	Diverse, creative, riskier output	Brainstorming, creative writing

Example: "The cat sat on the ___"

Candidate tokens: *mat* (p_1), *table* (p_2), *desk* (p_3), *floor* (p_4), ...

- **Top-K = 2**: only *mat* and *table* considered (highest two probabilities)
- Sample chosen from this reduced set

Top-K gives **explicit control** over how many alternatives the model considers.

Top-P (Nucleus) Sampling

Top-P selects the **smallest set of tokens** whose cumulative probability exceeds threshold p, then samples from that set.

P Value	Behaviour
Low (e.g., 0.7)	Focused, safe generation
High (e.g., 0.95)	Diverse, expressive generation

Example: Top-P = 0.5

Token	Probability	Cumulative
mat	0.20	0.20
table	0.15	0.35
floor	0.10	0.45
desk	0.08	0.53 (exceeds 0.5 — desk excluded if strict)

The candidate set **adapts dynamically** — unlike Top-K's fixed count. This flexibility makes Top-P the preferred strategy in industrial NLG systems.

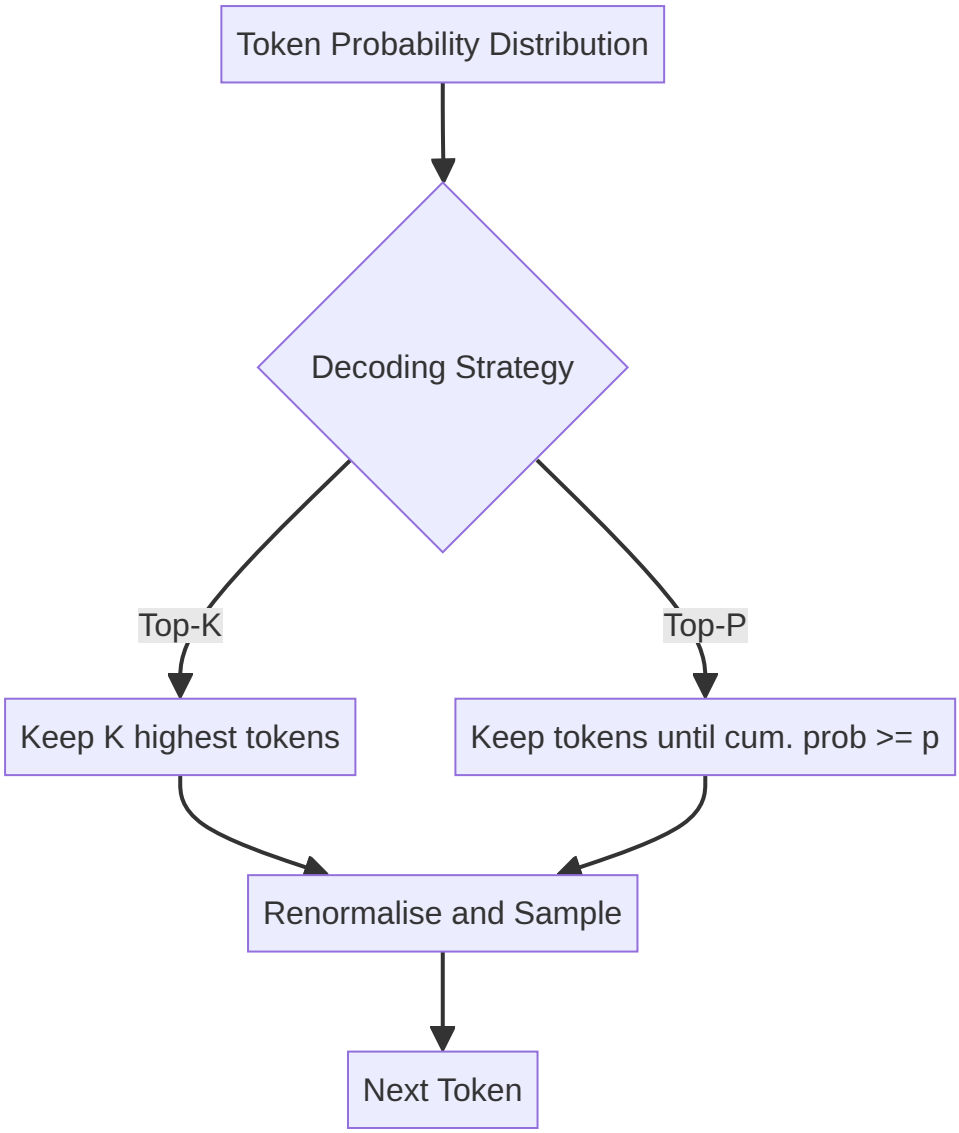

Temperature

Temperature scales the logits before softmax, controlling randomness:

P(token_i) = \frac{\exp(z_i / T)}{\sum_j \exp(z_j / T)}

Temperature	Effect	Recommended For
0.0 – 0.3	Deterministic, focused, conservative	Coding, data extraction, factual Q&A
0.4 – 0.7	Balanced	General-purpose applications
0.7 – 1.0	Creative, random, diverse	Brainstorming, poetry, creative writing
> 1.0 (up to 2.0)	Increasingly nonsensical	Not recommended

Some APIs allow T > 1; beyond 1.0, output becomes incoherent as the model loses control over token selection.

Strategy Comparison

Parameter	Controls	Fixed vs Adaptive
Top-K	Number of candidate tokens	Fixed count
Top-P	Cumulative probability mass	Adaptive set size
Temperature	Sharpness of distribution	Scales all probabilities

---

Common Pitfalls / Exam Traps

- **Confusing Top-K and Top-P** — K fixes the count; P fixes the cumulative probability threshold.
- **Using high temperature for factual tasks** — causes hallucination and inconsistency in code/data extraction.
- **Setting temperature > 1 for production** — produces nonsensical output; stay in 0–1 for most applications.
- **Assuming greedy decoding (K=1) is always best** — it maximises local probability but can produce repetitive text.
- **Ignoring that sampling introduces non-determinism** — same prompt can yield different outputs.

Quick Revision Summary

- LLM output variability comes from probabilistic sampling, not random model changes.
- Top-K: sample from the K most probable tokens; K=1 is greedy decoding.
- Top-P (nucleus): sample from smallest set with cumulative probability ≥ p; more flexible than Top-K.
- Temperature (0–1 typical): low = deterministic, high = creative; >1 causes nonsense.
- Use low temperature for coding/Q&A; high temperature for creative writing.
- Top-P with moderate values (e.g., 0.9) is widely used in production.

← **Previous** · **Index** · **Next** →
Choosing the Right Decoding Strategy

Choosing the Right Decoding Strategy

No Universal Best Settings

There is no single optimal decoding configuration for all applications. The right strategy depends on the task, acceptable risk of hallucination, and desired output diversity. The guidelines below provide starting points — always validate on your specific use case.

General Guidelines

Guideline	Recommendation	Rationale
Sampling method	**Top-P** over Top-K	Adaptive candidate set; balances control and flexibility
Top-P value	**Moderate** (e.g., 0.85–0.95)	Works well as a general default
Temperature for factual tasks	**Low** (0.0–0.3)	Minimises randomness in coding, extraction, Q&A
Temperature for creative tasks	**High** (0.7–1.0)	Enables diverse brainstorming and writing

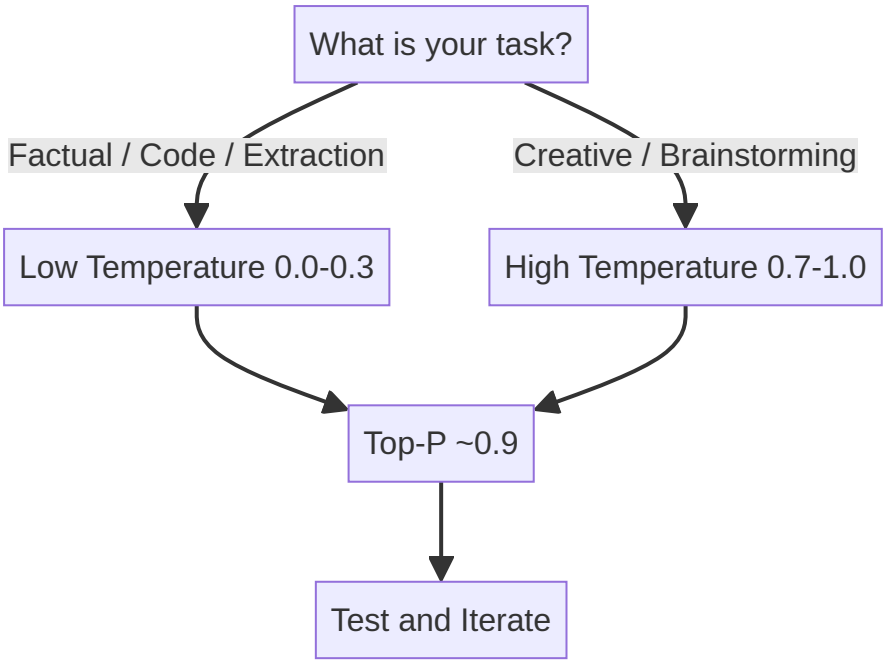

Task-Specific Starting Points

Application	Temperature	Top-P	Top-K
Code generation	0.0 – 0.2	0.9	Optional
Data extraction / JSON output	0.0 – 0.1	0.8 – 0.9	—
Factual Q&A	0.1 – 0.3	0.9	—
Marketing copy	0.7 – 0.9	0.95	—
Poetry / creative writing	0.8 – 1.0	0.95	—
Quiz generation (structured)	0.3 – 0.5	0.9	—

Why Top-P Is Preferred Industrially

Top-K always considers exactly K tokens, even when the probability distribution is sharply peaked (wasting slots on unlikely tokens) or flat (truncating valid options). Top-P adapts:

- **Peaked distribution**: few tokens exceed threshold → conservative output
- **Flat distribution**: many tokens qualify → diverse output

This adaptability makes Top-P the default in Google AI Studio, OpenAI Playground, and most production APIs.

The Experimental Mindset

1. Start with recommended defaults (Top-P ≈ 0.9, temperature matched to task)
2. Generate multiple outputs for the same prompt
3. Evaluate fluency, accuracy, and format compliance
4. Adjust one parameter at a time
5. Document settings per use case in your application config

Common Pitfalls / Exam Traps

- **Claiming one strategy works universally** — settings are application-specific.
- **Using high temperature for JSON/structured output** — causes format violations and hallucinated fields.
- **Choosing Top-K over Top-P without justification** — exams often expect Top-P as the industrial default.
- **Not testing after changing parameters** — small temperature shifts can dramatically alter output quality.
- **Setting both Top-K and Top-P without understanding interaction** — some APIs apply both; know your platform's behaviour.

Quick Revision Summary

- No universally best decoding strategy — tune per application.
- Top-P is generally preferred over Top-K for flexibility.
- Low temperature for informational/factual tasks; high for creative tasks.
- Moderate Top-P (≈0.9) works well as a general starting point.
- Always experiment and iterate — defaults are starting points, not final answers.
- Document hyperparameters per use case in production configs.

← **Previous** · **Index** · **Next** →
Prompts and Prompt Engineering

Prompts and Prompt Engineering

What Is a Prompt?

A **prompt** is the input provided to an LLM that initiates text generation. It is the sole interface for controlling model behaviour in most applications — there is no traditional API to "program" an LLM's logic directly.

A prompt can contain:

- **Instructions** — what task to perform
- **Context** — background information or source text
- **Examples** — demonstrations of desired input-output pairs
- **Constraints** — style, length, format, or tone requirements

There is no single "correct" prompt — only more or less effective ones for a given task.

What Is Prompt Engineering?

Prompt engineering is the practice of designing inputs that guide the model toward the desired output.

> Garbage in → garbage out.

No matter how powerful the model, poor input produces poor output. Because LLMs generate text entirely from the prompt:

- The same model can produce radically different outputs with different prompts
- Small wording changes can significantly alter results

- Unlike traditional software, behaviour is **described**, not **coded**

Prompt vs Traditional Programming

Aspect	Traditional Code	LLM + Prompt
Control mechanism	Explicit logic (if/else, functions)	Natural-language description
Behaviour change	Modify code	Modify prompt
Determinism	High (same input → same output)	Variable (sampling involved)
Flexibility	Rigid, precise	Flexible, probabilistic

Why Prompt Engineering Is a Core Skill

Industrial LLM applications — customer support bots, code assistants, document summarisers, quiz generators — all depend on prompt quality. A well-crafted prompt can make a smaller, cheaper model outperform a larger one with a vague prompt.

Key principles:

- Output depends entirely on input
- Specificity reduces ambiguity
- Iteration is expected — prompt design is experimental
- Constraints (format, length, tone) must be explicit

Common Pitfalls / Exam Traps

- **Assuming bigger models eliminate need for good prompts** — prompt quality always matters.
- **Treating prompts as one-time setup** — prompt engineering is iterative and task-specific.
- **Vague instructions expecting precise output** — "summarise this" vs "summarise in 3 bullet points under 100 words" produce very different results.
- **Confusing system prompt with user prompt** — system instructions set persistent behaviour; user prompt carries the task.

Quick Revision Summary

- A prompt is the LLM input that triggers text generation.
- Prompts can include instructions, context, examples, and constraints.
- Prompt engineering designs inputs for desired outputs.
- LLM behaviour is described through prompts, not coded explicitly.
- Small prompt changes can significantly affect results.
- No universal best prompt — effectiveness is task-dependent.

← **Previous** · **Index** · **Next** →
Core Components of a Prompt

Core Components of a Prompt

Why Structure Matters

Unstructured prompts produce inconsistent outputs. A well-structured prompt reduces ambiguity and gives the model clear signals about role, task, and expected response format.

The Four Components

Component	Purpose	Required?
Role	Who the model acts as; who the audience is	Recommended
Input	Context, data, or instructions for the task	Required
Output expectation	Format, tone, depth, length	Recommended
Examples	Demonstrations of desired behaviour	Optional

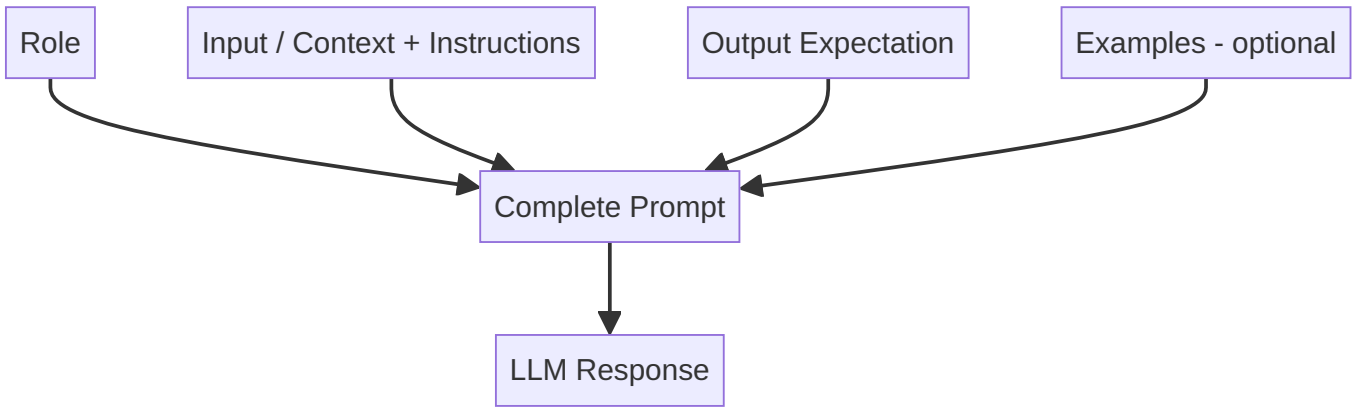

---

1. Role

Defines the persona the model should adopt and optionally the target audience.

Example:

> You are a math teacher explaining concepts to 3rd graders.

Element	Value
Model role	Math teacher
Audience	3rd graders

The role shapes vocabulary, analogies, and explanation depth.

---

2. Input

Provides the information the model uses to generate a response. Can be:

- **Context only** — a paragraph to summarise
- **Instruction only** — "Explain overfitting in machine learning"
- **Both** — context plus specific instructions about what to do with it

---

3. Output Expectation

Specifies how the response should look:

Constraint Type	Examples
Format	"Respond in 3 bullet points", "Return raw JSON"
Length	"Limit to 100 words"
Tone	Friendly, constructive, formal, sarcastic
Style	Simple language, poetic, literary, technical
Depth	Beginner overview vs expert analysis

Clear structure reduces ambiguity and improves consistency.

Worked Example: Teaching Quantum Physics

```
You are a physics teacher conducting a class for 3rd grade.

Explain the top 3 concepts of quantum physics.
Use analogies to explain each concept.
Tie back each concept to its respective analogy.
Use simple language and a friendly tone.
```

Component	Content
Role	Physics teacher for 3rd graders
Input	Explain top 3 quantum physics concepts with analogies
Output expectation	Simple language, friendly tone, analogy-linked explanations

---

4. Examples (Optional)

Demonstrations of input-output pairs — covered in detail under zero-shot, one-shot, and few-shot prompting. Examples anchor format and style when the task is complex or output structure matters.

Common Pitfalls / Exam Traps

- **Omitting output format specification** — leads to inconsistent structure across runs.
- **Confusing role with instruction** — "You are a lawyer" is role; "Summarise this contract" is instruction.
- **Vague output expectations** — "be concise" is weaker than "respond in exactly 3 bullet points, max 50 words each".
- **Skipping role for educational content** — audience-appropriate language requires explicit role/audience definition.
- **Listing only 3 components and forgetting examples** — exams may ask about all four.

Quick Revision Summary

- Four prompt components: Role, Input, Output expectation, Examples (optional).
- Role defines model persona and target audience.
- Input provides context, instructions, or both.
- Output expectation specifies format, tone, depth, and length.
- Clear structure reduces ambiguity and improves output quality.
- Examples (one-shot/few-shot) further anchor format when needed.

← **Previous** · **Index** · **Next** →
Zero-Shot, One-Shot, and Few-Shot Prompting

Zero-Shot, One-Shot, and Few-Shot Prompting

What "Shot" Means

In prompting terminology, **"shot"** refers to the number of **examples** provided in the prompt. More shots = stronger format guidance, but higher token cost.

Zero-Shot Prompting

Definition: No examples provided — only instructions.

```
Explain the concept of overfitting in machine learning.
```

Advantage	Limitation
Simple and fast	Output format may vary
Low token cost	Ambiguous for structured tasks
Often effective for open-ended tasks	Inconsistent for classification with specific labels

Best for: explanations, brainstorming, general Q&A where format flexibility is acceptable.

One-Shot Prompting

Definition: One example demonstrates the desired input-output pattern before the actual task.

```
Classify the sentiment of the following text as positive, negative, or neutral.

Example:
Text: "The movie was absolutely fantastic."
Sentiment: positive

Now classify:
Text: "The new restaurant's food was delicious but the service was terribly slow."
Sentiment:
```

Advantage	Limitation
Model infers desired format from example	Slightly higher token cost
More consistent output structure	One example may not cover edge cases
More control than zero-shot	Still not guaranteed correct

Best for: classification, formatting tasks where output structure matters.

Few-Shot Prompting

Definition: Multiple examples (typically 2–5) demonstrate the pattern.

```
Classify the sentiment as positive, negative, or neutral.

Example 1:
Text: "The movie was absolutely fantastic."
Sentiment: positive

Example 2:
Text: "I waited for 3 hours and nobody ever called me back."
Sentiment: negative

Example 3:
Text: "The package arrived on Tuesday as scheduled."
Sentiment: neutral

Now classify:
Text: "The new restaurant's food was delicious but the service was terribly slow."
Sentiment:
```

Advantage	Limitation
Strong guidance on style and format	Longer prompts, higher token cost
More reliable, less ambiguous	Still no correctness guarantee
Common in production NLG systems	Diminishing returns beyond ~5 examples

Best for: production classification, structured extraction, consistent formatting.

Comparison Table

Technique	Examples	Token Cost	Consistency	Control
Zero-shot	0	Lowest	Variable	Lowest
One-shot	1	Low	Better	Moderate
Few-shot	2+	Higher	Best	Highest

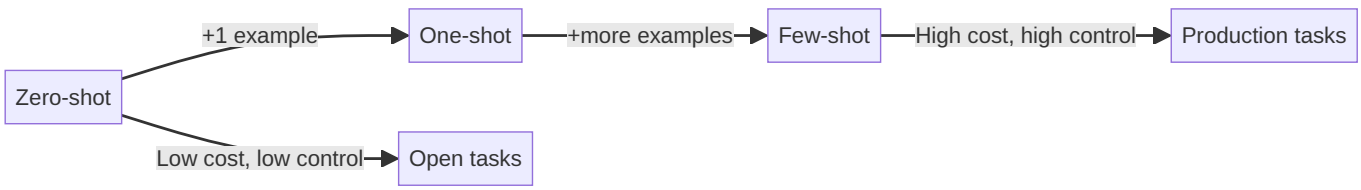

Trade-Off: Cost vs Reliability

Few-shot prompts consume more tokens per request → higher API cost. In production, balance example count against consistency requirements. For a sentiment API handling millions of requests, 3 well-chosen examples often suffice.

Common Pitfalls / Exam Traps

- **Confusing "shot" with "turn"** — shot = examples in prompt; turn = conversation exchange.
- **Claiming few-shot guarantees correctness** — it improves format consistency, not accuracy.
- **Using zero-shot for strict JSON output** — few-shot or explicit format constraints work better.
- **Too many examples wasting context window** — diminishing returns and higher cost beyond ~5 shots.
- **Inconsistent example format** — examples must match the exact output format expected.

Quick Revision Summary

- "Shot" = number of examples in the prompt.
- Zero-shot: no examples; simple, fast, but format may vary.
- One-shot: one example; better format consistency.
- Few-shot: multiple examples; strongest guidance, highest token cost.
- Few-shot is common in production but does not guarantee correctness.
- Choose based on format strictness, cost budget, and task complexity.

← **Previous** · **Index** · **Next** →
Practical Prompt Engineering Guidelines

Practical Prompt Engineering Guidelines

Prompting Is Experimental

No prompt works perfectly on the first attempt. Prompt engineering is an **iterative experimental process** — write, test, evaluate, refine, repeat. The guidelines below accelerate convergence but do not replace testing on your specific task and model.

Five Core Guidelines

1. Be Explicit About the Task

Describe exactly what you want the model to do. Vague tasks produce vague outputs.

Weak	Strong
"Analyse this text"	"Classify this product review as positive, negative, or neutral"
"Write about AI"	"Write a 200-word introduction to supervised learning for beginners"

2. Specify the Desired Output Format

State format requirements precisely:

- "Respond in 3 bullet points"
- "Return raw JSON with keys: title, summary, tags"
- "Use uppercase labels A, B, C, D only"

The more specific the format, the closer the output matches requirements.

3. Avoid Vague Instructions

Vague	Specific
"Be brief"	"Limit response to 100 words"
"Make it professional"	"Use formal tone, no contractions, third person"
"Summarise well"	"Summarise in 3 sentences covering cause, effect, and recommendation"

Specificity increases determinism and reduces unwanted variation.

4. Add Constraints When Needed

When output must meet strict requirements, enumerate every constraint:

- Difficulty level (easy / intermediate / hard)
- Number of items (exactly 3 questions)
- Option format (A, B, C, D uppercase only)
- Include explanations for correct answers
- JSON schema with required keys

5. Test and Iterate

- Create multiple prompt versions
- Run each against the same inputs
- Compare outputs for accuracy, format compliance, and tone
- Keep the best version; document what changed and why

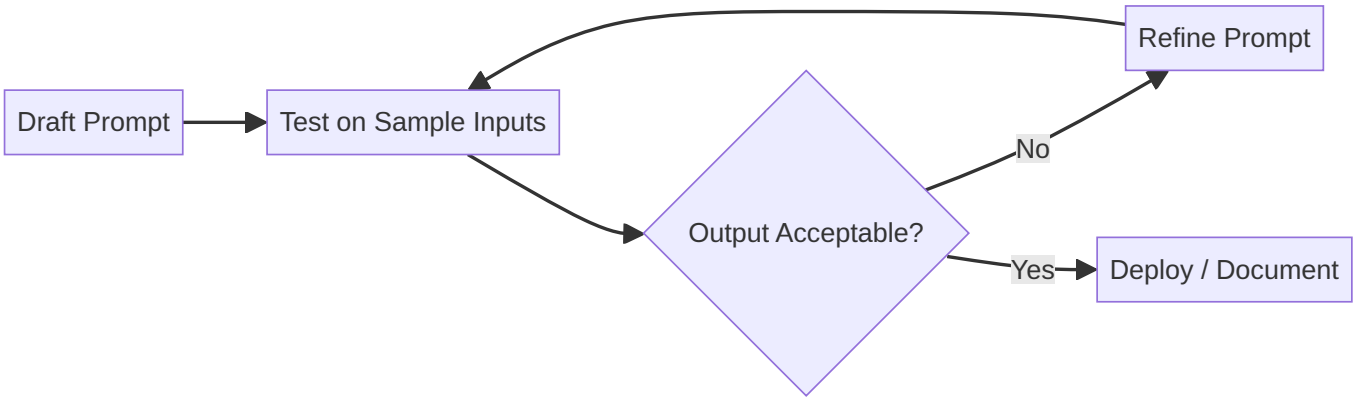

Checklist Before Finalising a Prompt

- [ ] Role defined (who the model is, who the audience is)
- [ ] Task stated explicitly
- [ ] Output format specified (structure, length, tone)
- [ ] Constraints listed (difficulty, count, labels)
- [ ] Examples included if format is complex (one-shot / few-shot)
- [ ] Tested on at least 3 diverse inputs
- [ ] Edge cases handled (empty input, off-topic requests)

Common Pitfalls / Exam Traps

- **Treating prompt design as a one-time step** — it requires iteration.
- **Adding constraints verbally but not in the prompt** — unstated constraints will not be followed.
- **Over-constraining creative tasks** — too many rules on brainstorming prompts stifles useful output.
- **Not testing edge cases** — empty inputs, very long inputs, and off-topic requests reveal prompt weaknesses.
- **Changing multiple variables at once during iteration** — change one element per iteration to understand what helps.

Quick Revision Summary

- Prompt engineering is iterative and experimental, not one-and-done.
- Be explicit about the task with detailed instructions.
- Specify output format precisely (structure, length, tone).
- Avoid vague language — specificity improves determinism.
- Add constraints when format or content must be exact.
- Test multiple prompt versions and refine based on results.

← **Previous** · **Index** · **Next** →
Google AI Studio Playground

Google AI Studio Playground

What Is Google AI Studio?

Google AI Studio is Google's platform for testing Gemini models, generating API keys, and prototyping LLM applications. Beyond text generation, it supports image/video generation, vibe-coding, and app building — but for NLG workflows, the **Playground** is the primary interface.

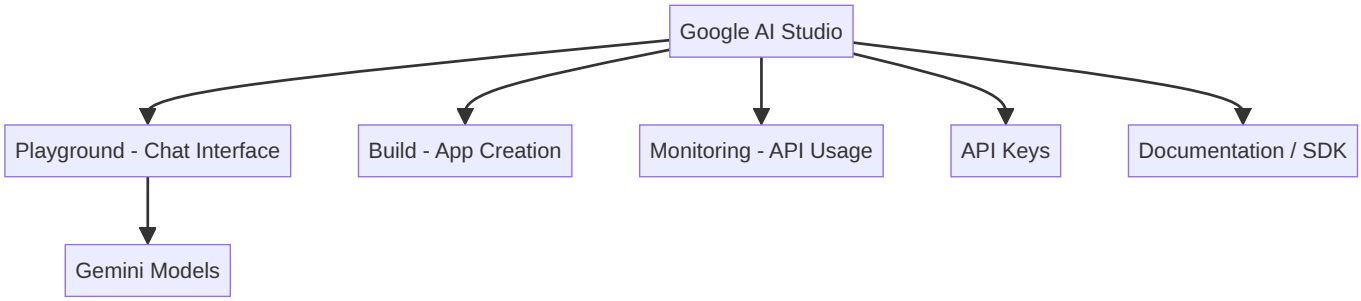

Playground Features

Chat Interface

- Enter prompts and receive generated responses (Ctrl+Enter to send)
- **Thoughts / Chain-of-thought**: expandable panel showing reasoning steps the model took (more steps for complex prompts)
- Model selection dropdown (e.g., Gemini Flash, Pro variants)

Single-Turn vs Multi-Turn

Mode	Behaviour
Multi-turn (default in Playground)	Prior messages retained; "What is his age?" resolves to the previously mentioned person
Single-turn (API design)	Each request is independent; no pronoun resolution

Knowledge cutoff: Models answer based on training data up to a cutoff date. A model trained before a leadership change may give outdated facts (e.g., answering "Joe Biden" when the current president has changed).

Configuration Panel

Setting	Purpose	Recommended
Model	Select Gemini variant	Flash for speed/cost; Pro for quality
System instructions	Persistent behaviour (same as system prompt)	e.g., "Always respond in a poem"
Temperature	Creativity control (0–2)	Stay 0–1; ~0.65 for balanced
Top-P	Nucleus sampling threshold	~0.95 default
Output length	Max response tokens	Set to control cost
Stop sequence	Halt generation at specific string	Optional
Structured outputs	Enforce JSON schema	Enable for API integrations
Thinking level	Reasoning depth for complex puzzles	High for multi-step problems
Grounding with Google Search	Fact-check via live search	Reduces hallucination
Safety settings	Block harassment, hate speech, etc.	Configure per use case

Tools and Integrations

- **Function calling** — attach external tools (weather API, database queries)
- **Code execution** — run generated code in sandbox
- **URL context** — browse URLs while generating
- **Google Maps grounding** — verify location data

Other Platform Sections

Section	Purpose
Build	Vibe-code full applications
Dashboard	Monitor API usage and costs
Documentation	SDK examples in Python, JavaScript
Get API Key	Generate keys for programmatic access

SDK Integration (Python)

```
from google import genai

client = genai.Client(api_key="YOUR_KEY")
response = client.models.generate_content(
    model="gemini-2.5-flash",
    contents="Explain how AI works in a few words"
)
print(response.text)
```

Documentation provides copy-paste code snippets for quick integration.

Common Pitfalls / Exam Traps

- **Trusting Playground answers as current facts** — knowledge cutoff causes outdated responses.
- **Setting temperature > 1 in production** — produces incoherent output.
- **Confusing system instructions with user prompt** — system instructions persist across turns; user prompt is per-message.
- **Ignoring API costs** — Playground usage counts toward billing once API key is linked.
- **Assuming single-turn in Playground** — the chat UI is multi-turn by default; API calls are typically single-turn unless history is passed.

Quick Revision Summary

- Google AI Studio: test Gemini models, get API keys, monitor usage.
- Playground provides chat UI with model, temperature, top-P, and safety settings.
- System instructions = persistent system prompt across the session.
- Multi-turn chat retains context; single-turn API calls do not.
- Knowledge cutoff causes outdated factual answers.
- Structured outputs, grounding, and function calling support production workflows.
- Python SDK: `google.genai` with `Client` and `generate_content`.

← **Previous** · **Index** · **Next** →
Gemini API Key Setup

Gemini API Key Setup

Why an API Key Is Required

Programmatic access to Gemini models requires authentication via an **API key**. Without it, SDK calls fail with:

```
ValueError: No API key was provided. Please pass a valid API key.
```

The key identifies your project and enables usage tracking and billing.

Step 1: Generate the API Key

1. Open Google AI Studio
2. Click **Get API key** (bottom-left)
3. Select or create a **GCP project**
4. Click **Create API key** → select **Gemini API**

5. Copy the key immediately (store securely)

Key Management

Action	Purpose
View usage	Monitor API requests and costs per key
Copy key	Use in applications
Revoke/regenerate	Rotate compromised keys

Step 2: Secure Storage

Never hardcode API keys in source code committed to version control.

Option A: Environment Variables (Production)

```
# .env file (add to .gitignore)
GEMINI_API_KEY=your_key_here
```

```
import os
from dotenv import load_dotenv

load_dotenv()
api_key = os.getenv("GEMINI_API_KEY")
```

Option B: Colab Secrets (Notebooks)

```
from google.colab import userdata
gemini_api_key = userdata.get("GEMINI_API_KEY")
```

Store the key in Colab's Secrets panel (key icon) with name `GEMINI_API_KEY`.

Step 3: Initialise the Client

```
from google import genai

client = genai.Client(api_key=gemini_api_key)

response = client.models.generate_content(
    model="gemini-2.5-flash",
    contents="Explain how AI works in a few words"
)
print(response.text)
# "AI learns patterns from data to make predictions."
```

Critical: passing the key to `genai.Client(api_key=...)` is required — importing the library alone is insufficient.

Security Best Practices

- Add `.env` to `.gitignore` — never commit secrets
- Use separate keys per environment (dev/staging/prod)
- Monitor usage to detect unexpected spend
- Rotate keys if exposed
- Use Colab secrets or cloud secret managers, not plaintext in notebooks shared publicly

Common Pitfalls / Exam Traps

- **Hardcoding API keys in notebooks shared on GitHub** — immediate security breach.
- **Forgetting to pass api_key to Client()** — most common integration error.
- **Confusing API key with OAuth credentials** — Gemini AI Studio uses simple API keys, not full OAuth for basic usage.
- **Not adding .env to .gitignore** — secrets get committed accidentally.
- **Using the same key across all environments** — prevents independent revocation and cost tracking.

Quick Revision Summary

- Gemini SDK requires an API key for all programmatic calls.
- Generate keys in Google AI Studio → Get API key → Create.
- Store keys in `.env` (production) or Colab Secrets (notebooks).
- Initialise: `genai.Client(api_key=your_key)`.
- Never commit API keys to version control.
- Monitor usage and costs via the AI Studio dashboard.

← **Previous** · **Index** · **Next** →
OpenAI Platform and API Setup

OpenAI Platform and API Setup

OpenAI Platform Overview

The OpenAI Platform (`platform.openai.com`) provides chat interfaces, agent builders, media generation tools, API management, and fine-tuning — all accessible from a single dashboard.

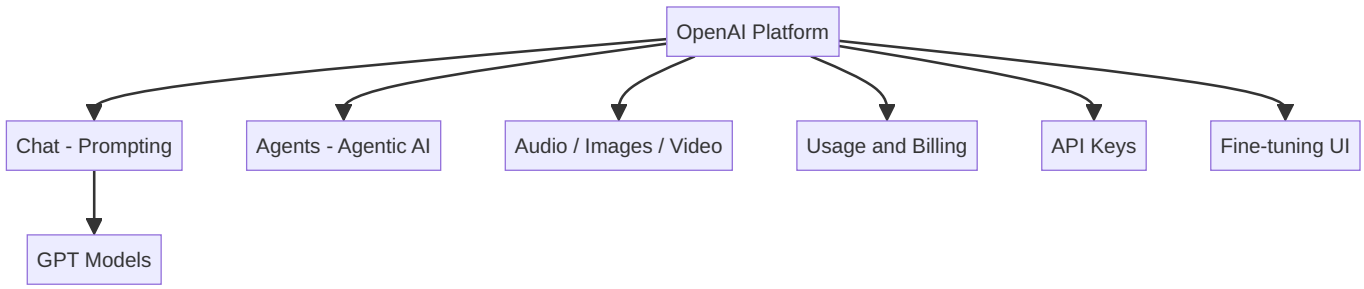

Dashboard Essentials

Section	Purpose
Chat	Interactive prompt testing with model selection
Agents	Build agentic AI workflows via UI
Audio / Images / Video	Multimodal generation tools
Usage	Track API consumption and costs
API Keys	Create and manage secret keys
Fine-tuning	Fine-tune base models without scripting
Logs / Storage	Data management for applications

Billing tip: Add a small credit balance (5–10) and monitor usage — overspending is easy with pay-per-token pricing.

Chat Playground Settings

Parameter	Purpose	Recommended
Model	GPT variant selection	Match task to model tier
System message	Persistent instructions (tone, format)	Set role and constraints
User prompt	Per-request task input	Your actual query
Temperature	Creativity (0–2)	0–1 for most tasks
Max tokens	Response length cap	Limits cost per request
Top-P	Nucleus sampling	~0.9 for balanced output
Tools	Function calling attachments	Optional per use case
Variables	Reusable template variables	e.g., `city = Mumbai`

API Key Generation

1. Navigate to **API Keys** in the left sidebar
2. Click **Create new secret key**
3. Select project, name the key (e.g., "demo")
4. **Copy immediately** — the key is shown only once
5. Store securely (`.env` or Colab Secrets)

Python SDK Integration

```
from openai import OpenAI
import os

client = OpenAI(api_key=os.getenv("OPENAI_API_KEY"))

response = client.responses.create(
    model="gpt-4.1",
    input="Write a one-sentence bedtime story about a unicorn"
)
print(response.output_text)
```

Error without key:

```
OpenAIError: The api_key client option must be set either by passing
api_key or by setting the OPENAI_API_KEY environment variable.
```

Colab Secrets Setup

```
from google.colab import userdata
openai_api_key = userdata.get("OPENAI_API_KEY")
```

Pass to `OpenAI(api_key=openai_api_key)`.

OpenAI vs Google AI Studio

Aspect	OpenAI Platform	Google AI Studio
Models	GPT family	Gemini family
Key storage	`OPENAI_API_KEY` env var	`GEMINI_API_KEY` env var
SDK	`openai` package	`google.genai` package
Playground	Chat with system/user messages	Playground with system instructions
Fine-tuning UI	Yes	Limited

Documentation

Full API reference: `developers.openai.com/api/docs`

Includes model pricing, endpoint specifications, and copy-paste code for Python and JavaScript.

Common Pitfalls / Exam Traps

- **Not saving the API key at creation** — it cannot be retrieved later; only regenerated.
- **Committing keys to GitHub** — use `.env` + `.gitignore`.
- **Ignoring max_tokens** — uncapped responses inflate API costs.
- **Confusing ChatGPT (consumer app) with OpenAI API** — different products, different billing.
- **Forgetting to set api_key in client constructor** — same error pattern as Gemini integration.

Quick Revision Summary

- OpenAI Platform: chat testing, agents, API keys, fine-tuning, usage monitoring.
- Playground settings: model, system message, temperature, max tokens, top-P.
- API keys shown once at creation — copy and store immediately.
- Python: `OpenAI(api_key=...)` then `client.responses.create(...)`.
- Set `OPENAI_API_KEY` environment variable or use Colab Secrets.
- Monitor credit balance to avoid unexpected charges.
- Documentation at developers.openai.com/api/docs.

← **Previous** · **Index** · **Next** →

Week 13

Week 13

QuizGenius AI: Capstone Project Overview

QuizGenius AI: Capstone Project Overview

Project Goal

QuizGenius AI is an LLM-powered application that dynamically generates multiple-choice quizzes on any user-specified topic. Built with the **Gemini API**, it demonstrates how generative AI integrates into interactive, educational applications — a template adaptable to any domain.

The capstone progresses from a Jupyter notebook proof-of-concept to a full-stack Streamlit application with proper API key management, structured JSON output, and a browser-based UI.

Key Features

Feature	Description
Dynamic quiz generation	3-question MCQ on any user-specified topic
Intermediate difficulty	Challenging yet accessible questions
Structured JSON output	Consistent, parseable quiz data from the LLM
Interactive UI	Real-time answer submission and immediate feedback
Secure API key management	Keys stored via Colab Secrets or `.env`, never hardcoded

---

Why Structured JSON Output Matters

Industrial LLM applications rarely consume free-form text. They need **machine-parseable output** to drive downstream logic — rendering UI components, storing in databases, or triggering workflows. Enforcing JSON schema via prompt design and API configuration is a high-demand production skill.

Technology Stack

Layer	Technology
LLM	Google Gemini (via `google.genai` SDK)
Notebook POC	Jupyter / Google Colab
Full-stack UI	Streamlit (`app.py`)
Backend logic	Python (`quiz_engine.py`)
Secrets	`.env` file / Colab `userdata`
Version control	GitHub

Learning Outcomes

- Integrate Gemini API into a Python application
- Design prompts that produce consistent structured output
- Parse and validate LLM JSON responses
- Build interactive NLG applications beyond notebooks
- Manage API keys securely in development and deployment

Extension Ideas

Adapt the template to other domains:

- **HR:** Policy comprehension quizzes for onboarding
- **Cloud:** AWS service knowledge checks for certification prep
- **Healthcare:** Patient education quizzes on treatment plans
- **Legal:** Clause understanding tests for contract review training

---

Common Pitfalls / Exam Traps

- **Hardcoding API keys in source code** — security violation; use secrets management.
- **Accepting free-form LLM text without JSON parsing** — breaks automated UI rendering.
- **Ignoring LLM factual errors in quiz answers** — generated questions may be incorrect; add disclaimers.
- **Treating notebook POC as production-ready** — full-stack deployment requires validation, secrets, and error handling.

Quick Revision Summary

- QuizGenius AI: Gemini-powered MCQ quiz generator on any topic.
- Features: dynamic generation, JSON output, interactive UI, secure API keys.
- Structured JSON from LLMs is critical for industrial applications.
- Template is domain-agnostic — adapt to education, cloud, HR, etc.
- Capstone builds from notebook POC to full-stack Streamlit app.
- Intermediate difficulty, 3 MCQs with A/B/C/D options and explanations.

← **Previous** · **Index** · **Next** →
Package Installation and API Setup

Package Installation and API Setup

Environment Setup Overview

Before building QuizGenius AI, install the Gemini SDK, configure imports, and securely load the API key. This setup pattern applies to any Gemini-based Python project.

Package Installation

```
pip install google-genai
```

Required Imports

```
import json
import re
from google import genai
from google.genai import types
```

Package	Purpose
`json`	Parse structured LLM responses
`re`	Regex cleanup of JSON strings
`google.genai`	Gemini SDK client and types
`google.genai.types`	Configuration objects (e.g., response MIME type)

API Key Configuration

Insecure (Demo Only)

```
api_key = "paste_key_here"  # NOT recommended for production
```

Secure: Colab Secrets

```
from google.colab import userdata
gemini_api_key = userdata.get("GEMINI_API_KEY")
```

Store the key in Colab Secrets panel with name `GEMINI_API_KEY`.

Secure: Environment Variables (Production)

```
# .env
GEMINI_API_KEY=your_key_here
MODEL_NAME=gemini-2.5-flash
```

```
import os
from dotenv import load_dotenv
load_dotenv()
api_key = os.getenv("GEMINI_API_KEY")
```

Client and Model Initialisation

```
client = genai.Client(api_key=gemini_api_key)
model_name = "gemini-2.5-flash"  # fast, cost-effective for POCs
```

Gemini 2.5 Flash is a lighter, faster, more cost-effective variant — ideal for demos and proof-of-concepts. Model names can also be stored in `.env` as `MODEL_NAME` to avoid hardcoding.

Fail-Safe Pattern

```
if not gemini_api_key:
    raise ValueError("GEMINI_API_KEY is missing. Please check your .env file.")
```

Stop early if the key is missing — prevents silent failures and cryptic API errors downstream.

Common Pitfalls / Exam Traps

- **Pasting API keys directly in notebook cells** — keys get committed to Git or shared notebooks.
- **Forgetting `genai.Client(api_key=...)`** — library import alone does not authenticate.
- **Using heavyweight models for simple demos** — Flash variants reduce cost and latency.
- **Not validating key presence before API calls** — add explicit checks with clear error messages.
- **Confusing `google-generativeai` (old) with `google-genai` (new SDK)** — use the current SDK per documentation.

Quick Revision Summary

- Install `google-genai`; import `genai` and `types`.
- Load API key via Colab Secrets or `.env` — never hardcode.
- Initialise `genai.Client(api_key=...)` before any API calls.
- Use `gemini-2.5-flash` for fast, cost-effective POCs.
- Add fail-safe: raise error if API key is missing.
- Also import `json` and `re` for response parsing.

← **Previous** · **Index** · **Next** →
Creating the Quiz Prompt Template

Creating the Quiz Prompt Template

Why a Template?

A **prompt template** is a reusable string with placeholders (`{topic}` , `{num_questions}` , `{difficulty}`) filled at runtime. It encodes role, constraints, and output schema in one place — ensuring every quiz request follows the same structure.

Template Structure

Role and Instruction

```
You are a quizmaster.
Generate a {num_questions}-question multiple choice quiz about {topic}.
```

Component	Value
Role	Quizmaster
Task	Generate N-question MCQ about user topic

Constraints

Rule	Specification
Difficulty	Intermediate (configurable: easy / intermediate / hard)
Options	Exactly 4 options: A, B, C, D (uppercase only)
Correct answer	Marked clearly per question
Explanation	Short explanation for each correct answer

Specificity prevents ambiguous output — the model knows exactly how many options, what labels, and what extras to include.

Output Format: JSON Schema

```
Return a raw JSON with this structure only:
```

```
{
  "quiz_title": "Catchy Title Here",
  "questions": [
    {
      "id": 1,
      "text": "Question text here?",
      "options": {
        "A": "Option A text",
        "B": "Option B text",
        "C": "Option C text",
        "D": "Option D text"
      },
      "correct_option": "B",
      "explanation": "Why B is correct."
    }
  ]
}
```

"Raw JSON" and **"this structure only"** are critical phrases — they reduce markdown wrapping and extraneous text around the JSON.

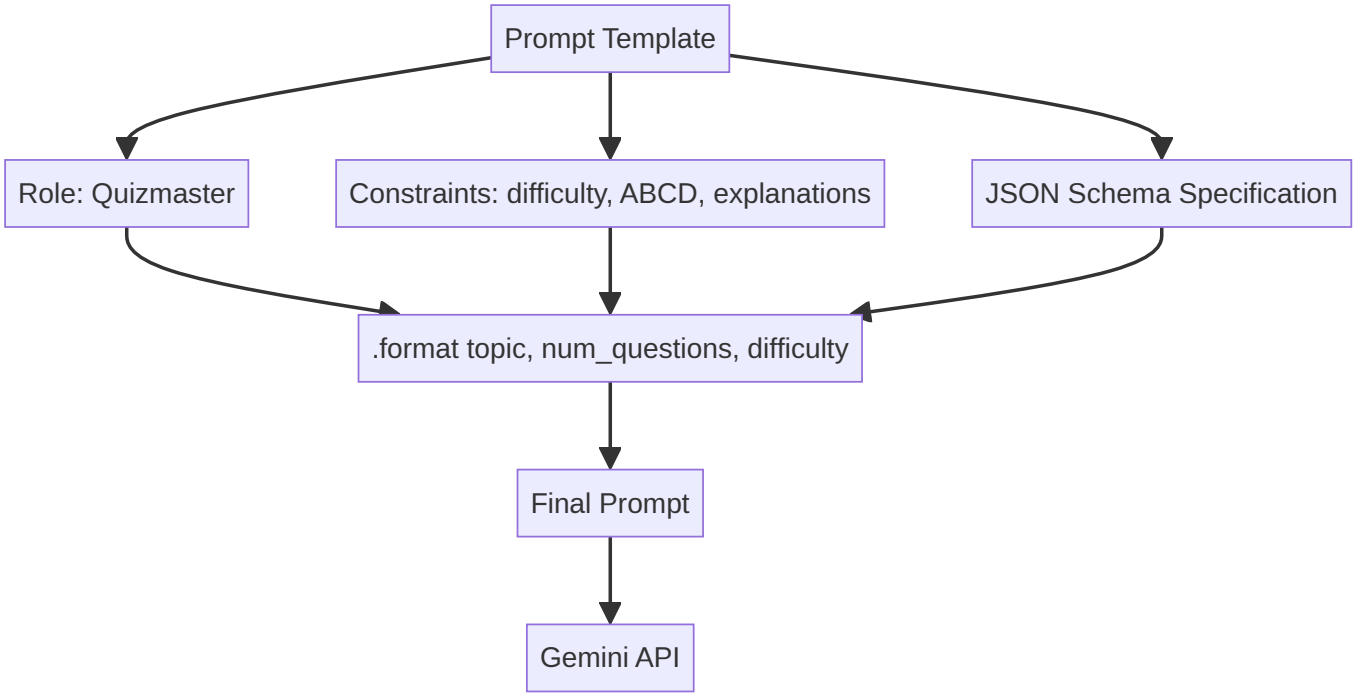

Runtime Formatting

```
prompt = prompt_template.format(
    topic=topic.strip(),
    num_questions=num_questions,
    difficulty=difficulty
)
```

User input is injected at call time. `.strip()` removes leading/trailing whitespace for consistency.

Design Principles Applied

Principle	How Applied
Explicit task	"Generate N-question MCQ about {topic}"
Format specification	Full JSON schema with sample keys
Constraints	Difficulty, option labels, explanations
Role	Quizmaster persona

Common Pitfalls / Exam Traps

- **Not specifying "raw JSON"** — model wraps output in markdown code fences, breaking `json.loads()`.
- **Vague option format** — "4 options" without "A, B, C, D uppercase" produces numbered or lowercase labels.
- **Omitting example schema** — LLM may invent different key names (`question` vs `text`, `answer` vs `correct_option`).
- **Hardcoding topic in template** — use placeholders for dynamic user input.

- **Skipping explanation requirement** — feedback screen needs explanations; must be in prompt.

Quick Revision Summary

- Prompt template: reusable string with `{topic}`, `{num_questions}`, `{difficulty}` placeholders.
- Role: quizmaster; constraints: difficulty, 4 uppercase options, explanations.
- Output: raw JSON with `quiz_title` and `questions` array.
- Each question: `id`, `text`, `options` (A–D), `correct_option`, `explanation`.
- Format with `.format()` at runtime after user input.
- Explicit JSON schema in prompt is essential for parseable industrial output.

← **Previous** · **Index** · **Next** →
Building the Quiz Generation Logic

Building the Quiz Generation Logic

The `generate_quiz()` Function

The core backend function accepts a topic, calls the Gemini API with the prompt template, and returns parsed JSON quiz data.

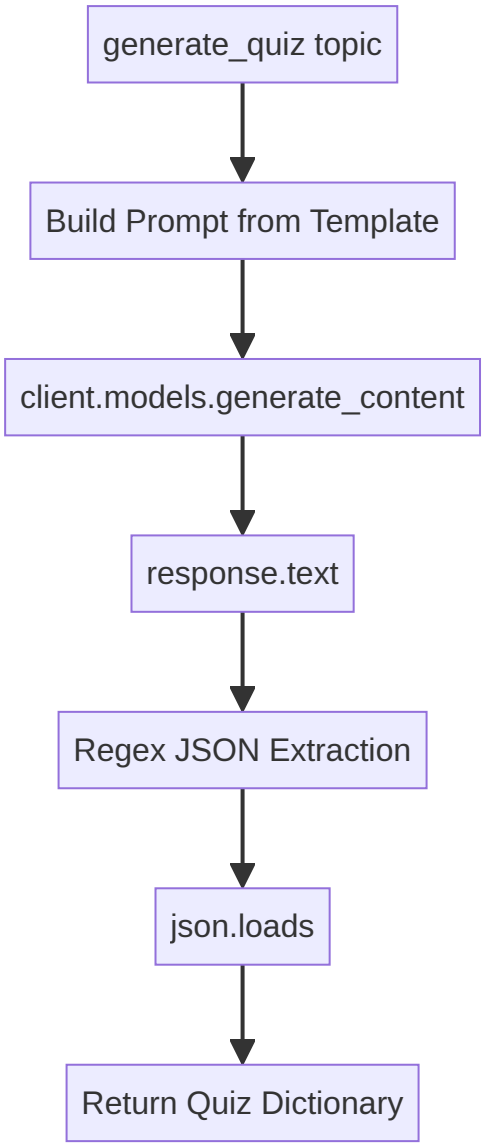

API Call with Structured Output

```
def generate_quiz(topic: str) -> dict:
    print(f"Generating a quiz on topic: {topic}")

    response = client.models.generate_content(
        model=model_name,
        contents=prompt_template.format(topic=topic),
        config=types.GenerateContentConfig(
            response_mime_type="application/json"
        )
    )
```

`response_mime_type="application/json"` instructs the API to return JSON-formatted output — a critical setting for industrial structured generation.

Response Extraction and Cleaning

```
raw_text = response.text

# Extract JSON block via regex
pattern = r'\{.*\}'
match = re.search(pattern, raw_text, re.DOTALL)
clean_json_string = match.group(0) if match else raw_text

# Remove markdown artifacts
clean_json_string = re.sub(r'```json|```', '', clean_json_string).strip()

quiz_data = json.loads(clean_json_string)
return quiz_data
```

Step	Purpose
`response.text`	Extract raw string from API response
Regex extraction	Isolate JSON object from surrounding text
Markdown cleanup	Remove ```` ```json ```` fences if present
`json.loads()`	Parse into Python dictionary

Error Handling

```
try:
    # ... generation logic ...
except json.JSONDecodeError as e:
    raise ValueError(f"Failed to parse quiz JSON: {e}")
except Exception as e:
    raise RuntimeError(f"Quiz generation failed: {e}")
```

Always wrap LLM calls in try/except — API failures, malformed JSON, and network errors are common in production.

Extended Version (Full-Stack App)

The production `quiz_engine.py` version adds:

- **Input validation** before API call (topic length, question count range, valid difficulty)
- **Dynamic prompt formatting** with `{topic}`, `{num_questions}`, `{difficulty}`
- **Schema validation** — verify `"questions"` key exists in parsed JSON
- **Fail-fast on missing API key**

Common Pitfalls / Exam Traps

- **Skipping** `response_mime_type="application/json"` — increases free-text wrapping and parsing failures.
- **Not using regex to extract JSON** — LLMs may add preamble text before the JSON object.
- **No** `json.JSONDecodeError` **handling** — malformed responses crash the app.
- **Trusting LLM JSON without schema validation** — verify required keys (`questions`, `quiz_title`) exist.
- **Calling API before input validation** — wastes tokens and cost on invalid inputs.

Quick Revision Summary

- `generate_quiz(topic)` calls Gemini with formatted prompt template.
- Set `response_mime_type="application/json"` for structured output.
- Extract text → regex JSON extraction → markdown cleanup → `json.loads()`.
- Wrap in try/except for `JSONDecodeError` and general exceptions.
- Production version adds input validation and schema checks.
- Never call the LLM before validating user inputs.

← **Previous** · **Index** · **Next** →
Building the Main Application Runner

Building the Main Application Runner

The `run_quiz_app()` Function

The application runner orchestrates the full quiz lifecycle: collect user input, generate quiz, run the interactive game loop, and display final results. This is the notebook POC entry point before the Streamlit full-stack version.

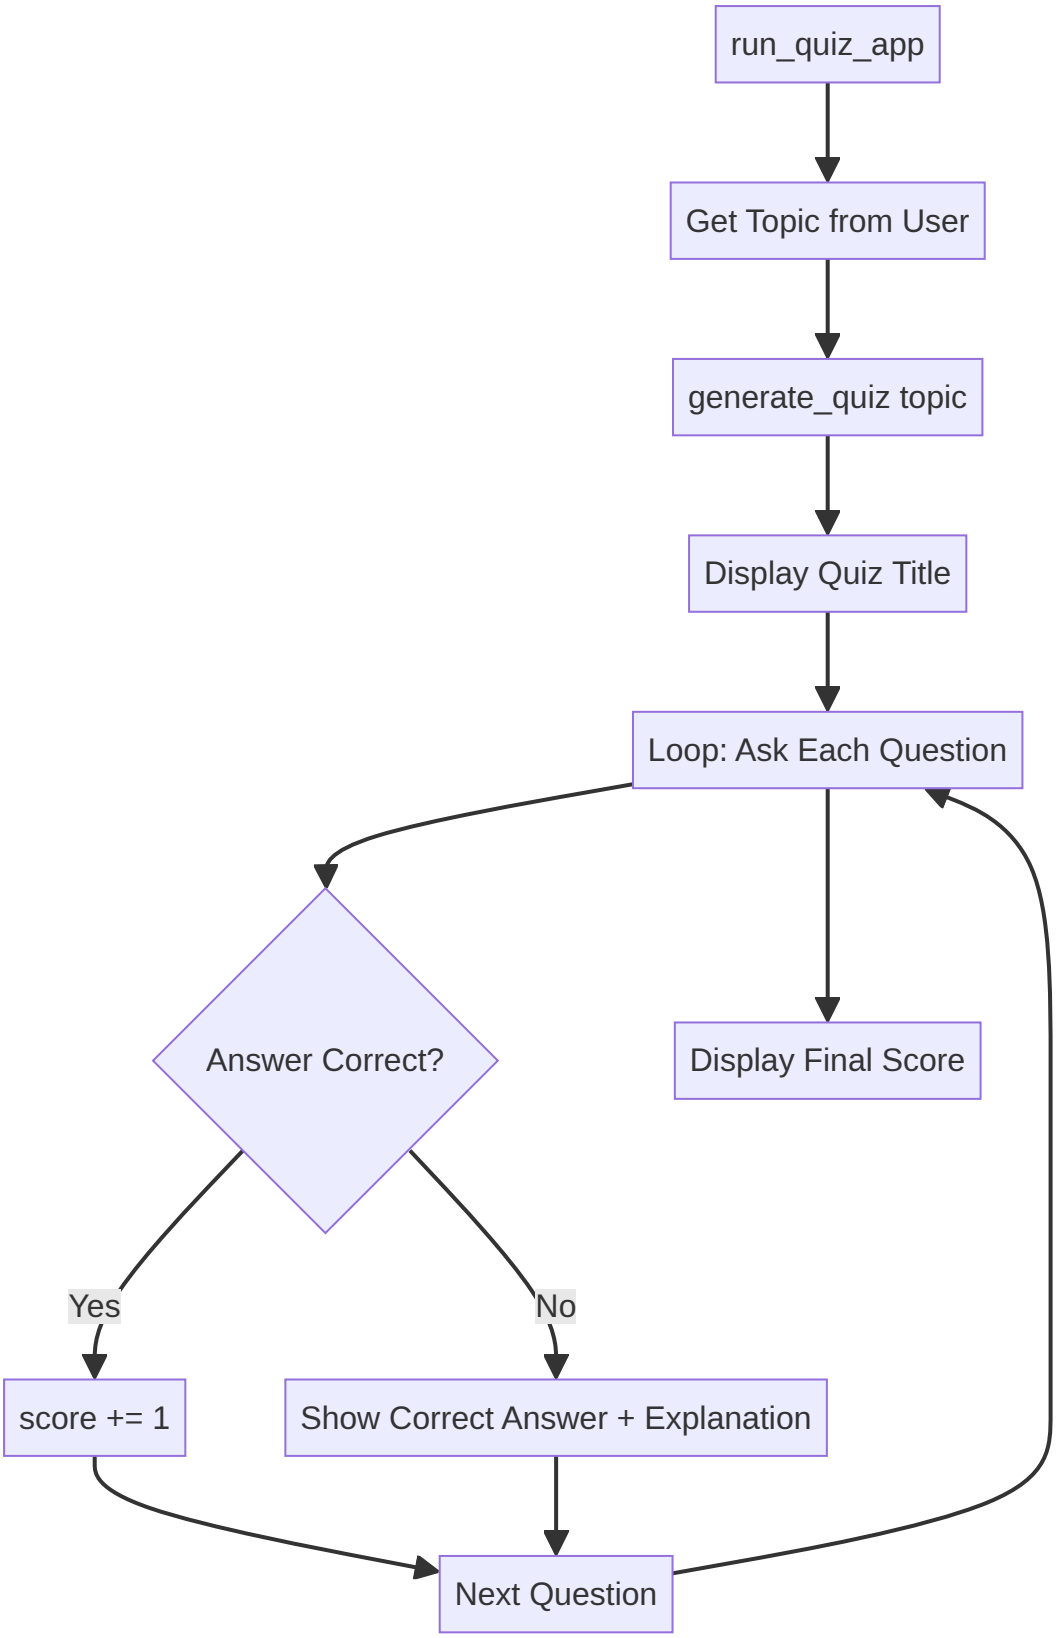
run_quiz_app
Get Topic from User
generate_quiz topic
Display Quiz Title
Loop: Ask Each Question
Answer Correct?
Display Final Score
Yes
No
score += 1
Show Correct Answer + Explanation
Next Question

Step-by-Step Logic

1. User Input

```
topic = input("What topic do you want to be quizzed on? ")
# topic = "mathematics"  # hardcoded alternative for testing
```

2. Generate Quiz

```
quiz_data = generate_quiz(topic)
if not quiz_data:
    print("Failed to generate quiz.")
    return
```

3. Display Quiz

```
print(quiz_data["quiz_title"])
score = 0
total = len(quiz_data["questions"])
```

`total` is computed dynamically — changing question count in the prompt does not require code changes.

4. Question Loop

```
for q in quiz_data["questions"]:
    print(q["text"])
    for key, value in q["options"].items():
        print(f"  {key}: {value}")

    user_answer = input("Your answer (A/B/C/D): ").strip().upper()

    if user_answer == q["correct_option"]:
        print("Correct!")
        score += 1
    else:
        print(f"Wrong. The answer was {q['correct_option']}.")
        print(f"Explanation: {q['explanation']}")

    time.sleep(1)  # pause between questions
```

5. Final Score

```
print(f"Final score: {score}/{total}")

if score == total:
    print("Perfect score! You are a master!")
elif score > 0:
    print("Good job! Keep studying.")
else:
    print("Better luck next time.")
```

JSON Structure Reference

Understanding the data shape is essential for writing the loop:

```
{
  "quiz_title": "Math Master's Challenge",
  "questions": [
    {
      "id": 1,
      "text": "What is the integral of x^2?",
      "options": {"A": "...", "B": "...", "C": "...", "D": "..."},
      "correct_option": "A",
      "explanation": "The integral of x^2 is x^3/3 + C."
    }
  ]
}
```

Access Pattern	Meaning
`quiz_data["quiz_title"]`	Quiz name
`quiz_data["questions"]`	List of question dicts
`q["text"]`	Question text
`q["options"]`	Dict of A/B/C/D options
`q["correct_option"]`	Correct label
`q["explanation"]`	Why the answer is correct

Common Pitfalls / Exam Traps

- **Hardcoding `total = 3`** — use `len(quiz_data["questions"])` for dynamic question counts.
- **Case-sensitive answer comparison** — normalise with `.strip().upper()`.
- **Not handling failed quiz generation** — check for `None` or empty response before looping.
- **Accessing wrong JSON keys** — `correct_option` not `answer` or `correct_answer` (depends on prompt schema).
- **Skipping explanation on wrong answers** — explanations aid learning; always display them.

Quick Revision Summary

- `run_quiz_app()`: input topic → generate → loop questions → score → feedback.
- `total = len(quiz_data["questions"])` keeps question count dynamic.
- Normalise user answers with `.strip().upper()`.
- Correct: increment score; wrong: show correct option and explanation.
- Final messages based on score/total ratio.
- `time.sleep(1)` adds pause between questions for readability.

← **Previous** · **Index** · **Next** →
Demo: Notebook Quiz Application

Demo: Notebook Quiz Application

Running the POC

The notebook proof-of-concept runs with a single function call:

```
run_quiz_app()
```

The terminal prompts for a topic, generates a quiz via Gemini, and runs the interactive question loop.

Demo Walkthrough: Mathematics

```
What topic do you want to be quizzed on? mathematics
Generating a quiz on topic: mathematics

Math Master's Challenge

Q1: What is the value of the integral of ...?
Your answer: D
Wrong. The answer was A. [explanation]

Q2: If matrix A has eigenvalues 1, 2, 3, ...?
Your answer: B
Correct! [explanation]

Q3: [question text]
Your answer: [option]
Correct!

Final score: 2/3
Good job! Keep studying.
```

Demo Walkthrough: Science

```
What topic do you want to be quizzed on? science

The Intermediate Science Challenge
[3 questions with answers and feedback]

Final score: 1/3
Good job! Keep studying.
```

What the Demo Validates

Capability	Evidence
Dynamic topic handling	Different subjects produce different quizzes
LLM-generated titles	"Math Master's Challenge", "Intermediate Science Challenge"
JSON parsing	Questions, options, and answers render correctly
Scoring logic	Correct/incorrect tracking with final ratio
Explanations	Shown after each answer
Real-time interaction	Input → generation → quiz → feedback in one session

---

Known Limitations

- **Factual accuracy not guaranteed** — LLM-generated questions may contain errors
- **Single-turn only** — no conversation memory between runs
- **CLI only** — requires Python familiarity to interact
- **No input validation** — notebook POC accepts any topic string

These limitations motivate the full-stack upgrade with validation, UI, and disclaimers.

Next Steps After Demo

1. Experiment with different topics (cloud computing, NLP, history)
2. Modify prompt template (difficulty, question count)
3. Adapt the template to your own project idea
4. Proceed to full-stack Streamlit implementation

Common Pitfalls / Exam Traps

- **Assuming LLM-generated answers are always correct** — verify critical content; add disclaimers in production.
- **Not testing multiple topics** — single-topic success does not prove robustness.
- **Ignoring JSON parse failures** — some topics may cause malformed output; handle gracefully.
- **Treating CLI demo as user-facing product** — full-stack UI needed for non-technical users.

Quick Revision Summary

- Run `run_quiz_app()` to start the interactive CLI quiz.
- User enters topic; Gemini generates titled MCQ quiz in real time.
- Each question: display options, accept A/B/C/D, show correct/incorrect + explanation.
- Final score: X/Y with motivational message based on performance.
- Demo validates end-to-end: prompt → API → JSON → game loop → scoring.
- LLM content may be factually incorrect — treat as educational demo, not authoritative source.

← **Previous** · **Index** · **Next** →
From Notebook POC to Full-Stack Application

From Notebook POC to Full-Stack Application

Why Upgrade Beyond the Notebook?

The Jupyter notebook POC works for Python practitioners but excludes non-technical users. Industrial LLM products — ChatGPT, Gemini, Claude — became viral only when wrapped in accessible frontends, not when exposed as scripting APIs.

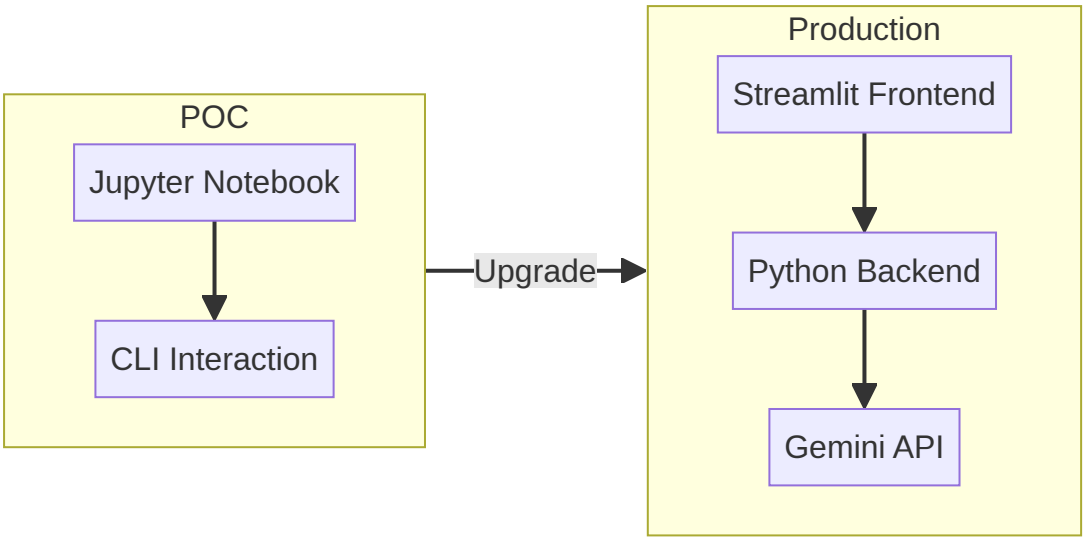

The Full-Stack Motivation

Notebook POC	Full-Stack App
Python users only	Any user via browser
Single file	Separated frontend + backend
Hardcoded secrets risk	`.env` secrets management
No input validation	Validated inputs before API calls
CLI text interface	Interactive forms, buttons, visual feedback

Portfolio impact: A full-stack LLM application demonstrates end-to-end capability — prompting, API integration, validation, UI, and deployment — far beyond notebook-only projects.

Architecture Components

Component	File	Technology
Frontend	`app.py`	Streamlit
Backend	`quiz_engine.py`	Python + Gemini SDK
Secrets	`.env`	`GEMINI_API_KEY`, `MODEL_NAME`
Dependencies	`requirements.txt`	Pinned versions
Docs	`README.md`	Project documentation
Ignore rules	`.gitignore`	Excludes `.env`, caches

Streamlit: Python Frontend Without HTML/CSS

Streamlit lets Python developers build browser UIs without HTML, CSS, or JavaScript knowledge. For NLP practitioners and AI engineers, this bridges the gap to user-facing applications. Frontend engineering teams can later replace Streamlit with React/Vue for polished production UIs.

Industry Workflow

1. **AI engineer** builds logic + Streamlit prototype
2. **Software engineering team** refines frontend for production
3. **DevOps** handles deployment, secrets, monitoring

You do not need to learn frontend frameworks to demonstrate a complete LLM product.

The Broader Lesson

GPT and transformer models existed years before ChatGPT's consumer launch. The **application layer** — accessible UI, session management, feedback loops — turned research into a product revolution. Building QuizGenius AI as full-stack follows the same pattern.

Common Pitfalls / Exam Traps

- **Stopping at notebook POC for portfolio** — full-stack demo is significantly more impressive.
- **Assuming users will run Python scripts** — non-technical users need browser UIs.
- **Learning React before shipping** — Streamlit is sufficient for prototypes and portfolios.
- **Ignoring the product wrapper** — LLM capability alone does not make a usable product.

Quick Revision Summary

- Notebook POC limits access to Python users; full-stack opens access to everyone.
- ChatGPT succeeded because of accessible UI, not just model capability.
- Full-stack: Streamlit frontend (`app.py`) + Python backend (`quiz_engine.py`).
- Streamlit enables Python-only frontend development without HTML/CSS.
- Industry pattern: AI engineer prototypes → frontend team polishes.
- Full-stack LLM apps are high-value portfolio pieces.

← **Previous** · **Index** · **Next** →
Application Flow and Architecture (v1)

Application Flow and Architecture (v1)

User Journey

QuizGenius AI follows a two-phase interaction loop:

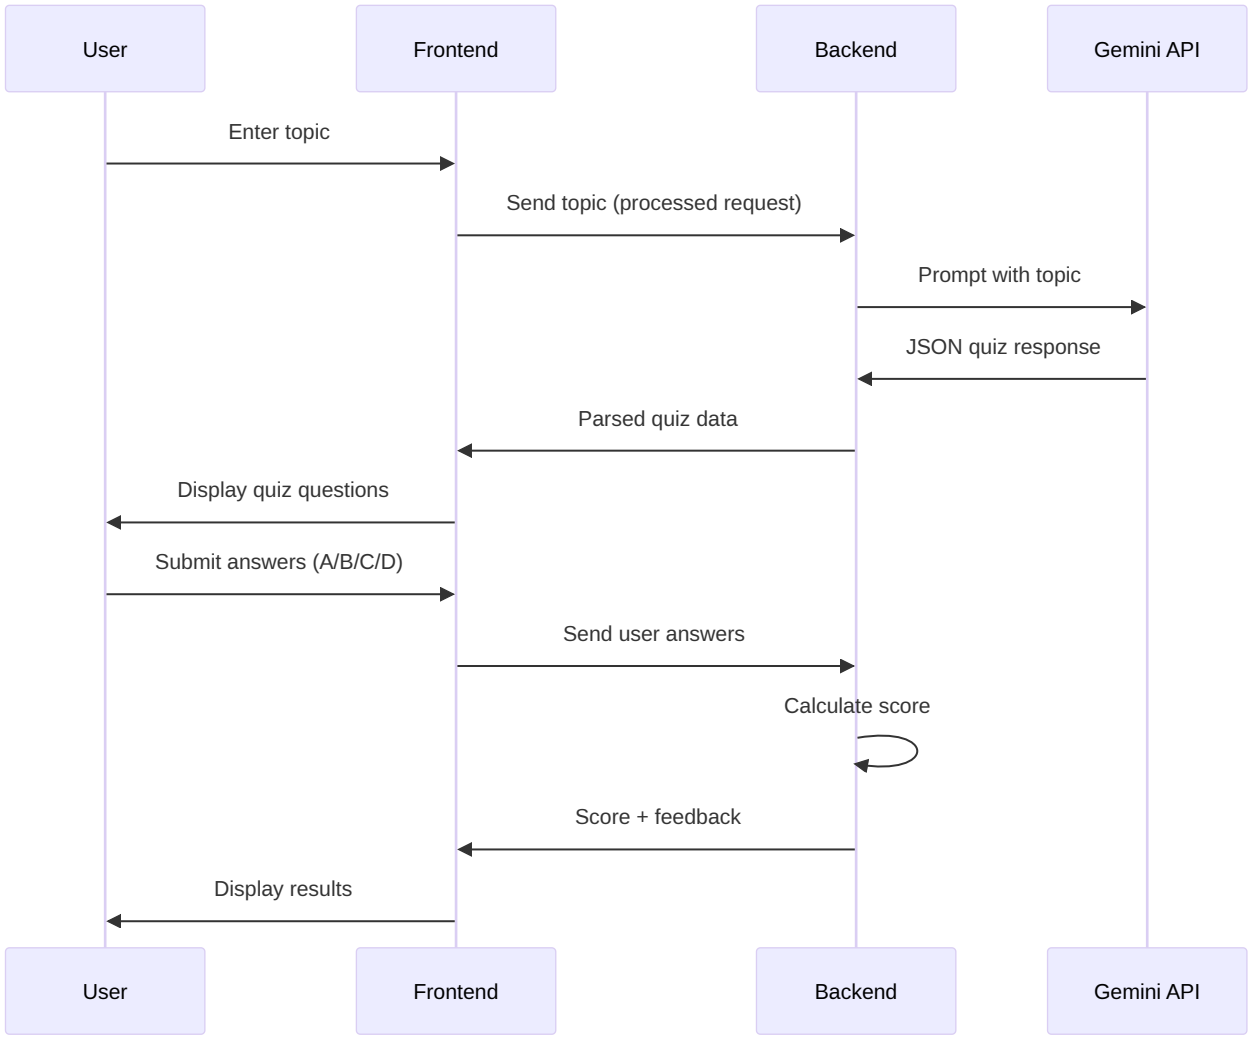

Phase 1: Quiz Generation

1. User enters a topic (e.g., "mathematics")
2. Frontend captures and forwards to backend
3. Backend calls Gemini LLM with the prompt template
4. LLM returns JSON quiz data
5. Backend parses and sends to frontend
6. Frontend displays quiz title and questions

Phase 2: Answer Evaluation

1. User selects answers for each question
2. Frontend sends answers to backend
3. Backend compares against `correct_option` fields
4. Backend calculates score and prepares feedback
5. Frontend displays score, per-question review, and explanations

The loop repeats when the user requests a new quiz.

Architecture Diagram (v1)

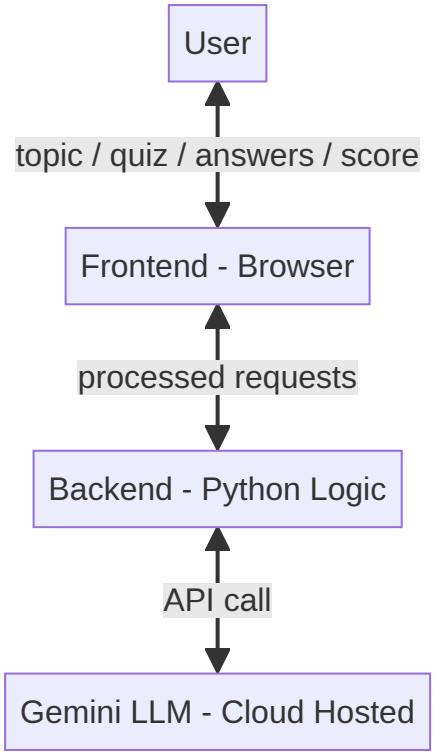

Important: The LLM is **not inside the backend**. It is hosted on Google's cloud and accessed via API. The backend is a client, not the model host.

Component Responsibilities

Component	Responsibility
Frontend	User interaction, form input, quiz display, answer collection, results rendering
Backend	Input validation, prompt construction, API call, JSON parsing, score calculation
Gemini API	Quiz content generation (questions, options, explanations)

Request/Response Flow Detail

Step	Direction	Payload
1	User → Frontend	Topic string
2	Frontend → Backend	Validated topic
3	Backend → LLM	Formatted prompt
4	LLM → Backend	JSON quiz
5	Backend → Frontend	Parsed quiz object
6	Frontend → User	Quiz UI
7	User → Frontend	Selected answers
8	Frontend → Backend	Answer dict
9	Backend → Frontend	Score + per-question feedback
10	Frontend → User	Results screen

Architecture Evolution

Real projects start with a base architecture (v1) and refine after implementation:

- Add input validation layers
- Introduce `.env` for secrets
- Split into `app.py` and `quiz_engine.py`
- Add session state management in Streamlit
- Update architecture diagram to v2 after build

Common Pitfalls / Exam Traps

- **Drawing LLM inside the backend box** — it is an external cloud API.
- **Skipping the two-phase flow** — generation and scoring are separate interactions.
- **Assuming frontend calls Gemini directly** — backend should mediate API calls to protect keys.
- **Not documenting architecture before coding** — leads to tangled responsibilities.
- **Treating v1 architecture as final** — iterate after implementation.

Quick Revision Summary

- Two-phase flow: (1) topic → quiz generation, (2) answers → scoring.
- Frontend handles UI; backend handles logic and API calls.
- Gemini LLM is external cloud service accessed via API, not embedded in backend.
- Loop repeats for each new quiz session.
- Architecture v1 is a starting point — update after implementation (v2).
- Backend parses JSON; frontend displays structured quiz data.

← **Previous** · **Index** · **Next** →
Project Setup and Repository Structure

Project Setup and Repository Structure

Repository Initialisation

Create a GitHub repository for version control, collaboration, and portfolio visibility.

Repository Settings

Setting	Value
Name	`quiz-genius-ai`
Visibility	Public (or private)
README	Yes — project documentation
`.gitignore`	Python template
Licence	MIT (common for open source)

```
git clone <repository-url>
cd quiz-genius-ai
```

Project Structure

```
quiz-genius-ai/
├── .env                    # API keys and secrets (NEVER commit)
├── .gitignore              # Excludes .env, caches, IDE files
├── LICENSE                 # MIT licence
├── README.md               # Documentation (updated as project evolves)
├── requirements.txt        # Pinned Python dependencies
├── quiz_engine.py          # Backend: quiz generation + validation
└── app.py                  # Frontend: Streamlit UI
```

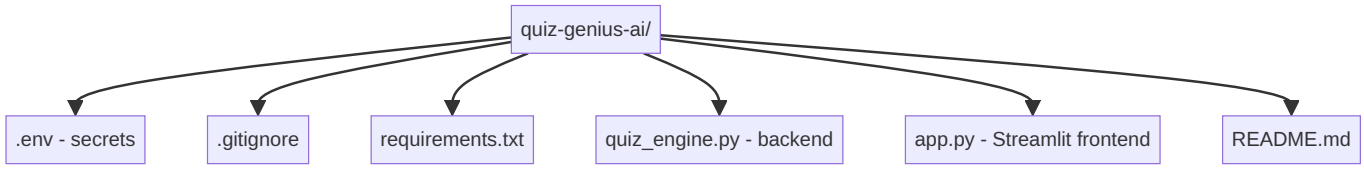

File Responsibilities

File	Purpose
`.env`	`GEMINI_API_KEY`, `MODEL_NAME` — loaded by `python-dotenv`
`.gitignore`	Prevents `.env`, `__pycache__/`, `.venv/` from being committed
`requirements.txt`	`streamlit`, `google-genai`, `python-dotenv` with pinned versions
`quiz_engine.py`	Validation, prompt template, `generate_quiz()`, constants
`app.py`	Streamlit pages: setup, quiz, results
`README.md`	Setup instructions, architecture diagram, usage guide

Secrets Management

```
# .env (add to .gitignore)
GEMINI_API_KEY=your_key_here
MODEL_NAME=gemini-2.5-flash
```

The `.gitignore` Python template already excludes `.env`. Never upload API keys to GitHub — revoked keys and billing surprises follow quickly.

Virtual Environment

```
conda create -n quiz-genius python=3.11 -y
conda activate quiz-genius
pip install -r requirements.txt
```

Pin versions in `requirements.txt` after installation:

```
streamlit==<version>
google-genai==<version>
python-dotenv==<version>
```

Anyone cloning the repo can reproduce the exact environment.

Running the Application

```
conda activate quiz-genius
cd quiz-genius-ai
streamlit run app.py
```

Common Pitfalls / Exam Traps

- **Committing `.env` to GitHub** — even private repos leak; use `.gitignore`.
- **No `requirements.txt`** — others cannot reproduce your environment.
- **Mixing frontend and backend in one file** — separation enables testing and maintenance.
- **Skipping README** — essential for portfolio and collaboration.
- **Not pinning dependency versions** — `pip install` without versions causes "works on my machine" failures.

Quick Revision Summary

- GitHub repo with README, `.gitignore` (Python), MIT licence.
- Key files: `.env`, `quiz_engine.py` (backend), `app.py` (Streamlit frontend).
- `.env` stores `GEMINI_API_KEY` and `MODEL_NAME` — never committed.
- `requirements.txt` with pinned versions for reproducibility.
- Virtual environment: `conda create -n quiz-genius python=3.11`.
- Run: `streamlit run app.py` after activating environment.

← Previous · Index · Next →
Building the Backend (`quiz_engine.py`)

Building the Backend (`quiz_engine.py`)

Backend Responsibilities

The backend module handles everything between user input and LLM response:

1. Load environment variables and initialise Gemini client
2. Validate user inputs before API calls
3. Format the prompt template
4. Call Gemini API with JSON response mode
5. Parse, clean, and validate the JSON response

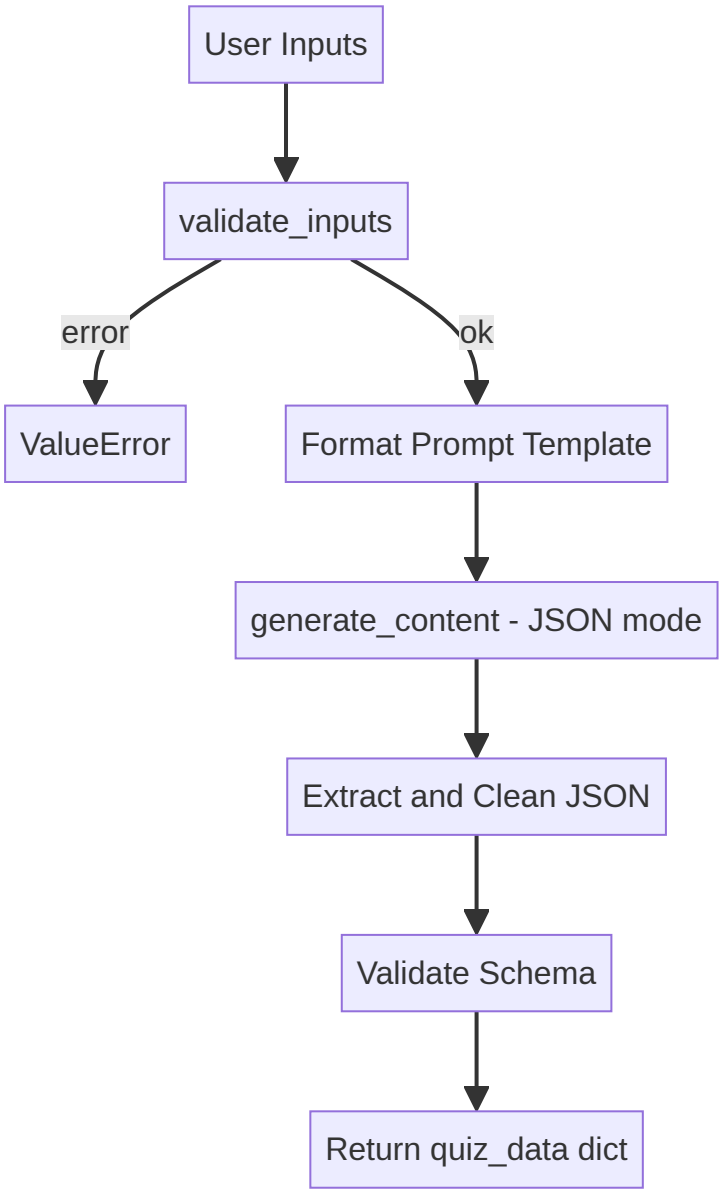

Environment and Client Setup

```
import json, re, os
from dotenv import load_dotenv
from google import genai
from google.genai import types

load_dotenv()
gemini_api_key = os.getenv("GEMINI_API_KEY")
model_name = os.getenv("MODEL_NAME", "gemini-2.5-flash")

if not gemini_api_key:
    raise ValueError("GEMINI_API_KEY is missing. Please check your .env file.")

client = genai.Client(api_key=gemini_api_key)
```

Input Validation

Validation runs **before** the LLM call to save API cost and ensure pipeline consistency.

Constants

```
VALID_DIFFICULTIES = ["easy", "intermediate", "hard"]
MAX_TOPIC_LENGTH = 200
MAX_QUESTIONS = 10
MIN_QUESTIONS = 1
```

`validate_inputs(topic, num_questions, difficulty)`

Check	Error Message
Empty topic	"Topic cannot be empty"
Topic > 200 chars	"Topic is too long. Please keep it under 200 characters"
Questions not in [1, 10]	"Number of questions must be between 1 and 10"
Invalid difficulty	"Difficulty must be one of: easy, intermediate, hard"
All valid	Returns `None`

`generate_quiz(topic, num_questions=3, difficulty="intermediate")`

```
def generate_quiz(topic: str, num_questions: int = 3,
                  difficulty: str = "intermediate") -> dict:
    error = validate_inputs(topic, num_questions, difficulty)
    if error:
        raise ValueError(error)

    prompt = prompt_template.format(
        topic=topic.strip(),
        num_questions=num_questions,
        difficulty=difficulty
    )

    try:
        response = client.models.generate_content(
            model=model_name,
            contents=prompt,
            config=types.GenerateContentConfig(
                response_mime_type="application/json"
            )
        )
        raw_text = response.text
        match = re.search(r'\{.*\}', raw_text, re.DOTALL)
        json_string = match.group(0) if match else raw_text
        json_string = re.sub(r'```json|```', '', json_string).strip()
        quiz_data = json.loads(json_string)

        if "questions" not in quiz_data:
            raise ValueError("LLM response missing 'questions' key")

        return quiz_data

    except json.JSONDecodeError as e:
        raise ValueError(f"Failed to parse quiz JSON: {e}")
    except Exception as e:
        raise RuntimeError(f"Quiz generation failed: {e}")
```

Why Validation Matters

Benefit	Explanation
Cost savings	Invalid inputs never reach the paid API
Consistency	Gibberish topics and bizarre question counts are rejected
Security	Length limits prevent prompt injection via oversized inputs
Reliability	Schema check catches malformed LLM responses early

---

Common Pitfalls / Exam Traps

- **Skipping validation and calling LLM directly** — wastes tokens and money.
- **No schema validation after JSON parse** — missing `questions` key crashes frontend.
- **Missing `except` block for try** — incomplete error handling causes unhandled crashes.
- **Hardcoding model name in code** — use `.env` with fallback default.
- **Not stripping topic whitespace** — leading spaces produce odd quiz content.

Quick Revision Summary

- Backend: env setup → validation → prompt format → API call → JSON parse → schema check.
- `validate_inputs()` checks topic, question count, and difficulty before API call.

- `generate_quiz()` returns dict with `quiz_title` and `questions` or raises errors.
- `response_mime_type="application/json"` enforces structured output.
- Validate `"questions"` key exists in parsed response.
- Handle `JSONDecodeError` and general exceptions separately.

← **Previous** · **Index** · **Next** →
Building the Frontend (`app.py`)

Building the Frontend (`app.py`)

Streamlit Overview

Streamlit (`st`) enables Python developers to build browser UIs without HTML/CSS. QuizGenius AI uses three screens managed by **session state** — because Streamlit reruns the entire script on every user interaction.

Page Configuration

```
st.set_page_config(
    page_title="QuizGenius AI",
    page_icon="🧠",
    layout="centered"
)
```

Setting	Effect
`page_title`	Browser tab title
`page_icon`	Favicon
`layout="centered"`	Content centred with side padding (vs `wide`)

Session State Management

Streamlit reruns the full script on every button click or input change. Session state persists data across reruns.

```
def init_session_state():
    defaults = {
        "quiz_data": None,       # Generated quiz JSON
        "user_answers": {},      # {question_index: "A"/"B"/"C"/"D"}
        "submitted": False,      # Whether answers submitted
        "current_step": "setup"  # "setup" | "quiz" | "results"
    }
    for key, value in defaults.items():
        if key not in st.session_state:
            st.session_state[key] = value
```

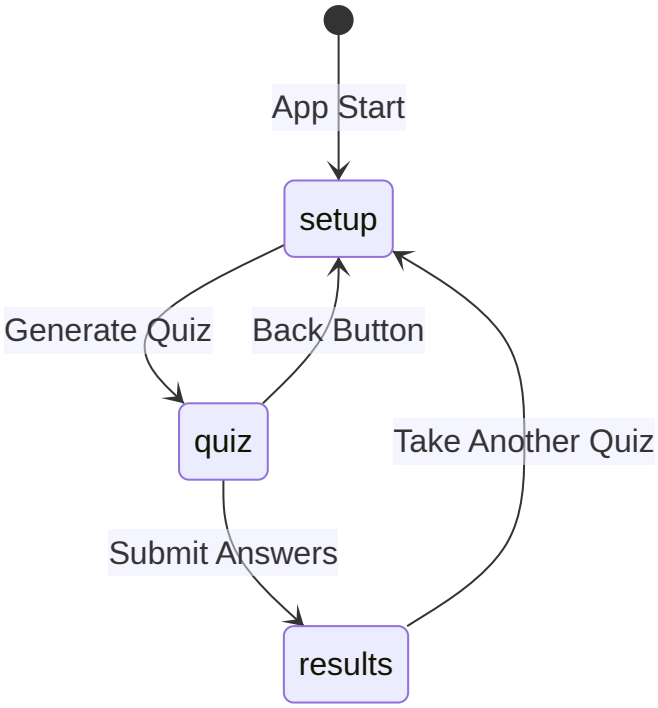

Screen 1: Setup Form (`render_setup_form`)

UI Element	Streamlit Widget	Purpose
Header	`st.header`	"Configure Your Quiz"
Topic	`st.text_input`	Free-text topic entry
Question count	`st.slider`	MIN_QUESTIONS to MAX_QUESTIONS (default 3)
Difficulty	`st.selectbox`	easy / intermediate / hard
Submit	`st.form_submit_button`	"Generate Quiz!"

On submit:

1. Validate topic is not empty
2. Show spinner: "Generating your quiz..."
3. Call `generate_quiz(topic, num_questions, difficulty)` from backend
4. Store result in `st.session_state.quiz_data`
5. Set `current_step = "quiz"` and `st.rerun()`

Screen 2: Quiz Page (`render_quiz`)

- Display `quiz_data["quiz_title"]` as header
- Loop through questions with `st.radio` for A/B/C/D selection
- Store selections in `st.session_state.user_answers`
- **Submit Answers** button (width ratio 3) — validates all questions answered
- **Back** button (width ratio 1) — calls `reset_quiz()` and returns to setup

---

Screen 3: Results Page (`render_results`)

Score Calculation

```
score = sum(
    1 for i, q in enumerate(questions)
    if user_answers.get(i) == q["correct_option"]
)
percentage = (score / total) * 100
```

Feedback Tiers

Percentage	Message
100%	Trophy emoji — "Perfect score! You're a master!"
> 70%	"Great job! You know your stuff!"
> 40%	"Good effort! Keep studying!"
≤ 40%	"Keep learning!"

Detailed Review

Per question:

- Correct answer: green checkmark (`st.success`)
- Wrong answer: red cross (`st.error`) + correct option shown
- Explanation from `q["explanation"]` via `st.info`

Take Another Quiz button resets state and returns to setup.

Main App Router (`main()`)

```
def main():
    init_session_state()
    st.title("QuizGenius AI")
    st.caption("Generate AI-powered quizzes on any topic in seconds.")

    if st.session_state.current_step == "setup":
        render_setup_form()
    elif st.session_state.current_step == "quiz":
        render_quiz()
    elif st.session_state.current_step == "results":
        render_results()
```

Common Pitfalls / Exam Traps

- **Not using session state** — quiz data lost on every rerun.
- **Forgetting `st.rerun()` after state change** — UI does not transition to next screen.
- **Hardcoding 3 questions in UI** — use `len(quiz_data["questions"])` dynamically.
- **Not validating all questions answered before submit** — partial submissions skew scoring.
- **Mixing backend logic in `app.py`** — import from `quiz_engine.py` only.

Quick Revision Summary

- Streamlit frontend: 3 screens (setup, quiz, results) routed by `current_step`.
- Session state persists `quiz_data`, `user_answers`, `submitted`, `current_step`.
- Setup: topic input, question slider, difficulty selectbox, generate button.
- Quiz: radio buttons per question, submit/back buttons.
- Results: score percentage, tiered feedback, per-question review with explanations.
- `main()` routes to correct screen based on session state.
- Always call `st.rerun()` after state transitions.

← **Previous** · **Index** · **Next** →
Running the Full-Stack Application

Running the Full-Stack Application

Launch Command

```
conda activate quiz-genius
cd quiz-genius-ai
streamlit run app.py
```

Streamlit starts a local web server (typically `http://localhost:8501`) and opens the browser automatically.

End-to-End Walkthrough

Screen 1: Setup

Element	Action
Topic input	Enter "science" (or any subject)
Question slider	Select 3 (or 1–10)
Difficulty dropdown	Select easy / intermediate / hard
Generate Quiz button	Triggers API call with spinner

Error case: empty topic → "Please enter a topic to generate a quiz."

Screen 2: Quiz

- Header: LLM-generated title (e.g., "The Intermediate Science Challenge")
- Caption: "3 questions — select your answers and submit"
- Per question: radio buttons for A, B, C, D
- Submit Answers (primary) | Back (secondary)

Screen 3: Results

- Score header: "1/3" with emoji
- Percentage and tiered message
- Detailed review per question:
  - Correct: green checkmark
  - Wrong: red cross + correct answer highlighted
  - Explanation for each question

- "Take Another Quiz" → resets to Screen 1

Code-to-UI Mapping

UI Element	Source in `app.py`
"QuizGenius AI" title	`main()` → `st.title()`
"Configure Your Quiz"	`render_setup_form()` → `st.header()`
Topic text box	`st.text_input("Topic", placeholder=...)`
Question slider	`st.slider()` with MIN/MAX from backend
Difficulty dropdown	`st.selectbox()` with VALID_DIFFICULTIES
"Generate Quiz!" button	`st.form_submit_button()` with `use_container_width=True`
Quiz title (dynamic)	`quiz_data["quiz_title"]` in `render_quiz()`
Question text (dynamic)	`f"Q{q['id']}: {q['text']}"`
Radio buttons	`st.radio()` with option labels from `q["options"]`
Score display	`render_results()` → `st.header(f"{emoji} Score: {score}/{total}")`

Testing Multiple Topics

Run the app with different topics to verify:

- Dynamic title generation
- Variable question counts via slider
- Difficulty level affecting question complexity
- Consistent JSON parsing across topics
- Score calculation accuracy

Troubleshooting

Issue	Fix
"GEMINI_API_KEY is missing"	Add key to `.env` file
Import error for `quiz_engine`	Run from project root directory
Streamlit not found	`pip install -r requirements.txt` in activated venv
JSON parse failure	Check prompt template; retry with simpler topic
Page does not transition	Verify `st.rerun()` called after state update

---

Common Pitfalls / Exam Traps

- **Running `python app.py` instead of `streamlit run app.py`** — Streamlit requires its own runner.
- **Not activating virtual environment** — wrong package versions or missing dependencies.
- **Forgetting `.env` file** — backend raises ValueError immediately.
- **Testing only one topic** — edge cases appear with unusual or very long topics.
- **Ignoring container width on buttons** — `use_container_width=True` improves visual layout.

Quick Revision Summary

- Launch: `streamlit run app.py` from project root with venv activated.
- Three screens: setup (form) → quiz (radio buttons) → results (score + review).
- Empty topic shows error; valid topic triggers spinner then quiz display.
- Results show score, percentage, tiered message, and per-question explanations.
- "Take Another Quiz" resets session state to setup.
- Map UI elements to `render_setup_form`, `render_quiz`, `render_results`, and `main()`.

← **Previous** · **Index** · **Next** →
Updated Project Architecture (v2)

Updated Project Architecture (v2)

From v1 to v2

After implementing the full-stack application, the architecture diagram is updated to reflect actual file names, screen flows, backend steps, and secrets management. Real projects routinely revise architecture post-implementation based on optimisation needs and changing requirements.

Architecture v2 Diagram

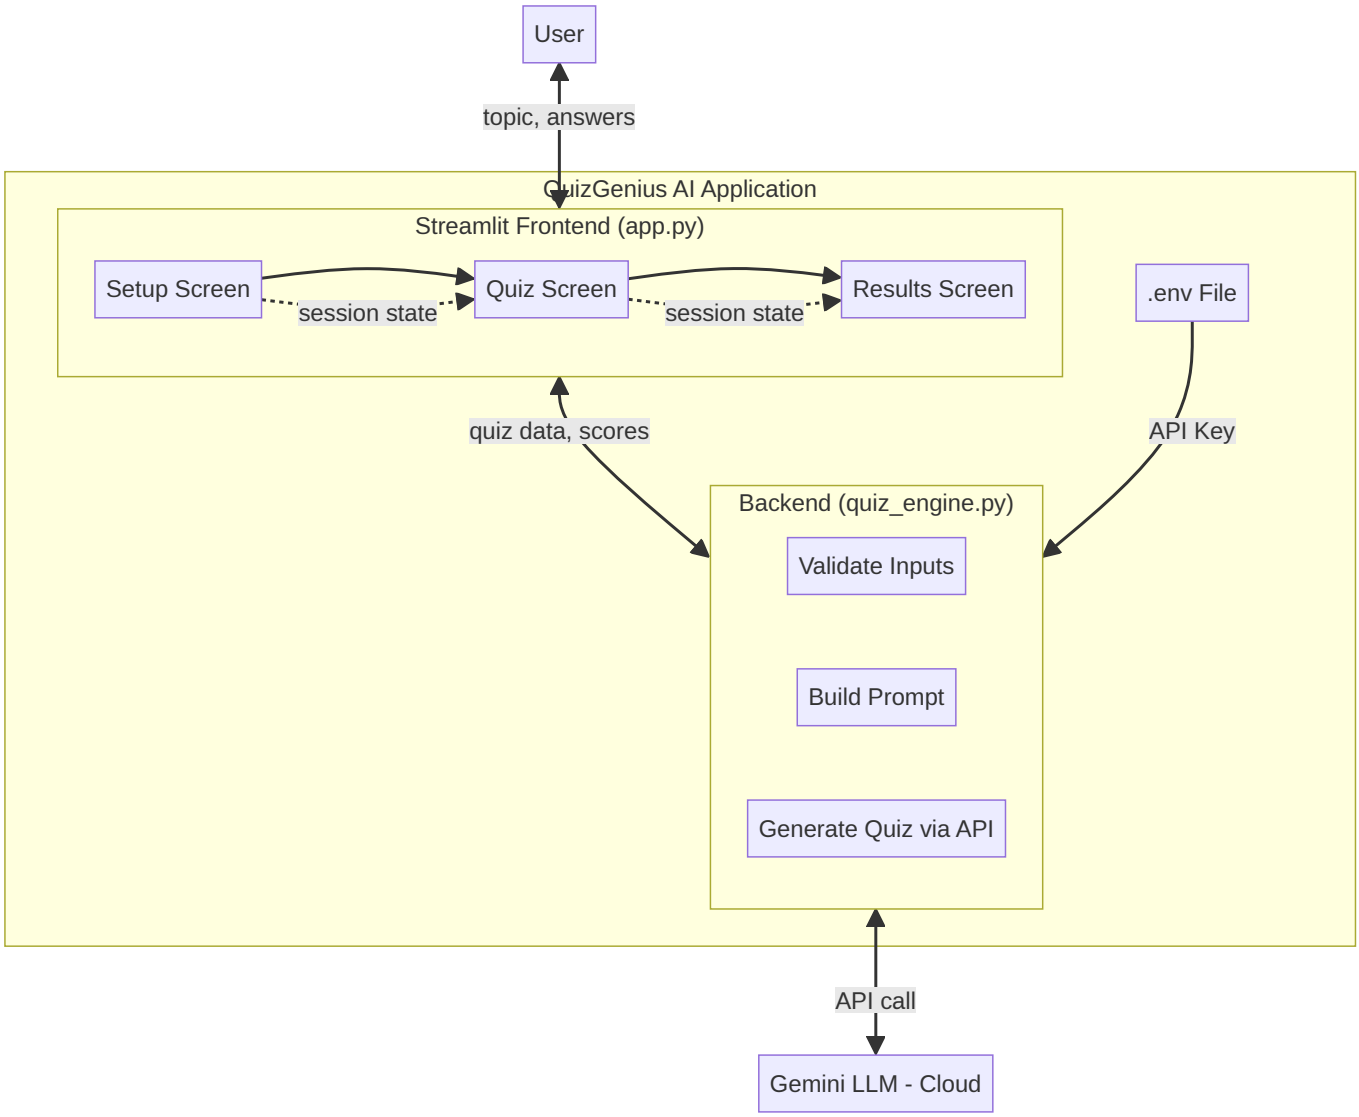

Component Details (v2)

Frontend: `app.py` (Streamlit)

Screen	Function	Session State
Setup	`render_setup_form()`	`current_step = "setup"`
Quiz	`render_quiz()`	`current_step = "quiz"`
Results	`render_results()`	`current_step = "results"`

All screens managed via `st.session_state` — persists data across Streamlit reruns.

Backend: `quiz_engine.py`

Step	Function
1	`validate_inputs()` — topic, question count, difficulty
2	Format `prompt_template` with user parameters
3	`generate_quiz()` — API call + JSON parse + schema validation

API Management: `.env`

```
GEMINI_API_KEY=your_key
MODEL_NAME=gemini-2.5-flash
```

Connected to backend via `python-dotenv` → `os.getenv()`. Never committed to Git.

---

v1 vs v2 Changes

Aspect	v1 (Conceptual)	v2 (Implemented)
Frontend detail	"Browser UI"	Streamlit, 3 named screens
Backend detail	"Python logic"	`quiz_engine.py` with validation pipeline
Secrets	Not shown	`.env` file with API key connection
Session management	Not shown	`st.session_state` with `current_step`
File names	Generic	`app.py`, `quiz_engine.py`

Complete Project Structure (v2)

```
quiz-genius-ai/
├── .env
├── .gitignore
├── LICENSE
├── README.md
├── requirements.txt
├── quiz_engine.py      # Backend logic
└── app.py              # Streamlit frontend (3 screens)
```

Portfolio and Next Steps

- Push completed project to GitHub with updated README and architecture diagram
- Add disclaimer: "AI-generated quizzes may contain inaccuracies"
- Experiment with prompt tweaks (difficulty descriptions, question styles)
- Adapt template to other domains (cloud certification, corporate training)
- Frontend teams can later replace Streamlit with production-grade React UI

Common Pitfalls / Exam Traps

- **Not updating architecture after implementation** — v1 diagrams become misleading documentation.
- **Omitting `.env` from architecture** — secrets management is a core production concern.
- **Drawing frontend as single box** — three screens with session routing is the actual design.
- **Forgetting session state in architecture** — critical for understanding Streamlit apps.
- **Skipping README architecture diagram** — essential for portfolio reviewers.

Quick Revision Summary

- Architecture v2 reflects actual implementation details post-build.
- Frontend: Streamlit `app.py` with 3 session-managed screens.
- Backend: `quiz_engine.py` — validate → prompt → generate → parse.
- `.env` provides API key to backend; never in version control.
- v2 adds file names, screen flow, validation pipeline, and secrets.
- Update architecture diagrams after implementation — industry standard practice.
- Project ready for GitHub portfolio with README and architecture snapshot.

← Previous · Index · Next →

Stories

Stories

Natural Language Processing and Understanding — Story-Based Learning

Natural Language Processing and Understanding — Story-Based Learning

The Complete Mental Model: The Language Detective Agency

> *"Every messy sentence is a crime scene. Every NLP pipeline is a detective turning chaos into evidence. This is your field manual."*

The Master Narrative: You work at **The Language Detective Agency**. Raw human language arrives as unstructured crime scenes — ambiguous, context-dependent, full of red herrings. Your job: **decode language into structured evidence** machines can act on. Each week adds a new division to the agency — from opening case files, to dusting for fingerprints (preprocessing), to tagging suspects (NER), to building semantic maps (embeddings), to deploying senior detectives with memory (RNNs), to the all-seeing spotlight team (Transformers), and finally the press office (LLMs).

PART 1: OPENING THE AGENCY (Week 0)

Course Overview and Practitioner Learning — The Agency Charter

The Story: Before any case lands on your desk, the agency director explains the mission: turn unstructured human language into actionable intelligence. The course arc mirrors a detective's career — crime-scene basics (preprocessing), evidence tagging (POS/NER), case archives (corpora), clue weighting (TF-IDF), semantic maps (embeddings), detectives with memory (RNNs), the spotlight team (Transformers), and the press office (LLMs). Industry practitioners run this agency daily; your job is to know *why* each tool exists, not just which API to call.

Course arc: Foundations → Preprocessing → Sparse vectors → Dense embeddings → Sequential models → Transformers → Applications → LLM capstone.

Tools in the toolkit: NLTK, spaCy, Flair, Hugging Face, Gensim, Gemini/OpenAI APIs.

Practitioner lens: Production trade-offs — latency, cost, maintainability — matter as much as benchmark accuracy.

Exam Tip: NLP is not only ChatGPT; classical preprocessing and sparse methods remain examinable foundations.

> **Course Overview** = The agency charter: NLP turns unstructured language into structured, actionable signals.

PART 2: THE NATURE OF THE CRIME SCENE (Week 1)

Introduction to NLP — Opening the First Case

The Story: The rookie detective's first briefing: human language is the most complex evidence you'll ever handle. This module introduces the agency's core divisions — understanding (NLU), generation (NLG), morphology, and ambiguity — before any tool is deployed. Rule-based lookup without context fails on real cases; every pipeline must account for structure, meaning, and surrounding sentences.

Core questions the module answers: What is NLP? How do NLU and NLG differ? Why is morphology relevant? Why does ambiguity break naive systems?

Exam Tip: NLP is not chatbots only; NLU and NLG are distinct design concerns even when LLMs blur the boundary.

> **NLP Module Introduction** = First briefing — language is complex evidence requiring structured pipelines.

What Is Natural Language Processing? — The Core Mission

The Story: Every day, billions of people type, speak, and message. Computers only see bits and numbers. The agency's core mission — **NLP** — is closing that gap: enabling machines to **read**, **extract**, and **respond** in human language. A crime scene (raw text) is unstructured, ambiguous, and context-dependent. No single regex solves it; you need a multi-stage pipeline.

Formal definition: NLP = AI branch enabling computers to (1) read and interpret text/speech, (2) extract structured information, (3) generate meaningful responses.

Property	Detective impact
Unstructured	No fixed schema — must parse and segment
Ambiguous	Context needed to disambiguate
Context-dependent	Prior sentences change meaning

Exam Tip: NLP includes understanding and extraction, not just generation. ASR converts audio to text; NLP operates on the resulting language.

> **NLP** = The agency's mission: bridge human communication and machine computation.

NLP, NLU, and NLG — Three Divisions in One Building

The Story: The agency has three divisions under one roof. **NLP** is the whole building — any computational work on language. **NLU** (Natural Language Understanding) is the analysis wing: "What does this witness mean?" — sentiment, NER, intent. **NLG** (Natural Language Generation) is the press office: "How do we express our findings?" — summaries, chatbots, reports. LLMs blur the lines, but the design distinction still matters when you architect a system.

Division	Core question	Direction
NLP	Process language computationally?	Umbrella
NLU	What does input **mean**?	Input → structured meaning
NLG	How to **express** information?	Data → human-readable text

Exam Tip: NLU and NLG are subsets of NLP, not disjoint fields. Design systems by asking: do I need to understand, generate, or both?

> **NLU / NLG** = Analysis wing vs press office — both inside the NLP building.

NLP Applications by Domain — Case Files Across Industries

The Story: The agency takes cases from every industry. Search engines need query understanding and document ranking. Voice assistants chain speech recognition with language understanding. Email filters classify spam. Keyboards predict the next word. Chatbots generate fluent replies. Each case combines sub-tasks — tokenisation, classification, generation — in a single pipeline.

Domain	Example	NLP task
Search	Google, Bing	Query understanding, ranking
Voice	Siri, Alexa	ASR + NLU
Email	Spam filter	Text classification
Keyboards	Autocomplete	Language modelling
Generative AI	ChatGPT, Gemini	Understanding + generation

Exam Tip: Link applications to underlying tasks (NER vs classification vs generation). Medical NER needs domain corpora, not general web text.

> **NLP Applications** = Every industry sends the agency a different kind of case file.

Morphology — Dissecting the Evidence

The Story: A detective dissects words like a forensic pathologist. **Morphology** is the study of word structure from **morphemes** — the smallest meaningful units. The **root** carries core meaning (*act* in *acting*). **Affixes** attach: prefixes (*un-*), suffixes (*-ing*). The **stem** is a rough chop (*stud* from *studying*); the **lemma** is the dictionary form (*study*). Stemming is fast but crude; lemmatisation is precise but needs POS tags. Modern transformers use **subword tokenisation** — a hybrid strategy.

Key terms: root, stem, affix, lemma, inflection, derivation, stemming, lemmatisation.

Exam Tip: Stem ≠ lemma. Over-stemming merges unrelated words (*university*/*universe*). Lemmatisation requires POS (*running* verb → *run*, not *running* noun).

> **Morphology** = Forensic word dissection — roots, affixes, stems, and lemmas.

Ambiguity — Polysemy and Synonymy — The Red Herrings

The Story: Every crime scene has red herrings. **Polysemy** is one word, multiple related meanings — *bank* (river vs financial), *run* (sprint vs operate). **Synonymy** is different words, same meaning — *buy*/*purchase*. Stemming cannot fix synonymy; static embeddings give *bank* one vector regardless of sense. **Contextual embeddings** (BERT) resolve polysemy by changing the vector per sentence. Sentence-level ambiguity (*"I saw her duck"*) needs full context.

Definitions:

- **Polysemy:** one word, multiple related senses
- **Synonymy:** different words, similar meaning

Exam Tip: Polysemy ≠ homonymy (unrelated senses). Static Word2Vec fails on polysemy; BERT handles it via context.

> **Polysemy / Synonymy** = Red herrings — one word many meanings, many words one meaning.

PART 3: CRIME-SCENE PROCESSING (Week 2)

Text Preprocessing Foundations — Cleaning the Scene

The Story: Raw text arrives covered in grime — HTML tags, URLs, inconsistent casing, filler words. **Cleaning** removes noise; **preprocessing** transforms text into machine-ready form. These are not the same job, and over-cleaning destroys evidence (negation in sentiment, numbers in finance). LLMs still benefit from corpus cleaning at scale. Every task needs its own pipeline — NER keeps capitalisation; topic modelling removes stopwords.

Exam Tip: Cleaning ≠ preprocessing. Over-cleaning loses negation and punctuation cues. Pipelines are task-specific, not universal.

> **Preprocessing** = Crime-scene cleaning — remove grime without destroying evidence.

Tokenisation — Bagging the Evidence

The Story: Before analysis, detectives bag individual pieces of evidence. **Tokenisation** splits text into **tokens** — words, sentences, or subwords. Whitespace splitting fails on contractions (*don't*). Word tokenisation is wrong for BERT, which uses subword pieces. NLTK and spaCy offer different tokenisers; choose based on downstream model.

Exam Tip: Tokenise before stopword removal. BERT needs subword tokenisation, not simple word split. Download NLTK `punkt` for sentence tokenisation.

> **Tokenisation** = Bagging evidence — split text into processable units.

Noise Removal with Regex — The Fine Brush

The Story: Some grime needs a fine brush, not a bulldozer. **Regular expressions** remove HTML tags, URLs, extra whitespace, and special characters. Pattern `\s+` collapses whitespace — but replacing with empty string destroys word boundaries. Case normalisation helps matching but loses named-entity cues. Domain matters: stripping numbers destroys financial and medical evidence.

Key patterns: `\s+` (whitespace), `<.*?>` (HTML tags), `https?://\S+` (URLs).

Exam Tip: Over-cleaning is as dangerous as under-cleaning. Never strip numbers in finance/medical domains without reason.

> **Regex cleaning** = Fine brush — precise noise removal without destroying boundaries.

Stopword Removal — Filtering the Noise

The Story: Function words (*the*, *is*, *and*) are like footprints everyone leaves — they tell you someone was here but not who. **Stopwords** carry little discriminative power. Removing them shrinks feature space for classification and topic modelling. But for sentiment, removing *not* destroys meaning. TF-IDF partially handles stopwords automatically; explicit removal still helps.

Exam Tip: Never remove negation words for sentiment analysis. Tokenise before removing stopwords. Stopword lists are language- and domain-specific.

> **Stopwords** = Common footprints — filter them when they add no clue.

Stemming and Lemmatisation — Normalising the Clues

The Story: Detectives normalise clues for comparison. **Stemming** chops suffixes with rules (Porter stemmer: *studying* → *studi*) — fast but produces invalid forms. **Lemmatisation** looks up dictionary forms (*studying* → *study*) — accurate but needs POS tags and is slower. Porter is English-centric. For short-text topic modelling (tweets), stemming helps; for NER, avoid aggressive normalisation.

Method	Speed	Accuracy	Needs POS
Stemming	Fast	Crude	No
Lemmatisation	Slower	Precise	Yes

Exam Tip: *Stems* may be invalid words. Lemmatisation needs POS (*running* verb → *run*). Download NLTK `wordnet` and `omw-1.4`.

> **Stemming / Lemmatisation** = Normalise word forms — rough chop vs dictionary lookup.

Preprocessing Pipeline Construction — The Standard Operating Procedure

The Story: Every detective follows a standard operating procedure. The canonical pipeline: **clean** (regex) → **tokenise** → **(remove stopwords)** → **(stem or lemmatise)**. Order matters: tokenise before cleaning destroys boundaries; remove stopwords before sentiment destroys negation. The same pipeline is wrong for NER (keep entities) vs topic modelling (remove stopwords).

```
Raw text → Clean → Tokenise → Stopwords → Stem/Lemmatise → Model input
```

Exam Tip: Pipelines are task-specific. Document your order; exam questions test sequencing errors.

> **Preprocessing pipeline** = Standard operating procedure — order matters.

PART 4: EVIDENCE TAGGING (Week 3)

Structural Text Analysis Overview — The Tagging Division

The Story: The Tagging Division assigns structure to raw evidence. **POS tagging** labels grammatical roles (noun, verb, adjective). **NER** identifies real-world entities (persons, organisations, locations). These are different jobs — POS tells you *what role* a word plays; NER tells you *who or what* it refers to. NLTK, spaCy, and Flair each have strengths; none is universally best.

Exam Tip: POS ≠ NER. POS is context-dependent (*book* as noun vs verb). Tool choice depends on speed vs accuracy trade-offs.

> **Structural analysis** = Tagging division — grammatical roles and named entities.

Part-of-Speech Tagging — Grammatical Role Labels

The Story: Every word at a crime scene plays a role — witness, suspect, location. **POS tagging** assigns grammatical categories: noun (NN), verb (VB), adjective (JJ). Tags are **context-dependent**: *book* is a noun in "read a book" but a verb in "book a flight." Penn Treebank and Universal POS tag sets differ across libraries. POS tags are required for accurate lemmatisation.

Exam Tip: POS ≠ NER. Tag sets differ (NLTK Penn vs spaCy Universal). Context determines the tag for ambiguous words.

> **POS tagging** = Assign grammatical roles — noun, verb, adjective, context-dependent.

POS Tagging in Practice — spaCy, NLTK, and Flair

The Story: Three junior detectives handle POS tagging differently. **spaCy** uses `token.pos_` (coarse) and `token.tag_` (fine-grained) — fast, production-ready. **NLTK** uses `pos_tag()` after tokenisation — modular, educational. **Flair** wraps text in `Sentence` objects and calls `SequenceTagger.load('pos')` — highest accuracy, slowest. All require loading models or downloading data first.

Key APIs:

- spaCy: `token.pos_`, `token.tag_`

- NLTK: `nltk.pos_tag(tokens)`
- Flair: `tagger.predict(sentence)`

Exam Tip: Load spaCy model before use. Download NLTK `averaged_perceptron_tagger`. Flair requires `Sentence` wrapper.

> **POS in practice** = Three detectives, three toolkits — spaCy for speed, Flair for accuracy.

Visualising POS and Dependency Trees — The Evidence Board

The Story: Detectives pin evidence on a board to see connections. **spaCy displacy** renders POS tags and **dependency trees** — who modifies whom (*nsubj*, *dobj*, *amod*). Use `style='dep'` for syntactic trees, `style='ent'` for named entities (not POS). Dependency labels show grammatical relationships beyond simple tags.

Exam Tip: `'ent'` style is for NER, not POS. Dependency labels (nsubj, dobj) ≠ POS tags.

> **Dependency trees** = Evidence board — visual grammar and entity connections.

Named Entity Recognition — Identifying the Suspects

The Story: NER is the suspect-identification unit. It finds **named entities** — real-world objects with proper names: PERSON, ORG, GPE, LOC, DATE, MONEY. Entities can span multiple tokens (*New York*). Not every noun is an entity. NER is the primary step in information extraction — pulling structured facts from unstructured text.

Common entity types: PERSON, ORG, GPE (geo-political), LOC, DATE, TIME, MONEY, PERCENT.

Exam Tip: NER ≠ POS. Tag sets differ across libraries (PER vs PERSON). Multi-token spans need special handling.

> **NER** = Suspect identification — find persons, organisations, locations in text.

NER in Practice — spaCy, NLTK, and Flair

The Story: Three tools, three accuracy levels. **spaCy** uses `doc.ents` and `ent.label_` — industry standard, fast. **NLTK** uses chunk-based `ne_chunk` — educational, fewer entity types (no MONEY/DATE). **Flair** loads `SequenceTagger.load('ner')` — highest accuracy, slowest. All require model loading; Flair needs `predict()` call.

Exam Tip: spaCy uses `doc.ents`, not `doc.entities`. NLTK accuracy gap vs spaCy. Call `predict()` on Flair tagger.

> **NER in practice** = Three suspect-ID teams — spaCy for production, Flair for accuracy.

Comparing NLTK, spaCy, and Flair — Choosing Your Detective

The Story: The agency maintains three detective ranks. **NLTK** is the trainee — slowest, lowest accuracy, best for learning linguistics. **spaCy** is the field agent — fast, good accuracy, production pipelines. **Flair** is the specialist — slowest, highest accuracy, research and critical cases. None is obsolete; choose by constraints.

Library	Speed	Accuracy	Best for
NLTK	Slowest	Lower	Education, linguistics
spaCy	Fast	Good	Production pipelines
Flair	Slowest	Highest	Research, accuracy-critical

Exam Tip: spaCy uses shallow neural nets, not BERT, for default POS/NER. Flair is slower but more accurate.

> **Tool comparison** = Trainee vs field agent vs specialist — match tool to constraints.

PART 5: THE CASE ARCHIVES (Week 4)

Text Corpora Overview — The Evidence Vault

The Story: Detectives don't work from single clues — they study **case archives** (corpora). A corpus is a collection of documents whose statistics shape every model trained on it. Model behaviour is a function of corpus statistics: \text{Model behaviour} = f(\text{corpus statistics}). Large does not mean unbiased. Always explore before modelling.

Exam Tip: Corpus ≠ labelled dataset. Single PDF ≠ corpus. Explore vocabulary and distribution before training.

> **Corpus linguistics** = The evidence vault — archives that shape every model.

What Is a Text Corpus? — Building the Archive

The Story: A **text corpus** is a structured collection of documents — books, articles, tweets, clinical notes. Models trained on a corpus reproduce its statistics. TF-IDF weights depend on document frequency across the archive: \text{TF-IDF}(t, d) \propto \text{count}(t \text{ in } d) \times \log\frac{N}{\text{df}(t)}. Garbage in, garbage out — the archive defines what the model "knows."

Exam Tip: Corpus ≠ dataset (which implies labels). Models inherit corpus biases and vocabulary.

> **Text corpus** = Case archive — documents whose statistics define model behaviour.

Types of Corpora — Filing the Archives

The Story: Archives come in different categories. **General corpora** (Gutenberg, Brown) cover broad language. **Domain-specific** corpora focus on medicine, law, finance. **Annotated corpora** (CoNLL) have human labels for training. **Task-specific** corpora target one NLP task. Training medical NER on Gutenberg fails — the vocabulary and entities differ completely.

Type	Example	Use
General	Gutenberg, Brown	Broad language study
Domain	PubMed, legal filings	Specialised models
Annotated	CoNLL NER	Supervised training
Task-specific	Sentiment reviews	Single-task models

Exam Tip: Don't train domain models on general literary corpora. Gutenberg is archaic literary English.

> **Corpus types** = Filing system — general, domain, annotated, task-specific archives.

Bias and Representation in Corpora — Skewed Archives

The Story: Archives are not neutral. A corpus over-representing one demographic produces models that fail on others: \text{Model bias} \leftarrow \text{Corpus bias}. Scale can amplify bias — bigger archives with skewed content make worse stereotypes. Annotation bias in labels propagates to predictions. Examine training data before blaming the algorithm.

Exam Tip: Examine corpus composition first. Scale amplifies existing bias. Annotation choices create label bias.

> **Corpus bias** = Skewed archives — models inherit what the vault contains.

Corpus Exploration — Gutenberg and Brown Analysis

The Story: Before opening a case, detectives survey the archive. **Types** (unique words) vs **tokens** (total word count) differ: \text{TTR} = \frac{|\text{types}|}{|\text{tokens}|}. Gutenberg reveals literary vocabulary and word frequency. **Brown corpus** compares genres (news, fiction, academic) — showing language varies by domain. TTR is not comparable across very different text lengths.

Key metrics: type-token ratio (TTR), word frequency distribution, genre comparison.

Exam Tip: Types ≠ tokens. TTR varies with text length. Brown corpus is 1960s American English — dated but genre-diverse.

> **Corpus exploration** = Survey the archive — types, tokens, frequency, genre.

PART 6: WEIGHING THE CLUES (Week 5)

Traditional Word Representations Overview — The Scoring Lab

The Story: Machines need numbers, not words. The Scoring Lab converts text to vectors. **Vectorisation** is not preprocessing — it happens after cleaning. "Embedding" broadly includes sparse methods (one-hot, BoW, TF-IDF) and dense methods (Word2Vec). Methods are not interchangeable; each makes different trade-offs.

Exam Tip: Vectorisation ≠ preprocessing. Sparse and dense embeddings serve different purposes. Know when each applies.

> **Word representations** = Scoring lab — turn words into numbers machines can compute.

One-Hot Encoding — The Identity Badge

The Story: Each word gets an identity badge — a vector of zeros with a single one. Vocabulary size |V| determines vector length. Document vector marks which words appear:

x_i =

$$\begin{cases} 1 & \text{if word } w_i \text{ appears in } d \\ 0 & \text{otherwise} \end{cases}$$

, \quad \mathbf{x}_d \in \{0, 1\}^{|V|}
Every word is orthogonal — *dog* and *puppy* have zero similarity. No word order, no frequency beyond presence.

Exam Tip: One-hot ≠ BoW (which counts frequency). No semantic similarity: \text{dog} \cdot \text{puppy} = 0.

> **One-hot encoding** = Identity badge — one slot per word, all others zero.

One-Hot Encoding Implementation — Printing the Badges

The Story: In code, build a sorted vocabulary list, then for each word in a sentence, set the corresponding index to 1. Lookup is O(n) with `list.index`. Case sensitivity matters — *The* and *the* are different badges. Sort vocabulary for deterministic output across runs.

Exam Tip: Case sensitivity creates duplicate vocabulary entries. Sort vocabulary for reproducibility.

> **One-hot implementation** = Print badges — vocabulary index lookup, watch case sensitivity.

Advantages and Limitations of One-Hot Encoding — Badge Pros and Cons

The Story: Identity badges are simple and interpretable but wasteful. Sparsity ratio \approx 1 - \frac{\text{words per sentence}}{|V|} — most entries are zero. High dimensionality grows with vocabulary. Orthogonality means linguistically related words appear unrelated. One-hot loses word order, frequency, and meaning.

Exam Tip: One-hot is not lossless — order, frequency, and semantics are all discarded. Orthogonality ≠ linguistic independence.

> **One-hot trade-offs** = Simple badges — interpretable but sparse, high-dimensional, no semantics.

Bag of Words — The Evidence Count

The Story: Instead of just marking presence, count how many times each word appears — a **multiset** of words. BoW preserves multiplicity but discards order: *"the man bit the dog"* and *"the dog bit the man"* produce identical vectors. Case and stemming choices affect counts. Still a sparse, high-dimensional representation.

Exam Tip: BoW preserves count, not order. Same words different order = same vector. Case and stemming matter.

> **Bag of Words** = Evidence count — word frequencies in a multiset, order discarded.

BoW Implementation with scikit-learn — Automated Counting

The Story: scikit-learn's `CountVectorizer` automates the counting pipeline. Fit on training data only to prevent data leakage. `binary=True` gives presence/absence, not frequency. BoW produces integer matrices; TF-IDF produces float weights — different downstream behaviour.

Exam Tip: Fit vectoriser on train only. `binary=True` ≠ frequency BoW. BoW = integers, TF-IDF = floats.

> **BoW implementation** = Automated counter — CountVectorizer, fit on train only.

Advantages and Limitations of Bag of Words — Count Pros and Cons

The Story: BoW is a workhorse — simple, fast, competitive for keyword-heavy tasks like spam detection. But it fails on negation (*not good* looks like *good*), ignores word order, and treats all non-zero counts equally without rarity weighting. Still widely used as a baseline.

Exam Tip: BoW fails on negation. Competitive for keyword tasks but lacks semantics and order.

> **BoW trade-offs** = Workhorse counter — fast baseline, blind to order and negation.

TF-IDF — Weighting Clues by Rarity

The Story: Not all clues are equal. **Term Frequency (TF)** asks: how often does this word appear in *this* document? **Inverse Document Frequency (IDF)** asks: how rare is it across the *entire archive*?

\text{TF}(x, d) = \text{count of } x \text{ in } d
\text{IDF}(x) = \log\frac{N}{\text{DF}(x)}
\text{TF-IDF}(x, d) = \text{TF}(x, d) \times \log\frac{N}{\text{DF}(x)}
Word in every document → IDF = 0 → downweighted. Rare domain term → high IDF → boosted.

Exam Tip: TF is per-document; IDF is corpus-wide. Log is required. TF-IDF does not capture semantics (*car* ≠ *automobile*).

> **TF-IDF** = Rarity-weighted clues — frequent locally, rare globally = high score.

TF-IDF Implementation — The Scoring Machine

The Story: scikit-learn's `TfidfVectorizer` builds sparse TF-IDF matrices. Sublinear TF scaling uses 1 + \log(\text{TF}). Smoothing parameter k prevents zero division. Fit on training data only. Same word gets different scores in different documents based on local TF.

Exam Tip: Fit on train only. Same word, different TF-IDF per document. Sublinear scaling dampens high frequencies.

> **TF-IDF implementation** = Scoring machine — TfidfVectorizer, sparse float matrix.

Advantages and Limitations of TF-IDF — Best Sparse Representation

The Story: TF-IDF is the best sparse representation for classification and retrieval — automatically downweighting *the*, *is*, *and* while boosting domain terms like *mitochondria*. But it remains high-dimensional, statistical (not semantic), and blind to word order. TF-IDF ≠ Word2Vec — different paradigm entirely.

Exam Tip: TF-IDF ≠ Word2Vec. High dimensionality remains. Statistical weighting, not semantic understanding.

> **TF-IDF trade-offs** = Best sparse scorer — auto-downweights common words, still no semantics.

PART 7: SEMANTIC MAPS (Week 6)

Dense Embeddings Overview — The Cartography Division

The Story: Sparse vectors tell you *which* words appear; dense embeddings tell you *what they mean*. The Cartography Division maps words into continuous space where similar meanings cluster. The famous analogy: $\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}} \approx \mathbf{v}_{\text{queen}}$. Word2Vec and GloVe produce **static** embeddings — one vector per word. BERT produces **contextual** ones.

Exam Tip: "Embedding" includes TF-IDF broadly. Word2Vec/GloVe are static; dense ≠ always better than sparse.

> **Dense embeddings** = Semantic maps — words as points in meaning-space.

Word2Vec — Predict to Compress

The Story: Word2Vec trains a shallow neural network on a prediction game. **CBOW** (Continuous Bag of Words): given context words, predict the center — fill in the blank. **Skip-gram**: given center word, predict context — opposite direction. To solve prediction efficiently, the network compresses meaning into hidden-layer vectors: $\mathbf{v}_w \in \mathbb{R}^{d}, d \in [50, 300]$.

Architecture	Direction	Speed	Rare words
CBOW	Context → target	Faster	Less effective
Skip-gram	Target → context	Slower	Better

Exam Tip: CBOW = many → one; Skip-gram = one → many. Shallow network, not deep. Learns by prediction, not co-occurrence counting (that's GloVe).

> **Word2Vec** = Prediction game — CBOW fills blanks, Skip-gram predicts neighbours.

Word2Vec Implementation — Training Your Own Map

The Story: Gensim's `Word2Vec` takes `List[List[str]]` — pre-tokenised sentences. Lowercase for consistency. `min_count` filters rare words — on small corpora, lower it. Pretrained models like `word2vec-google-news-300` provide 3M words in 300 dimensions ready for cosine similarity.

Exam Tip: Input must be list of token lists. Lowercase. `min_count` on small corpora. Pretrained models avoid training from scratch.

> **Word2Vec implementation** = Train the map — Gensim, tokenised sentences, min_count tuning.

GloVe — Global Co-occurrence Statistics

The Story: While Word2Vec learns locally from prediction windows, **GloVe** (Global Vectors) factorises a global co-occurrence matrix. X_{ij} = times word j appears in context of word i. The model learns: $\mathbf{w}_i^T \mathbf{w}_j \approx \log X_{ij}$. Matrix factorisation, not a prediction neural network. Still produces static embeddings.

Exam Tip: GloVe ≠ neural network — it's matrix factorisation on global statistics. Global vs Word2Vec's local windows.

> **GloVe** = Global statistics map — factorise co-occurrence matrix, not predict words.

GloVe Implementation — Loading Pretrained Maps

The Story: Gensim loads pretrained GloVe vectors (e.g., glove-wiki-gigaword). No training from scratch in standard workflow. Out-of-vocabulary words raise KeyError — handle gracefully. Small demo models (25-dim) are for illustration only; production uses 100–300 dimensions.

Exam Tip: Gensim loads pretrained only. OOV → KeyError. Small models for demos, not production.

> **GloVe implementation** = Load pretrained maps — Gensim, handle OOV gracefully.

Visualising Word Embeddings — Reading the Map

The Story: Semantic maps live in hundreds of dimensions — humans need 2D projections. **PCA** preserves global structure; **t-SNE** preserves local clusters better but distorts distances. Cosine similarity in original space is the reliable metric; 2D plots are for intuition only. Need 20+ words for meaningful visualisation. Never mix models in one plot.

Exam Tip: 2D distances are distorted — trust cosine similarity in original space. Don't mix embedding models.

> **Embedding visualisation** = Flatten the map — PCA/t-SNE for intuition, cosine for truth.

Static vs Contextual Embeddings — Fixed vs Adaptive Maps

The Story: Word2Vec and GloVe assign one fixed coordinate per word — *bank* gets the same vector whether river or financial. **Contextual embeddings** (ELMo, BERT) change the vector based on surrounding words. Vector arithmetic works on static embeddings: \mathbf{r}_{\text{gender}} = \mathbf{v}_{\text{woman}} - \mathbf{v}_{\text{man}}. Polysemy breaks static maps; context saves them.

Exam Tip: Word2Vec/GloVe = static (one vector per word). BERT = contextual. Vector arithmetic demonstrated on static embeddings only.

> **Static vs contextual** = Fixed map vs GPS that recalculates per sentence.

PART 8: DETECTIVES WITH MEMORY (Week 7)

Sequential Modelling Overview — Why Order Matters

The Story: A bag of evidence loses sequence — who spoke first, who reacted second. Language is a **sequence**, not a bag. TF-IDF vectors have fixed size regardless of sentence length. BoW destroys word order. Sequential models read evidence one piece at a time, carrying memory forward. Two problems to solve: variable-length input and preserved order.

Exam Tip: TF-IDF size is fixed per vocabulary, not per sentence length. BoW loses order — two problems drive sequential models.

> **Sequential modelling** = Detectives with notebooks — language is ordered, not a bag.

Sequence-to-Sequence Modelling — Translator Pairs

The Story: Some cases require transforming one sequence into another — translation, summarisation, dialogue. **Seq2Seq** uses an **encoder** that reads the input into a **context vector** and a **decoder** that generates output token by token. The context vector is an information bottleneck — long sequences lose detail. Attention (Part 9) solves this bottleneck.

Exam Tip: Encoder reads, decoder generates. Bottleneck limits long sequences. Attention addresses the bottleneck.

> **Seq2Seq** = Translator pair — encoder compresses, decoder generates.

Recurrent Neural Networks — The Memory Notebook

The Story: An RNN detective carries a **hidden state** notebook — at each time step, read new evidence X_t, update memory H_t, produce output Y_t. Weights are **shared across time steps** — same detective rules at every position.

H_t = \tanh(W_{hh} H_{t-1} + W_{xh} X_t + b_h)
Y_t = W_{hy} H_t + b_y
H_t = memory, Y_t = output. Activation is \tanh, not sigmoid.

Exam Tip: Weights shared across time. H_t is memory, Y_t is output. Vanishing gradients limit long-range memory.

> **RNN** = Memory notebook — hidden state updated at each time step.

RNN Architecture Patterns — Case Shapes

The Story: Same detective, different case shapes. **Many-to-one**: read entire sequence, one verdict (sentiment). **One-to-many**: one clue, generate sequence (image captioning). **Many-to-many**: input sequence → output sequence (translation). Sentiment is many-to-one, not one-to-many.

Pattern	Input	Output	Example
Many-to-one	Sequence	Single label	Sentiment
One-to-many	Single input	Sequence	Captioning
Many-to-many	Sequence	Sequence	Translation

Exam Tip: Sentiment = many-to-one. Translation = many-to-many (not synced step-by-step in basic form).

> **RNN patterns** = Case shapes — many-to-one, one-to-many, many-to-many.

Vanishing and Exploding Gradients — The Broken Telephone

The Story: Training an RNN is the telephone game — gradients multiply through time steps. If each multiplier < 1, the signal vanishes (0.9^{100} \approx 0). If > 1, it explodes (1.1^{100} \approx 13{,}780). Long sequences forget early evidence.

\frac{\partial \mathcal{L}}{\partial H_0} = \frac{\partial \mathcal{L}}{\partial H_T} \prod_{t=1}^{T} \frac{\partial H_t}{\partial H_{t-1}}
Gradient clipping caps exploding gradients: \mathbf{g} \leftarrow \frac{\text{threshold}}{\|\mathbf{g}\|} \cdot \mathbf{g} if \|\mathbf{g}\| > \text{threshold}. Clipping fixes exploding, not vanishing.

Exam Tip: Clipping fixes exploding, not vanishing. Solution for vanishing = LSTM/GRU gates.

> **Vanishing gradients** = Broken telephone — gradients fade or explode across time steps.

LSTM Networks — The Gated Memory Chip

The Story: LSTM detectives carry two notebooks: **cell state** C_t (long-term memory) and **hidden state** H_t (working memory). Three gates control flow:

f_t = \sigma(W_f \cdot [H_{t-1}, X_t] + b_f) \quad \text{(forget gate)}
i_t = \sigma(W_i \cdot [H_{t-1}, X_t] + b_i) \quad \text{(input gate)}
\tilde{C}_t = \tanh(W_C \cdot [H_{t-1}, X_t] + b_C) \quad \text{(candidate)}
o_t = \sigma(W_o \cdot [H_{t-1}, X_t] + b_o) \quad \text{(output gate)}
H_t = o_t \odot \tanh(C_t)
Forget gate uses sigmoid; cell update uses tanh. Three gates: forget, input, output.

Exam Tip: C_t = long-term, H_t = short-term. Forget gate = sigmoid, not tanh. Three gates, not two.

> **LSTM** = Gated memory chip — cell state with forget, input, output gates.

GRU — The Lean Memory Unit

The Story: GRU is LSTM's leaner sibling — no separate cell state, only two gates. **Update gate** z_t decides how much past memory to keep. **Reset gate** r_t decides how much past to forget when computing new candidate.

z_t = \sigma(W_z \cdot [H_{t-1}, X_t]) \quad \text{(update gate)}
r_t = \sigma(W_r \cdot [H_{t-1}, X_t]) \quad \text{(reset gate)}
H_t = (1 - z_t) \odot H_{t-1} + z_t \odot \tilde{H}_t
Fewer parameters, often comparable performance on smaller datasets.

Exam Tip: No separate cell state. Two gates (update, reset), not three. Often matches LSTM on smaller data.

> **GRU** = Lean memory unit — two gates, no separate cell state.

PART 9: THE SPOTLIGHT TEAM (Week 8)

From Sequential Models to Transformers — Breaking the Sequential Bottleneck

The Story: RNN detectives process evidence one piece at a time — slow, hard to parallelise, and LSTMs still struggle with very long sequences. The Spotlight Team processes **all tokens simultaneously** via self-attention. Parallel batching across sequences ≠ within-sequence parallelism. Positional information must still be injected explicitly.

Exam Tip: LSTMs don't solve all long-sequence problems. Transformers enable within-sequence parallelism. Positional encoding is mandatory.

> **Transformer motivation** = Spotlight team replaces sequential notebook — all tokens at once.

Attention Is All You Need — The 2017 Breakthrough

The Story: Before 2017, attention was an add-on to seq2seq models. The landmark paper made attention the entire architecture — no RNN needed. GPT is a decoder-only Transformer, not an LSTM. The original Transformer has both encoder and decoder stacks. Attention existed before 2017 but was not the core.

Exam Tip: GPT = decoder-only Transformer. Original Transformer = encoder + decoder. Attention predates 2017 but became central in Transformers.

> **Attention paper** = 2017 breakthrough — attention replaces recurrence entirely.

Transformer Architecture — The Spotlight Network

The Story: Every token shines a spotlight on every other token. **Self-attention** computes relevance scores; high scores dominate the final representation. Coreference (*it → animal*) is the canonical example. The attention equation:

\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V

Symbol	Role
Q (Query)	"What am I looking for?"
K (Key)	"What do I contain?"
V (Value)	"What information do I pass?"

Multi-head attention runs parallel specialised heads (syntax, semantics, coreference). **Positional encoding** (sin/cos waves) injects word order — without it, parallel processing loses sequence information.

Exam Tip: Don't omit \sqrt{d_k}. Encoder self-attention is bidirectional; decoder uses masking. Positional encoding is mandatory.

> **Transformer** = Spotlight network — every token attends to every other token.

Transformer Model Families — Three Detective Ranks

The Story: One architecture spawns three specialist ranks. **Encoder-only** (BERT, RoBERTa): read and understand — NLU, classification. **Decoder-only** (GPT, Gemini, Llama): generate autoregressively — NLG. **Encoder-decoder** (T5, BART): transform sequences — translation, summarisation.

Family	Examples	Best for
Encoder-only	BERT, RoBERTa	NLU, classification
Decoder-only	GPT, Gemini	NLG, chat
Encoder-decoder	T5, BART	Seq2seq tasks

Exam Tip: GPT for NER lacks left context in generation mode. BERT cannot generate paragraphs. T5/BART need both halves.

> **Transformer families** = Three ranks — encoders read, decoders write, both transform.

Hugging Face Hub — The Detective Supply Depot

The Story: Hugging Face is the agency's supply depot — pretrained models, datasets, tokenisers, and pipelines in one hub. Every model has a **model card** with license, bias notes, and limitations. Check commercial license before production. Not NLP-only — vision and audio models too.

Exam Tip: Check license for commercial use. Read bias/limitations on model cards. Hugging Face covers multiple ML domains.

> **Hugging Face** = Supply depot — models, tokenisers, pipelines, model cards.

Hugging Face Token and Model Implementation — Field Deployment

The Story: Access tokens authenticate API calls — never hardcode `hf_` tokens in notebooks committed to git. Tokeniser must match model checkpoint exactly. `truncation=True` for 512-token BERT limit. Pipeline task must match architecture (classification pipeline on BERT, not GPT).

Exam Tip: Never commit tokens. Match tokeniser to checkpoint. `truncation=True` for length limits. Task must match architecture.

> **HF implementation** = Field deployment — tokens, matching tokenisers, truncation.

PART 10: SENIOR DETECTIVE BERT (Week 9)

BERT and Contextual Embeddings — The Senior Analyst

The Story: Junior detectives (Word2Vec) assign one identity per word. **Senior Detective BERT** re-evaluates every word in context — *bank* gets different vectors in "river bank" vs "bank account." BERT is encoder-only Transformer, not GPT. Fine-tuning starts from pretrained weights, not from scratch.

Exam Tip: BERT embeddings are contextual. BERT ≠ GPT (encoder vs decoder). Fine-tuning uses pretrained weights.

> **BERT introduction** = Senior analyst — contextual embeddings from encoder-only Transformer.

BERT Architecture and Pre-Training — Two Training Drills

The Story: BERT trains on two self-supervised drills. **MLM (Masked Language Modelling):** hide ~15% of tokens, predict using bidirectional context — both left and right neighbours. **NSP (Next Sentence Prediction):** given sentences A and B, predict if B truly follows A. Input format: `[CLS] sentence_A [SEP] sentence_B [SEP]`.

Task	What it teaches
MLM	Token meaning from full context
NSP	Sentence-pair coherence

BERT-base = 12 layers, 768 dims. BERT-large = 24 layers. RoBERTa later removed NSP.

Exam Tip: MLM is bidirectional; GPT is left-to-right only. BERT-base = 12 layers, not 24. RoBERTa dropped NSP.

> **BERT pre-training** = Two drills — MLM (bidirectional fill-in) and NSP (sentence pairs).

BERT Applications — Case Specialisations

The Story: Senior Detective BERT excels at understanding tasks: sentiment classification (using `[CLS]` token), NER (with subword merging), extractive QA (span prediction), and natural language inference. BERT is not for open-ended generation — that's the press office's job (GPT).

Exam Tip: BERT not for open-ended generation. Subword tokens need merging for NER spans. QA is extractive, not generative.

> **BERT applications** = Understanding specialist — sentiment, NER, QA, NLI.

BERT Variants — Specialised Units

The Story: The agency deploys specialised BERT units. **RoBERTa:** more data, no NSP, better training recipe. **ALBERT:** parameter sharing for efficiency. **DistilBERT:** ~40% smaller, ~60% faster, ~97% of BERT performance via knowledge distillation. **Domain fine-tunes:** BioBERT, FinBERT for specialised corpora.

Exam Tip: RoBERTa changes training, not architecture entirely. Distillation ≠ fine-tuning. Domain models need domain data.

> **BERT variants** = Specialised units — RoBERTa, ALBERT, DistilBERT, domain fine-tunes.

PART 11: MOOD ANALYSIS DIVISION (Week 10)

Text Classification Overview — Sorting the Case Files

The Story: The Mood Analysis Division sorts incoming case files into categories: y \in \{c_1, c_2, \ldots, c_k\} — mapping text x to label y. Sentiment is one type of text classification. Class imbalance makes accuracy misleading; use precision, recall, F1. Topic modelling discovers themes; topic labelling assigns names — different jobs.

Exam Tip: Sentiment is text classification. Class imbalance distorts accuracy. Topic modelling ≠ topic labelling.

> **Text classification** = Sort case files — map text to discrete category labels.

Sentiment Analysis — Reading the Room

The Story: Sentiment analysis reads emotional tone at three levels: **document** (whole review), **sentence** (one statement), **aspect-based** (ABSA — sentiment toward a specific feature). Labels: positive, negative, neutral. Traps abound: *"not bad"* is positive, sarcasm breaks rules, *long* is context-dependent. Subjectivity (opinion vs fact) ≠ polarity (positive vs negative).

Exam Tip: "Not bad" = positive. Sarcasm breaks rule-based systems. Subjectivity ≠ polarity.

> **Sentiment analysis** = Reading the room — document, sentence, or aspect-level tone.

Sentiment Implementations — VADER, BERT, Flair, spaCy

The Story: Four mood-readers, four philosophies. **VADER**: rule/lexicon, compound score \in [-1, 1], thresholds at \pm 0.05 — fast, explainable, social-media tuned. **BERT**: contextual deep learning via Hugging Face pipeline — accurate on complex language, 512 token limit. **Flair**: neural tagger with `Sentence` + `predict()`. **spaCy/TextBlob**: pattern-based polarity [-1, 1] and subjectivity [0, 1].

Contrast: *"I usually hate waiting, but this was worth it"* → VADER: neutral; BERT/Flair: positive; TextBlob: negative.

Exam Tip: Use VADER `compound`, not `pos` alone. BERT needs `truncation=True`. Flair requires `predict()`.

> **Sentiment tools** = Four mood-readers — VADER (rules), BERT (context), Flair (neural), TextBlob (patterns).

Comparing Sentiment Approaches — Choosing the Mood Reader

The Story: Use **VADER** when speed, cost, and explainability matter — social media, real-time dashboards. Use **BERT** when accuracy on complex language matters — negation, contrast, domain nuance. BERT is not always better; regulated industries may require explainable VADER scores. Black-box constraints are a real production concern.

Exam Tip: BERT not always better. VADER is ML-adjacent, not "non-computational." Match tool to constraints.

> **Sentiment comparison** = Match mood-reader to case — speed vs accuracy vs explainability.

PART 12: THEME CLUSTERING UNIT (Week 11)

Introduction to Topic Modelling — Finding Hidden Themes

The Story: The Theme Clustering Unit discovers hidden themes in unsorted case archives — no labels needed. A **topic** is a probability distribution over words; a **document** is a mixture of topics. Topic modelling ≠ classification. Documents are mixtures, not single assignments. Word order is ignored (bag-of-words assumption).

Exam Tip: Topic modelling is unsupervised. Documents are mixtures, not single-topic. Word order ignored.

> **Topic modelling** = Theme clustering — discover hidden topics as word distributions.

Latent Dirichlet Allocation — The Probabilistic Theme Finder

The Story: LDA is the classic theme finder. Two core ideas: (1) documents are **mixtures** of topics — a tech-health article might be 70% technology, 30% healthcare; (2) topics are **distributions over words** — technology weights *software*, *algorithm*; healthcare weights *patient*, *treatment*. Output is proportions like (0.70, 0.30), not a single label. Number of topics K must be set manually.

Exam Tip: LDA assigns proportions, not one topic per document. K is manual. Topics are unlabelled — interpret via top words.

> **LDA** = Probabilistic theme finder — documents as topic mixtures.

Key Assumptions of LDA — The Rulebook

The Story: LDA detectives follow strict rules: bag-of-words input (no word order), documents are mixtures of topics, topics are distributions over words, fixed K topics, and **exchangeability** (word order doesn't matter). Short text (tweets) violates the mixture assumption. Unsupervised ≠ no preprocessing —

stopword removal is critical.

Exam Tip: No context preserved. Short text violates assumptions. Unsupervised still needs preprocessing.

> **LDA assumptions** = Strict rulebook — bag-of-words, mixtures, fixed K, no word order.

LDA Implementation — Running the Theme Finder

The Story: Gensim pipeline: preprocess (tokenise, remove stopwords) → build BoW dictionary → `LdaModel(corpus, num_topics=K)`. Inspect top words per topic for interpretation. Proportions are not classification probabilities — they are topic mixture weights.

Exam Tip: Set `num_topics=K` explicitly. Proportions ≠ classification probabilities. Stopword removal critical.

> **LDA implementation** = Run the theme finder — Gensim, BoW dictionary, tune K.

Why Alternatives to LDA Are Needed — When Classic Fails

The Story: LDA hits a ceiling: no semantic understanding (*car* ≠ *automobile*), poor on short text (tweets), manual K tuning, incoherent topics on sparse data. Alternatives: **GSDMM** for short text (one-topic-per-document assumption), **BERTopic** for semantic clustering.

Exam Tip: LDA fails on short text and lacks semantics. Know when to switch to GSDMM or BERTopic.

> **LDA limitations** = Classic theme finder hits ceiling — short text, no semantics, manual K.

BERTopic — Semantic Theme Clustering

The Story: BERTopic is not "LDA with BERT." Pipeline: embed documents (sentence transformer) → **UMAP** dimensionality reduction → **HDBSCAN** clustering (no preset K) → **c-TF-IDF** for interpretable topic words → optional MMR for diversity. Embedding model choice matters significantly.

Exam Tip: Not "LDA with BERT." c-TF-IDF uses IDF per cluster, not global. HDBSCAN discovers topic count.

> **BERTopic** = Semantic theme clustering — embed, cluster, extract words.

BERTopic Architecture — Five-Stage Pipeline

The Story: The full BERTopic pipeline has five stages: (1) sentence embeddings, (2) UMAP reduction (non-linear, preserves local structure), (3) HDBSCAN clustering (automatic K), (4) c-TF-IDF topic representation (cluster-level IDF), (5) optional MMR for diverse topic words. UMAP vs PCA: UMAP is non-linear and better for clustering.

Exam Tip: Five stages, not three. HDBSCAN discovers topic count automatically. UMAP is non-linear.

> **BERTopic architecture** = Five-stage pipeline — embed, UMAP, HDBSCAN, c-TF-IDF, MMR.

GSDMM for Short Text — The Tweet Specialist

The Story: GSDMM (Gibbs Sampling Dirichlet Multinomial Mixture) assumes **one topic per document** — perfect for short text like tweets where LDA's mixture assumption fails. Uses Gibbs sampling with hyperparameters \alpha and \beta. Stemming helps on short text. Choose GSDMM over LDA when documents are very short.

Exam Tip: GSDMM for tweets, not LDA. One-topic-per-document vs mixture is the key difference. Stemming helps.

> **GSDMM** = Tweet specialist — one topic per short document via Gibbs sampling.

PART 13: THE PRESS OFFICE (Week 12)

Why Decoder-Only Models Dominate NLG — The Press Release Machine

The Story: The press office writes fluent releases one word at a time: $P(t_n \mid t_1, t_2, \ldots, t_{n-1})$. **Decoder-only** models (GPT, Gemini) excel because autoregressive generation is their native training objective. BERT cannot write press releases — it reads, not writes. No planning ahead — each token conditions on all prior tokens.

Exam Tip: BERT not for open-ended NLG. Autoregressive = one token at a time. Decoder-only dominates generation.

> **Decoder-only NLG** = Press release machine — autoregressive, one token at a time.

What Is a Large Language Model? — The Senior Press Officer

The Story: An **LLM** is a neural network with billions of parameters (10^9+) pretrained on massive text via self-supervised next-token prediction — no labels needed for pre-training. **Fine-tuning** adapts to labelled tasks afterward. No fixed parameter threshold defines "large" — scale, data, and capability matter.

Exam Tip: No fixed parameter cutoff for "large." Pre-training is self-supervised (no labels). Fine-tuning uses labelled data.

> **LLM** = Senior press officer — billions of parameters, pretrained on next-token prediction.

Why LLMs Excel at NLG — Pattern Mastery

The Story: LLMs excel at NLG because they model $P(\text{token} \mid \text{context})$ over vast training data — learning grammar, style, facts, and reasoning patterns statistically. But they mimic patterns, not true reasoning. Prompt quality shapes output. Without RAG, they don't retrieve — they generate. Output is non-deterministic.

Exam Tip: Pattern mimicry, not guaranteed reasoning. Prompt quality matters. Fluency ≠ accuracy.

> **LLM NLG strength** = Pattern mastery — statistical fluency from massive pre-training.

Natural Language Generation — From Templates to Probabilistic Writing

The Story: NLG transforms structured data into human-readable text. Two paradigms: **deterministic** (templates, mail-merge) and **probabilistic** (LLMs sampling from learned distributions). The **prompt** is the control mechanism — it steers the LLM without changing weights. Prompts must be domain-adapted; generic prompts produce generic output.

Exam Tip: LLMs use statistical patterns, not explicit rules. Prompts must be domain-adapted.

> **NLG** = Probabilistic writing — templates replaced by learned language generation.

How LLMs Generate a Response — The Writing Process

The Story: Generation loop: tokenise input → embed tokens → pass through Transformer layers → softmax over vocabulary → sample next token → append and repeat until stop condition. Each token is generated, not retrieved. Single-turn has no memory of prior conversations. Knowledge cutoff limits factual currency. Fluency does not guarantee accuracy.

Exam Tip: Generated, not retrieved. No cross-turn memory in single-turn. Knowledge cutoff applies. Fluency ≠ accuracy.

> **LLM generation** = Writing loop — tokenise, embed, predict, sample, repeat.

Controlling LLM Output — The Editor's Knobs

The Story: The editor controls output via decoding strategies. **Temperature** T scales logits before softmax:

P(token_i) = \frac{\exp(z_i / T)}{\sum_j \exp(z_j / T)}
Low T (0–0.3) = deterministic, factual. High T (0.7–1.0) = creative, risky. **Top-K**: keep K highest-probability tokens. **Top-P (nucleus)**: keep smallest set whose cumulative probability exceeds p — adapts dynamically.

Exam Tip: Top-K ≠ Top-P. High temperature causes hallucination on factual tasks. T > 1 incoherent in production.

> **Decoding strategies** = Editor's knobs — temperature, Top-K, Top-P control randomness.

Choosing the Right Decoding Strategy — Match Settings to Task

The Story: No universal best settings. Code/JSON extraction: low temperature (0–0.3), Top-P. Creative writing: medium-high temperature, Top-P. Factual Q&A: low temperature, greedy or low Top-P. High temperature breaks JSON structure. Top-P is the industrial default for flexibility.

Task	Temperature	Strategy
Code/JSON	0–0.3	Top-P
Creative	0.7–1.0	Top-P
Factual Q&A	Low	Greedy or low Top-P

Exam Tip: No universal settings. High temperature breaks structured output. Top-P is production default.

> **Decoding selection** = Match editor settings to task — factual needs low T, creative needs high T.

Prompts and Prompt Engineering — The Briefing Document

The Story: A **prompt** is the briefing document sent to the press officer. Core components: **role** (who you are), **instruction** (what to do), **context/input** (the evidence), **output format** (JSON, bullets, prose), **constraints** (length, tone), and **examples** (shots). Role ≠ instruction. Bigger model does not eliminate need for good prompts.

Exam Tip: Role ≠ instruction. Specify output format explicitly. Good prompts required even for large models.

> **Prompt engineering** = Briefing document — role, instruction, context, format, constraints, examples.

Zero-Shot, One-Shot, and Few-Shot Prompting — Example Count

The Story: Zero-shot: no examples, just instruction. **One-shot**: one example in the prompt. **Few-shot**: several examples. "Shot" = example demonstrations, not conversation turns. Few-shot improves output format adherence but not guaranteed accuracy. Zero-shot weak for strict JSON schemas. Too many shots waste context window.

Exam Tip: Shot = examples, not turns. Few-shot helps format, not always accuracy. Zero-shot weak for strict JSON.

> **Shot prompting** = Example count — zero, one, or few demonstrations in the briefing.

Practical Prompt Engineering Guidelines — The Editor's Playbook

The Story: Prompting is experimental — iterate, test edge cases, state constraints explicitly. Balance creativity vs constraints. Unstated constraints are ignored. One-time prompt design fails in production. Document what works and what breaks.

Exam Tip: Iterate and test edge cases. Unstated constraints are ignored. Prompting is experimental, not one-shot design.

> **Prompt guidelines** = Editor's playbook — iterate, constrain, test edge cases.

Google AI Studio Playground — The Testing Desk

The Story: Google AI Studio is the testing desk for Gemini models — experiment with prompts before API integration. Watch knowledge cutoff dates. Temperature above 1.0 produces incoherent output. System instructions vs user messages serve different roles. API billing applies beyond free tier.

Exam Tip: Knowledge cutoff limits facts. Temperature > 1 incoherent. System vs user instruction roles differ.

> **Google AI Studio** = Testing desk — experiment with Gemini prompts before deployment.

LLM API Setup — Gemini and OpenAI Keys

The Story: API keys authenticate calls — never commit keys to git. Use `.env` files and `.gitignore`. OpenAI keys shown once at creation — store immediately. Gemini requires `genai.Client(api_key=...)`. ChatGPT app subscription ≠ API access — separate products.

Exam Tip: Never commit API keys. `.env` + `.gitignore`. OpenAI key shown once. ChatGPT app ≠ API.

> **API setup** = Secure the keys — .env files, never commit, separate products.

PART 14: CAPSTONE CASE — QUIZGENIUS (Week 13)

QuizGenius Capstone Project Overview — The Final Case

The Story: The capstone case: build **QuizGenius AI** — a quiz generator powered by Gemini API, Python backend, Streamlit frontend, JSON quiz schema. No hardcoded keys. Parse structured JSON, not free text. LLM-generated answers may be wrong — validate. Notebook proof-of-concept ≠ production application.

Stack: Gemini API, Python backend, Streamlit frontend, JSON schema.

Exam Tip: No hardcoded keys. Parse JSON, not free text. LLM answers need validation. Notebook ≠ production.

> **QuizGenius overview** = Final case — LLM-powered quiz generator with structured JSON output.

Package Installation and API Setup — Equipping the Team

The Story: Install dependencies from `requirements.txt`. Store API keys in Colab Secrets or `.env` — never in notebook cells. Validate key before API calls. Use Flash models for demos to control cost and latency.

Exam Tip: Colab Secrets, not notebook cells. Validate key first. Flash models for demos.

> **Project setup** = Equip the team — requirements, secrets, key validation.

Quiz Prompt Template Design — The Quiz Briefing

The Story: The prompt template is the briefing for quiz generation. Requirements: raw JSON only (no markdown fences), schema with `questions`, `options`, `correct_option`, placeholders for topic. Include a schema example in the prompt. Vague option format produces inconsistent output.

Exam Tip: Specify JSON schema in prompt. No markdown fences. Use placeholders for dynamic topic.

> **Prompt template** = Quiz briefing — JSON schema, options format, topic placeholder.

Quiz Generation Logic and Application Runner — The Engine Room

The Story: The engine room: `generate_quiz()` calls Gemini with `response_mime_type="application/json"`, extracts JSON via regex if needed, validates schema, checks answers case-insensitively. Handle `JSONDecodeError`. Use dynamic `len(questions)`, not hardcoded count. Key naming: `correct_option` must match schema exactly.

Exam Tip: Handle JSON parse errors. `correct_option` key naming matters. Dynamic question count.

> **Quiz generation logic** = Engine room — API call, JSON parse, schema validation, scoring.

Application Architecture — The Agency Floor Plan

The Story: User flow: User → Streamlit frontend → Backend (`quiz_engine.py`) → Gemini API → JSON response → UI scoring. LLM is external — not inside the backend box. Backend mediates API to protect keys. Session state preserves data across multi-screen flow (v1 notebook → v2 full-stack).

Exam Tip: LLM is external service. Backend protects API keys. Session state for multi-screen apps.

> **App architecture** = Agency floor plan — frontend, backend, external LLM, session state.

Project Setup and Repository Structure — Filing the Case

The Story: Repository structure: `app.py` (frontend), `quiz_engine.py` (backend), `.env` (secrets), `requirements.txt`, `.gitignore`. Never commit `.env`. Avoid monolithic single-file apps. Separate concerns for maintainability.

Exam Tip: Never commit `.env`. Include `requirements.txt`. Separate frontend and backend files.

> **Repository structure** = Filing the case — app.py, quiz_engine.py, .env, requirements.txt.

Streamlit Frontend — The Client Desk

The Story: Streamlit builds the client-facing desk. Key APIs: `st.session_state` preserves quiz state across reruns, `st.rerun()` refreshes after submission, dynamic question rendering from JSON. Launch with `streamlit run app.py`, not `python app.py`.

Exam Tip: Data lost without session state. Use `streamlit run app.py`, not `python app.py`.

> **Streamlit frontend** = Client desk — session state, dynamic rendering, correct launch command.

Running the Full-Stack Application — Opening for Business

The Story: Launch checklist: activate virtual environment, ensure `.env` with valid API key, run `streamlit run app.py`, test multiple topics. Verify JSON parsing and scoring across edge cases.

Exam Tip: Activate venv. `.env` required. Test multiple topics and edge cases.

> **Running the app** = Opening for business — venv, .env, streamlit run, multi-topic test.

Demo Notebook Quiz Application — The Proof of Concept

The Story: The notebook POC validates the core loop before building the full app. Verify LLM-generated answers manually — models hallucinate. Handle JSON parse failures gracefully. CLI notebook demo ≠ user-facing product.

Exam Tip: Verify LLM answers manually. Handle JSON failures. POC ≠ production product.

> **Notebook demo** = Proof of concept — validate loop before full-stack build.

From Notebook POC to Full-Stack Application — Leveling Up

The Story: Why upgrade beyond the notebook? Product wrapper matters — session state, error handling, clean UI, key protection. Streamlit is sufficient for portfolio projects vs React. The LLM capability is one layer; the application architecture is the deliverable.

Exam Tip: Product wrapper matters beyond LLM capability. Streamlit sufficient for portfolio. Architecture is the deliverable.

> **POC to production** = Level up — wrapper, error handling, and architecture matter.

Updated Project Architecture — The Revised Floor Plan

The Story: Version 2 architecture adds session state, separated backend, environment-based config, and improved error handling. Update architecture diagrams after implementation changes. Include `.env` and session state in documentation — they are not optional extras.

Exam Tip: Update diagrams after implementation. Document `.env` and session state in architecture.

> **Updated architecture** = Revised floor plan — v2 adds session state and separated concerns.

Implementation Overview and Integration — Closing the Case

The Story: Integration summary: prompt template → generation logic → JSON validation → Streamlit UI → scoring. Cross-cutting patterns: never hardcode keys, always validate LLM output, handle errors at every layer. Stopping at notebook POC misses the production lessons — error handling and validation are the examinable architecture skills.

Exam Tip: Don't stop at notebook POC. Error handling and validation at every layer. Integration patterns are examinable.

> **Integration summary** = Closing the case — prompt, generate, validate, display, score.

ONE-LINE SUMMARIES — The Complete Set

Course Overview = Agency charter: NLP turns unstructured language into structured signals. **NLP Module Introduction** = First briefing — language is complex evidence requiring structured pipelines. **NLP** = Core mission: bridge human communication and machine computation. **NLU / NLG** = Analysis wing vs press office inside the NLP building. **NLP Applications** = Case files across search, voice, email, keyboards, chatbots. **Morphology** = Forensic word dissection — roots, affixes, stems, lemmas. **Polysemy / Synonymy** = Red herrings — one word many meanings, many words one meaning. **Preprocessing** = Crime-scene cleaning without destroying evidence. **Tokenisation** = Bagging evidence into processable units. **Regex cleaning** = Fine brush for precise noise removal. **Stopwords** = Common footprints to filter when they add no clue. **Stemming / Lemmatisation** = Rough chop vs dictionary lookup for word forms. **Preprocessing pipeline** = Standard operating procedure — order matters. **Structural analysis** = Tagging division for POS and NER. **POS tagging** = Grammatical role labels — context-dependent. **POS in practice** = spaCy for speed, Flair for accuracy. **Dependency trees** = Evidence board for syntax and entities. **NER** = Suspect identification — persons, orgs, locations. **NER in practice** = spaCy for production, Flair for accuracy. **Tool comparison** = Trainee (NLTK) vs field agent (spaCy) vs specialist (Flair). **Corpus linguistics** = Evidence vault shaping every model. **Text corpus** = Case archive whose statistics define behaviour. **Corpus types** = General, domain, annotated, task-specific filing. **Corpus bias** = Skewed archives produce skewed models. **Corpus exploration** = Survey types, tokens, frequency, genre. **Word representations** = Scoring lab — words to numbers. **One-hot encoding** = Identity badge — one slot per word. **One-hot implementation** = Vocabulary index lookup, watch case sensitivity. **One-hot trade-offs** = Simple but sparse, high-dimensional, no semantics. **Bag of Words** = Evidence count — frequencies, order discarded. **BoW implementation** = CountVectorizer, fit on train only. **BoW trade-offs** = Fast baseline, blind to order and negation. **TF-IDF** = Rarity-weighted clues — local frequency × global rarity. **TF-IDF implementation** = TfidfVectorizer sparse float matrix. **TF-IDF trade-offs** = Best sparse scorer, still no semantics. **Dense embeddings** = Semantic maps in continuous space. **Word2Vec** = Prediction game — CBOW fills blanks, Skip-gram predicts neighbours. **Word2Vec implementation** = Gensim on tokenised sentences. **GloVe** = Global co-occurrence matrix factorisation. **GloVe implementation** = Load pretrained vectors, handle OOV. **Embedding visualisation** = PCA/t-SNE for intuition, cosine for truth. **Static vs contextual** = Fixed map vs GPS recalculating per sentence. **Sequential modelling** = Language is ordered, not a bag. **Seq2Seq** = Encoder compresses, decoder generates. **RNN** = Memory notebook updated each time step. **RNN patterns** = Many-to-one, one-to-many, many-to-many shapes. **Vanishing gradients** = Broken telephone — gradients fade or explode. **LSTM** = Gated memory chip with forget, input, output gates. **GRU** = Lean memory — two gates, no cell state. **Transformer motivation** = Spotlight team replaces sequential processing. **Attention paper** = 2017 — attention replaces recurrence. **Transformer** = Every token attends to every other token. **Transformer families** = Encoders read, decoders write, both transform. **Hugging Face** = Supply depot for models and pipelines. **HF implementation** = Tokens, matching tokenisers, truncation. **BERT introduction** = Senior analyst with contextual embeddings. **BERT pre-training** = MLM (bidirectional) + NSP (sentence pairs). **BERT applications** = Sentiment, NER, QA, NLI — not generation. **BERT variants** = RoBERTa, ALBERT, DistilBERT, domain fine-tunes. **Text classification** = Sort case files into category labels. **Sentiment analysis** = Read tone at document, sentence, or aspect level. **Sentiment tools** = VADER (rules), BERT (context), Flair, TextBlob. **Sentiment comparison** = Match tool to speed, accuracy, explainability. **Topic modelling** = Discover hidden themes as word distributions. **LDA** = Documents as topic mixtures, topics as word distributions. **LDA assumptions** = Bag-of-words, mixtures, fixed K, no word order. **LDA implementation** = Gensim pipeline with BoW dictionary. **LDA limitations** = No semantics, poor on short text, manual K. **BERTopic** = Embed, UMAP, HDBSCAN, c-TF-IDF pipeline. **BERTopic architecture** = Five stages with automatic topic count. **GSDMM** = One topic per short document via Gibbs sampling. **Decoder-only NLG** = Press release machine — autoregressive generation. **LLM** = Billions of parameters, next-token pre-training. **LLM NLG strength** = Statistical fluency from massive pre-training. **NLG** = Probabilistic writing steered by prompts. **LLM generation** = Tokenise, embed, predict, sample, repeat. **Decoding strategies** = Temperature, Top-K, Top-P control randomness. **Decoding selection** = Low T for facts, high T for creativity. **Prompt engineering** = Briefing with role, instruction, format, examples. **Shot prompting** = Zero, one, or few examples in the briefing. **Prompt guidelines** = Iterate, constrain, test edge cases. **Google AI Studio** = Testing desk for Gemini prompts. **API setup** = Secure keys in .env, never commit. **QuizGenius overview** = LLM quiz generator with JSON schema. **Project setup** = Requirements, secrets, key validation. **Prompt template** = JSON schema briefing for quiz generation. **Quiz generation logic** = API call, JSON parse, validation, scoring. **App architecture** = Frontend, backend, external LLM, session state. **Repository structure** = app.py, quiz_engine.py, .env, requirements.txt. **Streamlit frontend** = Session state, dynamic rendering, streamlit run. **Running the app** = venv, .env, multi-topic testing. **Notebook demo** = POC before full-stack build. **POC to production** = Wrapper and architecture matter beyond LLM. **Updated architecture** = v2 with session state and separated backend. **Integration summary** = Prompt, generate, validate, display, score.

Last compiled: 2026-08-01 | BITS Pilani — NLP and Understanding