Question 1 (10 marks)

You are working on a credit-risk prediction model for a bank. The dataset contains the following columns:

- **Loan_ID** - Unique identifier for each loan
- **Applicant_Income** - Monthly income of the borrower
- **Loan_Purpose** - Category describing why the loan was taken (17 unique categories)
- **Default_Flag** - Target variable (1 = default, 0 = non-default)

Answer the following:

Q1.1 (2 marks)

Should Loan_ID be included as a feature in the model? Explain your reasoning.

Q1.2 (3 marks)

Loan_Purpose contains 17 categories with no ordinal meaning.
(a) Should you use Ordinal Encoding? Why or why not?
(b) Suggest a more suitable encoding method and briefly describe how it works.

Q1.3 (3 marks)

Applicant_Income is extremely right-skewed.
(a) Name one transformation technique you can apply.
(b) Explain why this transformation helps the model.

Q1.4 (2 marks)

Give one reason why feature scaling is important if using Logistic Regression or KNN for credit risk.

Question 2 (10 marks)

Dataset: Product reviews from an e-commerce website

Review_ID	Review_Text	Stars
R101	“Delivered quickly but packaging was damaged.”	3

R102	“Excellent quality! Worth every rupee.”	5
R103	“Product stopped working after two days.”	1
R104	“Good value for money but not very durable.”	4
R105	“Terrible experience, completely disappointed.”	1

Goal: Predict the **Stars** rating (1–5) from the review text.

Q2.1 (6 marks)

List **three feature engineering techniques** to convert Review_Text into numeric representations for modelling.
Explain *why* each technique is useful.

Q2.2 (4 marks)

Explain why **dimensionality reduction** is often necessary when working with text features.
Name one suitable technique and describe how it works.

Question 3 (10 marks)

Dataset: IoT sensor data from a smart building

Timestamp	Room_ID	Temperature	Humidity	Motion_Flag
2025-07-01 08:00	R12	26.3	48	1
2025-07-01 08:05	R12	26.8	47	1
2025-07-01 08:10	R12	27.1	46	0
2025-07-01 08:15	R12	27.4	45	0
2025-07-01 08:20	R12	27.9	44	1

Goal: Predict **Temperature** at the next timestamp.

Q3.1 (5 marks)

Suggest **three time-derived features** you can extract from Timestamp.
For each, explain *why* it is useful for temperature prediction.

Q3.2 (5 marks)

Humidity and Motion_Flag may carry more information than appears.
Propose **two feature engineering strategies** to make these features more predictive and explain how they help.

Question 4 (10 marks)

Answer the following:

Q4.1 (5 marks)

Explain what feature engineering involves when working with **audio (.wav)** files.

Q4.2 (5 marks)

Is it possible to train a model directly on raw audio waveforms without explicit feature extraction?
Explain when this is possible and when it is not.